




Introduction to AI Platform Pipelines

Nitin Aggarwal
Technical PM, Cloud AI




Agenda



- Overview
- Introduction to AI Platform Pipelines
- AI Platform Pipelines: Concepts
- AI Platform Pipelines: When to Use?
- AI Platform Pipelines: Ecosystem



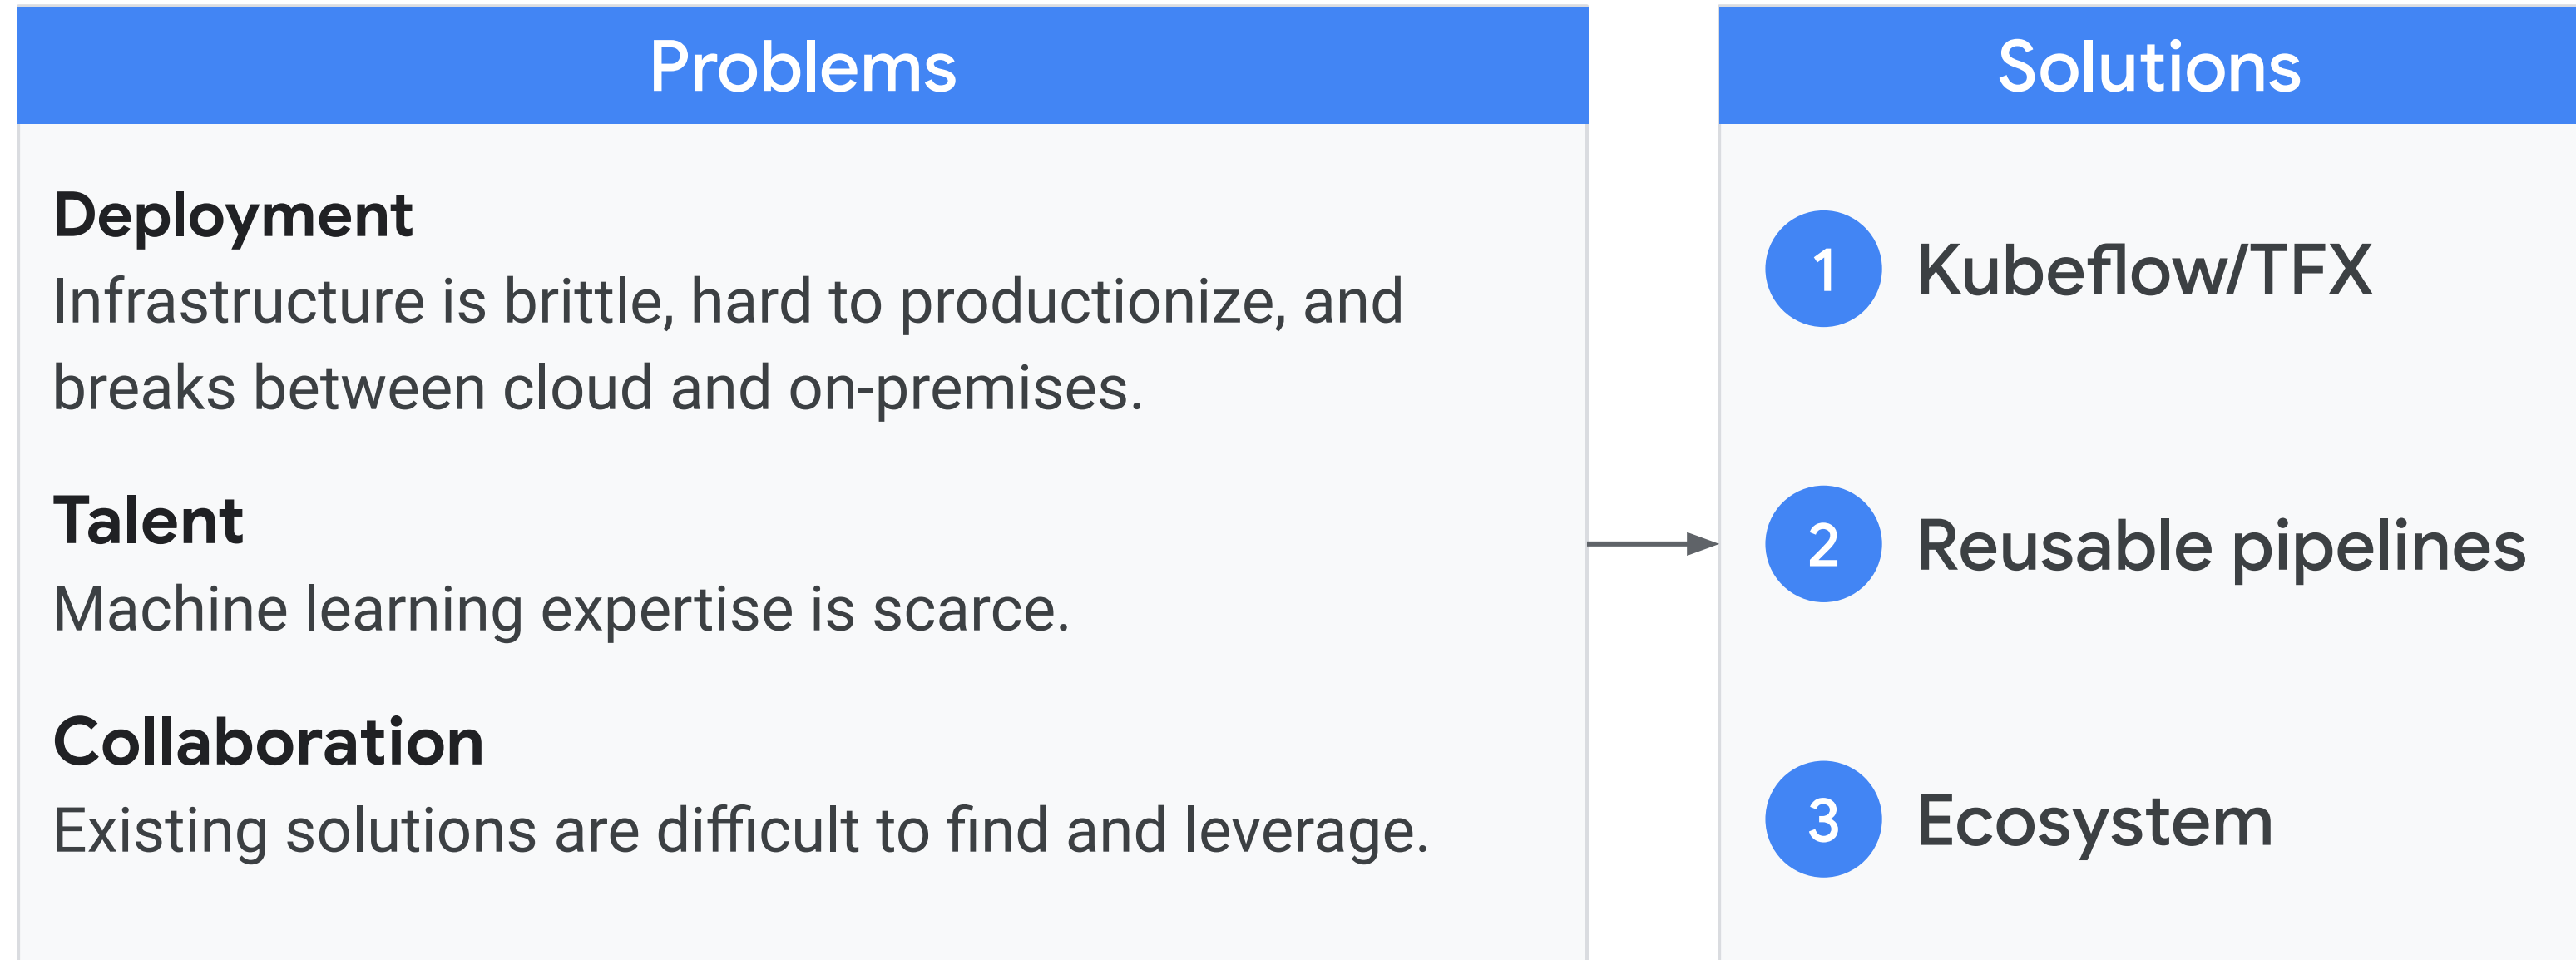
Agenda

- Overview
- Introduction to AI Platform Pipelines
- AI Platform Pipelines: Concepts
- AI Platform Pipelines: When to Use?
- AI Platform Pipelines: Ecosystem





Few common AI problems



1 Kubeflow/TFX **scalable ML services** on Kubernetes

Easy to get started

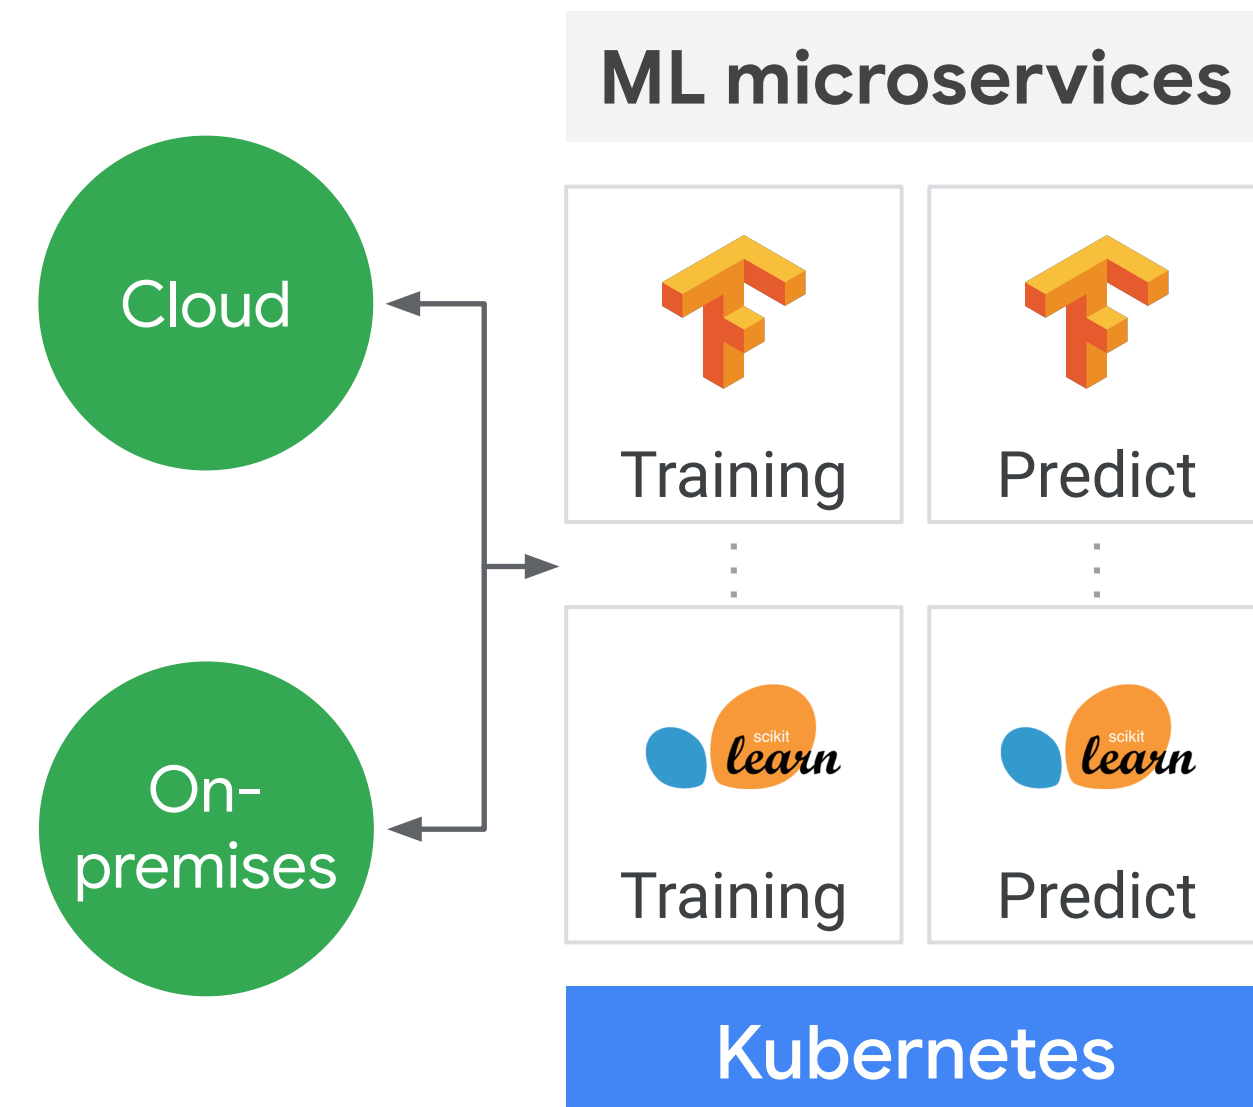
- Kubeflow: Out-of-box support for top frameworks
 - pytorch, caffe, tf, and xgboost
- TFX: Google best practices on TF
- Kubernetes manages dependencies and resources

Swappable and scalable

- Library of ML services
- GPU support
- Massive scale

Meet customers where they are


- Google Cloud
- On-premises



2 Reusable pipelines

Enable developers to build custom ML applications by easily “stitching” and connecting various components, like LEGO.

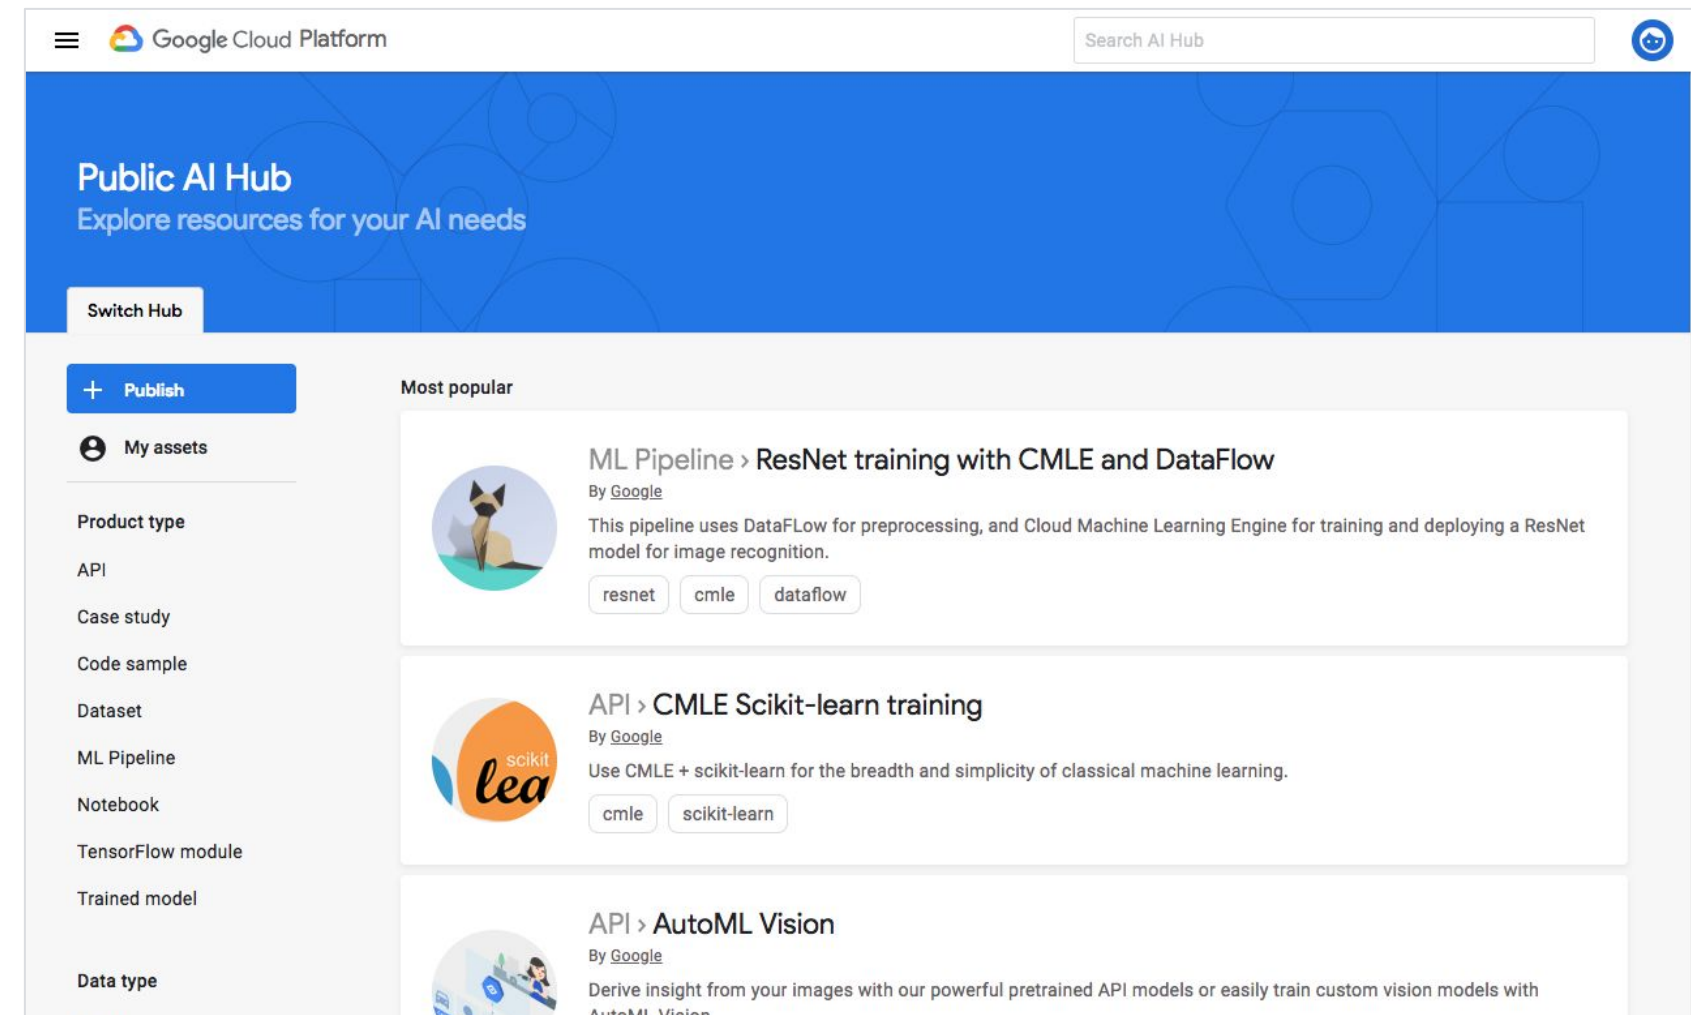
- Reuse instead of reimplement or reinvent.
- Discover, learn, and replicate successful pipelines.



3 Ecosystem

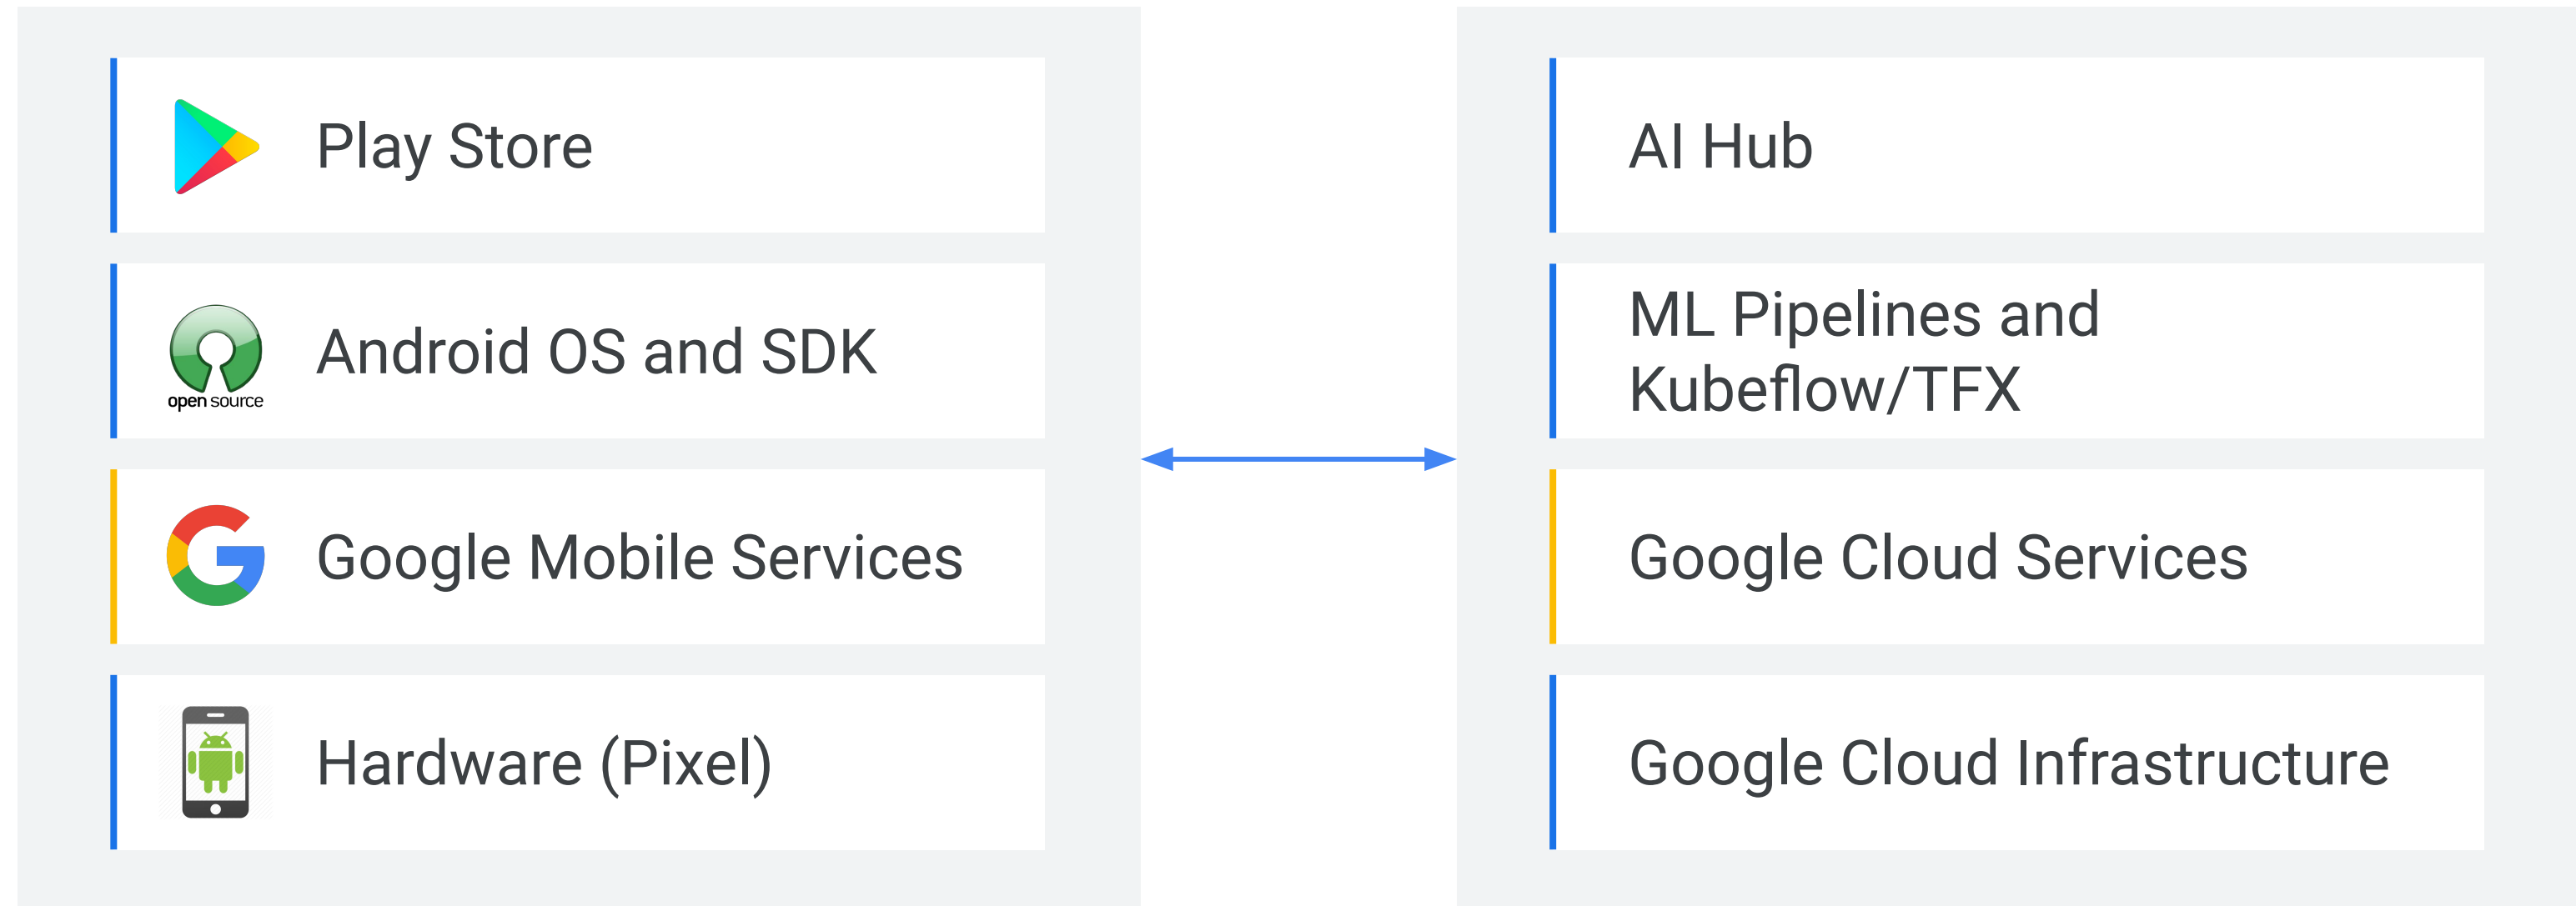
AI Hub at a glance

- 1 All AI content in one place
 - Quick discovery of plug and play AI pipelines and other content build by teams across Google and by partners and customers
- 2 Fast and simple implementation of AI on Google Cloud
 - One-click deployment of AI pipelines via Kubeflow on Google Cloud




- 3 Enterprise-grade internal & external sharing
 - Foster reuse by sharing deployable AI pipelines and other content privately within organizations and publicly

Android ecosystem analogy



Agenda

- Overview
- Introduction to AI Platform Pipelines
- AI Platform Pipelines: Concepts
- AI Platform Pipelines: When to Use?
- AI Platform Pipelines: Ecosystem



What is an ML pipeline?

ML pipelines are portable, scalable ML workflows based on containers.

You can use ML pipelines to:

- Apply MLOps strategies to automate repeatable processes.
- Experiment by running an ML workflow with different sets of hyperparameters, number of training steps or iterations, etc.
- Reuse a pipeline's workflow to train a new model.

Understanding pipeline components

- Self-contained sets of code to perform various operations
- Composed of set of input/output parameters and location of container image
- **Container image**: Package that includes the component's executable code and a definition of the environment that the code runs in

Example pipeline component

Container image location

gcr.io/example-components/...

Container image registry

Component container image

Inputs

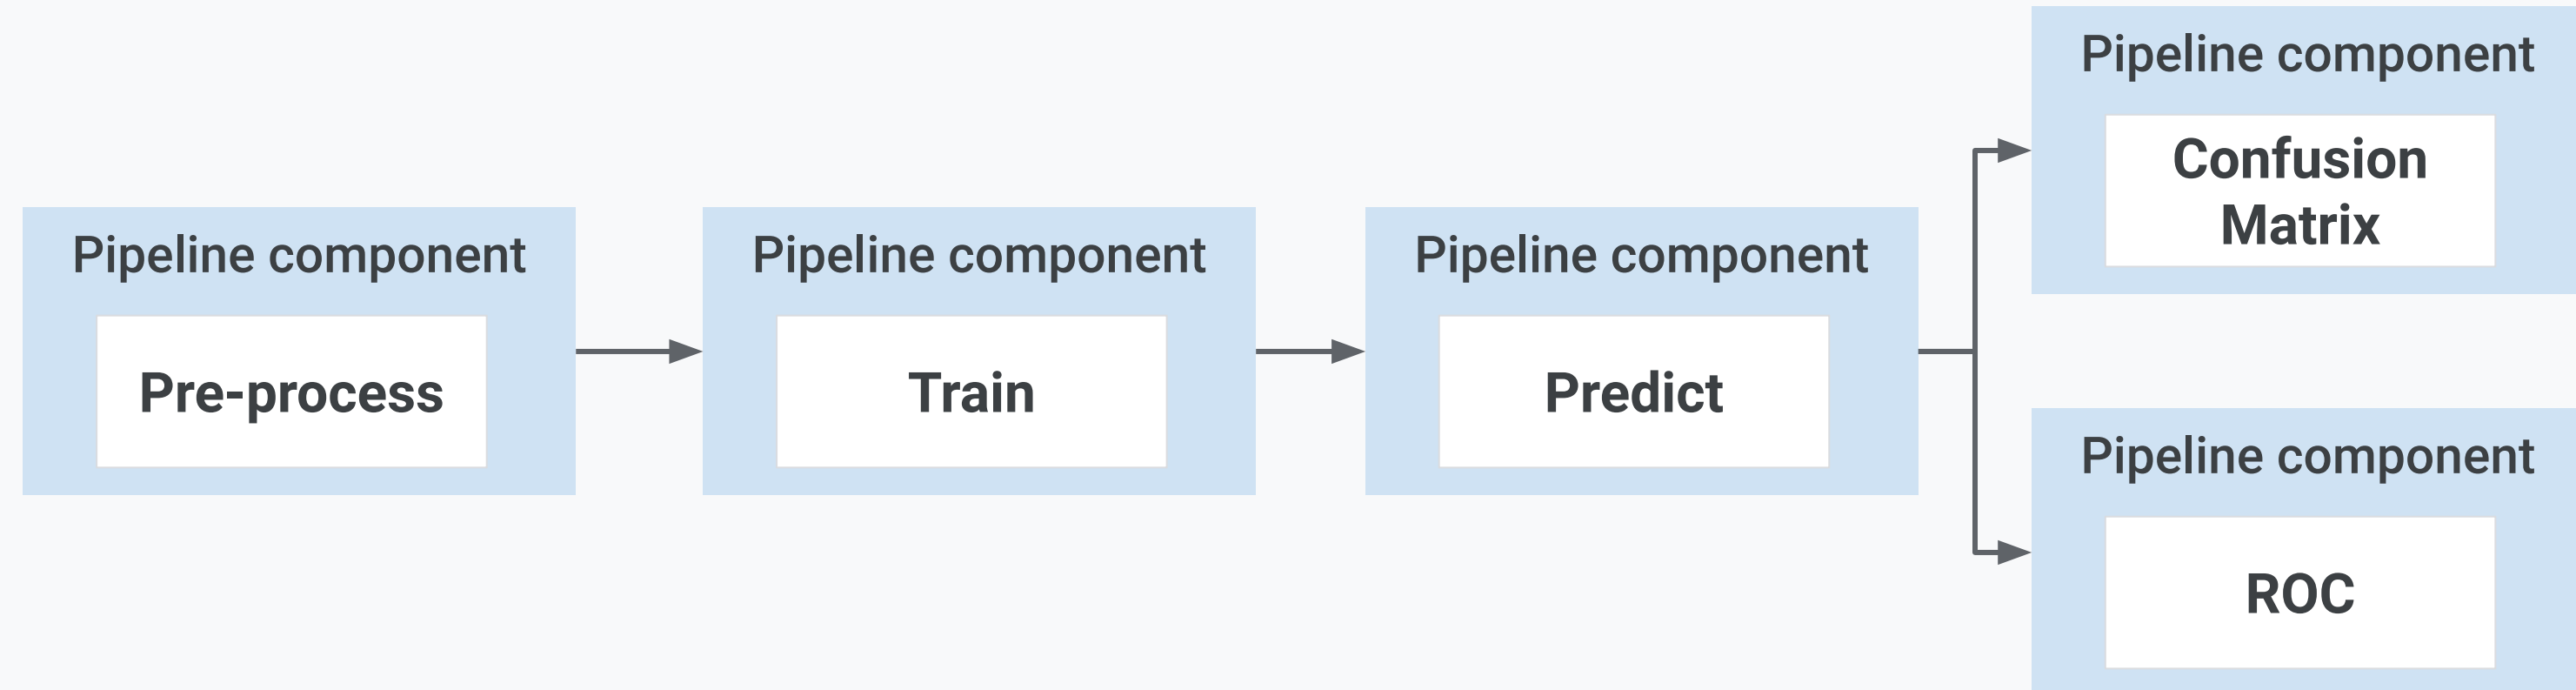
train_data=gs:examples/train.csv,...
eval_data=gs:examples/train.csv,...

Outputs

output=output.txt

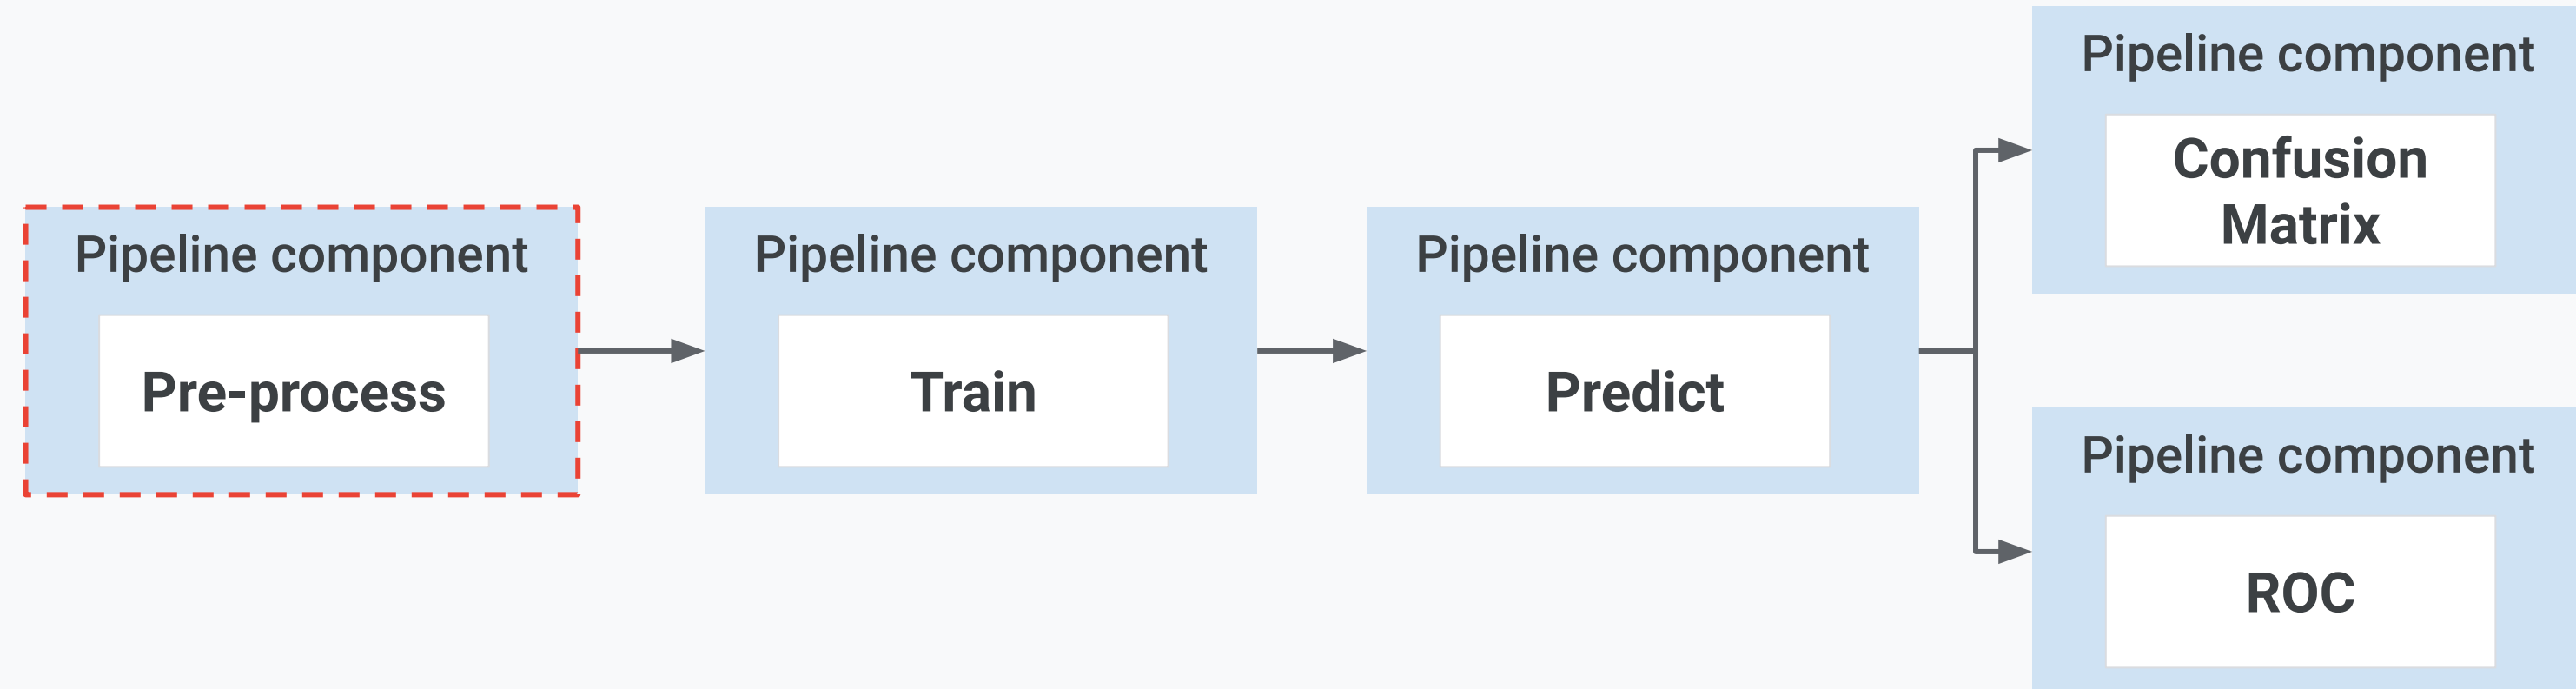
Understanding a pipeline workflow

Example pipeline



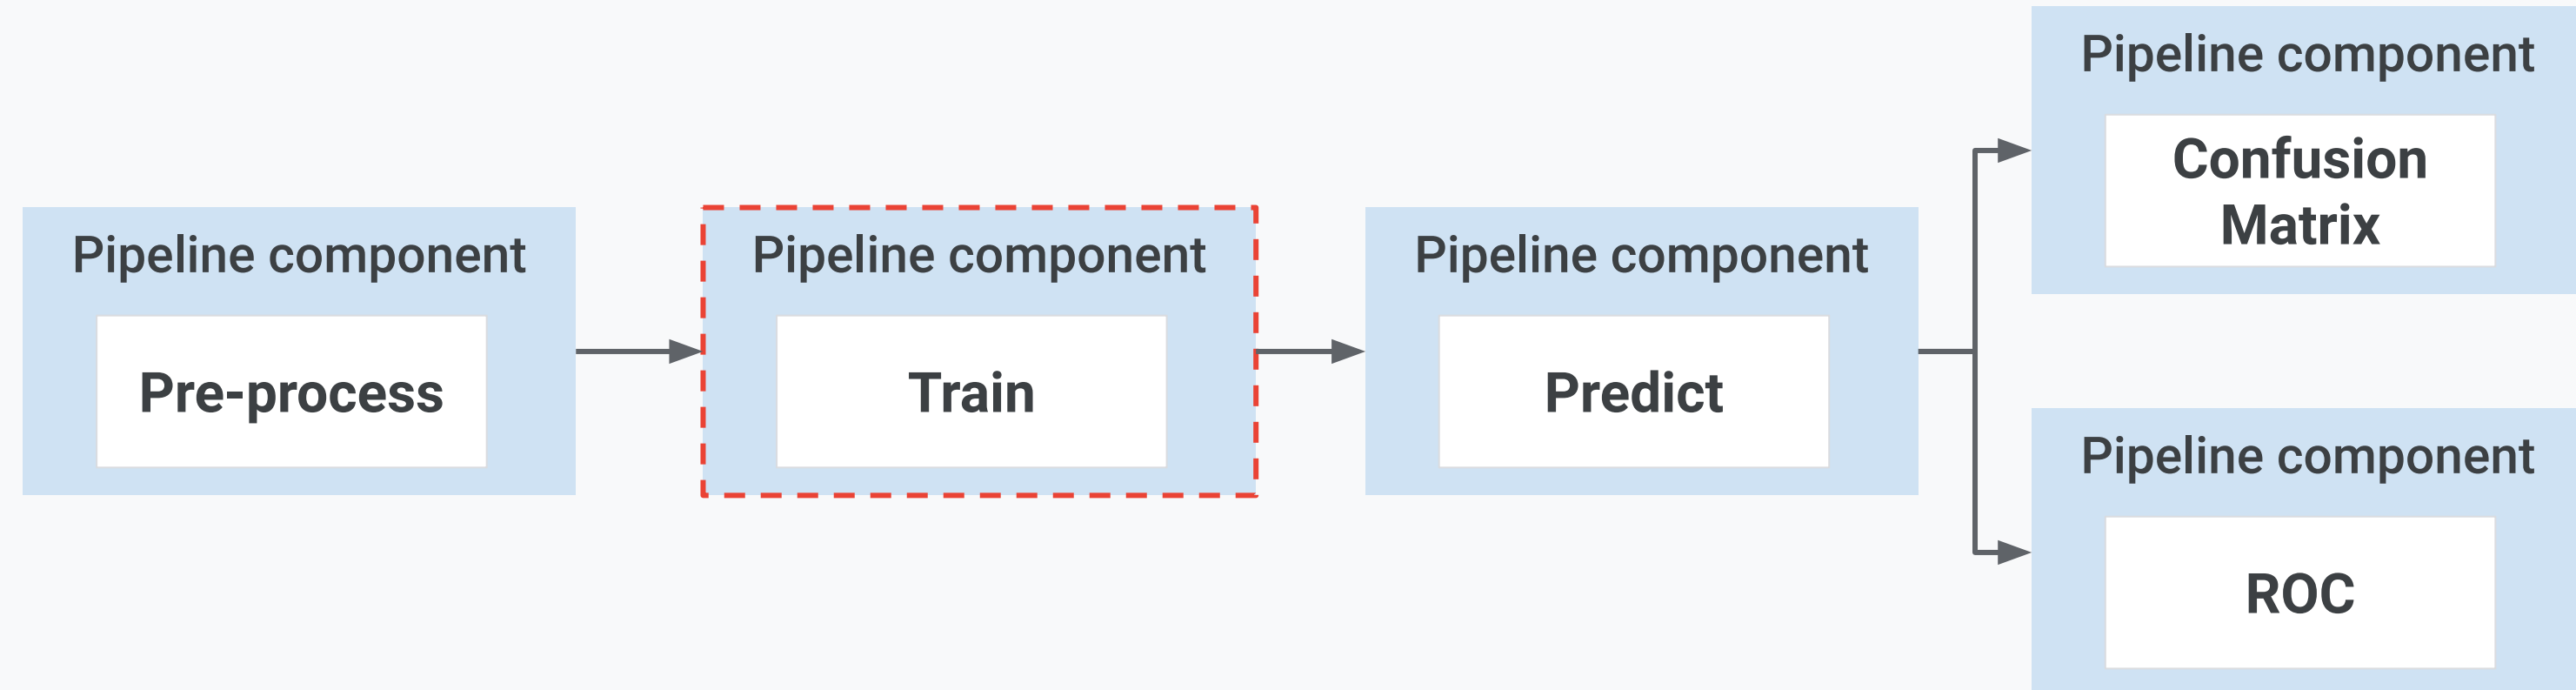
Understanding a pipeline workflow

Example pipeline



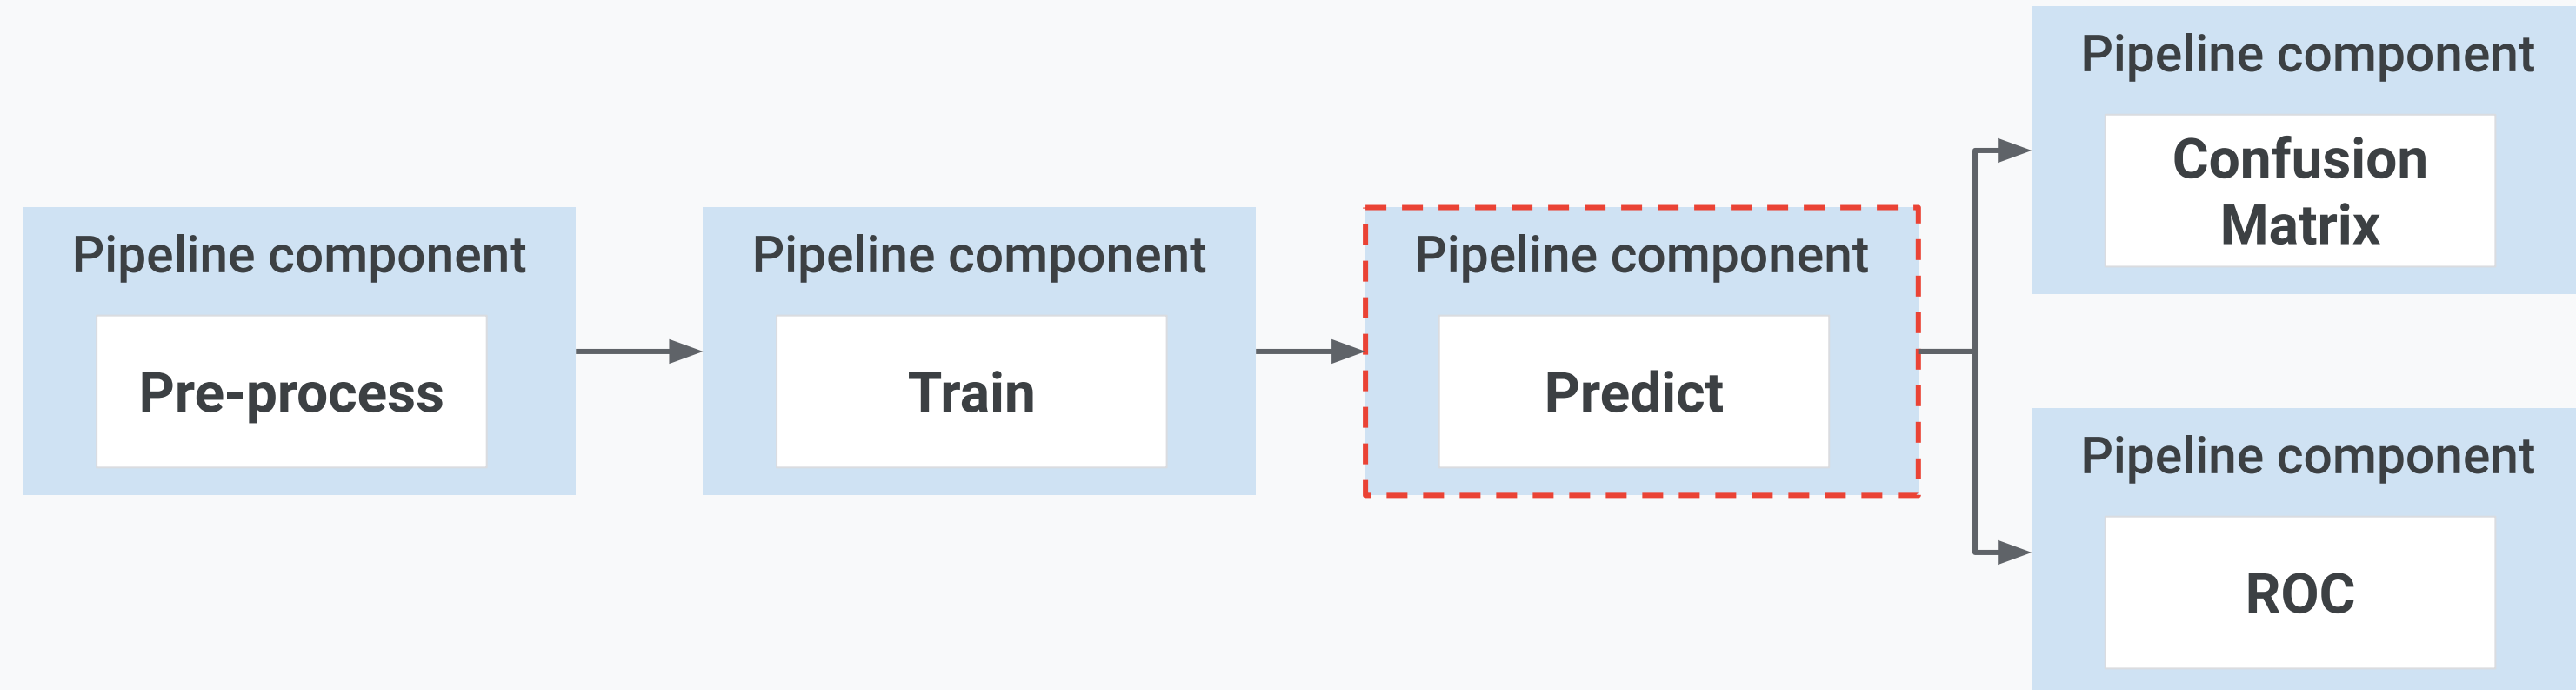
Understanding a pipeline workflow

Example pipeline



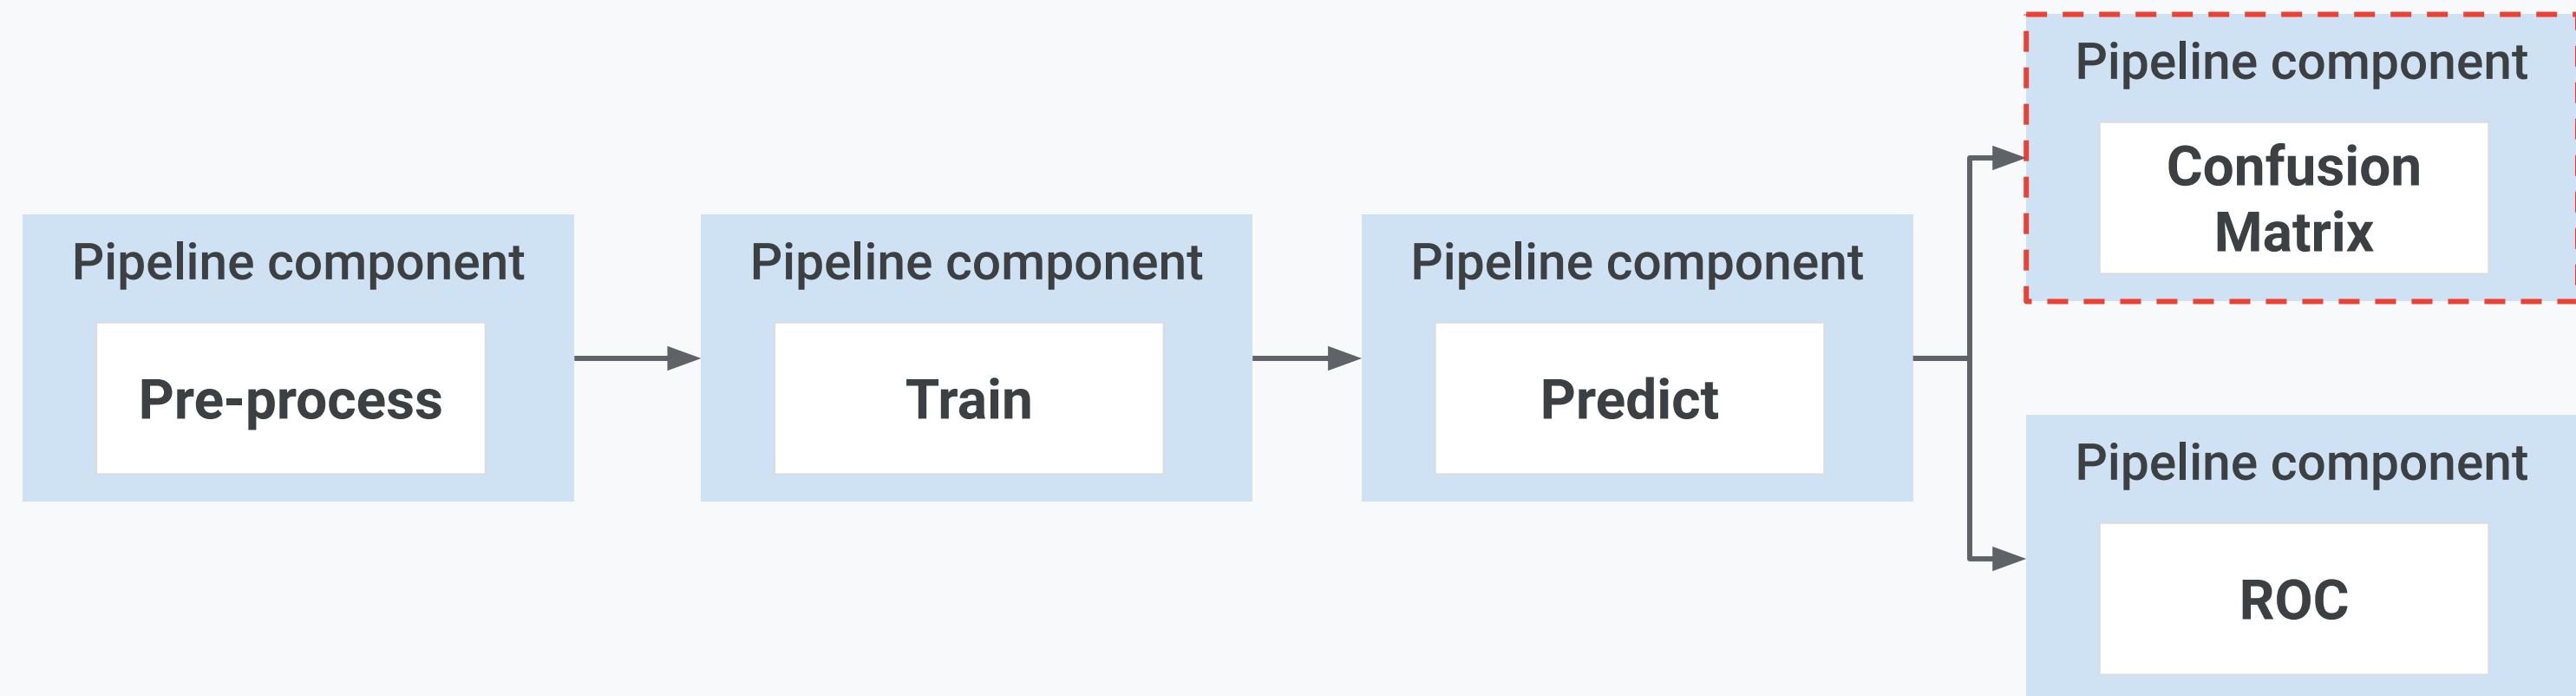
Understanding a pipeline workflow

Example pipeline



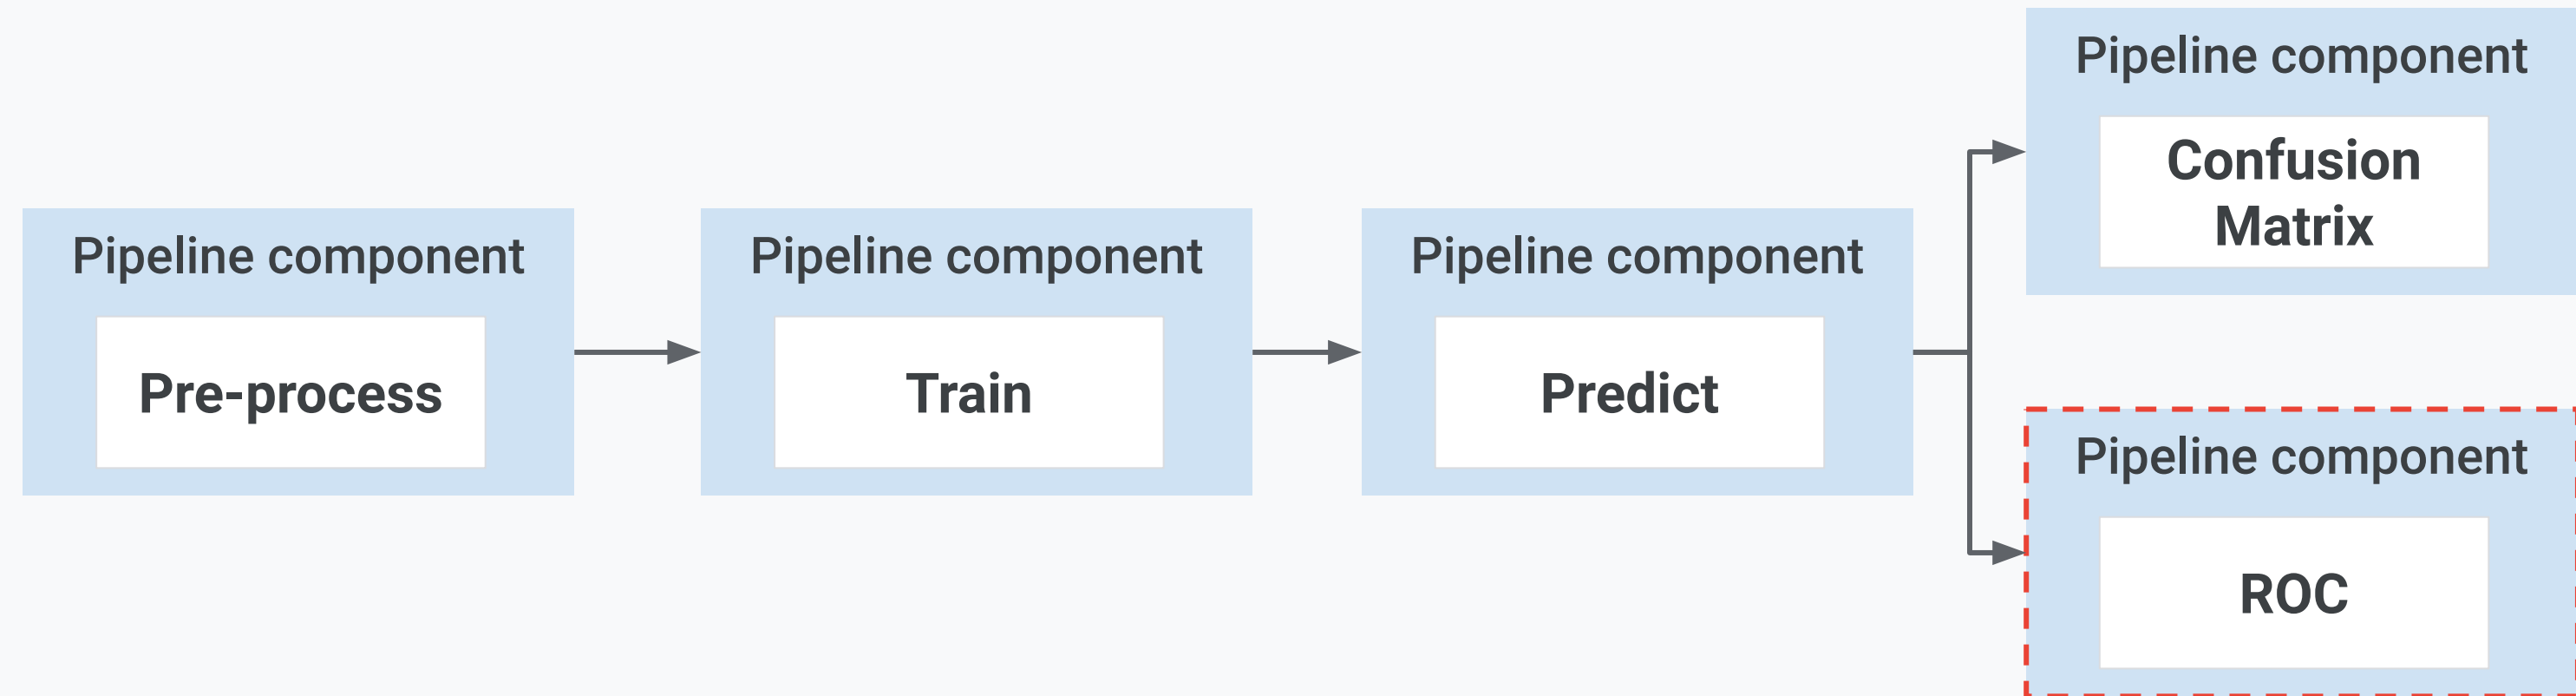
Understanding a pipeline workflow

Example pipeline



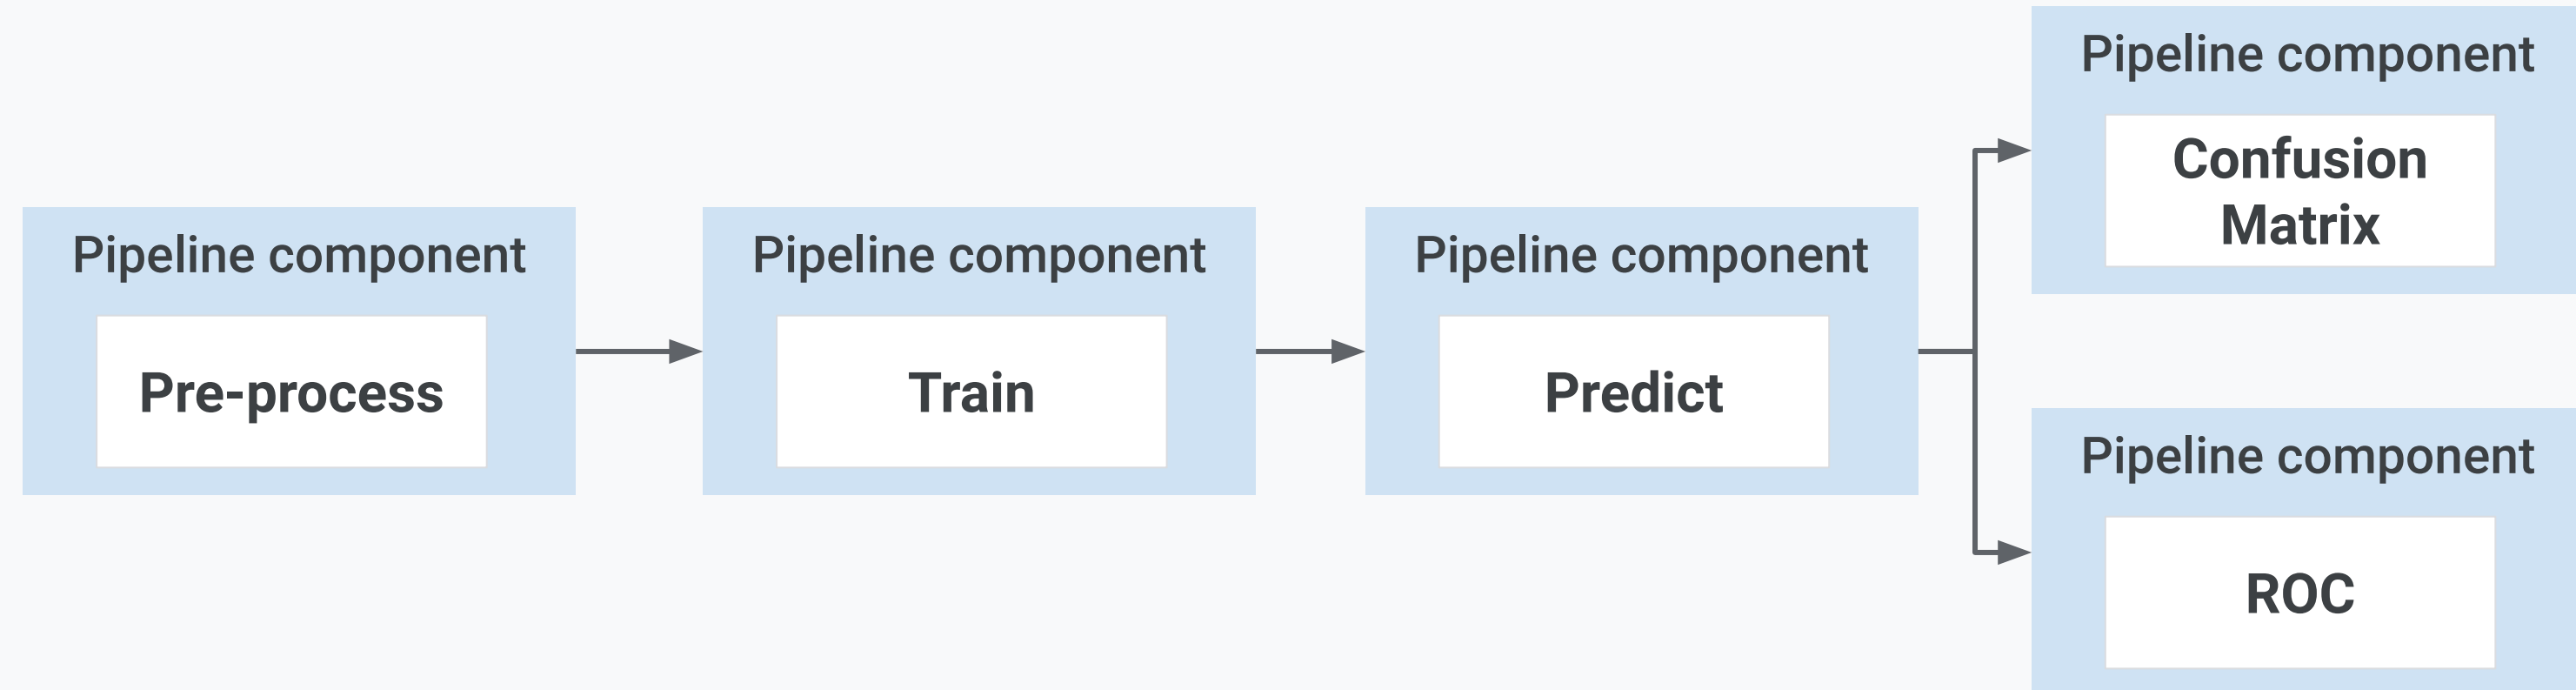
Understanding a pipeline workflow

Example pipeline



Understanding a pipeline workflow


Example pipeline




Current process for building an MLOps pipeline

- Set up a Google Kubernetes Engine (GKE) cluster.
- Create a Cloud Storage bucket for storing data.
- Install Kubeflow pipelines.
- Set up port forwarding.
- Create a process to share the pipeline with your team.

Can we automate these processes to make MLOps a seamless and easy experience?



A new product was needed to deploy robust, repeatable machine learning pipelines along with monitoring, auditing, version tracking, and reproducibility and deliver an enterprise-ready, easy-to-install, secure execution environment for your ML workflows.



What is AI Platform Pipelines?

- One-click installation via the Google Cloud Console
- Enterprise features for running MLOps workloads
- Seamless integration with Google Cloud managed services
- Prebuilt pipeline *components* (pipeline steps) for ML workflows
- Easy customization for new components


Two major parts of an AI Platform Pipelines instance

Enterprise-ready infrastructure

For deploying and running
ML workflows with tight
integration with Google Cloud
services

Pipeline ecosystem

For building, debugging, and
sharing the pipeline and
components



Key benefits of using AI Platform Pipelines

1. Easy installation

- Click one button on the Google Cloud Console
- Automated creation of Google Kubernetes Engine (GKE) cluster
- Customization of existing GKE cluster for deployment
- Easy to delete and reinstall to retain persisted previous state while updating the pipeline version

AI Platform Pipelines



BETA

+ New Instance

Refresh

Delete


Filter

	Status	Name ↑		Zone	Version	Cluster	Namespace	
<input type="checkbox"/>	✓	Kuberflow-pipeline-1	OPEN PIPELINES DASHBOARD	us-central1-a	0.2.4	cluster-5	default	<div> SETTINGS</div>
<input type="checkbox"/>	✓	Kuberflow-pipeline-2	OPEN PIPELINES DASHBOARD	us-west1-a	0.2.5	mpdev-verify	kfp	<div> SETTINGS</div>

Key benefits of using AI Platform Pipelines


2. Easy authentication process

- Fully secure and provides authenticated access to Pipelines UI
- No need to set up port forwarding
- Easy to share with team members
- Easy to access through REST API service
- Seamless performance of using Pipelines SDK from Notebooks
 - Define pipeline
 - Schedule run job

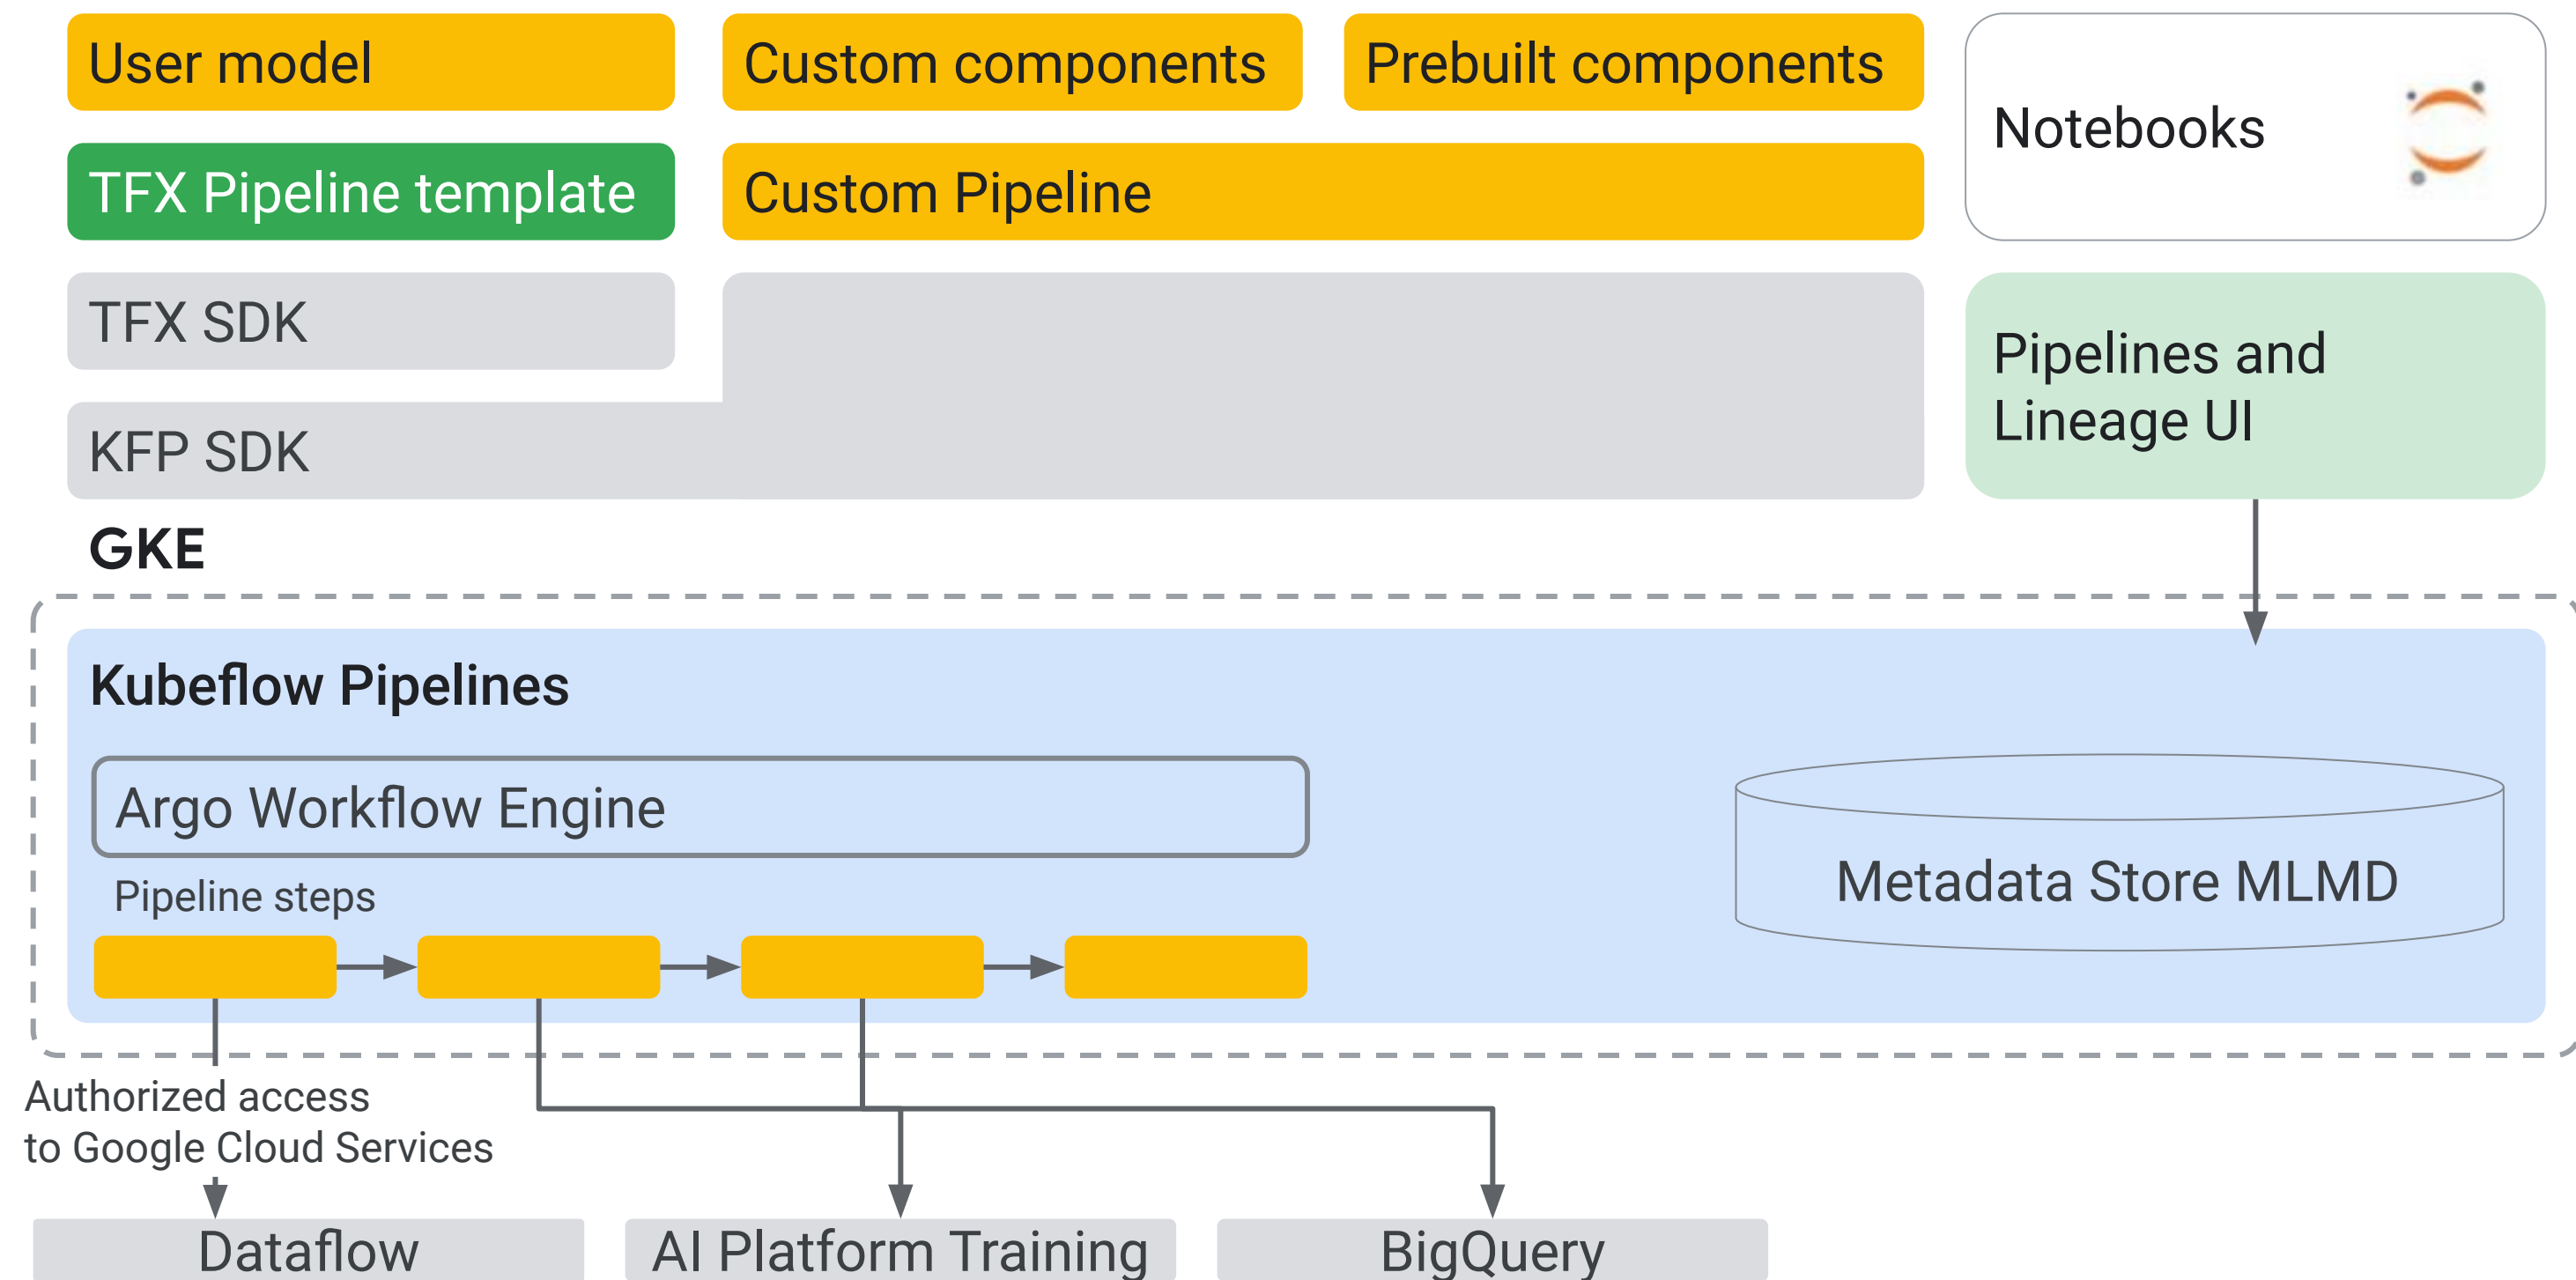


Agenda

- Overview
- Introduction to AI Platform Pipelines
- **AI Platform Pipelines: Concepts**
- AI Platform Pipelines: When to Use?
- AI Platform Pipelines: Ecosystem



AI Platform Pipelines tech stack



AI Platform Pipelines implementation strategy

Kubeflow Pipelines

Through Kubeflow Pipelines SDK

- Lower-level ML framework–agnostic implementation
- Enables direct control of Kubernetes resource control
- Simple sharing of containerized components
- Use it for fully custom pipelines

TensorFlow Extended (TFX)

Through TFX SDK

- Higher-level abstraction
- Prescriptive but customizable components with pre-defined ML types
- Brings Google best practices for robust/scalable ML workloads
- Use it for E2E TF-based pipeline with customizable data pre-processing and training code

Some features of AI Platform Pipelines

1. Build your own ML pipeline with TFX examples

Pipelines

Experiments

Notebooks

Graph

Run output

Config

csvexa

statis

sche

examplevalidator

tra

evaluator

push

Artifacts

Input/Output

Volumes

Manifest

Logs

Sort by

Feature

Reverse order

Feature search (...)

Feature:

int (8)

float (7)

string (2)

Numeric Features (15)

count	missing	mean	Std dev
dropoff_census_tract			
3,618	28.93%	17.0B	333K
dropoff_community_area			
4,905	3.65%	21.2	17.85
dropoff_latitude			
4,915	3.46%	41.9	0.04

Some features of AI Platform Pipelines

2. Pipeline versioning

Pipelines

+ Upload pipeline

Refresh

Delete

	Pipeline name	Description	Uploaded on	↓
<input type="checkbox"/>	▶ [Tutorial] DSL - Control struct...	Source code Shows how to use conditional execution and exit...	2/20/2020, 3:28:12 PM	
<input type="checkbox"/>	▶ [Tutorial] Data passing in pyt...	Source code Shows how to pass data between python comp...	2/20/2020, 3:28:11 PM	
<input type="checkbox"/>	▼ [Demo] TFX - Taxi Tip Predict...	Source code GCP Permission Requirements . Example pipeline that does classification...	2/20/2020, 3:28:10 PM	
Version name			Uploaded on	↓
<input type="checkbox"/>	TFX - Taxi Tip Prediction Model Trainer_version_at_2020-03-03T15:44:30.197Z		3/3/2020, 7:55:03 AM	
<input type="checkbox"/>	[Demo] TFX - Taxi Tip Prediction Model Trainer		2/20/2020, 3:28:10 PM	
1 of 200				
<input type="checkbox"/>	▶ [Demo] XGBoost - Training...	Source code GCP Prmission requirements . A trainer that does end-to-end distributed tra...	2/20/2020, 3:28:09 PM	

Some features of AI Platform Pipelines

3. Artifact tracking

Artifacts					
<div></div>					
Pipeline / Workspace	Name	ID	Type	URL	Created at
		1	ExternalArtifact	gs://ml-pipeline-play...	2/20/2020, 5:1...
▼ taxi_pipeline_with_parameters	examples	2	Examples	gs://ml-pipeline-play...	2/20/2020, 5:1...
	statistics	3	ExampleStatistics	gs://ml-pipeline-play...	2/20/2020, 5:1...
	schema	4	Schema	gs://ml-pipeline-play...	2/20/2020, 5:1...
	anomalies	5	ExampleAnomalies	gs://ml-pipeline-play...	2/20/2020, 5:1...
	transform_gra...	6	TransformGraph	gs://ml-pipeline-play...	2/20/2020, 5:1...
	transformed_e...	7	Examples	gs://ml-pipeline-play...	2/20/2020, 5:1...

Some features of AI Platform Pipelines

3. Artifact tracking

Artifacts

← model

Overview

Lineage Explorer

Type: Model

URI

gs://aju-pipelines/tfx_taxi_simple/85265540-6a06-4969-a49e-1f65741878be/Trainer/model/7

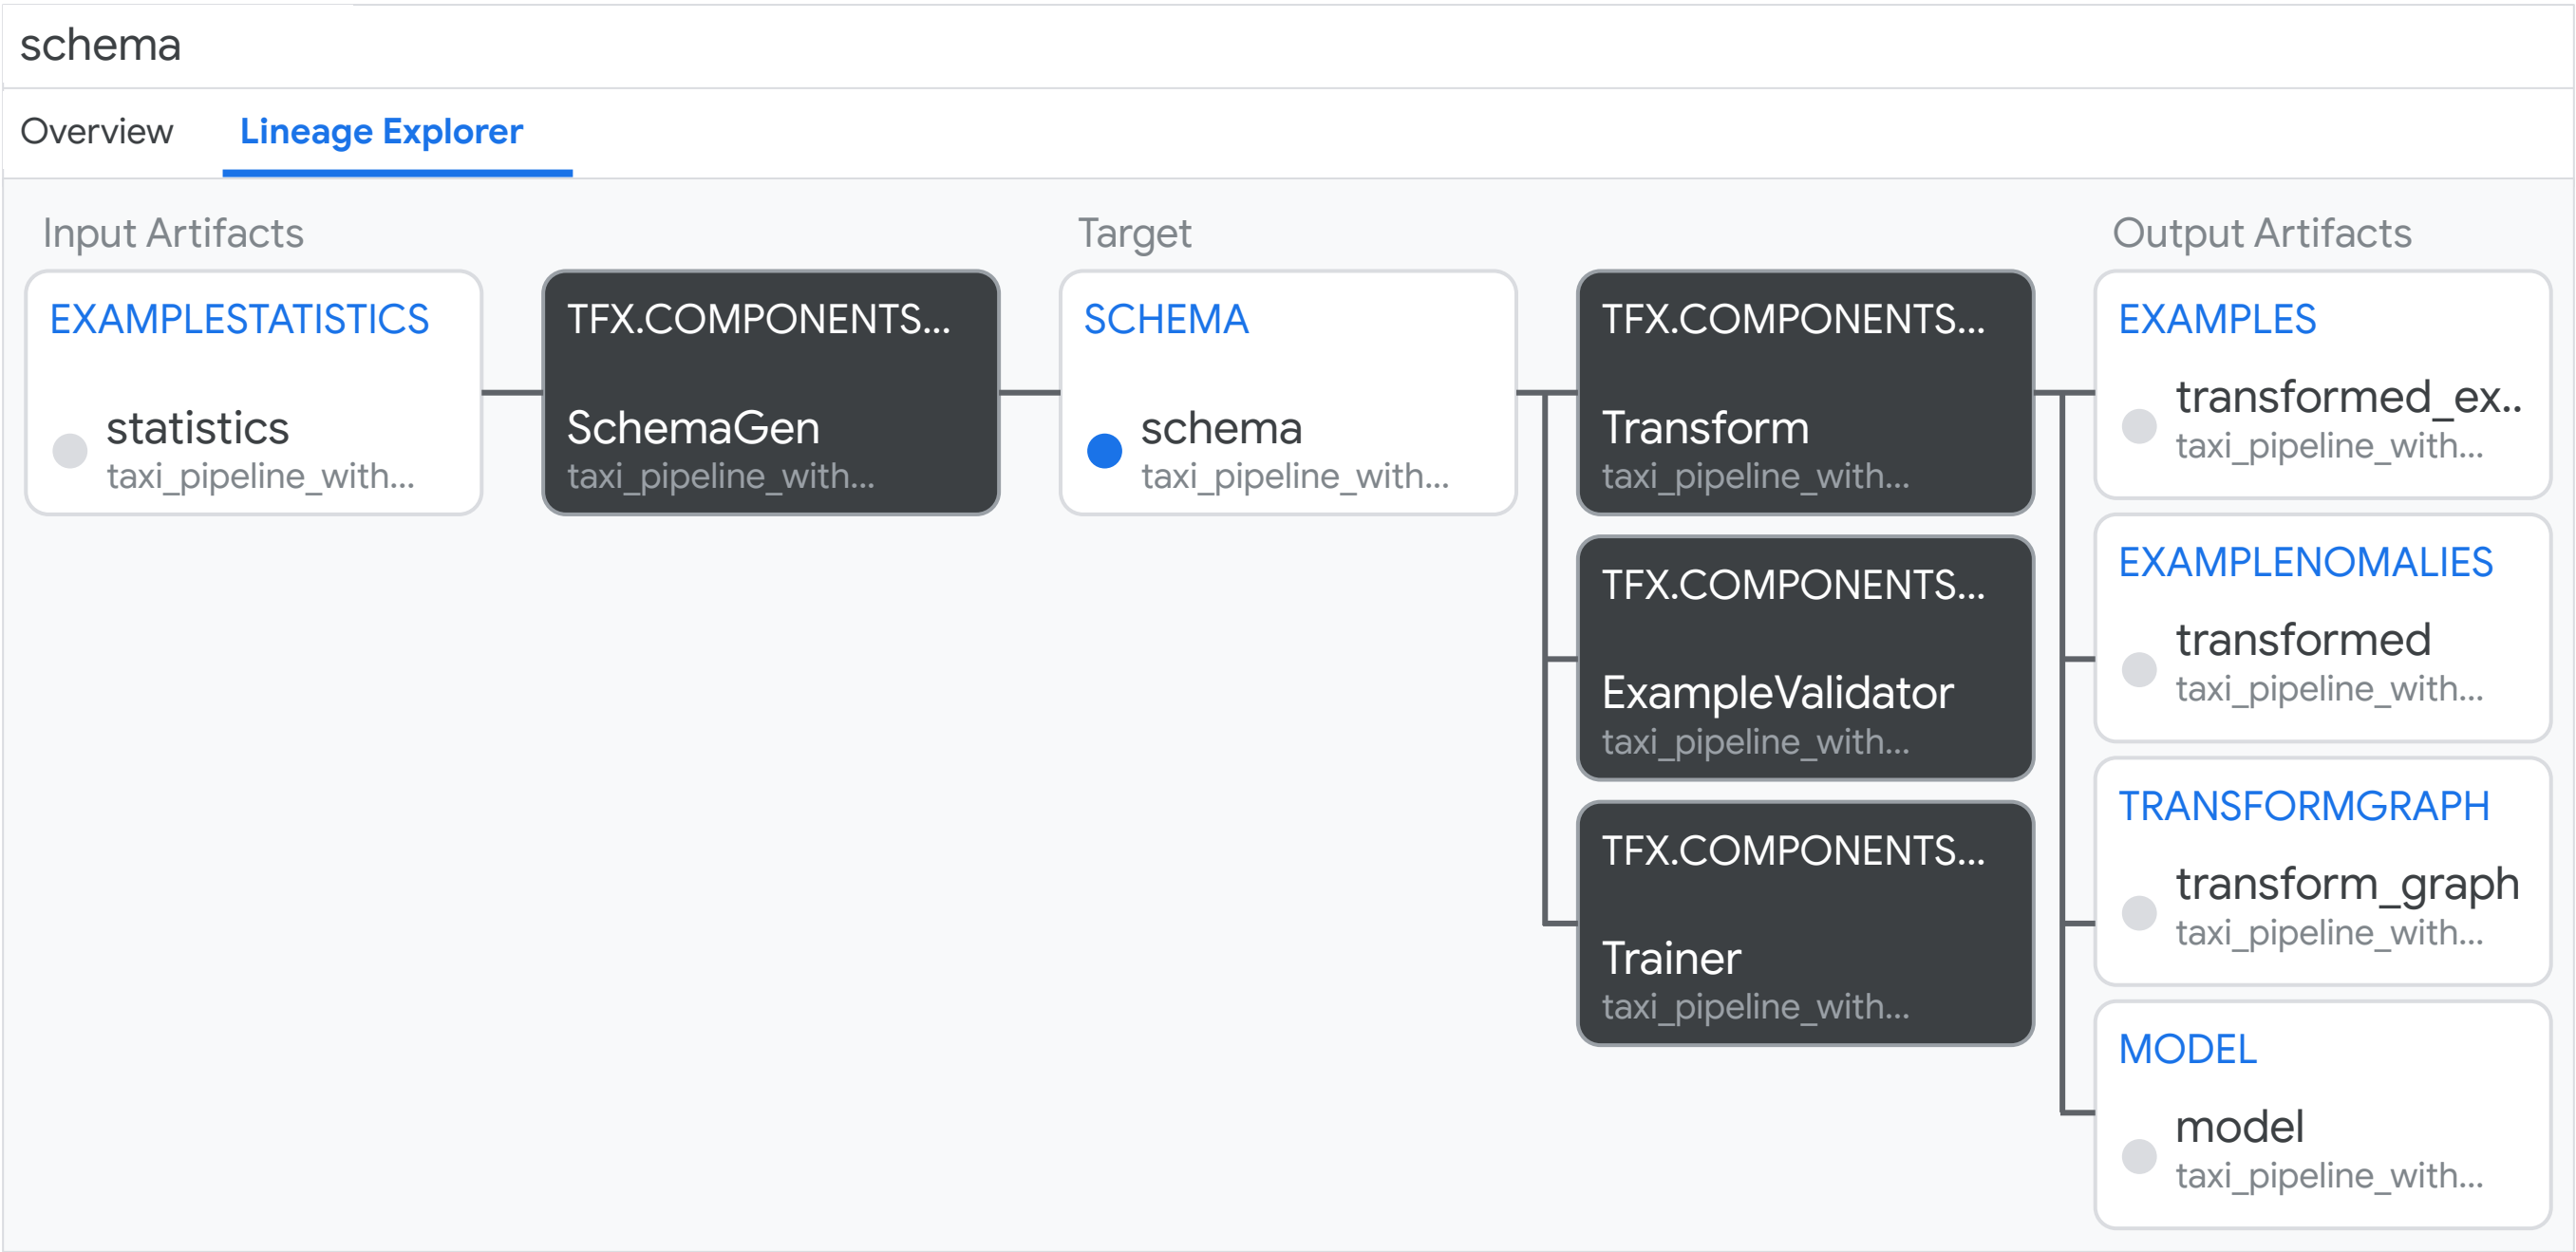
Properties

Custom Properties

name	pipeline_name	producer_component	state
model	taxi_pipeline_with_parameters	Trainer	published


Some features of AI Platform Pipelines

4. Lineage tracking




Agenda

- Overview
- Introduction to AI Platform Pipelines
- AI Platform Pipelines: Concepts
- AI Platform Pipelines: When to Use?
- AI Platform Pipelines: Ecosystem




What does AI Platform Pipelines enable?



**Workflow
orchestration**




**Rapid, reliable,
repeatable
experimentation**



**Share, re-use,
and compose**


What does AI Platform Pipelines enable?



**Workflow
orchestration**

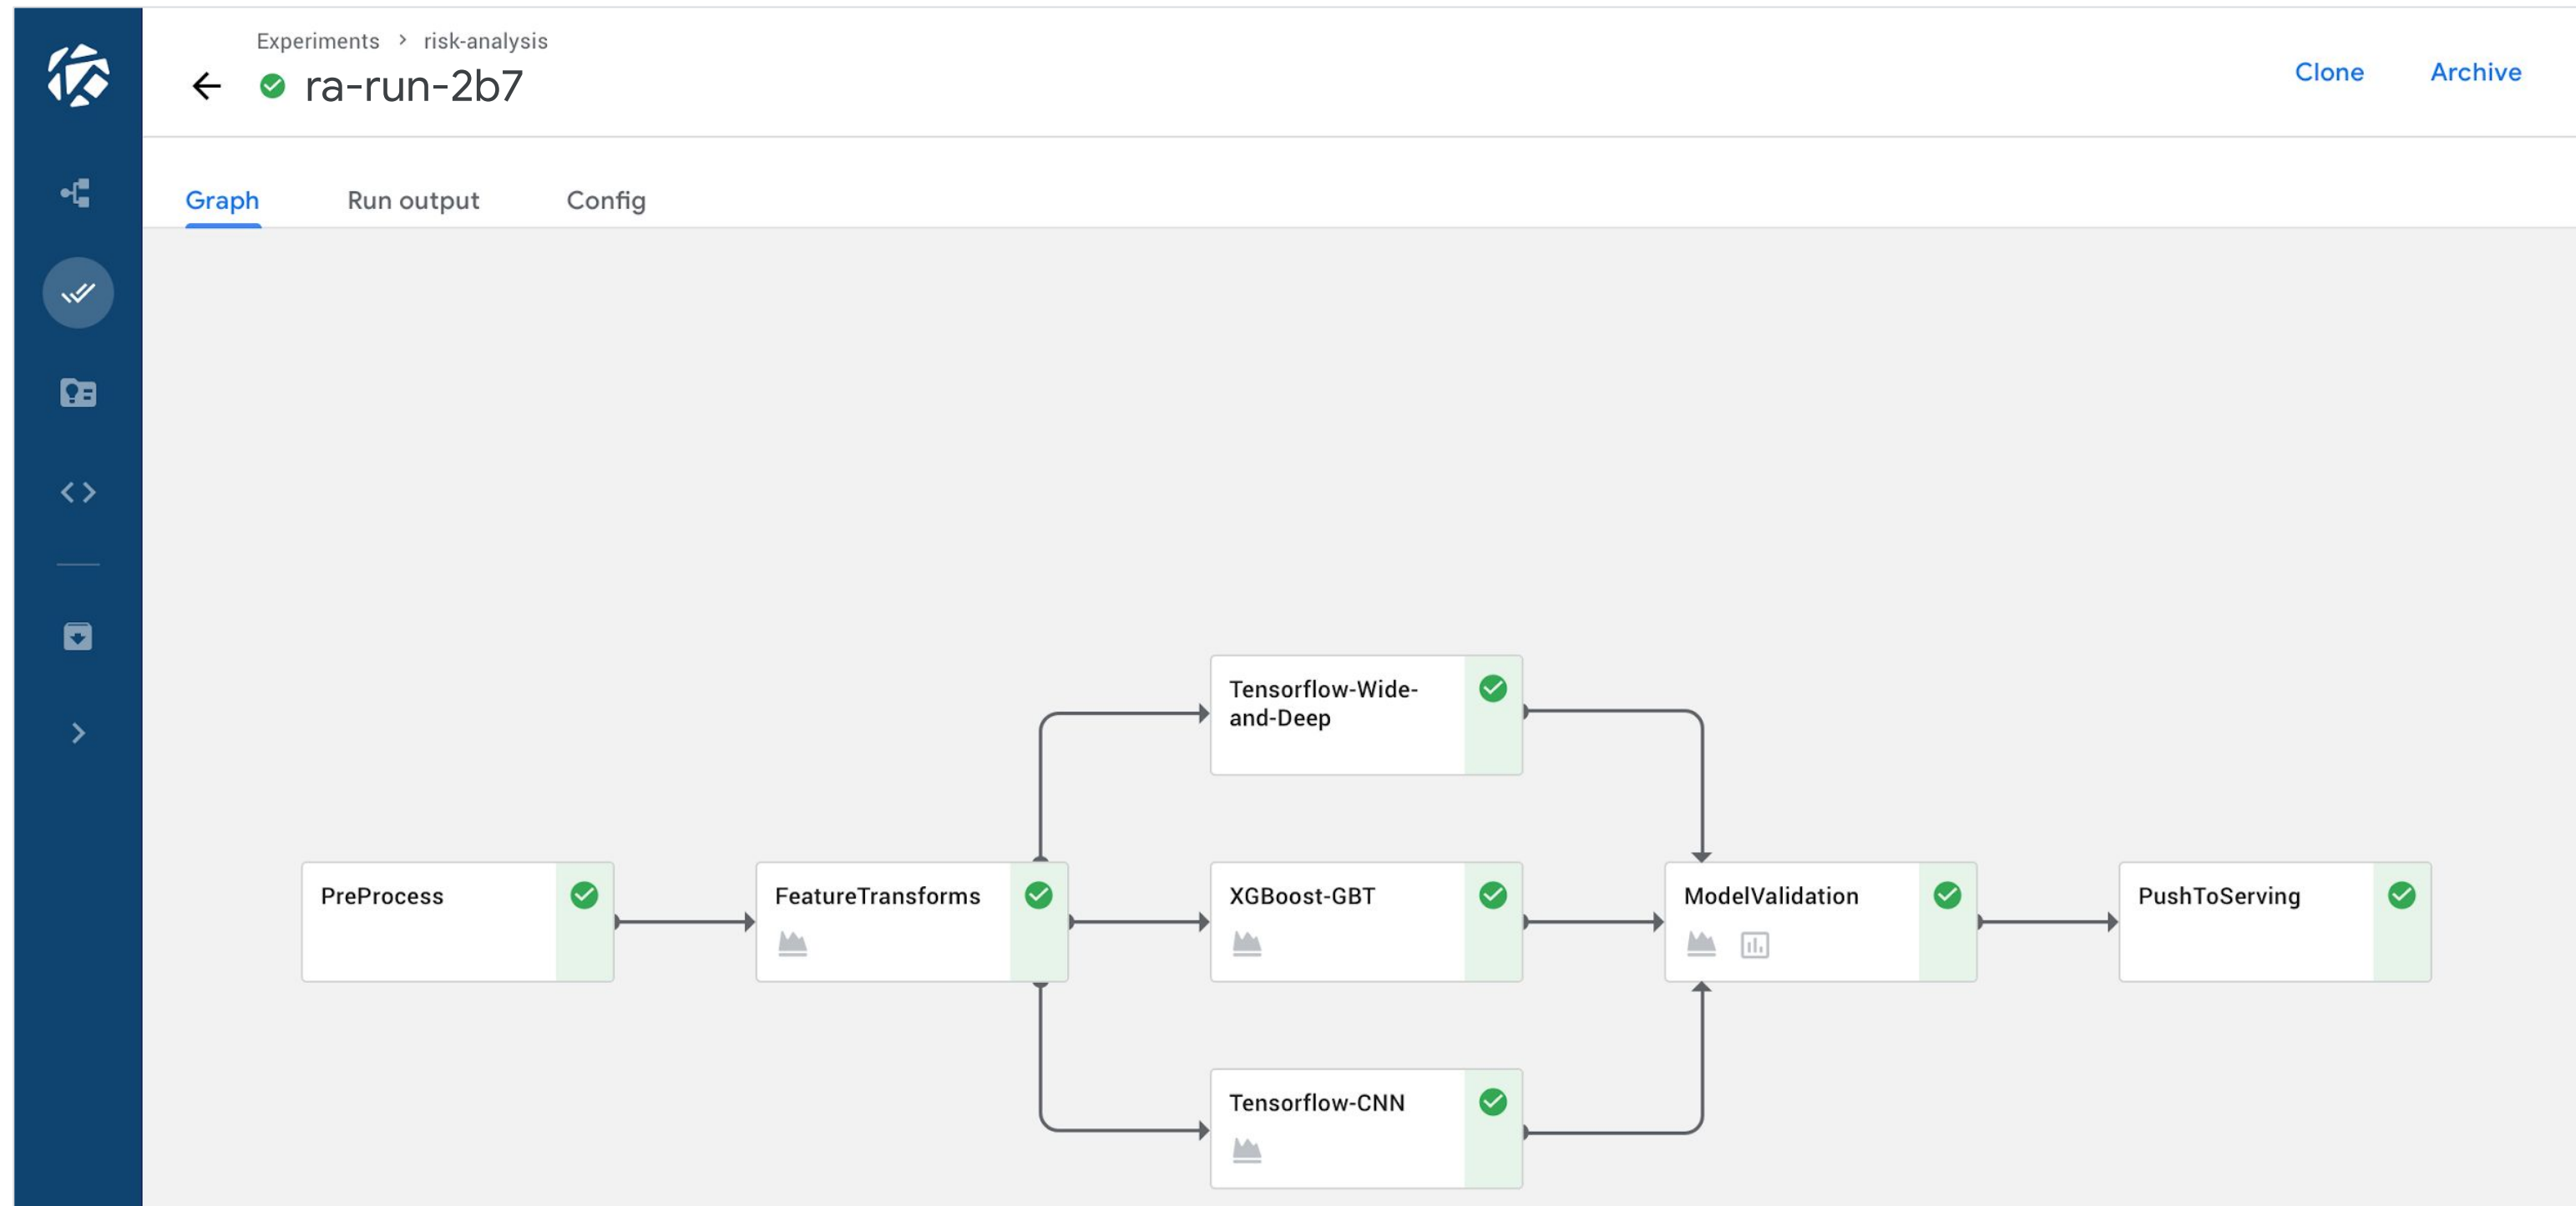


**Rapid, reliable,
repeatable
experimentation**

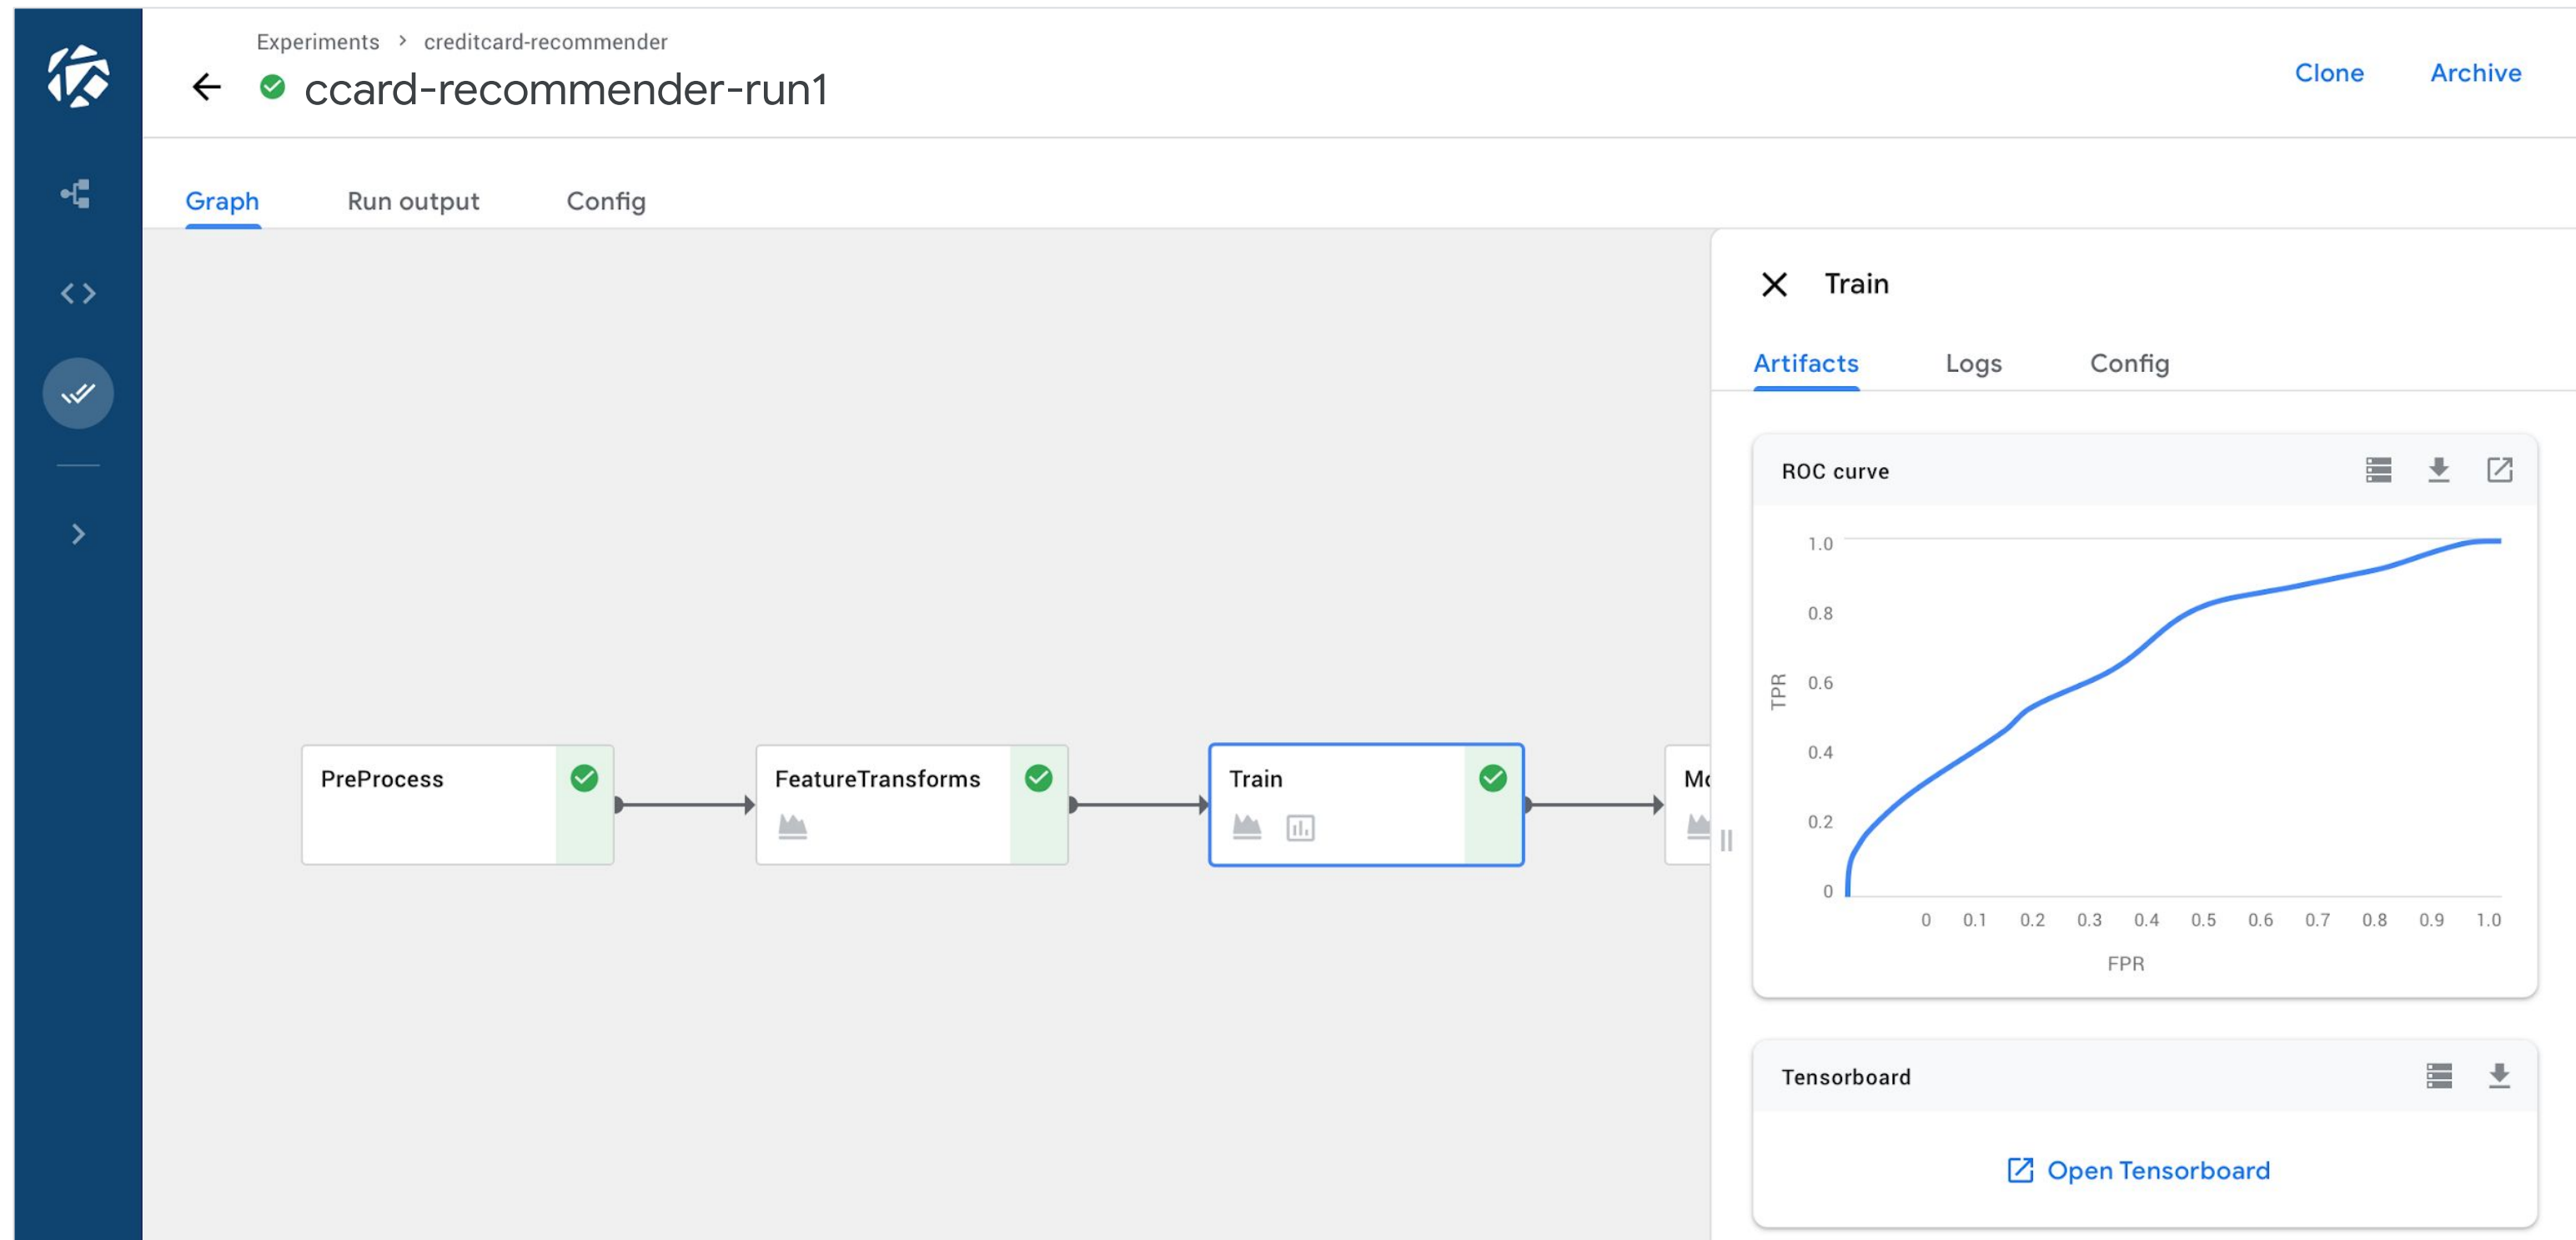


**Share, re-use,
and compose**






Visual depiction of pipeline topology



Rich visualizations of metrics



Access to all
config params,
inputs, and
outputs for each
run



Experiments > Product Image Classification

←

✔ Simple XGBoost Classifier

Graph

[Config](#)

Run details


Status	Succeeded
Description	
Created at	11/25/2018, 12:56:44 PM
Started at	11/25/2018, 12:56:44 PM
Finished	11/25/2018, 12:16:37 PM
Duration	0:19:53

Run parameters

output	gs://mipipelines
project	foo2thebar
region	us-central1
train-data	gs://ml-pipeline-playground/sfpd/train.csv
eval-data	gs://ml-pipeline-playground/sfpd/eval.csv
schema	gs://ml-pipeline-playground/sfpd/schema.json
target	resolution
rounds	200

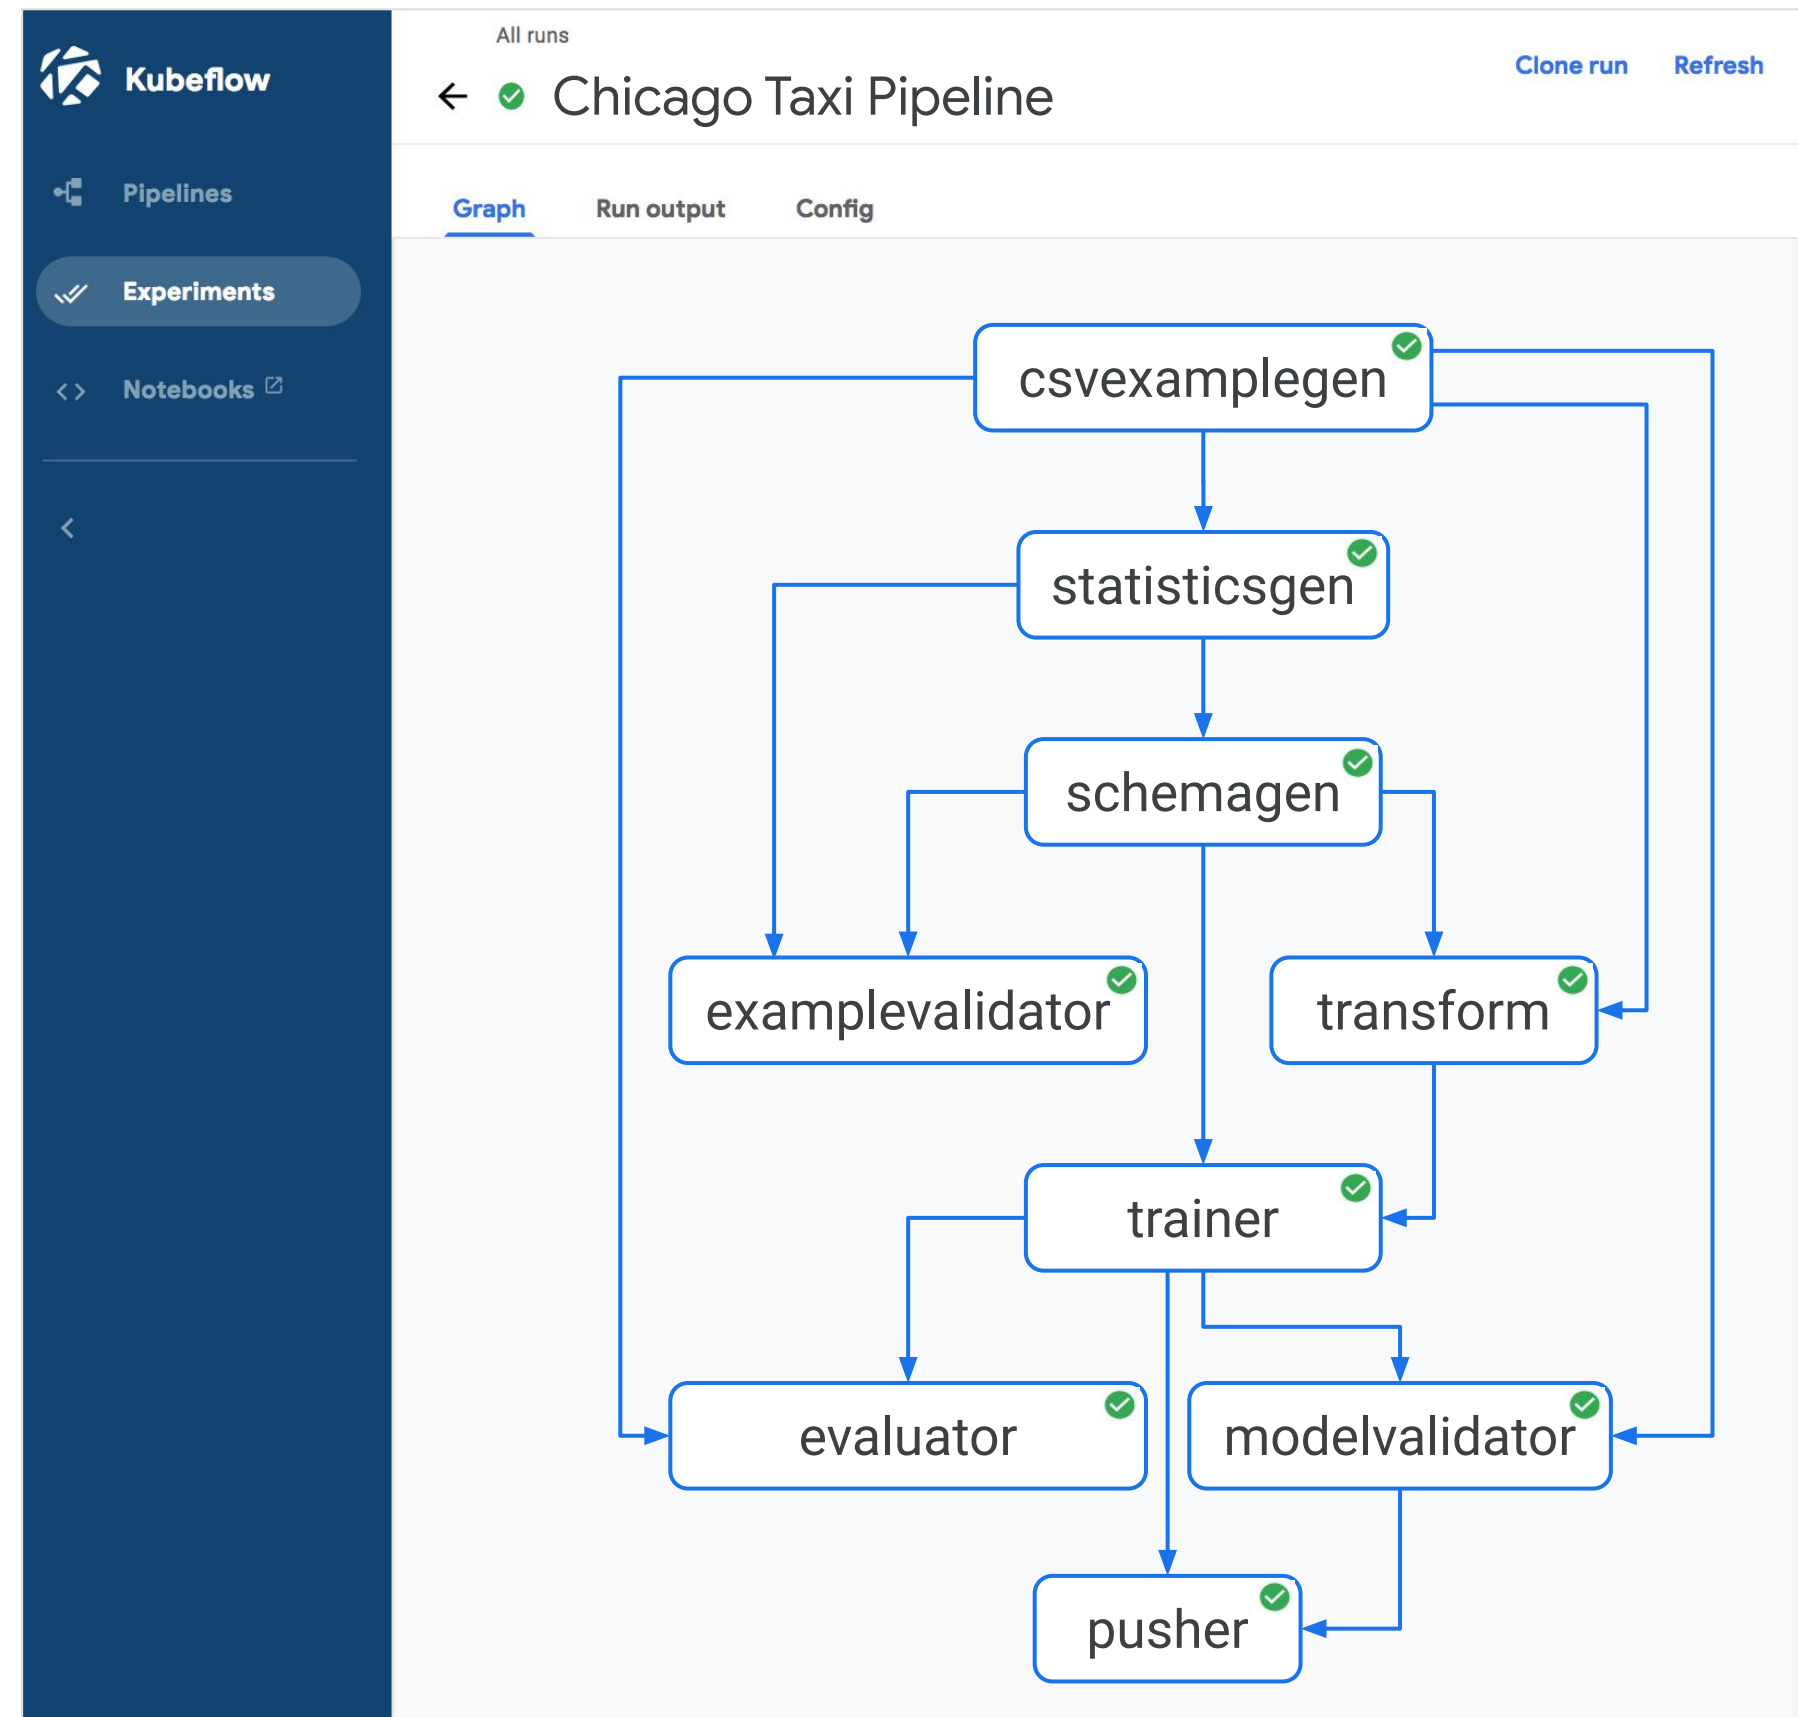
What constitutes an AI Platform Pipelines instance?

- Containerized implementations of ML tasks
 - Containers provide portability, repeatability, and encapsulation.
 - A task can be single node or *distributed*.
 - A containerized task can invoke other services, like CAIP, Dataflow, or Dataproc.
- Specification of the sequence of steps
 - Specified via Python SDK
- Input parameters
 - A “job” = A pipeline invoked with specific parameters




TFX + Kubeflow Pipelines

Robust extensible platform and [best-in-class](#) components for typical stages of good ML workflow




What does AI Platform Pipelines enable?



**Workflow
orchestration**


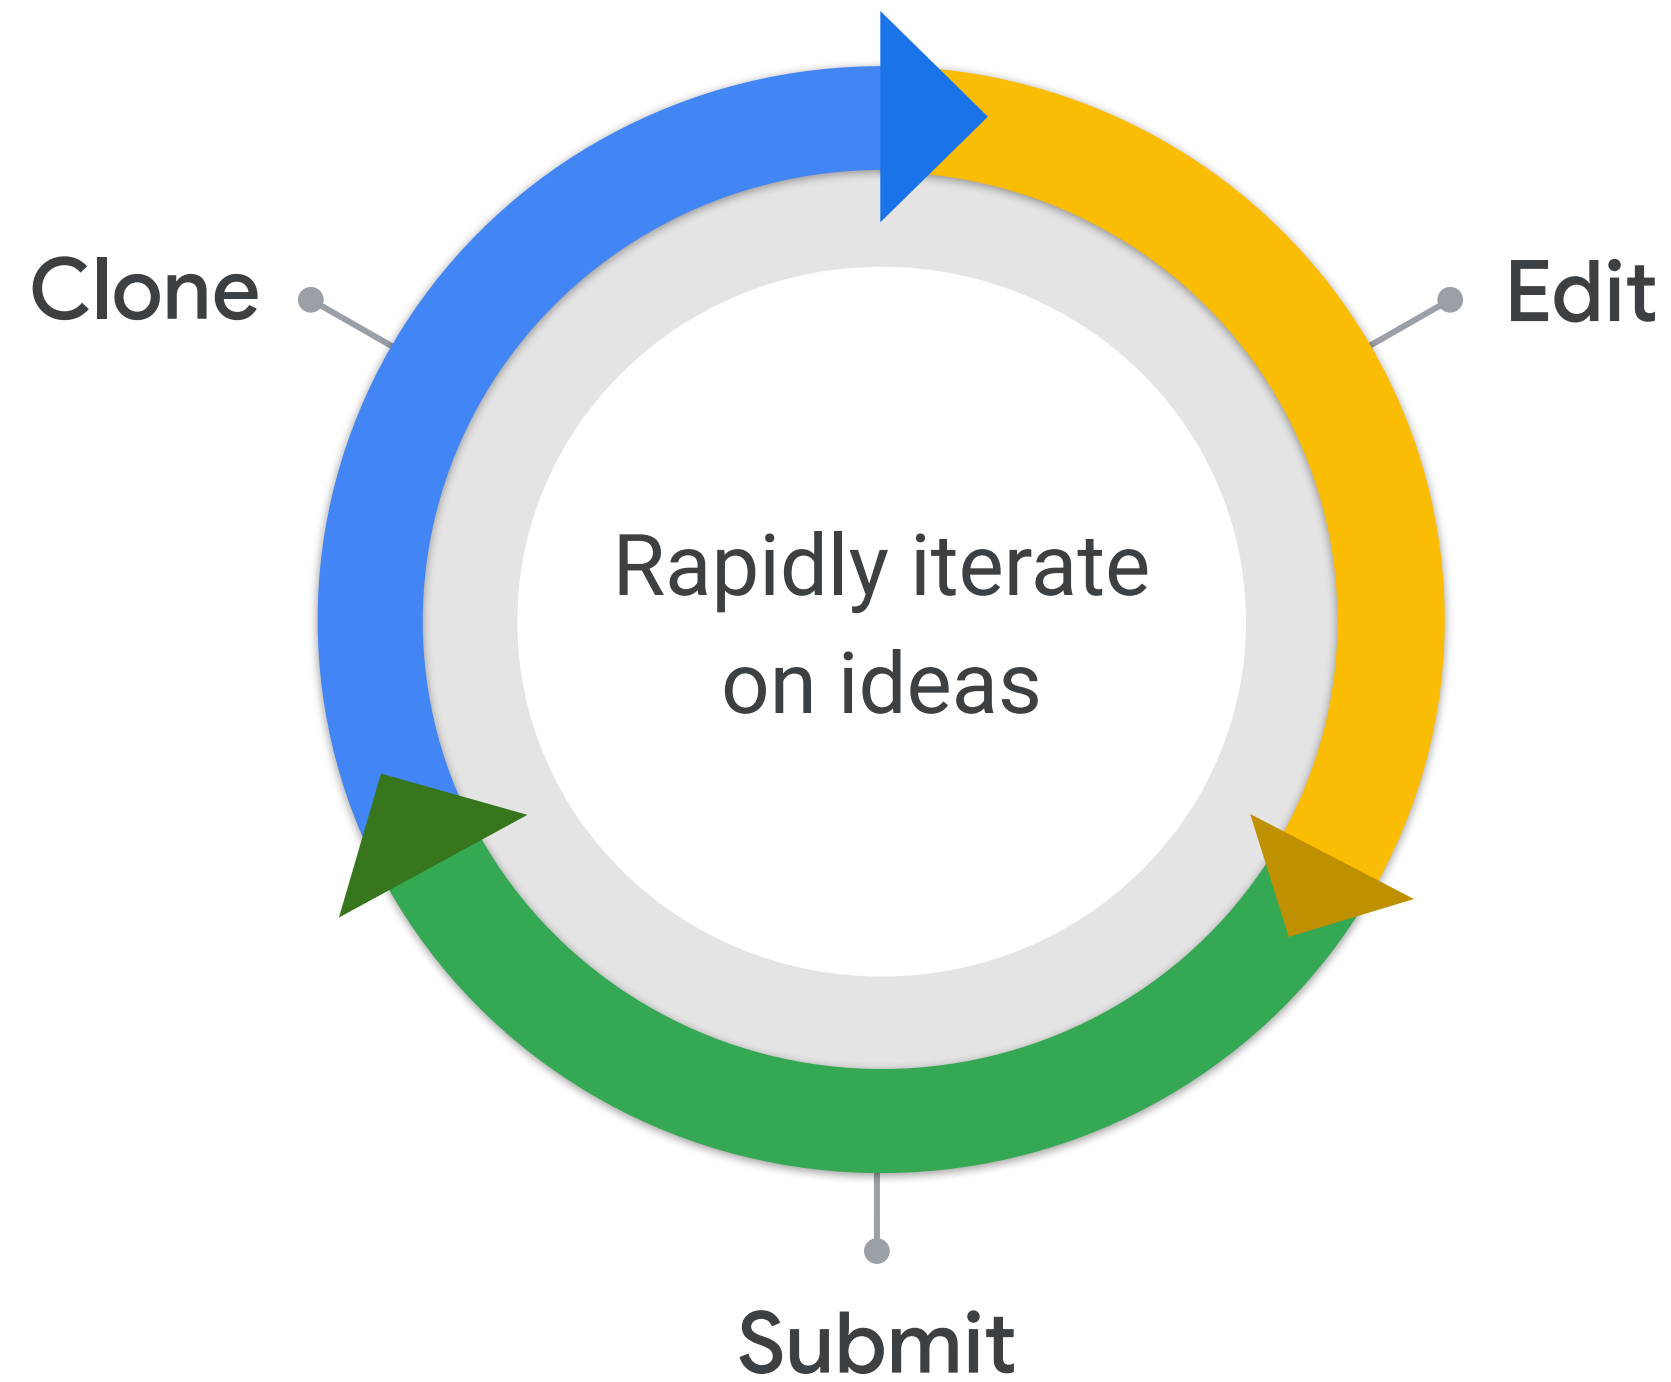


**Rapid, reliable,
repeatable
experimentation**



**Share, re-use,
and compose**

Rapid, reliable experimentation



View all current and historical runs grouped as “Experiments”

Experiments

+ Create experiment

Compare runs

Archive

All experiments

All runs


Filter experiments

<input type="checkbox"/>	Experiment name	Last 5 runs	Created on	Created by
<input type="checkbox"/>	▶ tfma-experiment		6:17 PM, Aug 24, 2018	John Doe
<input type="checkbox"/>	▶ xgboost-train		6:17 PM, Aug 24, 2018	John Doe
<input type="checkbox"/>	▶ promo-email		6:17 PM, Aug 24, 2018	Walter Fisher
<input type="checkbox"/>	▶ data-prep		6:17 PM, Aug 24, 2018	Walter Fisher
<input type="checkbox"/>	▶ tf-preprocessing		6:17 PM, Aug 24, 2018	John Doe
<input type="checkbox"/>	▶ tf-training		6:17 PM, Aug 24, 2018	Walter Fisher
<input type="checkbox"/>	▶ tf-serving		6:17 PM, Aug 24, 2018	Walter Fisher

Rows per page: 10

1–10 of 241

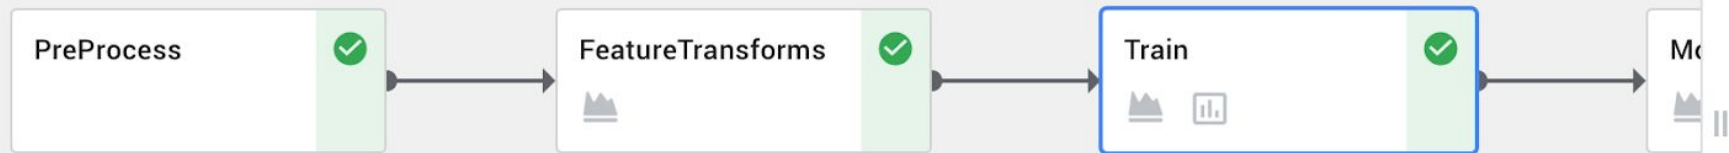
Select any Run to see all params and metrics



Experiments > creditcard-recommender

[←](#) ✓ ccard-recommender-run1 [Clone](#) [Archive](#)

[Graph](#) [Run output](#) [Config](#)

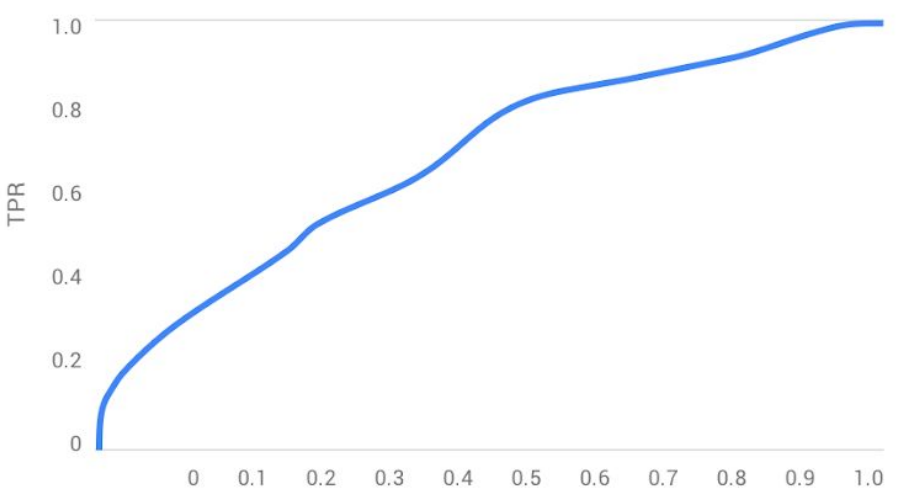


×

Train

[Artifacts](#) [Logs](#) [Config](#)


ROC curve



Tensorboard

[Open Tensorboard](#)

Clone an existing pipeline

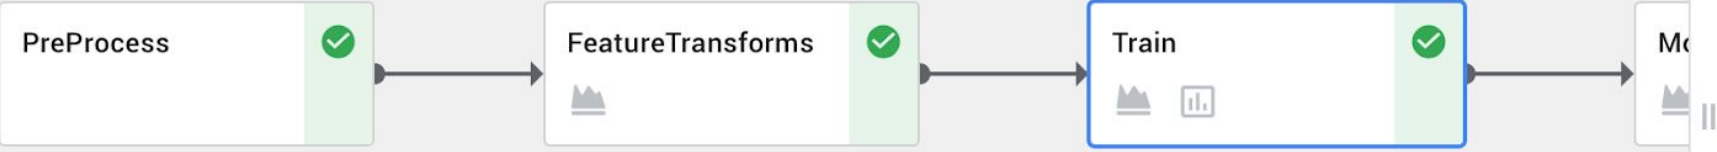


Experiments > creditcard-recommender

[←](#) ✓ ccard-recommender-run1

[Clone](#) [Archive](#)

[Graph](#) [Run output](#) [Config](#)

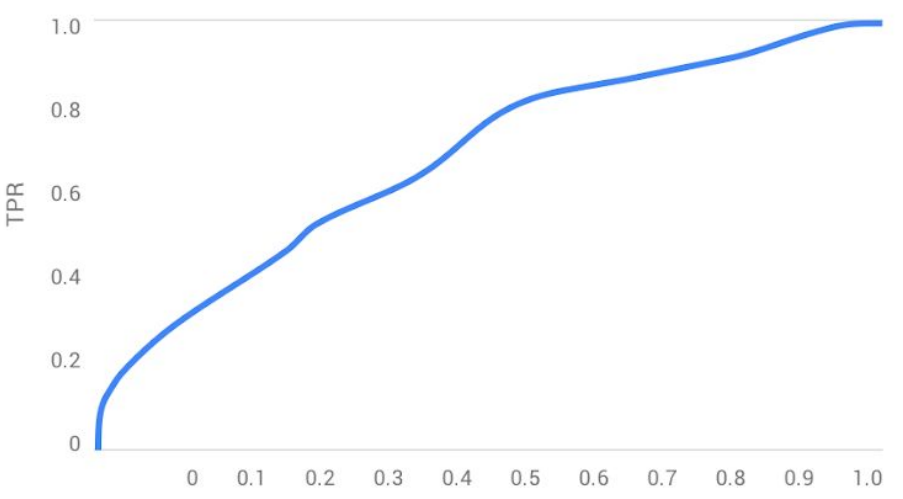


×

Train

[Artifacts](#) [Logs](#) [Config](#)


ROC curve



Tensorboard

[Open Tensorboard](#)


What does AI Platform Pipelines enable?



Workflow
orchestration



Rapid, reliable,
repeatable
experimentation



Share, re-use,
and compose




Share, re-use, and compose

- Re-use instead of re-implement
- Easy to compose or swap
 - Standard artifact types
 - Standard interfaces
 - Conventions and best practices

A lot of reusable
components


+

Composable
pipelines



Agenda

- Overview
- Introduction to AI Platform Pipelines
- AI Platform Pipelines: Concepts
- AI Platform Pipelines: When to Use?
- AI Platform Pipelines: Ecosystem



AI problems today

Talent


Lack of expertise in ML

Ecosystem

Difficult to find and leverage existing solutions

Flexibility

Brittle, opinionated infrastructure that breaks between cloud and on-premises



AI problems today

Talent

Lack of expertise in ML


Ecosystem

Difficult to find and leverage existing solutions


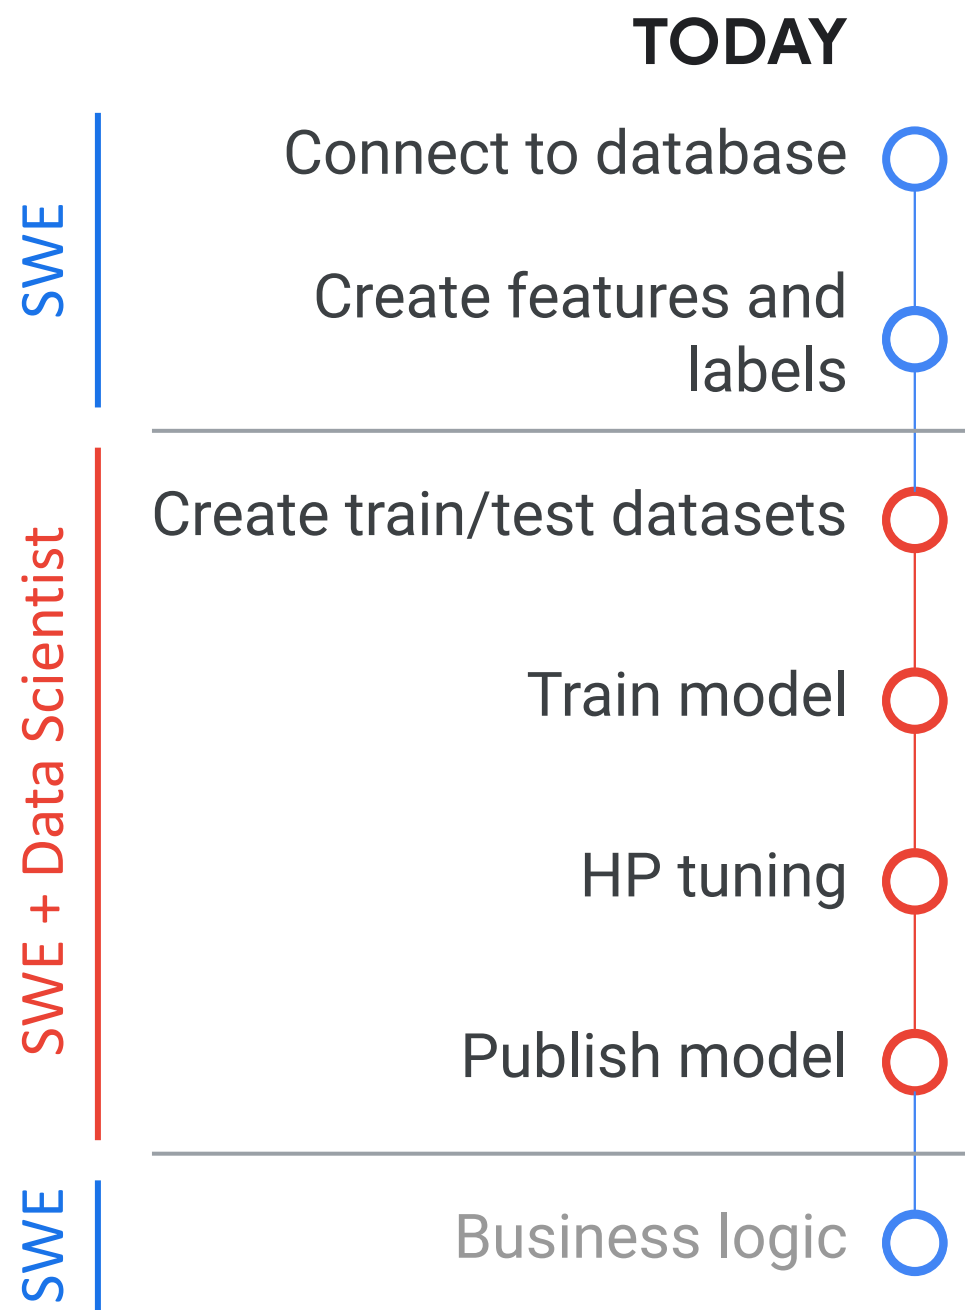
Flexibility

Brittle, opinionated infrastructure that breaks between cloud and on-premises

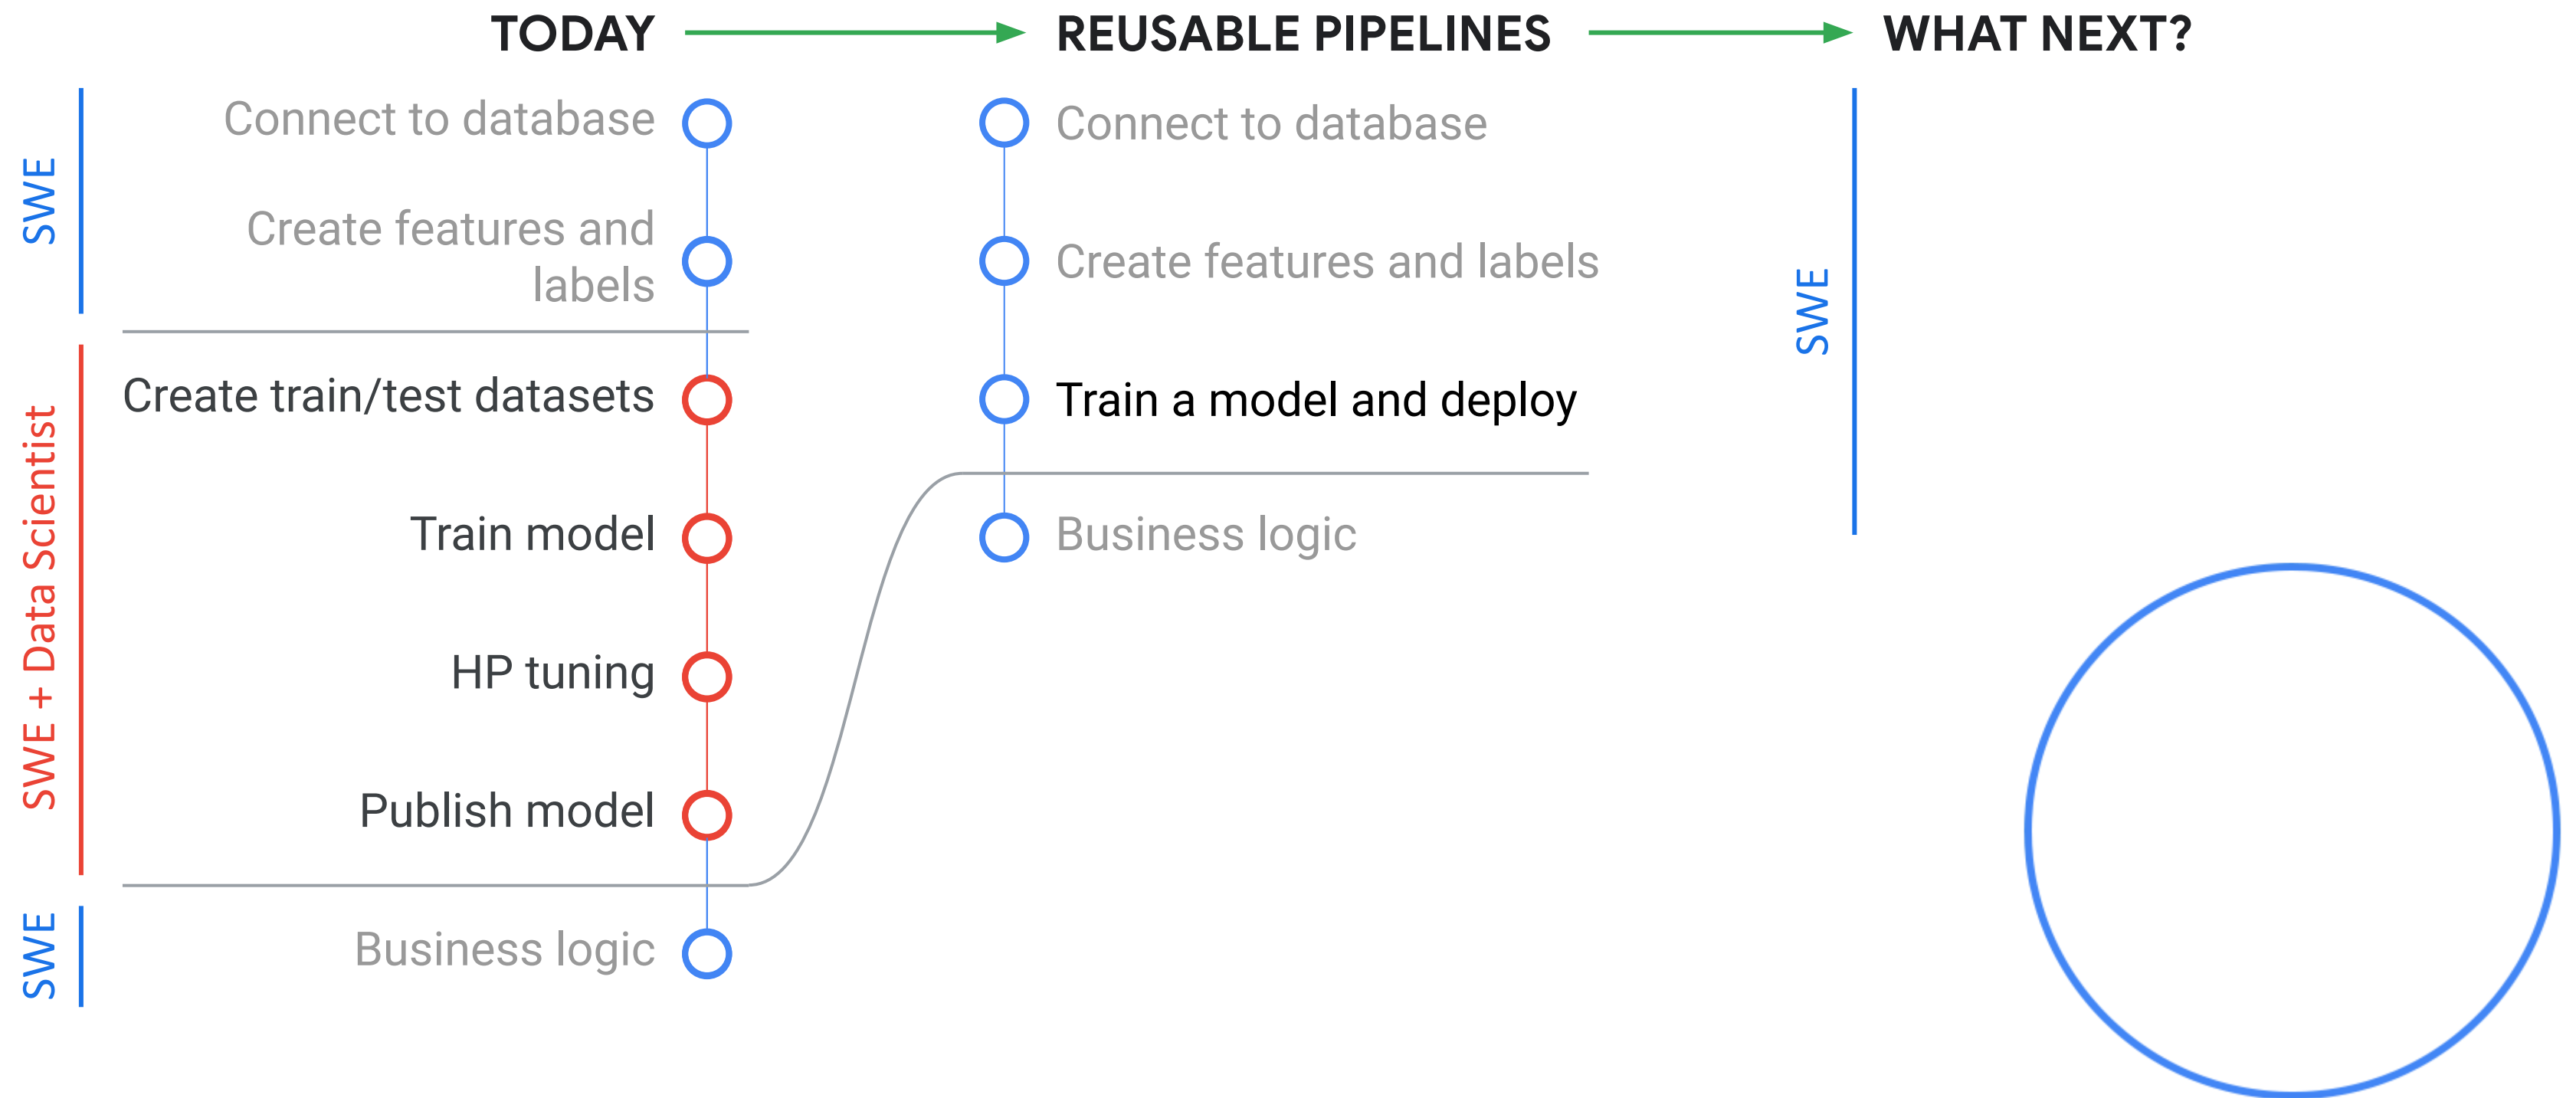
One of the most important factors to address for better collaboration and business impact



Reusable pipelines: Force multiplier for data scientists



Reusable pipelines: Force multiplier for data scientists



Mission

The one place for everything AI,
from experimentation to production



0.8x
Google Cloud AI Hub


AI Hub

Public Content

+ Private Content

By Google

Unique AI assets by Google

 AutoML, TPUs, [kaggle](#)
Cloud AI Platform, etc.

 Research at Google


 DeepMind

By Partners

Created, shared, and
monetized by anyone

By Customers

Content shared securely within and
with other organizations




AI Hub and Pipelines: Fast and simple adoption of AI

The flywheel of AI adoption

1. SEARCH and DISCOVER: Find best-of-breed pipelines on the hub that leverage Cloud AI solutions (AutoML, GPU, TPU, CMLE, etc.).

5. PUBLISH: Upload and share pipelines that run best on Google Cloud with your org or publicly.



2. DEPLOY: Quick 1-click implementation of ML pipelines onto Google Cloud/GKE.

3. CUSTOMIZE: Experiment and adjust out-of-the-box pipelines to custom use cases via pipelines UI.

4. RUN IN PRODUCTION: Deploy customized pipelines in production on Google Cloud.


Let's start building an AI Platform Pipeline

[Table of Contents](#)



Agenda

- Start building an AI Platform Pipelines Instance
- Run an example pipeline




Start with your GCP project

1. Open Cloud shell



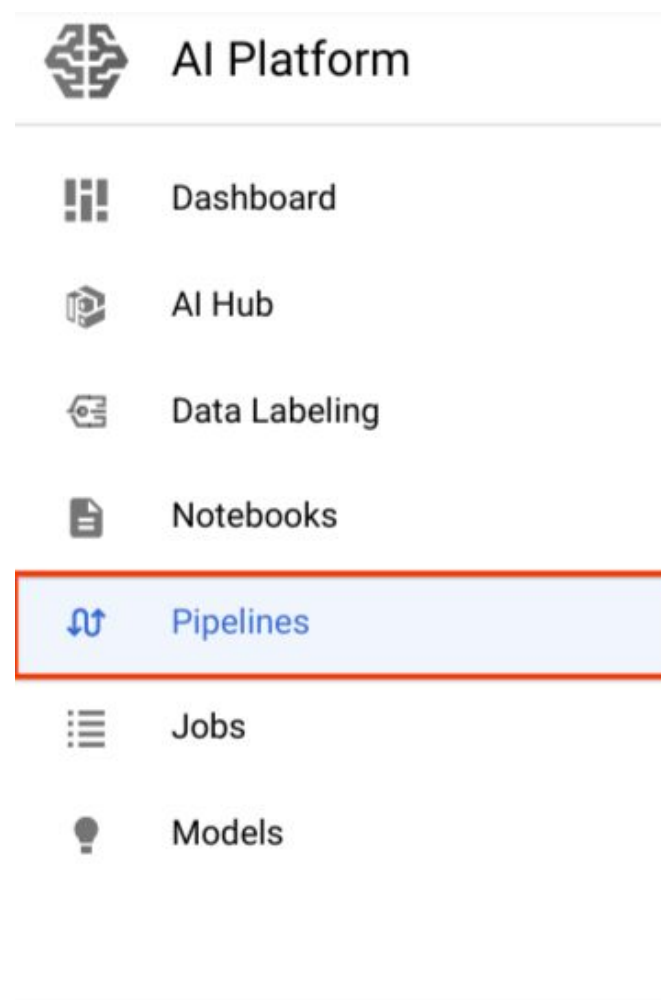
2. Edit IAM policy

```
gcloud projects get-iam-policy <you-project-name> --flatten="bindings[].members"  
--format="table(bindings.role, bindings.members)" --filter="bindings.role:roles/container.admin OR  
bindings.role:roles/viewer"
```




Start with your GCP project


3. Open AI platform Pipelines



4. Create a new instance



5. This will open Kubeflow pipeline page. Click “*Configure*”.



Start with your GCP project

6. Select any region for the cluster and check “access to Cloud APIs” for default access. Click “*Create cluster*”. This will automatically create cluster for you in the selected region. This may take a moment to create an instance.

Your app will use compute instances managed in a logical grouping called a "cluster", which will be configured in a way that's great for getting started with Kubernetes. For more options, visit the Kubernetes engine [cluster creation page](#).

Cluster zone ?

us-central1-a


☒ Allow access to the following Cloud APIs ?

<https://www.googleapis.com/auth/cloud-platform>

i All applications in cluster will have the access without additional settings. [Learn more](#) [↗](#)

Create cluster

[or select an existing cluster](#)



Start with your GCP project

7. After the cluster is created, Select “default” namespace and your App instance name **my-first-pipeline**. This step will take some time.

✔ Created "cluster-2" in zone "us-central1-a".

Namespace ?

default

App instance name ?

my-first-pipeline

☒ Use managed storage ?

Artifact storage Cloud Storage bucket (Managed storage only) ?

Cloud SQL instance connection name (Managed storage only) ?

Database username (Managed storage only) ?


Database password (Managed storage only) ?

Database name prefix (Managed storage only) ?

ⓘ This app has permission to modify resources at the cluster scope.

⌵ More

Deploy




Start with your GCP project

8. This will create App instance as well as AI Platform pipeline instance. Click on “*Open Pipelines Dashboard*”


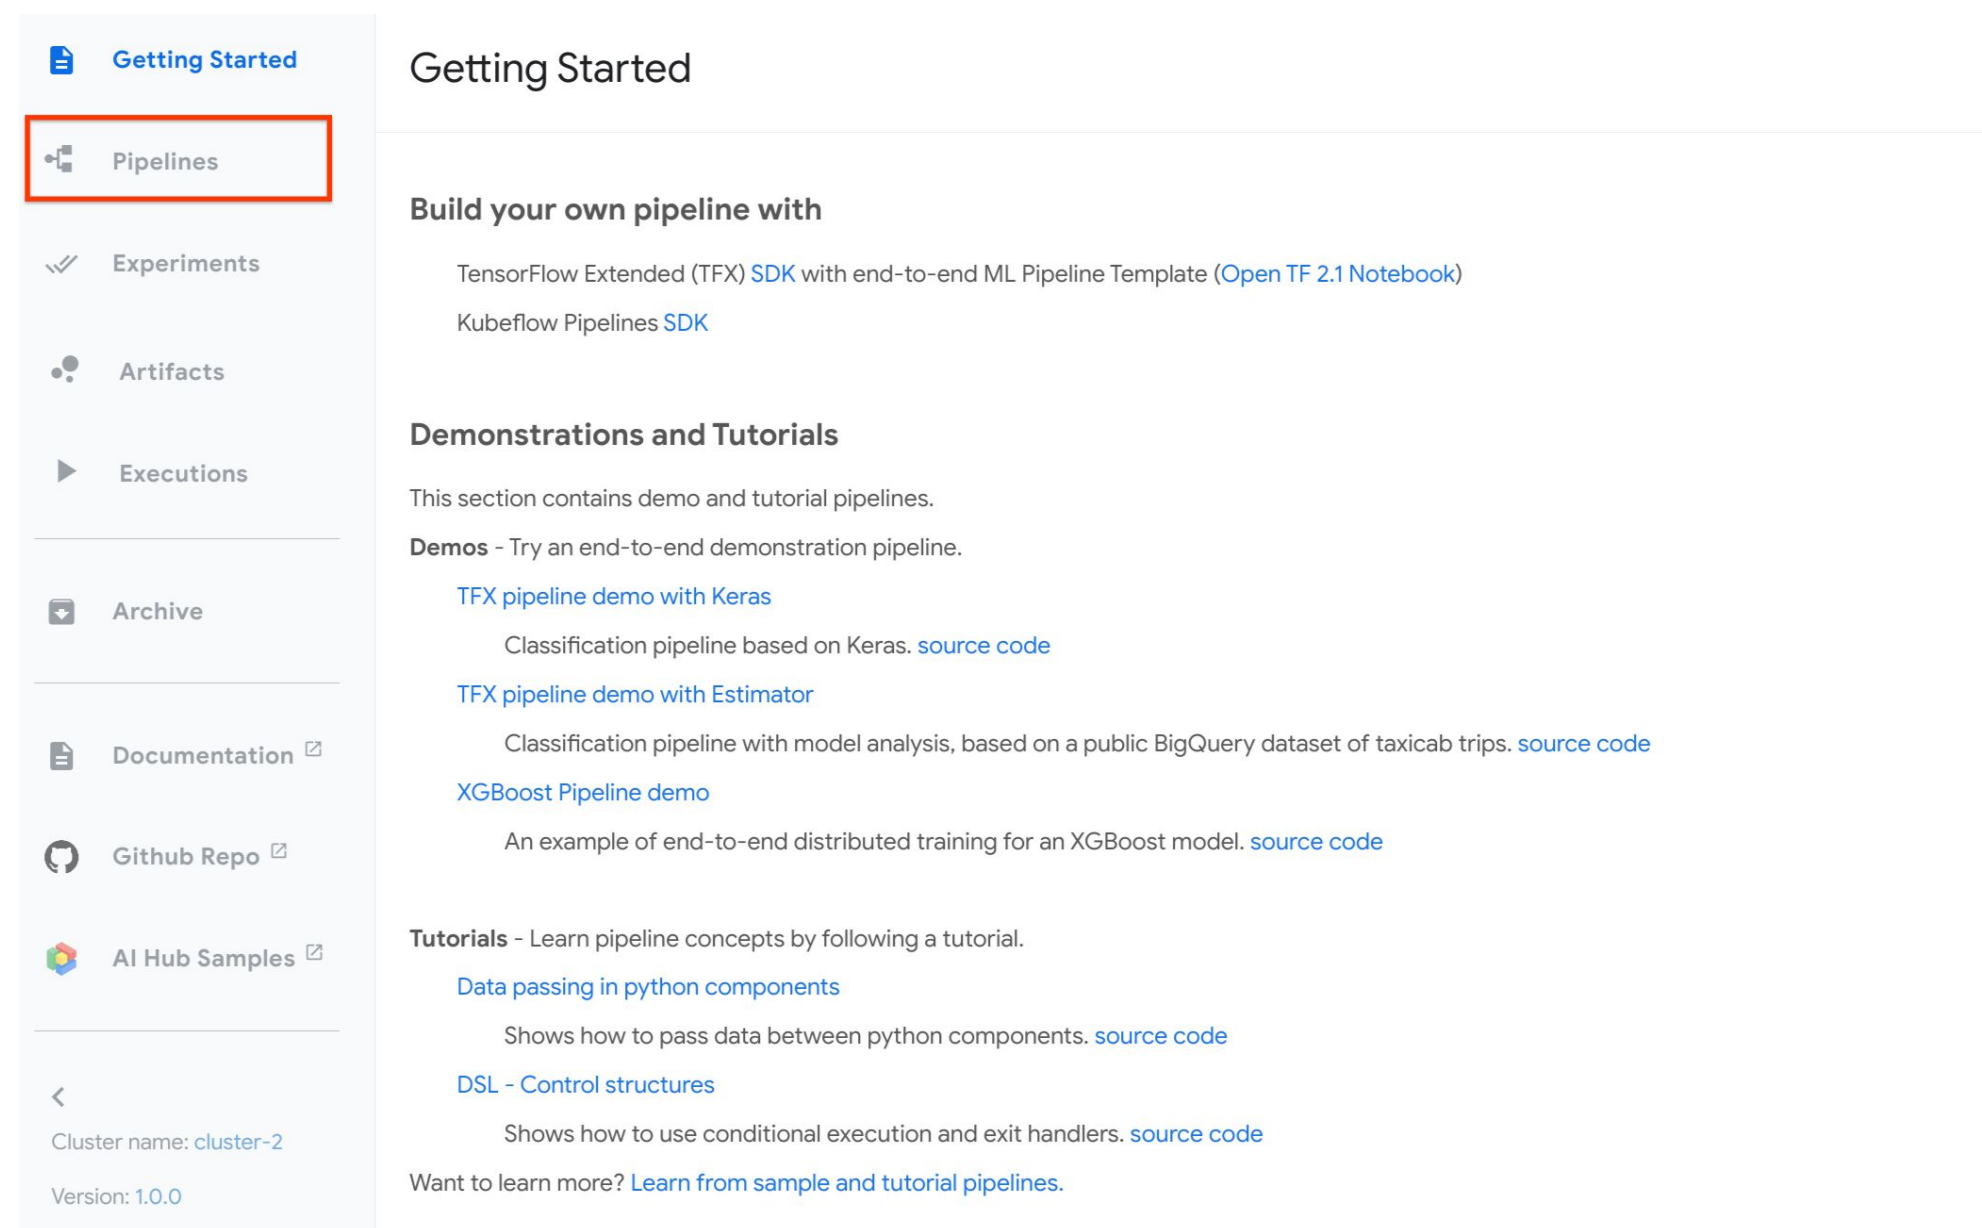
AI Platform Pipelines BETA NEW INSTANCE REFRESH DELETE LEARN MORE

<input type="checkbox"/>	Status	Name ↑		Zone	Version	Cluster	Namespace	
<input type="checkbox"/>	✓	my-first-pipeline	<a>OPEN PIPELINES DASHBOARD	us-central1-a	1.0.0	cluster-2	default	<a>SETTINGS




Start with your GCP project

9. This will open AI Platform Pipelines Dashboard. Click “*Pipelines*”




Agenda

- Start building an AI Platform Pipelines Instance
- Run an example pipeline



Start with an example pipeline

1. Let's start with an example pipeline “[*Demo*] TFX - Taxi Tip Prediction Model Trainer”
2. To run or schedule the pipeline, click Create run. A form where you can enter the run details opens.
3. You will need to add Run details that includes:
 - a. Pipeline
 - b. Pipeline version
 - c. Run name
 - d. Description
 - e. Experiment



Start with an example pipeline

Getting Started

Pipelines

Experiments

Artifacts

Executions

Archive

Documentation

Github Repo

AI Hub Samples

Cluster name: cluster-3

Version: 1.0.0

Pipelines

+ Upload pipeline


Refresh

Delete

Filter pipelines

<input type="checkbox"/>	Pipeline name	Description	Uploaded on ↓
<input type="checkbox"/>	▶ [Tutorial] DSL - Control str...	source code Shows how to use conditional execution and exit handlers. This pipelin...	10/7/2020, 8:02:10 PM
<input type="checkbox"/>	▶ [Tutorial] Data passing in p...	source code Shows how to pass data between python components.	10/7/2020, 8:02:08 PM
<input type="checkbox"/>	▶ [Demo] TFX - Iris classifica...	source code . Example pipeline that classifies Iris flower subspecies and how to use...	10/7/2020, 8:02:07 PM
<input type="checkbox"/>	▶ [Demo] TFX - Taxi tip predi...	source code GCP Permission requirements . Example pipeline that does classificati...	10/7/2020, 8:02:06 PM
<input type="checkbox"/>	▶ [Demo] XGBoost - Training ...	source code GCP Permission requirements . A trainer that does end-to-end distribut...	10/7/2020, 8:02:05 PM

Rows per page: 10 < >



Start with an example pipeline

Pipelines

← [Demo] TFX - Taxi tip prediction model trai...

+ Create run

+ Upload version

+ Create experiment

Delete

Graph

YAML

csvexamplegen

statisticsgen

schemagen

Summary

Hide

ID

b781b2d1-8b6a-4195-b61c-2895a56e96ac

Version

[Demo] TFX - Taxi tip prediction model trainer

Version source


Uploaded on

10/7/2020, 8:02:06 PM

Description

[source code](#) [GCP Permission requirements](#). Example pipeline that does classification with model analysis based on a public tax cab dataset.

Static pipeline graph



Final output from the experiment pipeline

← ✓ Run of [Demo] TFX - Taxi tip prediction model trainer (202fa)

