Training, Tuning, and Serving on AI Platform

Priscilla Moraes
Learning Portfolio Manager, ML and AI

Agenda

System and Concepts Overview

Create a Reproducible Dataset

Implement a Tunable Model

Build and Push a Training Container

Train and Tune a Model

Serve and Query a Model

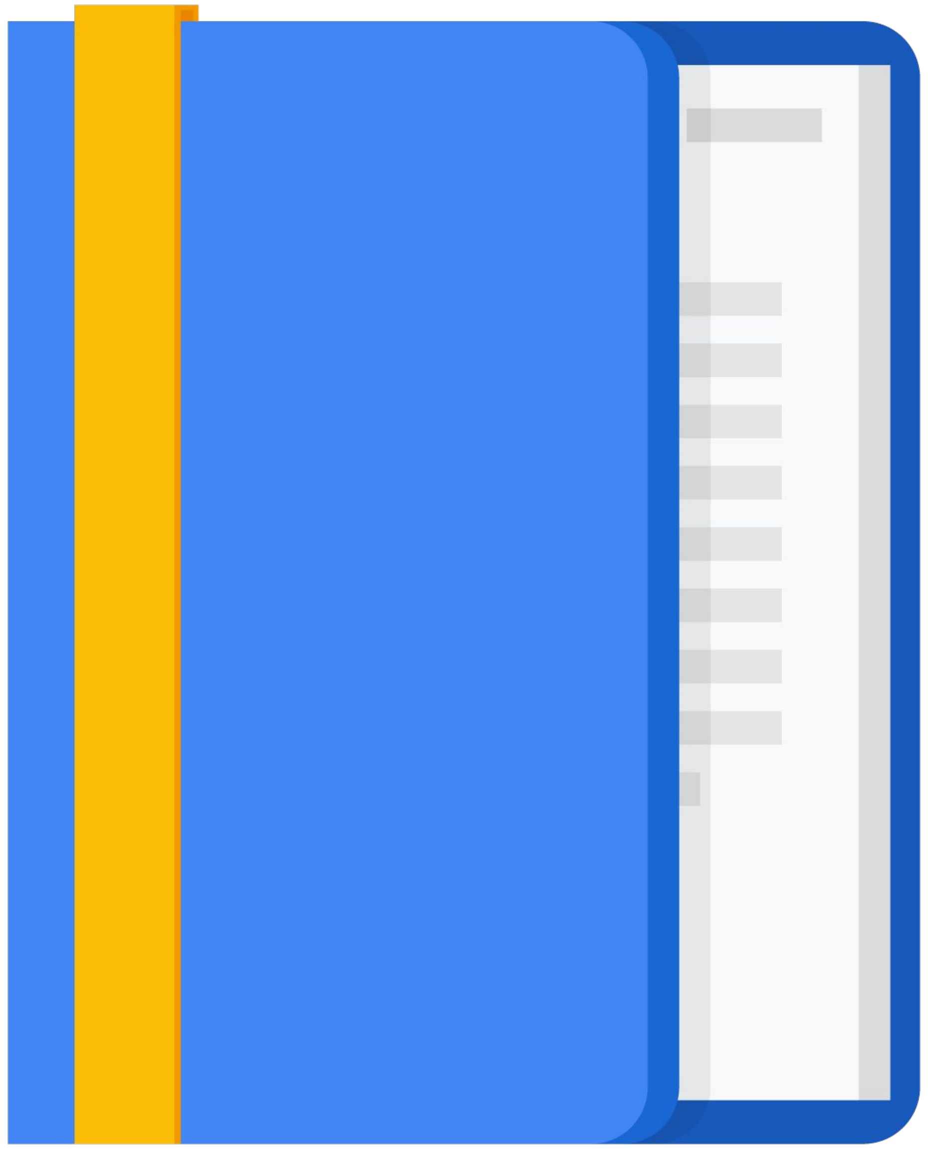

ML model building process

Create
the dataset

Build
the model

Operationalize
the model

Building and operationalizing the model

train.py

Implement a tunable training application

Step 1

Dockerfile

Package your training application

Step 2

config.yaml

Configure and start an AI Platform job

Step 3

MLOps building blocks on Google Cloud in this module

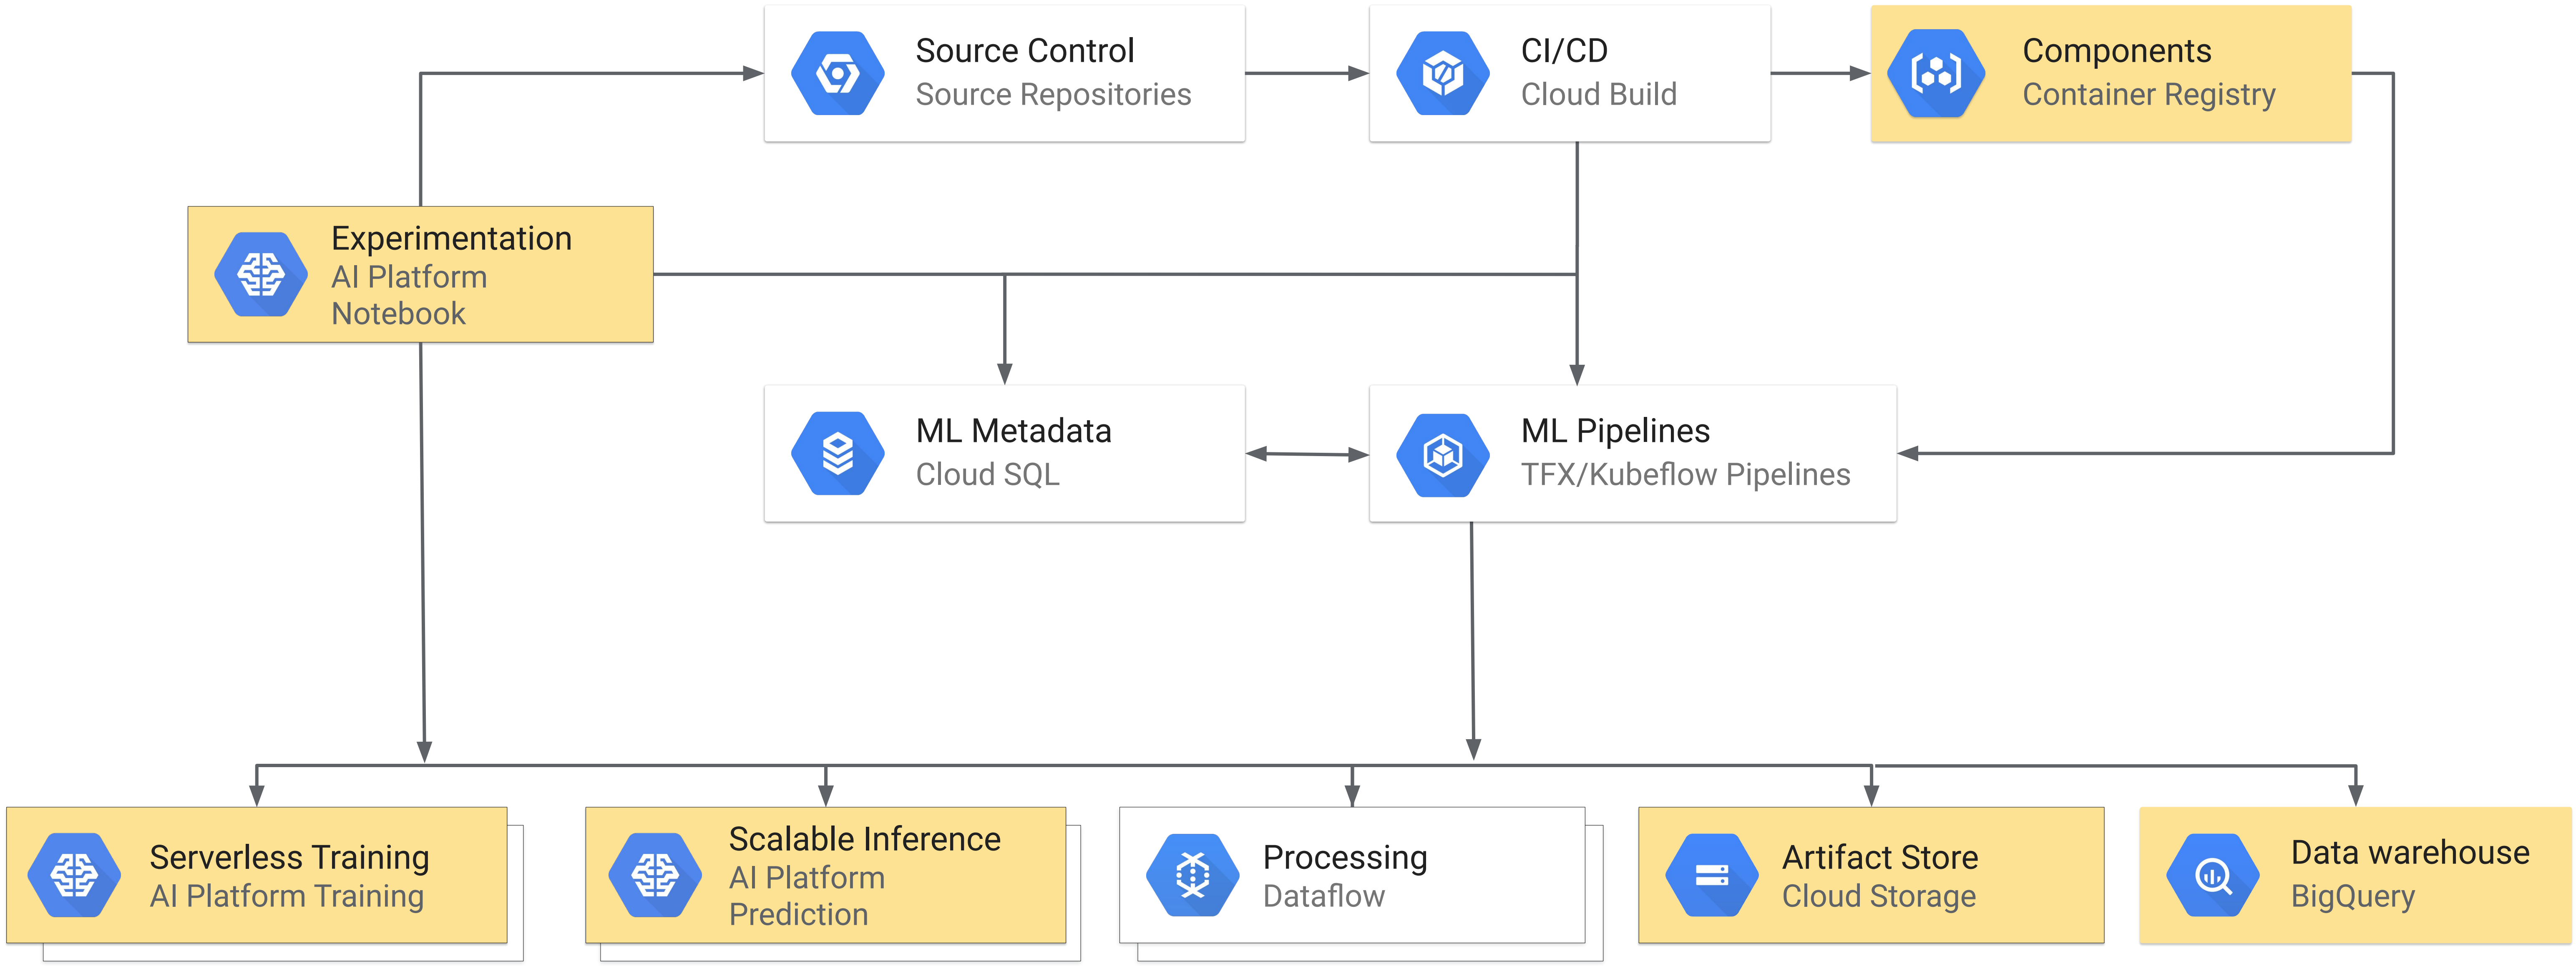

System overview

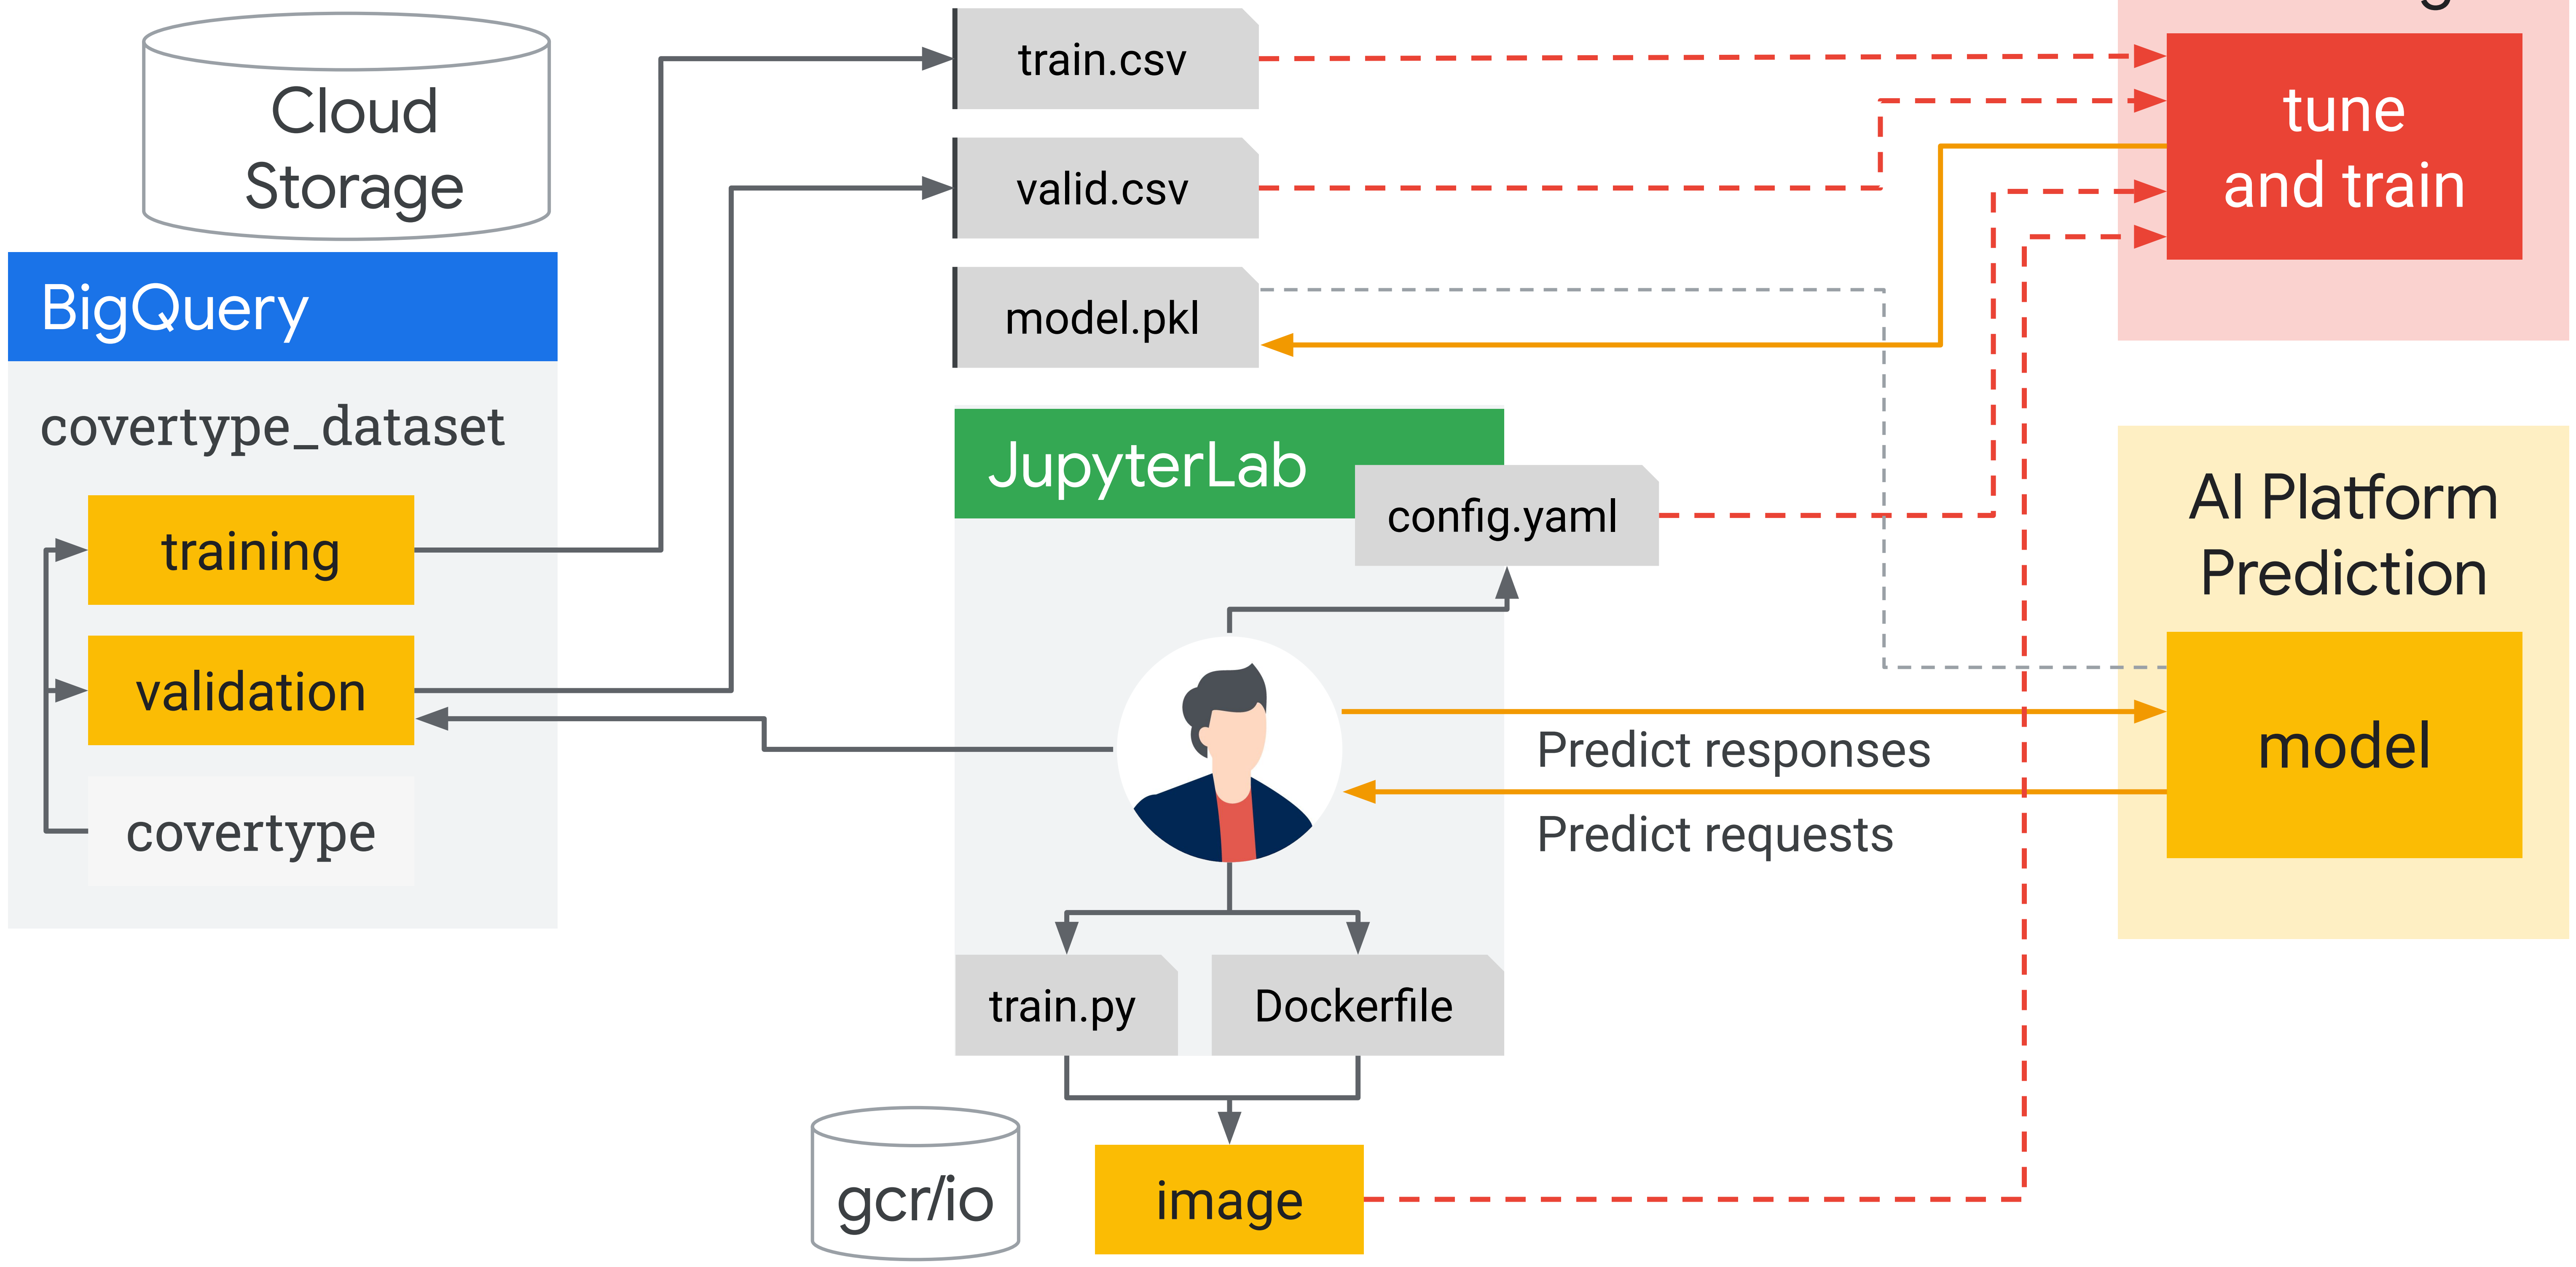

Where we are going next

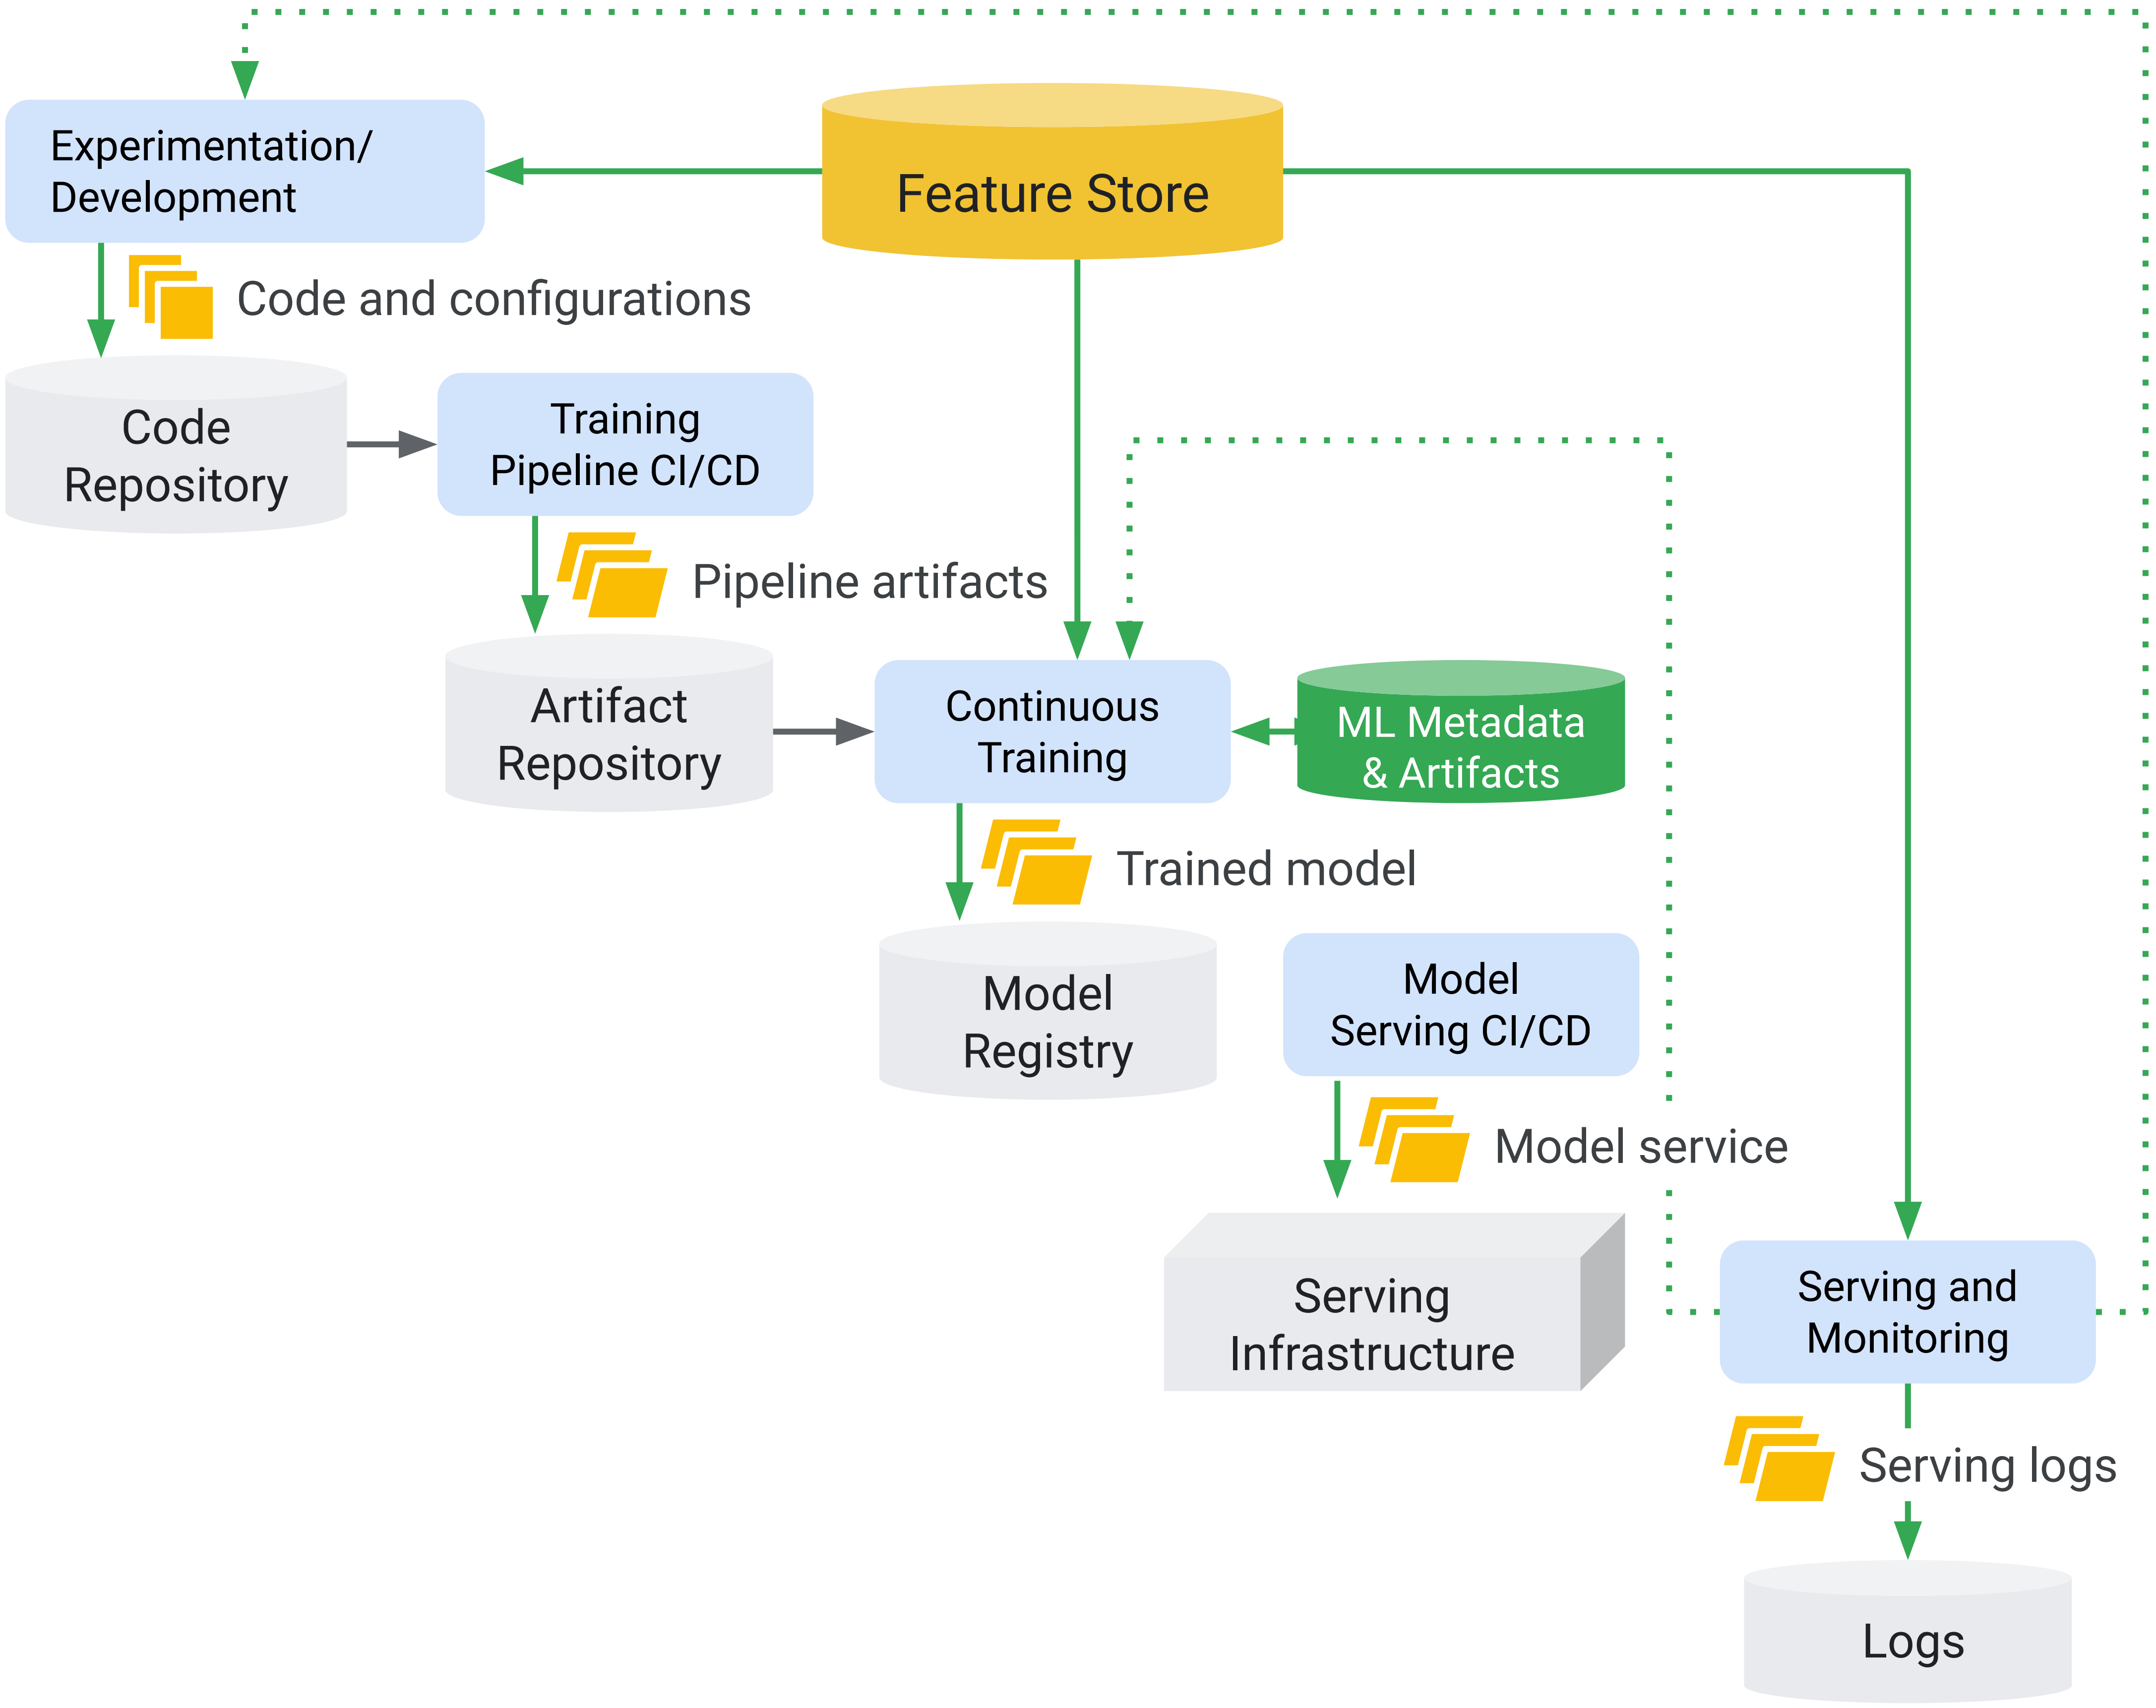

Agenda

System and Concepts Overview

Create a Reproducible Dataset

Implement a Tunable Model

Build and Push a Training Container

Train and Tune a Model

Serve and Query a Model

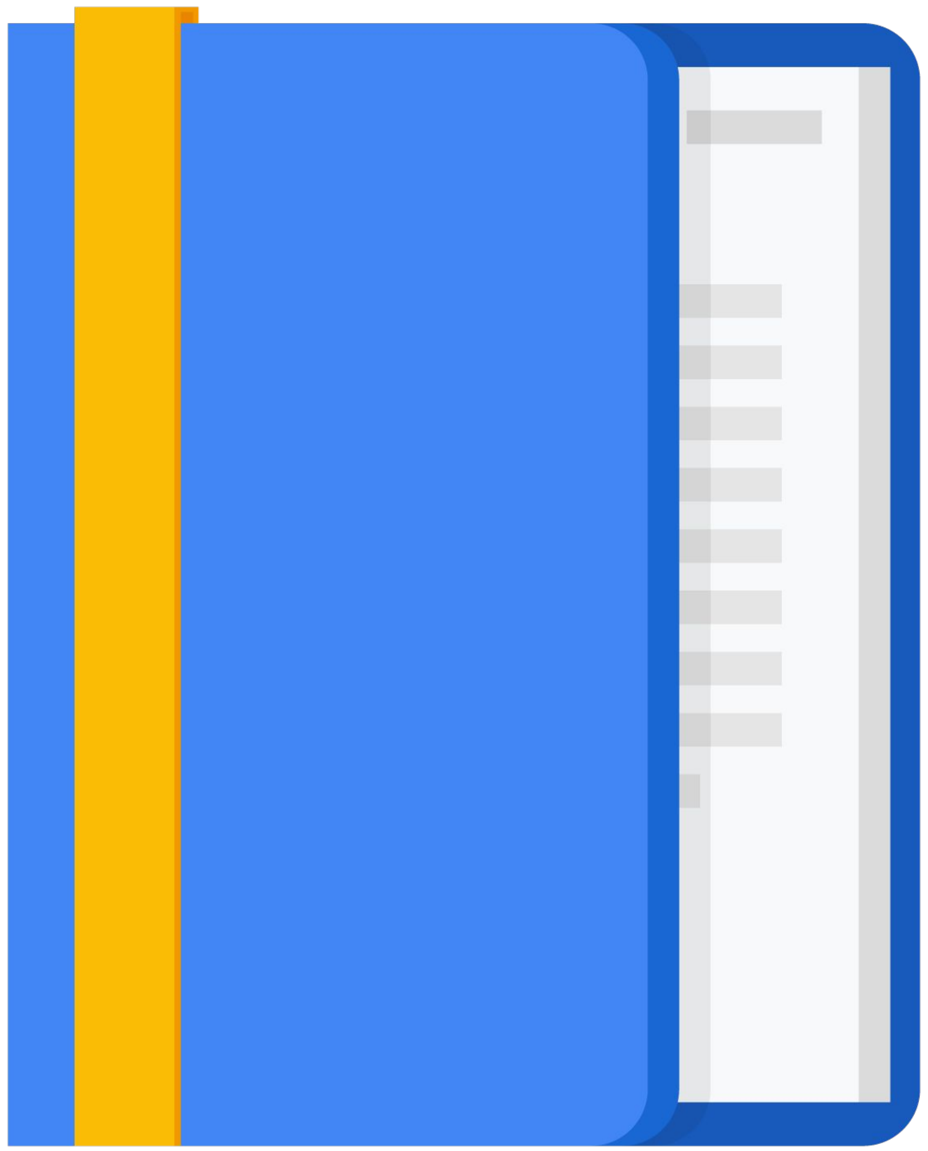

System overview

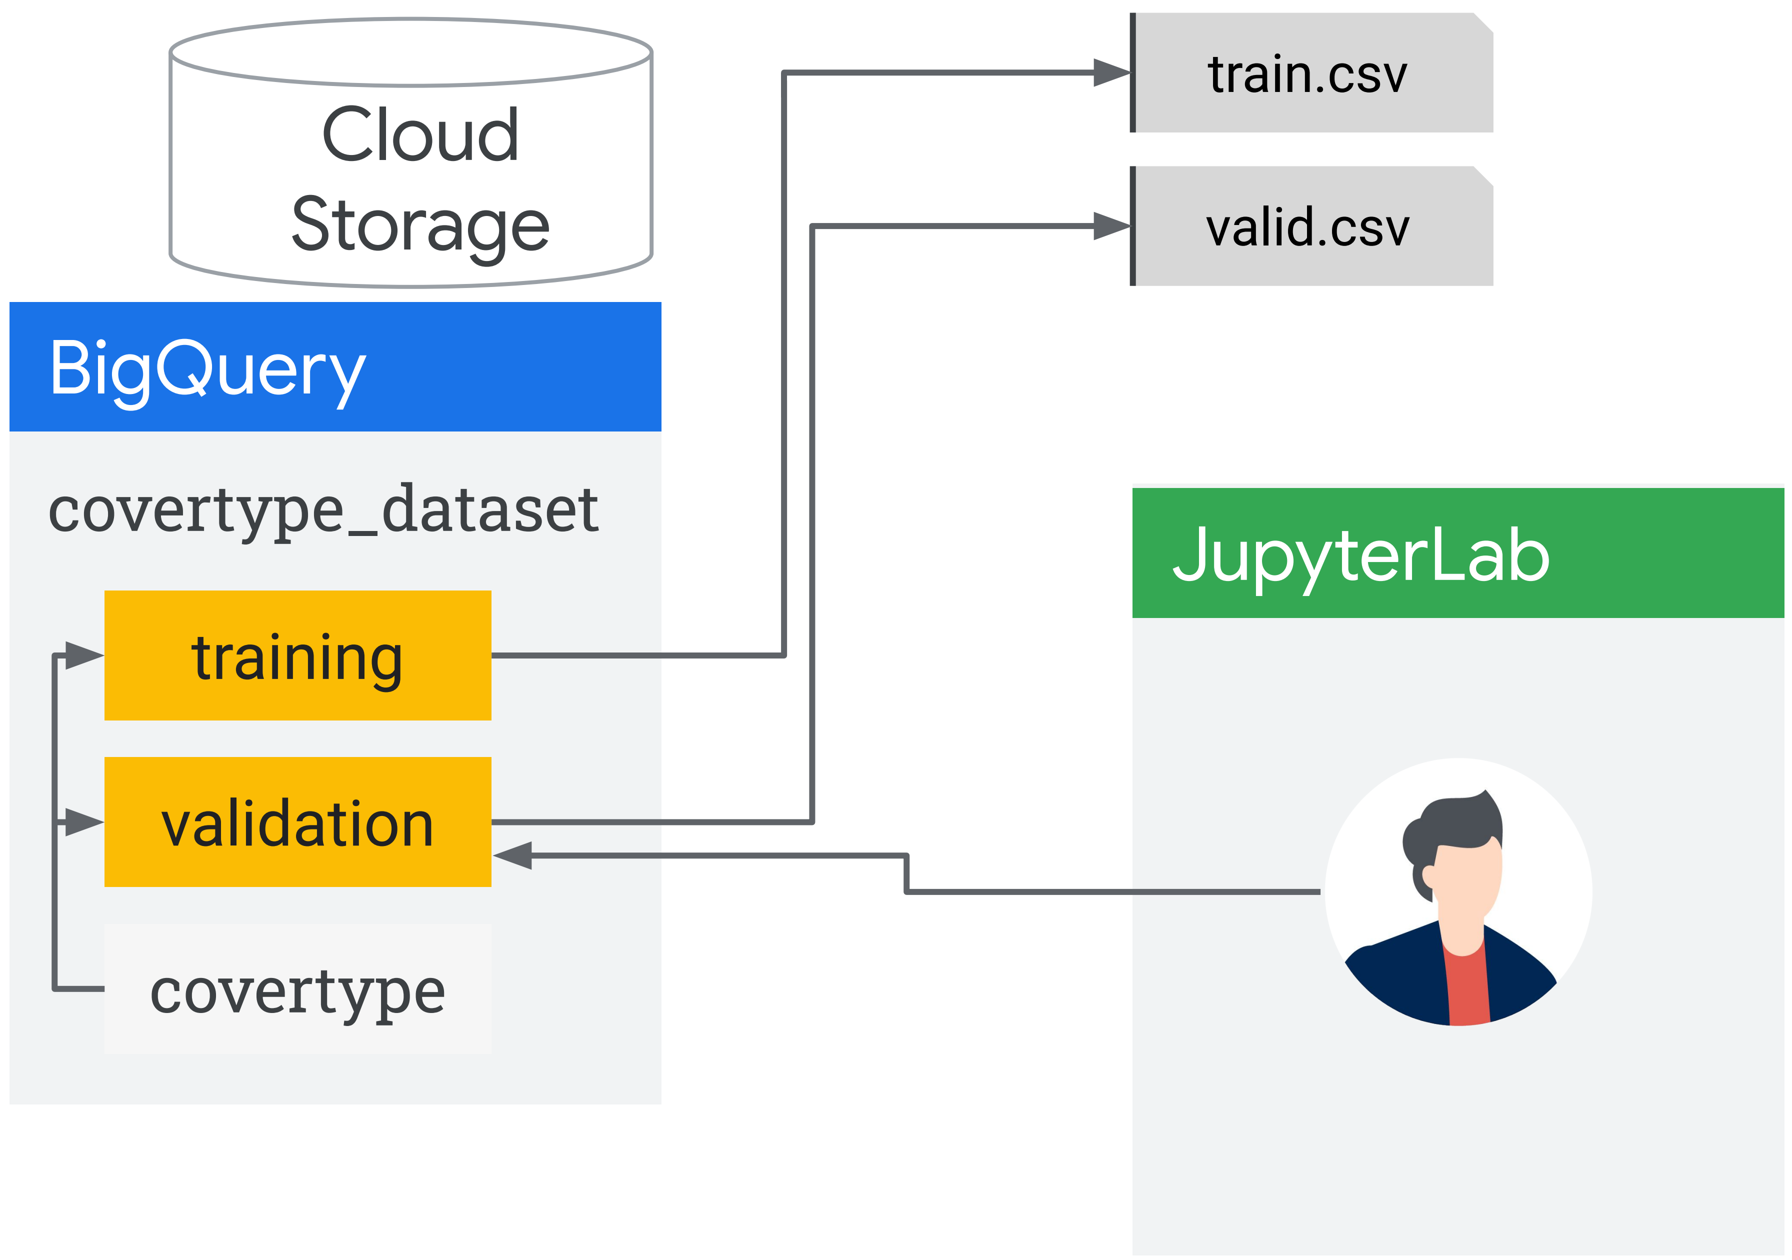

Field name	Type
Elevation	INTEGER
Aspect	INTEGER
Slope	INTEGER
Horizontal_Distance_To_Hydrology	INTEGER
Vertical_Distance_To_Hydrology	INTEGER
Horizontal_Distance_To_Roadways	INTEGER
Hillshade_Noon	INTEGER
Hillshade_3pm	INTEGER
Horizontal_Distance_To_Fire_Points	INTEGER
Wilderness_Area	STRING
Soil_Type	STRING
Cover_Type	INTEGER

UCI

Machine Learning Repository

Center for Machine Learning and Intelligent System

Covertype Data Set

Download: Data Folder, Data Set Description

Abstract: Forest CoverType dataset

Data Set Characteristics:	Multivariate	**Number of Instances:**	581012	**Area:**	Life
Attribute Characteristics:	Categorical, Integer	**Number of Attributes:**	54	**Date Donated**	1998-08-01
Associated Tasks:	Classification	**Missing Values?**	No	**Number of Web Hits:**	289499

https://archive.ics.uci.edu/ml/datasets/covertype

Features

Target

Elevation	Aspect	Slope	Horizontal_Distance_To_Hydrology	Vertical_Distance_To_Hydrology	Horizontal_Distance_To_Roadways	Hillshade_9am	Hillshade_Noon	Hillshade_3pm	Horizontal_Distance_To_Fire_Points	Wilderness_Area	Soil_Type	Cover_Type
2067	0	21	270	9	755	184	196	145	900	Cache	C2702	5
2574	0	2	319	20	1419	216	235	156	1595	Commanche	C2703	4
2559	0	0	510	16	1113	218	238	156	1332	Commanche	C2703	2
2647	0	6	402	94	641	212	229	155	1104	Commanche	C2703	2
2651	0	3	335	103	488	215	233	156	1381	Commanche	C2703	2
2647	0	6	417	94	648	212	229	155	1082	Commanche	C2703	2
2639	0	10	366	80	589	206	222	154	1041	Commanche	C2703	2
2590	0	2	201	13	1200	216	235	156	1719	Commanche	C2703	1
2447	0	4	0	0	631	213	232	156	711	Commanche	C2705	5
2501	0	6	228	31	1012	211	228	155	930	Commanche	C2705	1
2500	0	4	30	3	1746	213	232	156	886	Commanche	C2705	5
2641	0	1	90	15	1518	217	236	156	182	Commanche	C2705	2

Split the dataset and experiment with models

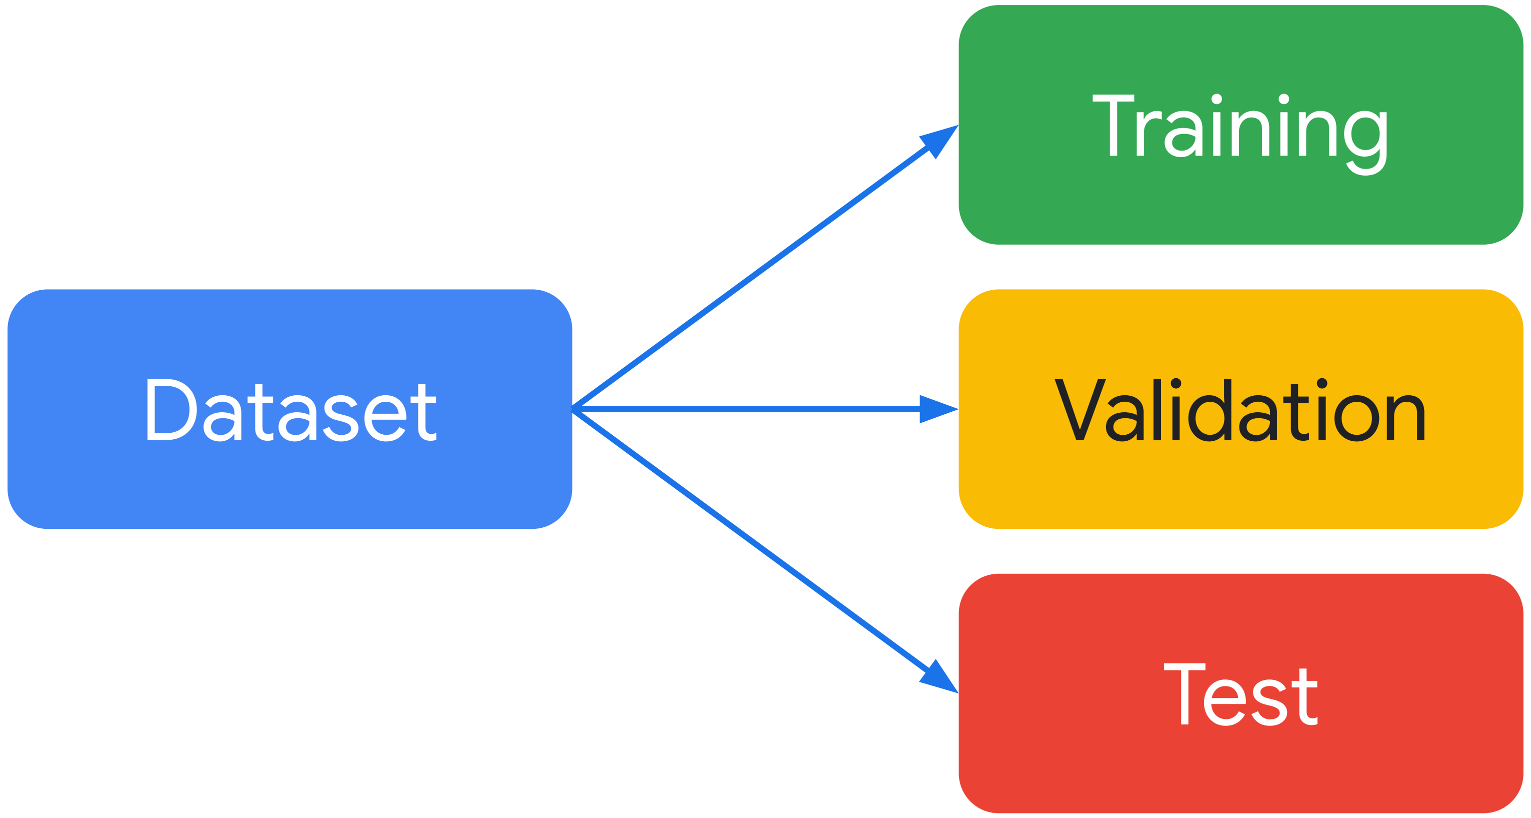

Getting a random 80% of your dataset for training is easy

```
#standardSQL
SELECT
  date,
  airline,
  departure_airport,
  departure_schedule,
  arrival_airport,
  arrival_delay
FROM
  `bigquery-samples.airline_ontime_data.flights`
WHERE
  RAND() < 0.8
```

RAND will return a number between 0 and 1.

However, experimentation requires repeatability

You need to know which specific data was involved in training, validation, and testing.

Naive random splitting is not repeatable

The order of rows in BigQuery is not certain without ORDER BY.

Identifying and splitting the remaining 20% of data for validation and testing is difficult.

RAND() will return different results each time →

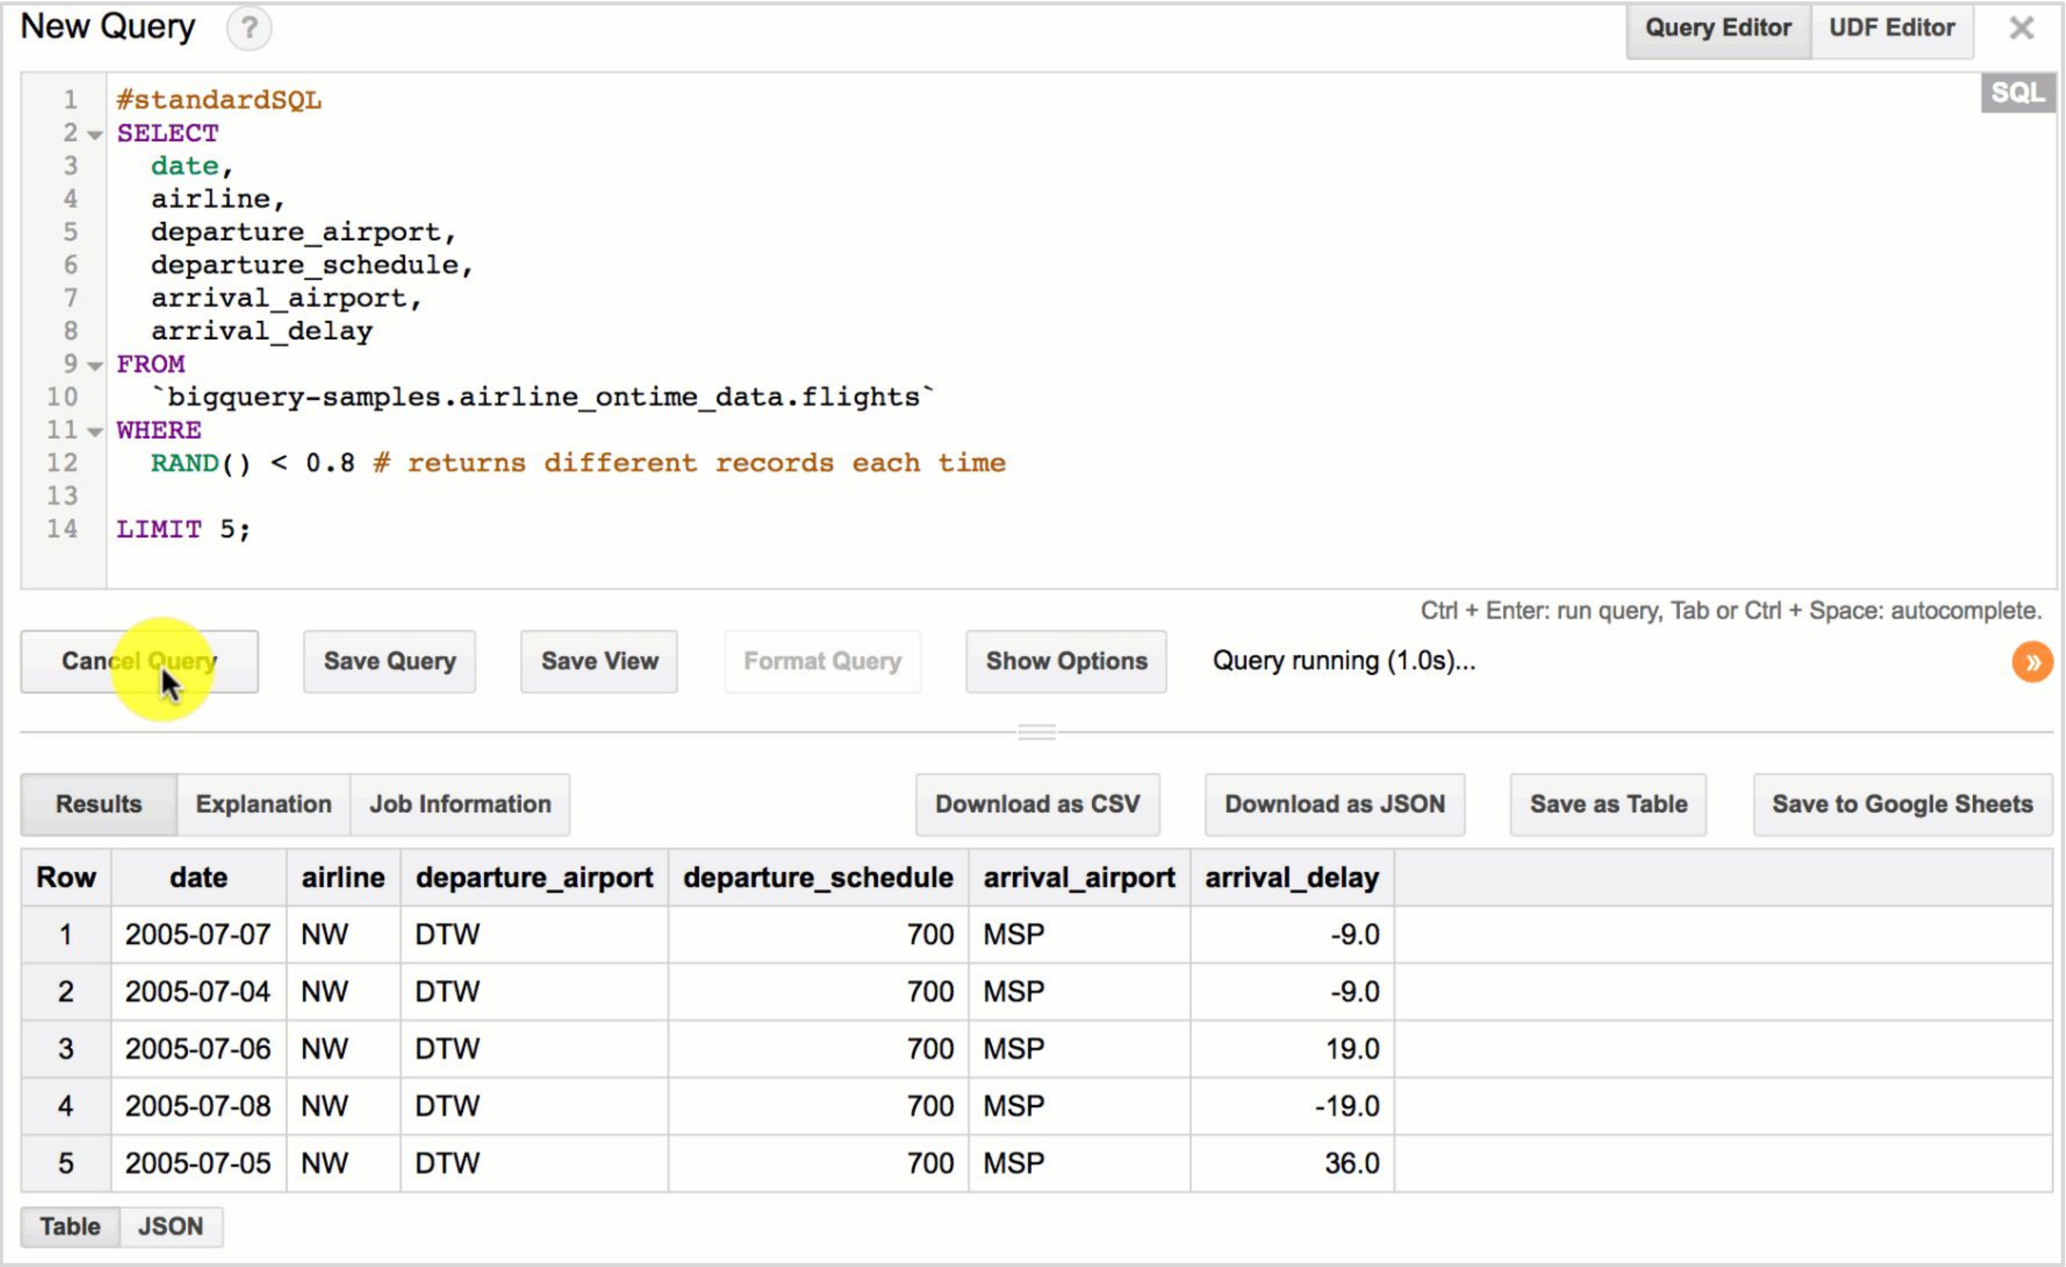

Solution: Use hashing and modulo operators to split a dataset into training/validation

Input		Hash sum
Fox	Hash function	DFCD3454
The red fox **runs** across the ice	Hash function	52ED879E
The red fox **walks** across the ice	Hash function	46042841

Solution: Use hashing and modulo operators to split a dataset into training/validation

```
#standardSQL
SELECT
  date,
  airline,
  departure_airport,
  departure_schedule,
  arrival_airport,
  arrival_delay
FROM
 `bigquery-samples.airline_ontime_data.flights`

WHERE
  MOD(ABS(FARM_FINGERPRINT(date)),10) < 8
```

Solution: Use hashing and modulo operators to split a dataset into training/validation

```
#standardSQL
SELECT
  date
  airline,
  departure_airport,
  departure_schedule,
  arrival_airport,
  arrival_delay
FROM
 `bigquery-samples.airline_ontime_data.flights`

WHERE
  MOD(ABS(FARM_FINGERPRINT(date)),10) < 8
```

Note: Even though we select date, our model wouldn't actually use it during training.

Hash value on the date will always return the same value.

Then we use a modulo operator to pull only 80% of that data based on the last few hash digits.

Solution: Use hashing and modulo operators to split a dataset into training/validation

```
#standardSQL
SELECT
  date
  airline,
  departure_airport,
  departure_schedule,
  arrival_airport,
  arrival_delay
FROM
 `bigquery-samples.airline_ontime_data.flights`

WHERE
  MOD(ABS(FARM_FINGERPRINT(date)),10) == 8
```

Note: Even though we select date, our model wouldn't actually use it during training.

Training

Solution: Use hashing and modulo operators to split a dataset into training/validation

```
#standardSQL
SELECT
  date,
  airline,
  departure_airport,
  departure_schedule,
  arrival_airport,
  arrival_delay
FROM
 `bigquery-samples.airline_ontime_data.flights`

WHERE
  MOD(ABS(FARM_FINGERPRINT(date)),10) == 9
```

Note: Even though we select date, our model wouldn't actually use it during training.

Testing

Solution: Use hashing and modulo operators to split a dataset into training/validation

```
#standardSQL
SELECT
  date,
  airline,
  departure_airport,
  departure_schedule,
  arrival_airport,
  arrival_delay
FROM
 `bigquery-samples.airline_ontime_data.flights`

WHERE
  MOD(ABS(FARM_FINGERPRINT(date)),10) == 9
```

Note: Even though we select date, our model wouldn't actually use it during training.

Which field to hash on?

1. Not correlated to label *(otherwise, you'll leave valuable information out of the training set)*

Which field to hash on?

1. Not correlated to label *(otherwise, you'll leave valuable information out of the training set)*

2. Granular enough for your desired module split

Which field to hash on?

Which field to hash on?

Possible solution: Concatenate all the fields as a JSON string, and hash on that.

```
TO_JSON_STRING(cover)
```

Create a training split

```
bq query \
  -n 0 \
  --destination_table covertype_dataset.training \
  --replace \
  --use_legacy_sql=false \
    'SELECT * \
     FROM `covertype_dataset.covertype` AS cover \
     WHERE \
     MOD(ABS(FARM_FINGERPRINT(TO_JSON_STRING(cover))), 10) IN (1, 2, 3, 4)'
```

Create the training table in BigQuery.

Create a training split

```
bq query \
  -n 0 \
  --destination_table covertype_dataset.training \
  --replace \
  --use_legacy_sql=false \
    'SELECT * \
     FROM `covertype_dataset.covertype` AS cover \
     WHERE \
     MOD(ABS(FARM_FINGERPRINT(TO_JSON_STRING(cover))), 10) IN (1, 2, 3, 4)'
```

Create the training table in BigQuery.

Create a training split

```
bq query \
  -n 0 \
  --destination_table covertype_dataset.training \
  --replace \
  --use_legacy_sql=false \
    'SELECT * \
     FROM `covertype_dataset.covertype` AS cover \
     WHERE \
     MOD(ABS(FARM_FINGERPRINT(TO_JSON_STRING(cover))), 10) IN (1, 2, 3, 4)'
```

Create the training table in BigQuery.

```
bq extract \
  --destination_format CSV \
  covertype_dataset.training \
  $TRAINING_FILE_PATH
```

Export it to Cloud Storage as a CSV file.

Create a training split

```
bq query \
  -n 0 \
  --destination_table covertype_dataset.training \
  --replace \
  --use_legacy_sql=false \
    'SELECT * \
     FROM `covertype_dataset.covertype` AS cover \
     WHERE \
     MOD(ABS(FARM_FINGERPRINT(TO_JSON_STRING(cover))), 10) IN (1, 2, 3, 4)'
```

Create the training table in BigQuery.

```
bq extract \
  --destination_format CSV \
  covertype_dataset.training \
  $TRAINING_FILE_PATH
```

Export it to Cloud Storage as a CSV file.

Do the same for the validation split

```
bq query \
  -n 0 \
  --destination_table covertype_dataset.validation \
  --replace \
  --use_legacy_sql=false \
    'SELECT * \
     FROM `covertype_dataset.covertype` AS cover \
     WHERE \
     MOD(ABS(FARM_FINGERPRINT(TO_JSON_STRING(cover))), 10) IN (8)'

bq extract \
  --destination_format CSV \
  covertype_dataset.validation \
  $VALIDATION_FILE_PATH
```

Do the same for the validation split

```
bq query \
  -n 0 \
  --destination_table covertype_dataset.validation \
  --replace \
  --use_legacy_sql=false \
    'SELECT * \
     FROM `covertype_dataset.covertype` AS cover \
     WHERE \
     MOD(ABS(FARM_FINGERPRINT(TO_JSON_STRING(cover))), 10) IN (8)'

bq extract \
  --destination_format CSV \
  covertype_dataset.validation \
  $VALIDATION_FILE_PATH
```

Agenda

System and Concepts Overview

Create a Reproducible Dataset

Implement a Tunable Model

Build and Push a Training Container

Train and Tune a Model

Serve and Query a Model

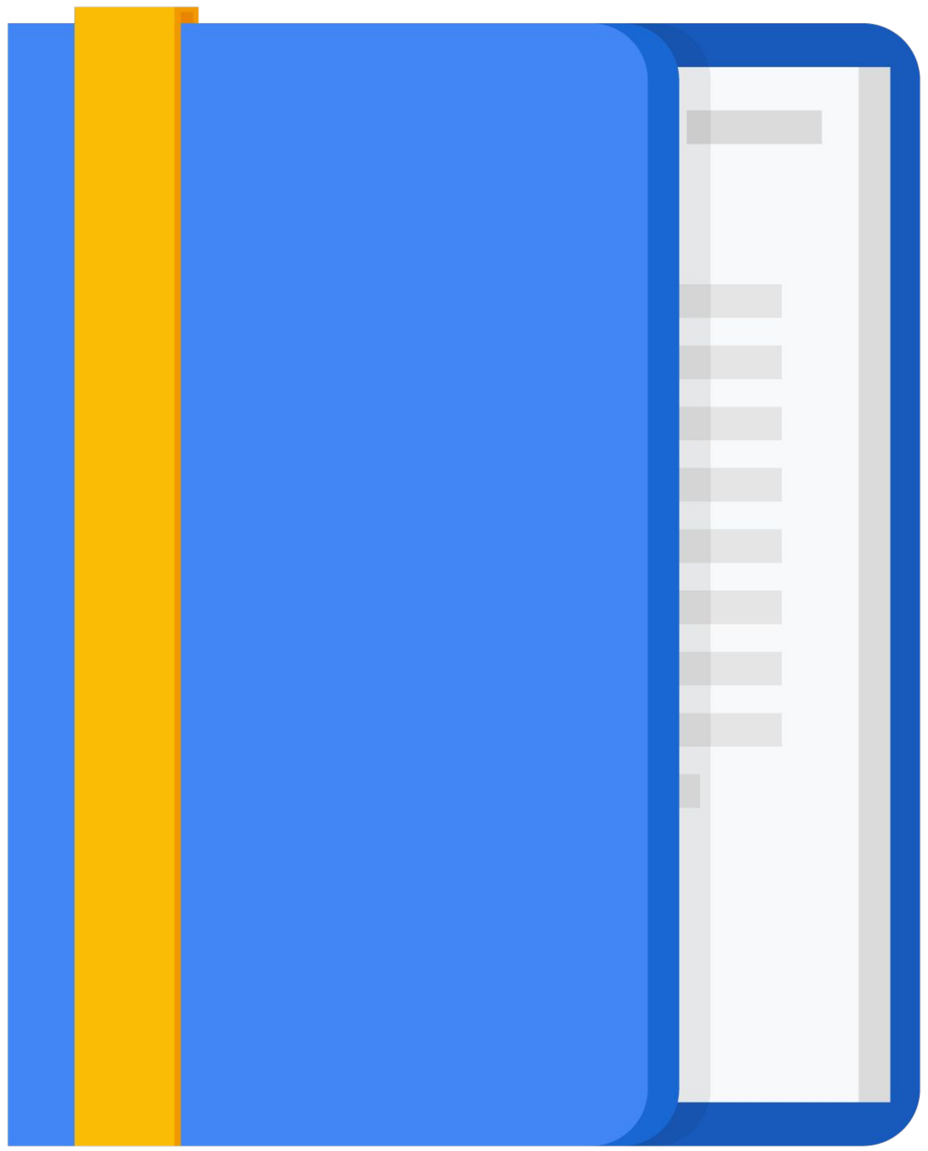

System overview

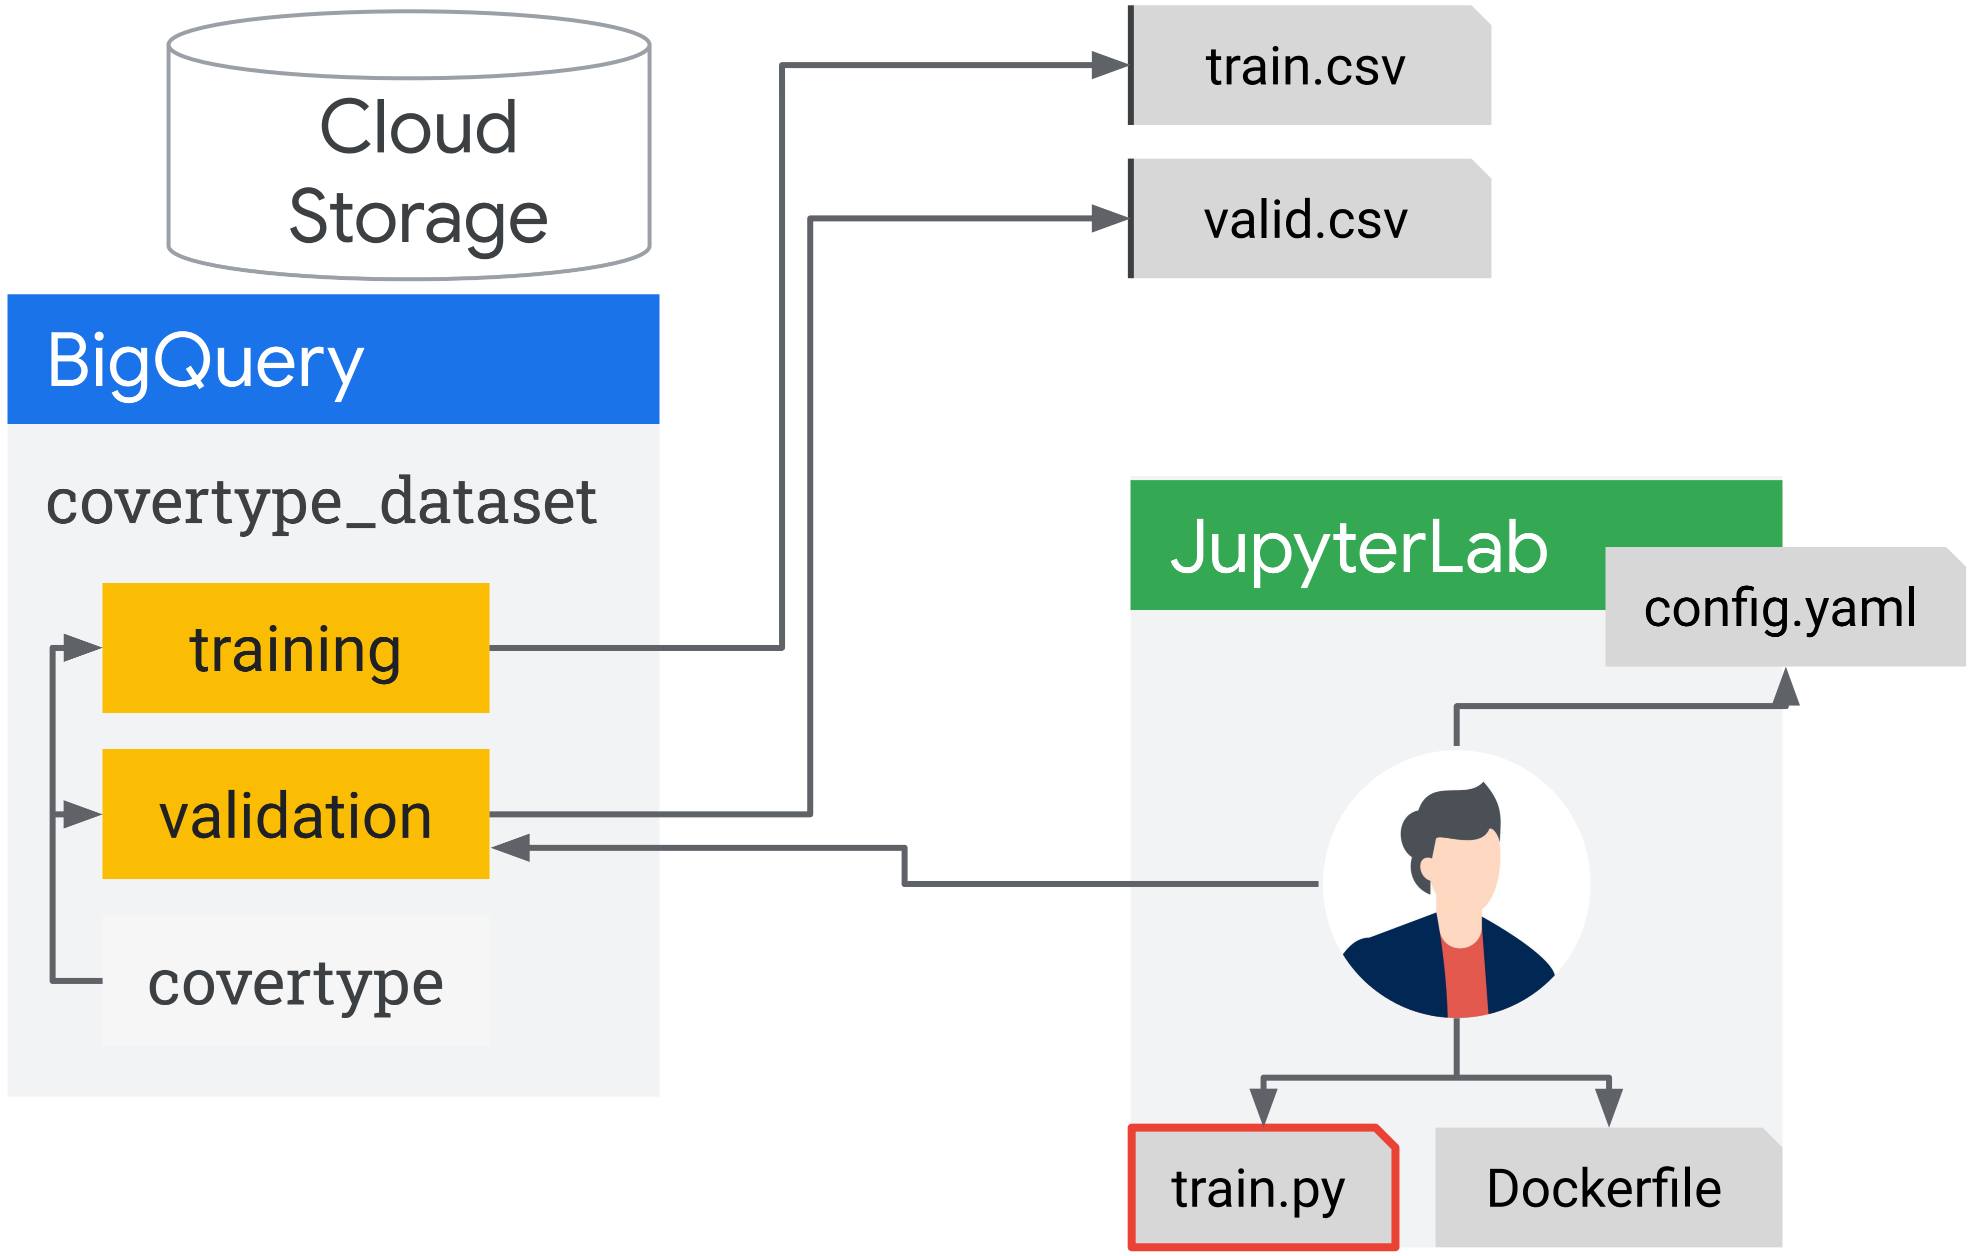

System overview

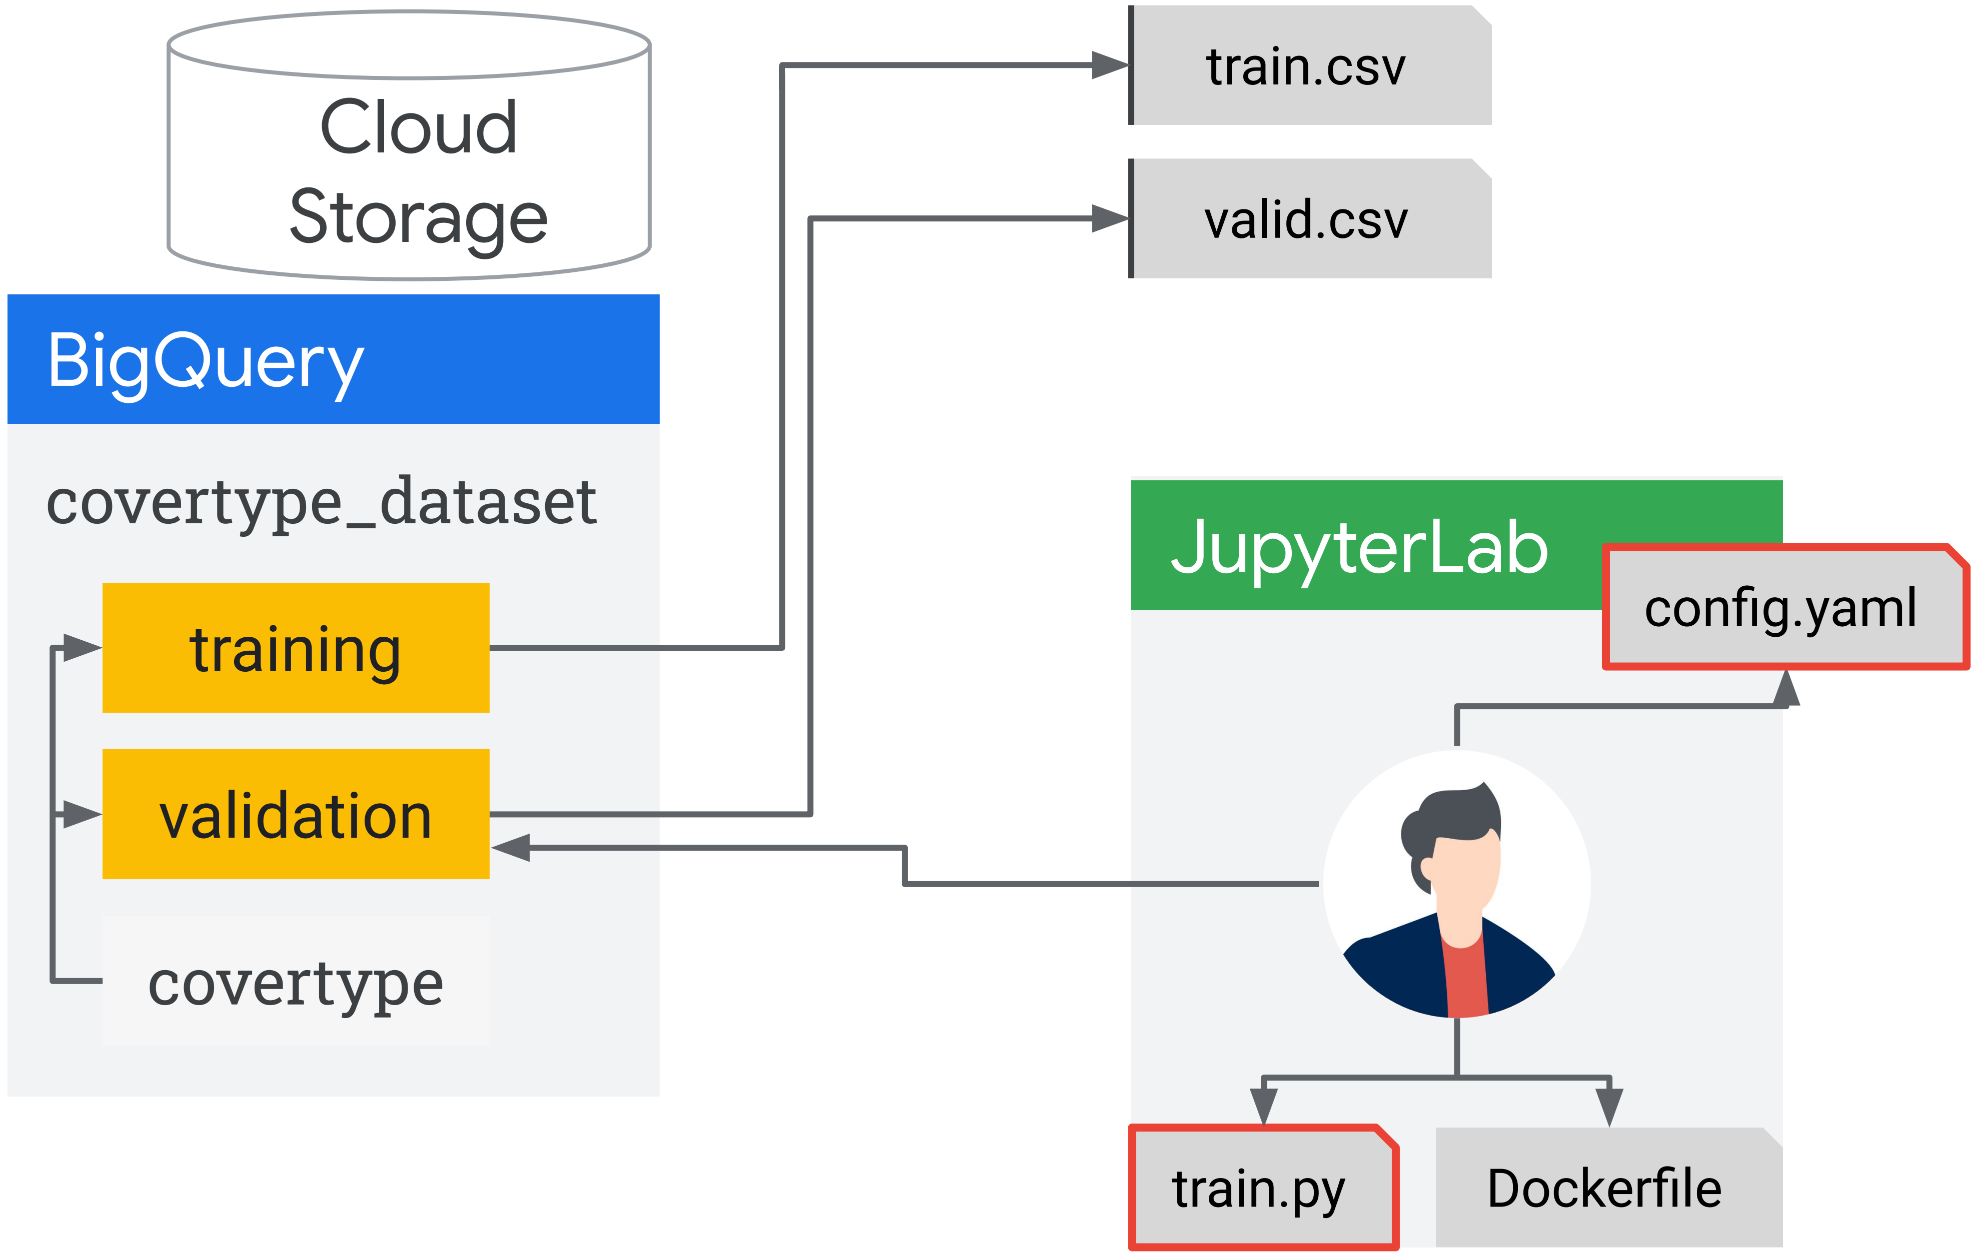

ML models are functions with parameters and hyperparameters

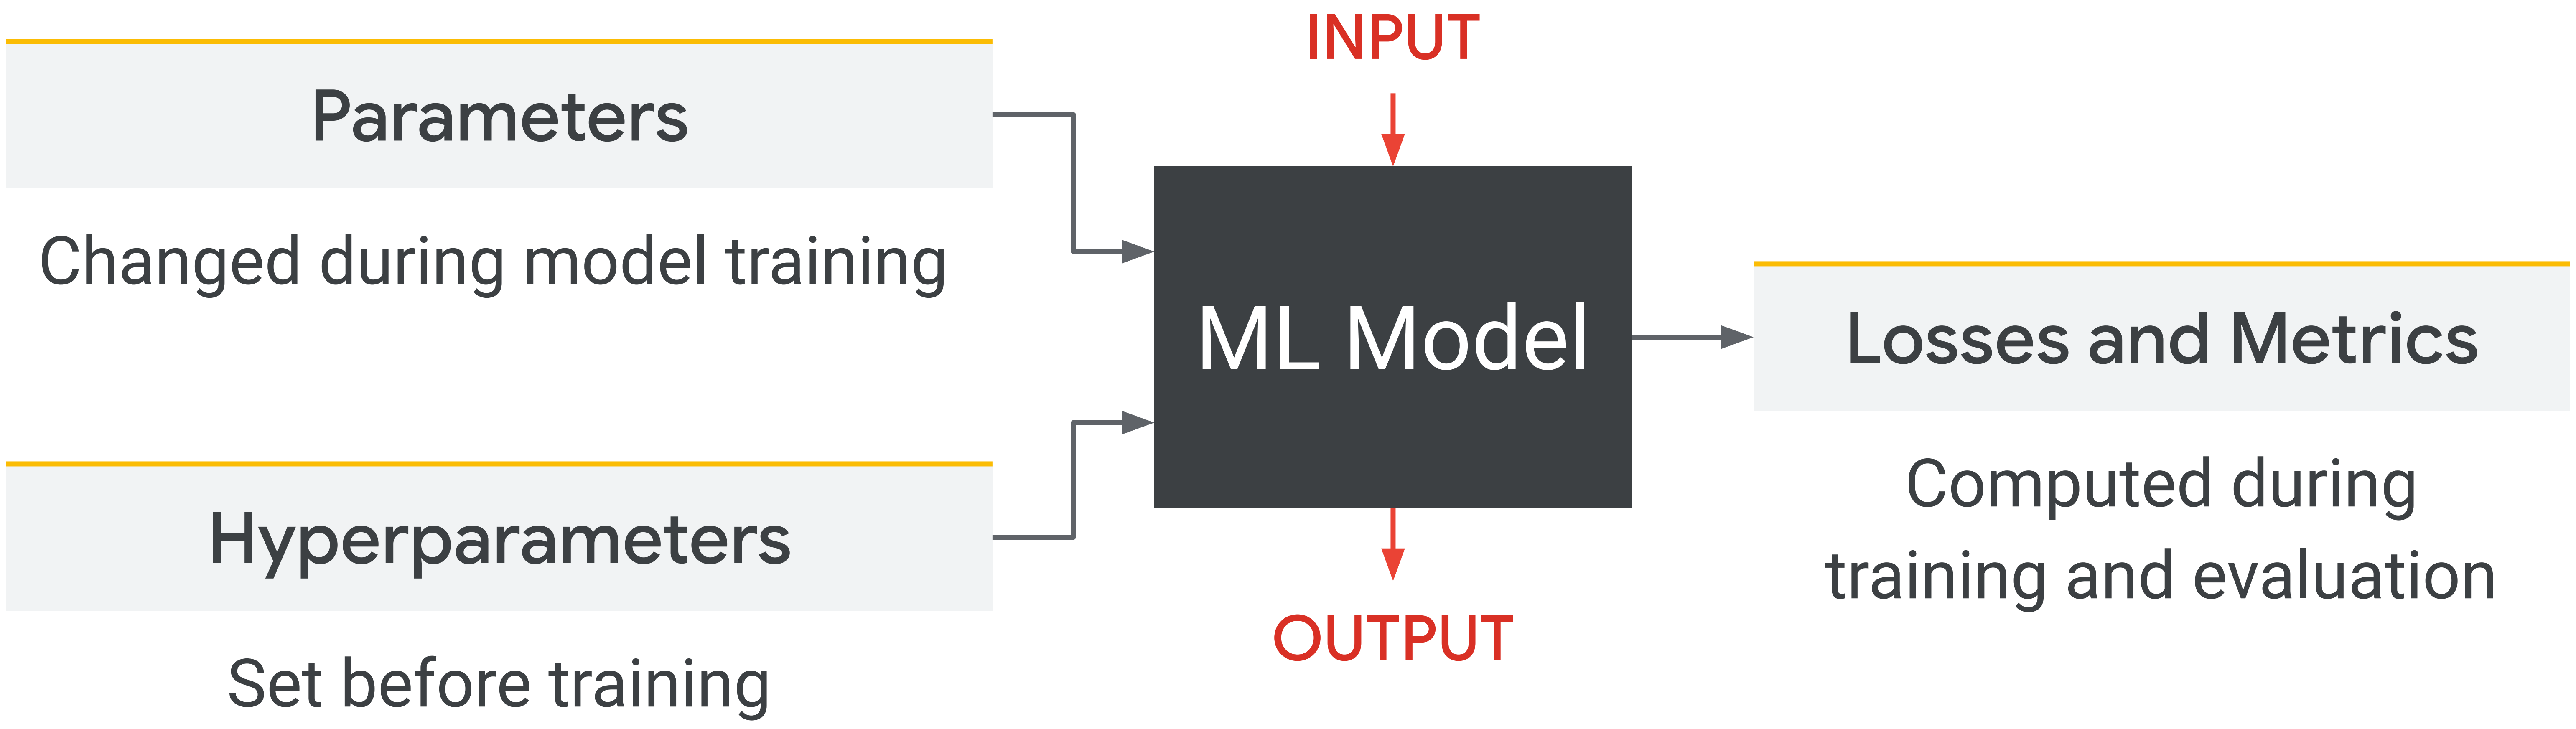

ML model: Sklearn pipeline

```
preprocessor = ColumnTransformer(
    transformers=[
        ('num', StandardScaler(), numeric_feature_indexes),
        ('cat', OneHotEncoder(), categorical_feature_indexes)
    ])

pipeline = Pipeline([
    ('preprocessor', preprocessor),
    ('classifier', SGDClassifier(loss='log', tol=1e-3))
])

pipeline.set_params(classifier__alpha=0.001, classifier__max_iter=200)
pipeline.fit(X_train, y_train)

accuracy = pipeline.score(X_validation, y_validation)
```

ML model: Sklearn pipeline

train.py

```
preprocessor = ColumnTransformer(
    transformers=[
        ('num', StandardScaler(), numeric_feature_indexes),
        ('cat', OneHotEncoder(), categorical_feature_indexes)
    ])

pipeline = Pipeline([
    ('preprocessor', preprocessor),
    ('classifier', SGDClassifier(loss='log', tol=1e-3))
])

pipeline.set_params(classifier__alpha=0.001, classifier__max_iter=200)
pipeline.fit(X_train, y_train)

accuracy = pipeline.score(X_validation, y_validation)
```

ML model: Sklearn pipeline

train.py

```
preprocessor = ColumnTransformer(
    transformers=[
        ('num', StandardScaler(), numeric_feature_indexes),
        ('cat', OneHotEncoder(), categorical_feature_indexes)
    ])

pipeline = Pipeline([
    ('preprocessor', preprocessor),
    ('classifier', SGDClassifier(loss='log', tol=1e-3))
])

pipeline.set_params(classifier__alpha=0.001, classifier__max_iter=200)
pipeline.fit(X_train, y_train)

accuracy = pipeline.score(X_validation, y_validation)
```

ML model: Sklearn pipeline

train.py

```
preprocessor = ColumnTransformer(
    transformers=[
        ('num', StandardScaler(), numeric_feature_indexes),
        ('cat', OneHotEncoder(), categorical_feature_indexes)
    ])

pipeline = Pipeline([
    ('preprocessor', preprocessor),
    ('classifier', SGDClassifier(loss='log', tol=1e-3))
])

pipeline.set_params(classifier__alpha=0.001, classifier__max_iter=200)
pipeline.fit(X_train, y_train)

accuracy = pipeline.score(X_validation, y_validation)
```

ML model: Sklearn pipeline

train.py

```
preprocessor = ColumnTransformer(
    transformers=[
        ('num', StandardScaler(), numeric_feature_indexes),
        ('cat', OneHotEncoder(), categorical_feature_indexes)
    ])

pipeline = Pipeline([
    ('preprocessor', preprocessor),
    ('classifier', SGDClassifier(loss='log', tol=1e-3))
])

pipeline.set_params(classifier_alpha=0.001, classifier_max_iter=200)
pipeline.fit(X_train, y_train)

accuracy = pipeline.score(X_validation, y_validation)
```

ML model: Sklearn pipeline

train.py

```
preprocessor = ColumnTransformer(
    transformers=[
        ('num', StandardScaler(), numeric_feature_indexes),
        ('cat', OneHotEncoder(), categorical_feature_indexes)
    ])

pipeline = Pipeline([
    ('preprocessor', preprocessor),
    ('classifier', SGDClassifier(loss='log', tol=1e-3))
])

pipeline.set_params(classifier__alpha=0.001, classifier__max_iter=200)
pipeline.fit(X_train, y_train)

accuracy = pipeline.score(X_validation, y_validation)
```

ML model: Sklearn pipeline

train.py

```
preprocessor = ColumnTransformer(
    transformers=[
        ('num', StandardScaler(), numeric_feature_indexes),
        ('cat', OneHotEncoder(), categorical_feature_indexes)
    ])

pipeline = Pipeline([
    ('preprocessor', preprocessor),
    ('classifier', SGDClassifier(loss='log', tol=1e-3))
])

pipeline.set_params(classifier__alpha=0.001, classifier__max_iter=200)
pipeline.fit(X_train, y_train)

accuracy = pipeline.score(X_validation, y_validation)
```

ML model: Sklearn pipeline

train.py

```
preprocessor = ColumnTransformer(
    transformers=[
        ('num', StandardScaler(), numeric_feature_indexes),
        ('cat', OneHotEncoder(), categorical_feature_indexes)
    ])

pipeline = Pipeline([
    ('preprocessor', preprocessor),
    ('classifier', SGDClassifier(loss='log', tol=1e-3))
])

pipeline.set_params(classifier__alpha=0.001, classifier__max_iter=200)
pipeline.fit(X_train, y_train)

accuracy = pipeline.score(X_validation, y_validation)
```

ML model: Sklearn pipeline

```
preprocessor = ColumnTransformer(
    transformers=[
        ('num', StandardScaler(), numeric_feature_indexes),
        ('cat', OneHotEncoder(), categorical_feature_indexes)
    ])

pipeline = Pipeline([
    ('preprocessor', preprocessor),
    ('classifier', SGDClassifier(loss='log', tol=1e-3))
])

pipeline.set_params(classifier__alpha=0.001, classifier__max_iter=200)
pipeline.fit(X_train, y_train)

accuracy = pipeline.score(X_validation, y_validation)
```

How to use AI Platform for hyperparameter tuning

1. Make the hyperparameter a command-line argument.

2. Set up `cloudml-hypertune` to record training metrics.

3. Export the final trained model.

4. Supply hyperparameters to the training job.

Hyperparameter tuning

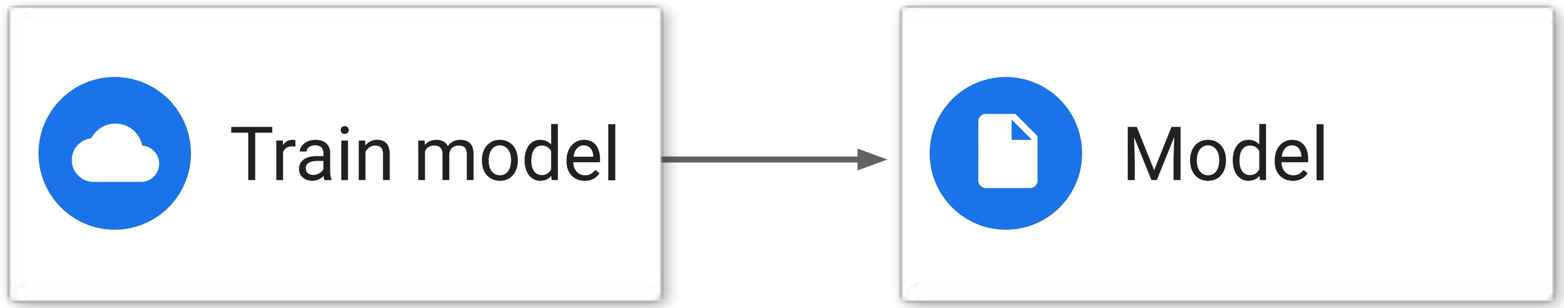

How to use AI Platform for hyperparameter tuning

1. Make the hyperparameter a command-line argument.

2. Set up `cloudml-hypertune` to record training metrics.

3. Export the final trained model.

4. Supply hyperparameters to the training job.

Hyperparameter tuning

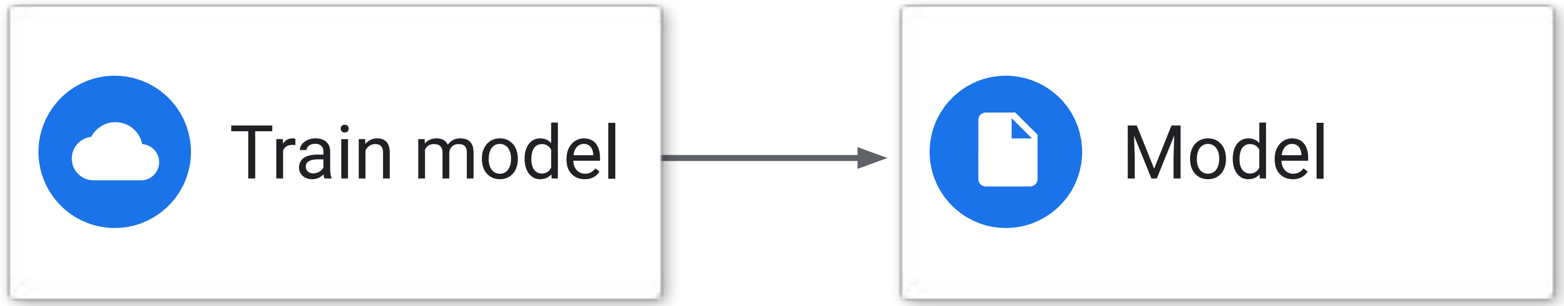

How to use AI Platform for hyperparameter tuning

1. Make the hyperparameter a command-line argument.
2. Set up `cloudml-hypertune` to record training metrics.
3. Export the final trained model.
4. Supply hyperparameters to the training job.

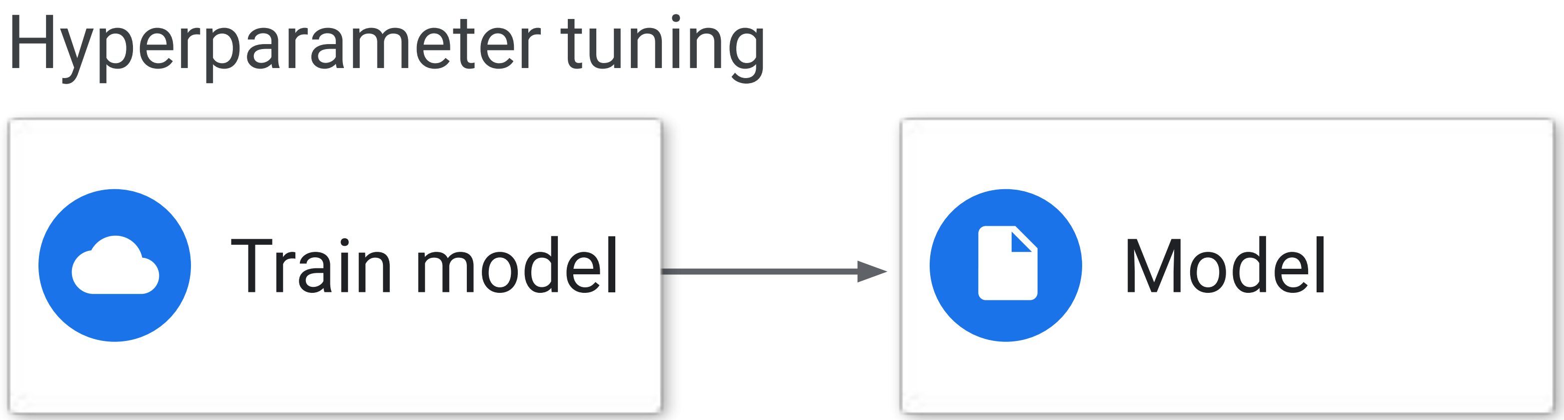

How to use AI Platform for hyperparameter tuning

1. Make the hyperparameter a command-line argument.

2. Set up `cloudml-hypertune` to record training metrics.

3. Export the final trained model.

4. Supply hyperparameters to the training job.

Hyperparameter tuning

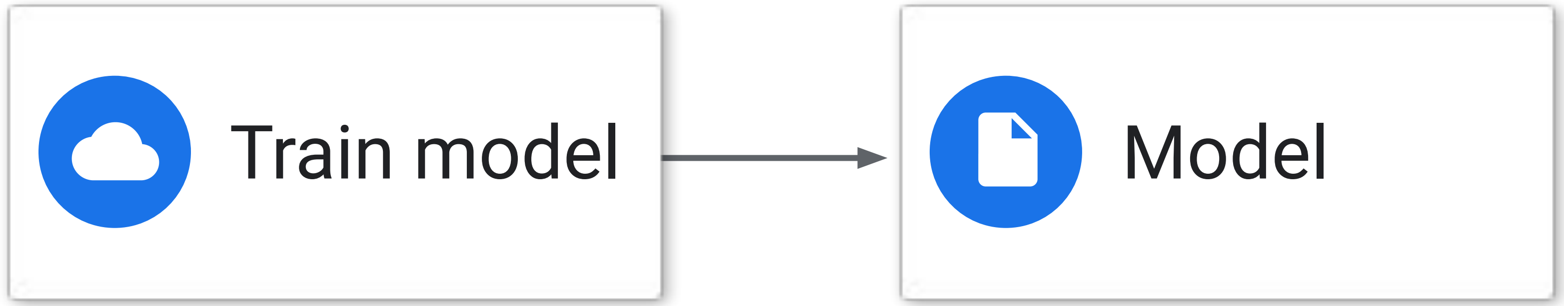

1. Make the hyperparameter a command-line argument

train.py

```
import fire

def train_evaluate(job_dir,
                   training_dataset_path,
                   validation_dataset_path,
                   alpha, max_iter, hptune):


    # [...]


if __name__ == "__main__":
    fire.Fire(train_evaluate)
```

```
Python train.py \
  --job_dir $JOBDIR \
  --training_dataset_path $TRAINING_PATH \
  --validation_dataset_path $VALID_PATH \
  --alpha \
  --maxiter \
  --hptune
```

2. Set up `cloudml-hypertune` to record training metrics

train.py

```
import hypertune

def train_evaluate(job_dir,
                   training_dataset_path,
                   validation_dataset_path,
                   alpha, max_iter, hptune):
    # [...]

    if hptune:
        accuracy = pipeline.score(X_validation, y_validation)

        hpt = hypertune.HyperTune()

        hpt.report_hyperparameter_tuning_metric(
          hyperparameter_metric_tag='accuracy',
          metric_value=accuracy
        )

if __name__ == "__main__":
    fire.Fire(train_evaluate)
```

Import cloudml-hypertune.

2. Set up `cloudml-hypertune` to record training metrics

train.py

```
import hypertune

def train_evaluate(job_dir,
                   training_dataset_path,
                   validation_dataset_path,
                   alpha, max_iter, hptune):
    # [...]

    if hptune:
        accuracy = pipeline.score(X_validation, y_validation)

        hpt = hypertune.HyperTune()

        hpt.report_hyperparameter_tuning_metric(
          hyperparameter_metric_tag='accuracy',
          metric_value=accuracy
        )

if __name__ == "__main__":
    fire.Fire(train_evaluate)
```

Capture the metrics.

3. Export the final trained model

train.py

```
import pickle

def train_evaluate(job_dir,
                   training_dataset_path,
                   validation_dataset_path,
                   alpha, max_iter, hptune):
    # [...]

    if not hptune:
        model_filename = 'model.pkl'
        with open(model_filename, 'wb') as model_file:
            pickle.dump(pipeline, model_file)
        gcs_model_path = "{}/{}".format(job_dir, model_filename)
        subprocess.check_call(['gsutil', 'cp', model_filename, gcs_model_path],
        stderr=sys.stdout)

if __name__ == "__main__":
    fire.Fire(train_evaluate)
```

3. Export the final retrain model when not tuning

train.py

```
import pickle

def train_evaluate(job_dir,
                   training_dataset_path,
                   validation_dataset_path,
                   alpha, max_iter, hptune):
    # [...]

    if not hptune:
        model_filename = 'model.pkl'
        with open(model_filename, 'wb') as model_file:
            pickle.dump(pipeline, model_file)
        gcs_model_path = "{}/{}".format(job_dir, model_filename)
        subprocess.check_call(['gsutil', 'cp', model_filename, gcs_model_path],
        stderr=sys.stdout)

if __name__ == "__main__":
    fire.Fire(train_evaluate)
```

3. Export the final retrain model when not tuning

train.py

```
import pickle

def train_evaluate(job_dir,
                   training_dataset_path,
                   validation_dataset_path,
                   alpha, max_iter, hptune):
    # [...]

    if not hptune:
        model_filename = 'model.pkl'
        with open(model_filename, 'wb') as model_file:
            pickle.dump(pipeline, model_file)
        gcs_model_path = "{}/{}".format(job_dir, model_filename)
        subprocess.check_call(['gsutil', 'cp', model_filename, gcs_model_path],
        stderr=sys.stdout)

if __name__ == "__main__":
    fire.Fire(train_evaluate)
```

3. Export the final retrain model when not tuning

train.py

```
import pickle

def train_evaluate(job_dir,
                   training_dataset_path,
                   validation_dataset_path,
                   alpha, max_iter, hptune):
    # [...]

    if not hptune:
        model_filename = 'model.pkl'
        with open(model_filename, 'wb') as model_file:
            pickle.dump(pipeline, model_file)
        gcs_model_path = "{}/{}".format(job_dir, model_filename)
        subprocess.check_call(['gsutil', 'cp', model_filename, gcs_model_path],
        stderr=sys.stdout)

if __name__ == "__main__":
    fire.Fire(train_evaluate)
```

4. Supply hyperparameters to the training job

config.yaml

```
trainingInput:
  hyperparameters:
    goal: MAXIMIZE
    maxTrials: 4
    maxParallelTrials: 4
    hyperparameterMetricTag: accuracy
    enableTrialEarlyStopping: TRUE
    params:
    - parameterName: max_iter
      type: DISCRETE
      discreteValues: [
          200,
          500
         ]
    - parameterName: alpha
      type: DOUBLE
      minValue:  0.00001
      maxValue:  0.001
      scaleType: UNIT_LINEAR_SCALE
```

4. Supply hyperparameters to the training job

config.yaml

```
trainingInput:
  hyperparameters:
    goal: MAXIMIZE
    maxTrials: 4
    maxParallelTrials: 4
    hyperparameterMetricTag: accuracy
    enableTrialEarlyStopping: TRUE
    params:
    - parameterName: max_iter
      type: DISCRETE
      discreteValues: [
          200,
          500
          ]
    - parameterName: alpha
      type: DOUBLE
      minValue:  0.00001
      maxValue:  0.001
      scaleType: UNIT_LINEAR_SCALE
```

4. Supply hyperparameters to the training job

config.yaml


```yaml
trainingInput:
  hyperparameters:
    goal: MAXIMIZE
    maxTrials: 4
    maxParallelTrials: 4
    hyperparameterMetricTag: accuracy
    enableTrialEarlyStopping: TRUE
    params:
    - parameterName: max_iter
      type: DISCRETE
      discreteValues: [
          200,
          500
         ]
    - parameterName: alpha
      type: DOUBLE
      minValue:  0.00001
      maxValue:  0.001
      scaleType: UNIT_LINEAR_SCALE
```

4. Supply hyperparameters to the training job

config.yaml

```
trainingInput:
  hyperparameters:
    goal: MAXIMIZE
    maxTrials: 4
    maxParallelTrials: 4
    hyperparameterMetricTag: accuracy
    enableTrialEarlyStopping: TRUE
    params:
    - parameterName: max_iter
      type: DISCRETE
      discreteValues: [
          200,
          500
         ]
    - parameterName: alpha
      type: DOUBLE
      minValue:  0.00001
      maxValue:  0.001
      scaleType: UNIT_LINEAR_SCALE
```

4. Supply hyperparameters to the training job

config.yaml

```
trainingInput:
  hyperparameters:
    goal: MAXIMIZE
    maxTrials: 4
    maxParallelTrials: 4
    hyperparameterMetricTag: accuracy
    enableTrialEarlyStopping: TRUE
    params:
    - parameterName: max_iter
      type: DISCRETE
      discreteValues: [
          200,
          500
          ]
    - parameterName: alpha
      type: DOUBLE
      minValue:  0.00001
      maxValue:  0.001
      scaleType: UNIT_LINEAR_SCALE
```

4. Supply hyperparameters to the training job

config.yaml

```
trainingInput:
  hyperparameters:
    goal: MAXIMIZE
    maxTrials: 4
    maxParallelTrials: 4
    hyperparameterMetricTag: accuracy
    enableTrialEarlyStopping: TRUE
    params:
    - parameterName: max_iter
      type: DISCRETE
      discreteValues: [
          200,
          500
          ]
    - parameterName: alpha
      type: DOUBLE
      minValue:  0.00001
      maxValue:  0.001
      scaleType: UNIT_LINEAR_SCALE
```

```
gcloud ai-platform jobs submit training $JOB_NAME \
  -- [...]
  --config config.yaml \
  -- \
  --training_dataset_path=$TRAINING_FILE_PATH \
  --validation_dataset_path=$VALIDATION_FILE_PATH \
  --hptune
```

Agenda

System and Concepts Overview

Create a Reproducible Dataset

Implement a Tunable Model

Build and Push a Training Container

Train and Tune a Model

Serve and Query a Model

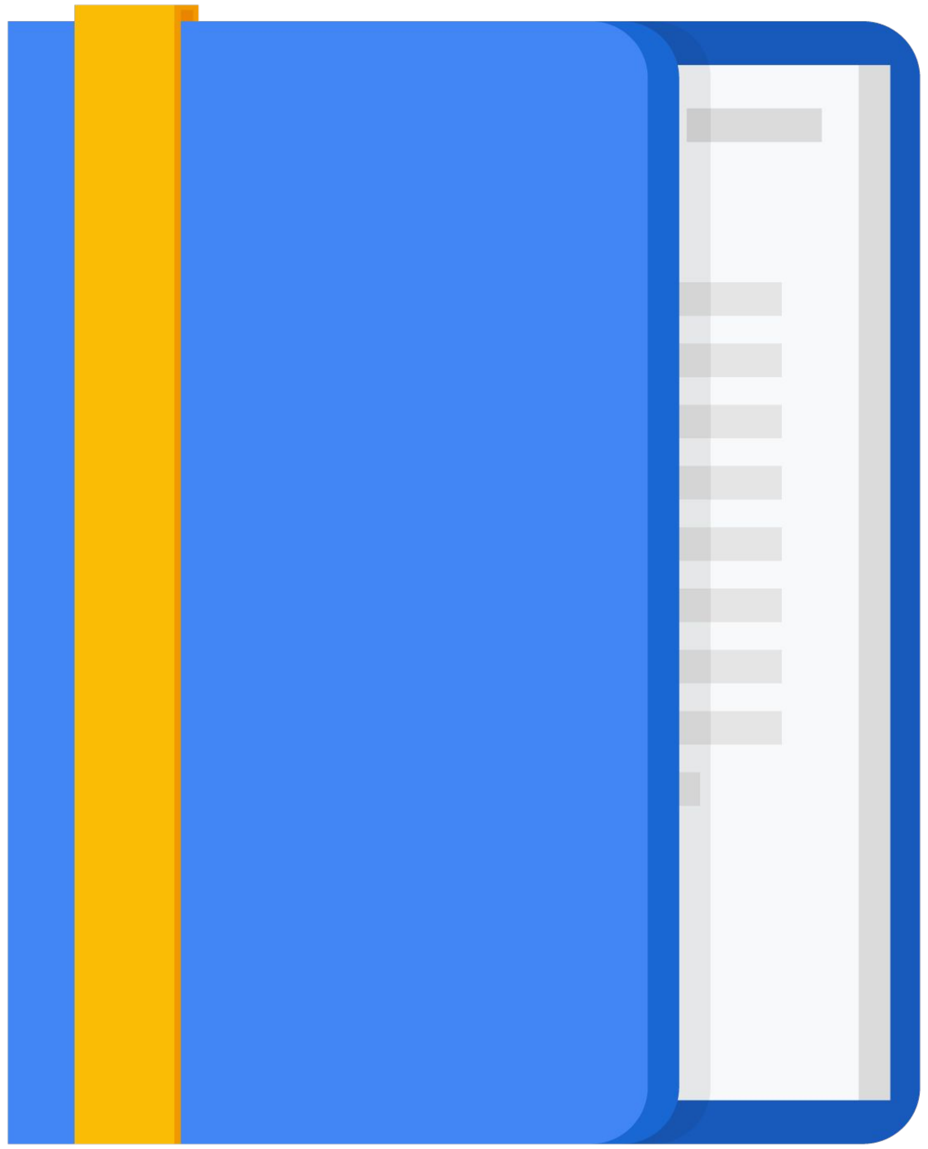

System overview

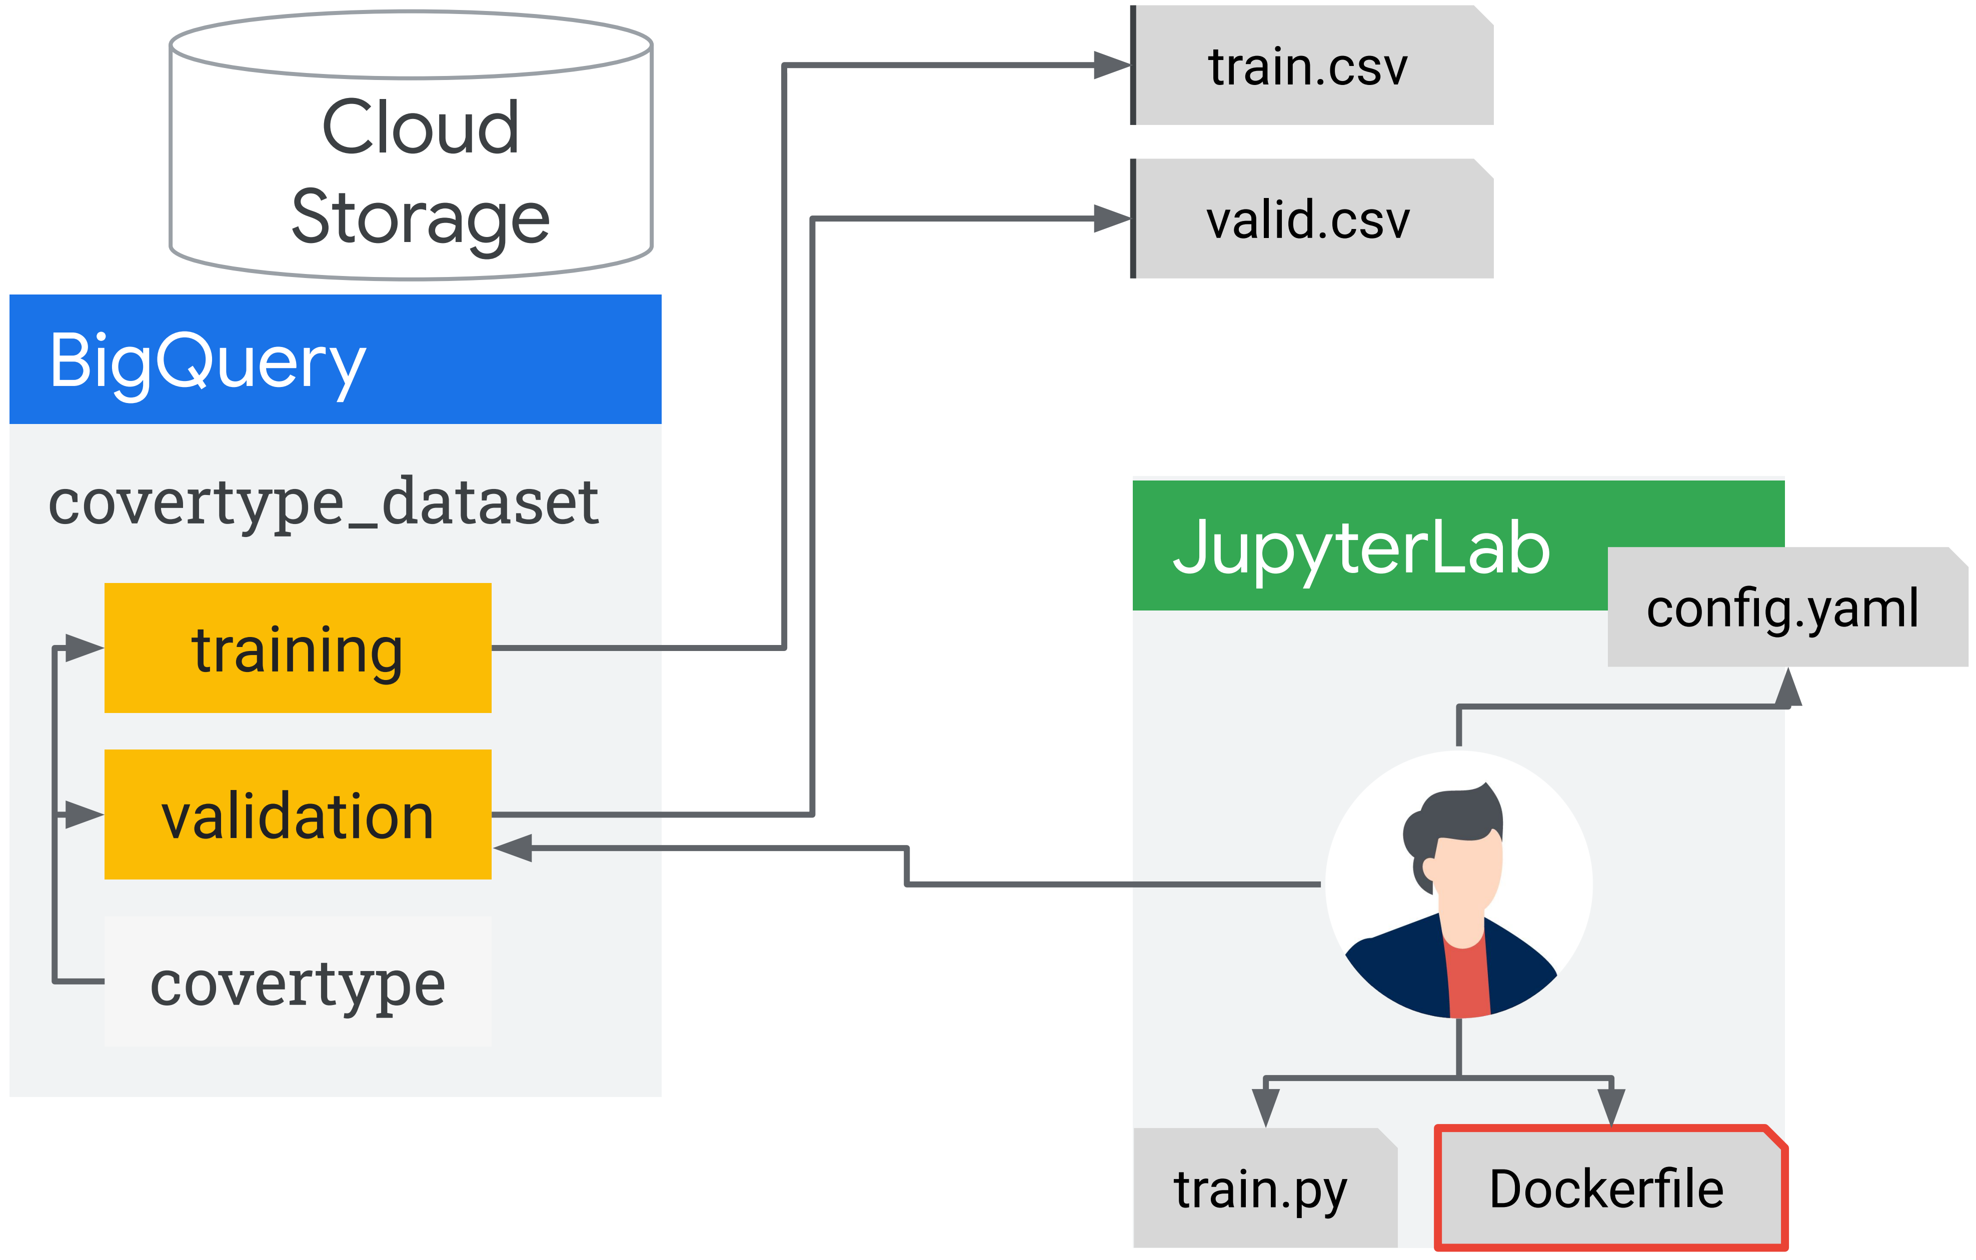

Create the training Docker container

Dockerfile

```
FROM gcr.io/deeplearning-platform-release/base-cpu

RUN pip install -U fire cloudml-hypertune scikit-learn==0.20.4 pandas==0.24.2

WORKDIR /app

COPY train.py .

ENTRYPOINT ["python", "train.py"]
```

Create the training Docker container

Dockerfile

```
FROM gcr.io/deeplearning-platform-release/base-cpu

RUN pip install -U fire cloudml-hypertune scikit-learn==0.20.4 pandas==0.24.2

WORKDIR /app

COPY train.py .

ENTRYPOINT ["python", "train.py"]
```

Create the training Docker container

Dockerfile

```
FROM gcr.io/deeplearning-platform-release/base-cpu

RUN pip install -U fire cloudml-hypertune scikit-learn==0.20.4 pandas==0.24.2

WORKDIR /app

COPY train.py .

ENTRYPOINT ["python", "train.py"]
```

Create the training Docker container

Dockerfile

```
FROM gcr.io/deeplearning-platform-release/base-cpu

RUN pip install -U fire cloudml-hypertune scikit-learn==0.20.4 pandas==0.24.2

WORKDIR /app

COPY train.py .

ENTRYPOINT ["python", "train.py"]
```

Create the training Docker container

Dockerfile

```
FROM gcr.io/deeplearning-platform-release/base-cpu

RUN pip install -U fire cloudml-hypertune scikit-learn==0.20.4 pandas==0.24.2

WORKDIR /app

COPY train.py .

ENTRYPOINT ["python", "train.py"]
```

Create the training Docker container

Dockerfile

```
FROM gcr.io/deeplearning-platform-release/base-cpu

RUN pip install -U fire cloudml-hypertune scikit-learn==0.20.4 pandas==0.24.2

WORKDIR /app

COPY train.py .

ENTRYPOINT ["python", "train.py"]
```

```
gcloud builds submit --tag gcr.io/$PROJECT/$IMAGE:$TAG $TRAINING_APP_FOLDER
```

Create the training Docker container

Dockerfile

```
FROM gcr.io/deeplearning-platform-release/base-cpu

RUN pip install -U fire cloudml-hypertune scikit-learn==0.20.4 pandas==0.24.2

WORKDIR /app

COPY train.py .

ENTRYPOINT ["python", "train.py"]
```

```
gcloud builds submit --tag gcr.io/$PROJECT/$IMAGE:$TAG $TRAINING_APP_FOLDER
```

Create the training Docker container

Dockerfile

```
FROM gcr.io/deeplearning-platform-release/base-cpu

RUN pip install -U fire cloudml-hypertune scikit-learn==0.20.4 pandas==0.24.2

WORKDIR /app

COPY train.py .

ENTRYPOINT ["python", "train.py"]
```

```
gcloud builds submit --tag gcr.io/$PROJECT/$IMAGE:$TAG $TRAINING_APP_FOLDER
```

Agenda

System and Concepts Overview

Create a Reproducible Dataset

Implement a Tunable Model

Build and Push a Training Container

Train and Tune a Model

Serve and Query a Model

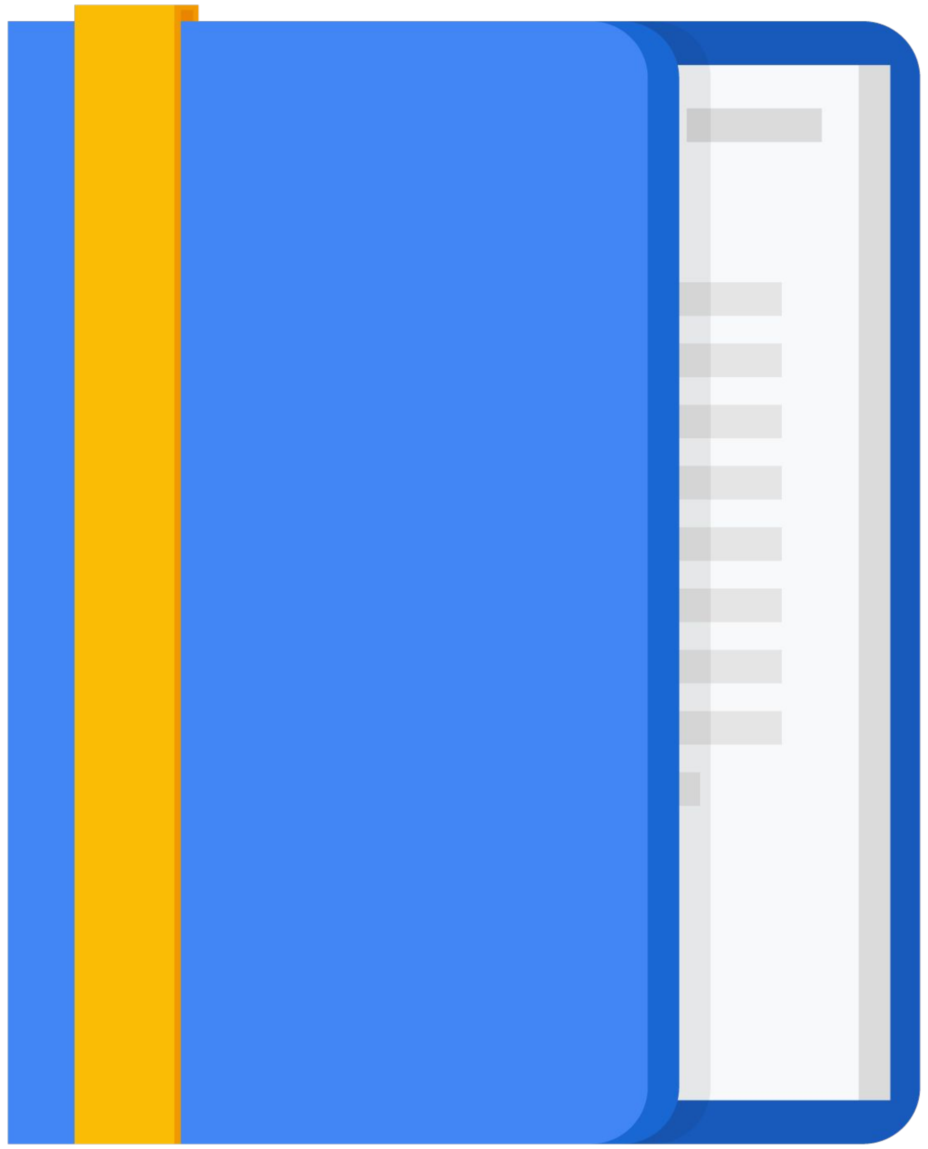

System overview

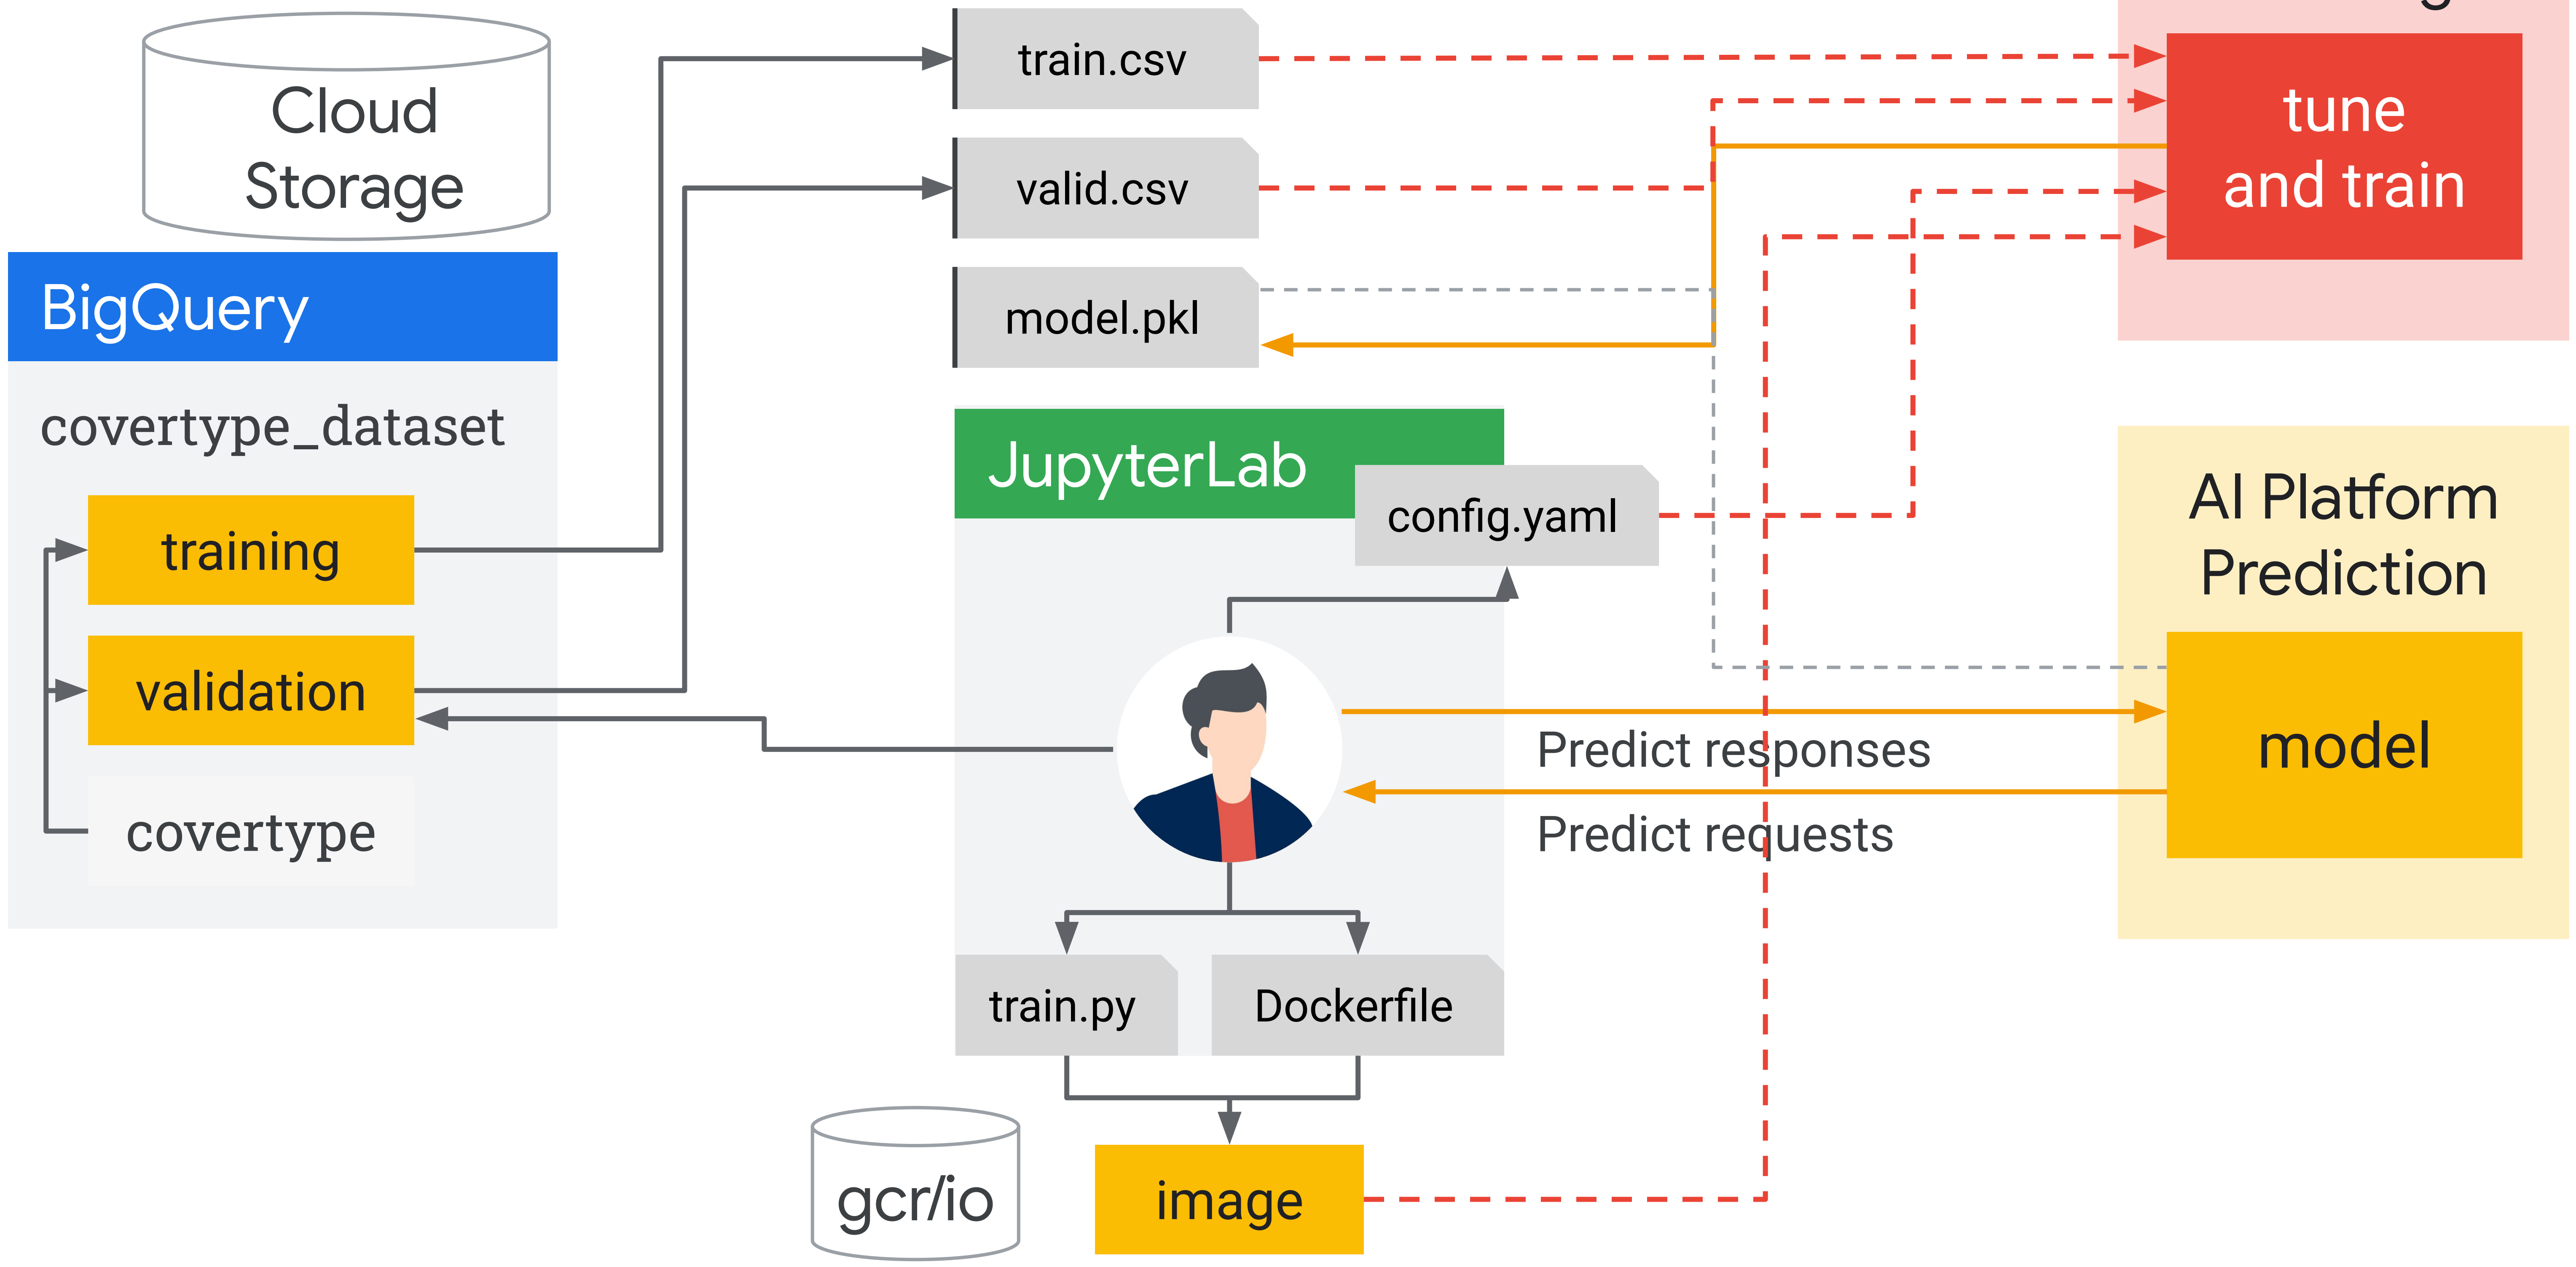

Start the hyper tuning job on AI Platform

```
gcloud ai-platform jobs submit training $JOB_NAME \
   --region=$REGION \
   --job-dir=$JOB_DIR \
   --master-image-uri=$IMAGE_URI \
   --scale-tier=$SCALE_TIER \
   --config $TRAINING_APP_FOLDER/hptuning_config.yaml \
   -- \
   --training_dataset_path=$TRAINING_FILE_PATH \
   --validation_dataset_path=$VALIDATION_FILE_PATH \
   --hptune
```

Start the hyper tuning job on AI Platform

```
gcloud ai-platform jobs submit training $JOB_NAME \
    --region=$REGION \
    --job-dir=$JOB_DIR \
    --master-image-uri=$IMAGE_URI \
    --scale-tier=$SCALE_TIER \
    --config $TRAINING_APP_FOLDER/hptuning_config.yaml \
    -- \
    --training_dataset_path=$TRAINING_FILE_PATH \
    --validation_dataset_path=$VALIDATION_FILE_PATH \
    --hptune
```

Start the hyper tuning job on AI Platform

```
gcloud ai-platform jobs submit training $JOB_NAME \
    --region=$REGION \
    --job-dir=$JOB_DIR \
    --master-image-uri=$IMAGE_URI \
    --scale-tier=$SCALE_TIER \
    --config $TRAINING_APP_FOLDER/hptuning_config.yaml \
    -- \
    --training_dataset_path=$TRAINING_FILE_PATH \
    --validation_dataset_path=$VALIDATION_FILE_PATH \
    --hptune
```

Start the hyper tuning job on AI Platform

```
gcloud ai-platform jobs submit training $JOB_NAME \
   --region=$REGION \
   --job-dir=$JOB_DIR \
   --master-image-uri=$IMAGE_URI \
   --scale-tier=$SCALE_TIER \
   --config $TRAINING_APP_FOLDER/hptuning_config.yaml \
   -- \
   --training_dataset_path=$TRAINING_FILE_PATH \
   --validation_dataset_path=$VALIDATION_FILE_PATH \
   --hptune
```

Start the hyper tuning job on AI Platform

```
gcloud ai-platform jobs submit training $JOB_NAME \
   --region=$REGION \
   --job-dir=$JOB_DIR \
   --master-image-uri=$IMAGE_URI \
   --scale-tier=$SCALE_TIER \
   --config $TRAINING_APP_FOLDER/hptuning_config.yaml \
   -- \
   --training_dataset_path=$TRAINING_FILE_PATH \
   --validation_dataset_path=$VALIDATION_FILE_PATH \
   --hptune
```

Start the hyper tuning job on AI Platform

```
gcloud ai-platform jobs submit training $JOB_NAME \
    --region=$REGION \
    --job-dir=$JOB_DIR \
    --master-image-uri=$IMAGE_URI \
    --scale-tier=$SCALE_TIER \
    --config $TRAINING_APP_FOLDER/hptuning_config.yaml \
    -- \
    --training_dataset_path=$TRAINING_FILE_PATH \
    --validation_dataset_path=$VALIDATION_FILE_PATH \
    --hptune
```

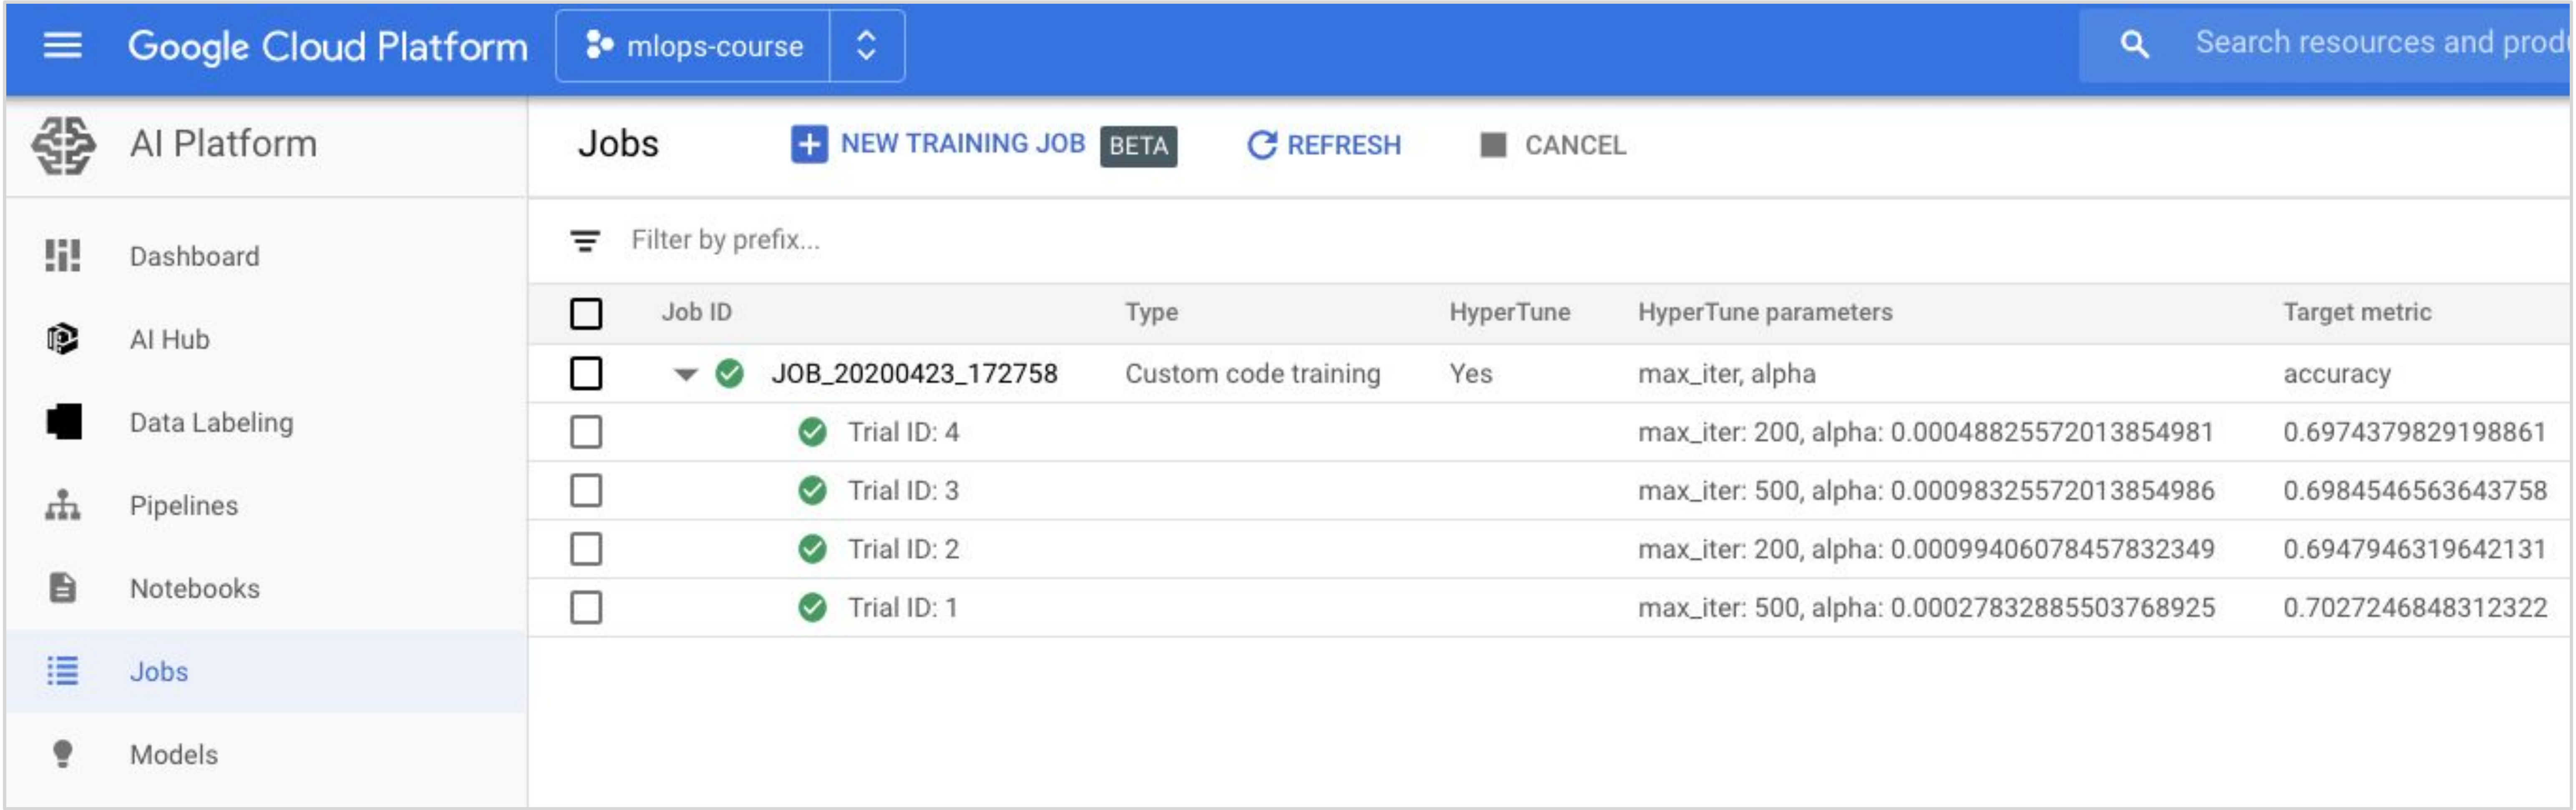
Google Cloud Platform
mlops-course
Search resources and prod
AI Platform
Jobs
NEW TRAINING JOB
BETA
REFRESH
CANCEL
Dashboard
AI Hub
Data Labeling
Pipelines
Notebooks
Jobs
Models
Filter by prefix...
Job ID
Type
HyperTune
HyperTune parameters
Target metric
JOB_20200423_172758
Custom code training
Yes
max_iter, alpha
accuracy
Trial ID: 4
max_iter: 200, alpha: 0.00048825572013854981
0.6974379829198861
Trial ID: 3
max_iter: 500, alpha: 0.00098325572013854986
0.6984546563643758
Trial ID: 2
max_iter: 200, alpha: 0.00099406078457832349
0.6947946319642131
Trial ID: 1
max_iter: 500, alpha: 0.00027832885503768925
0.7027246848312322

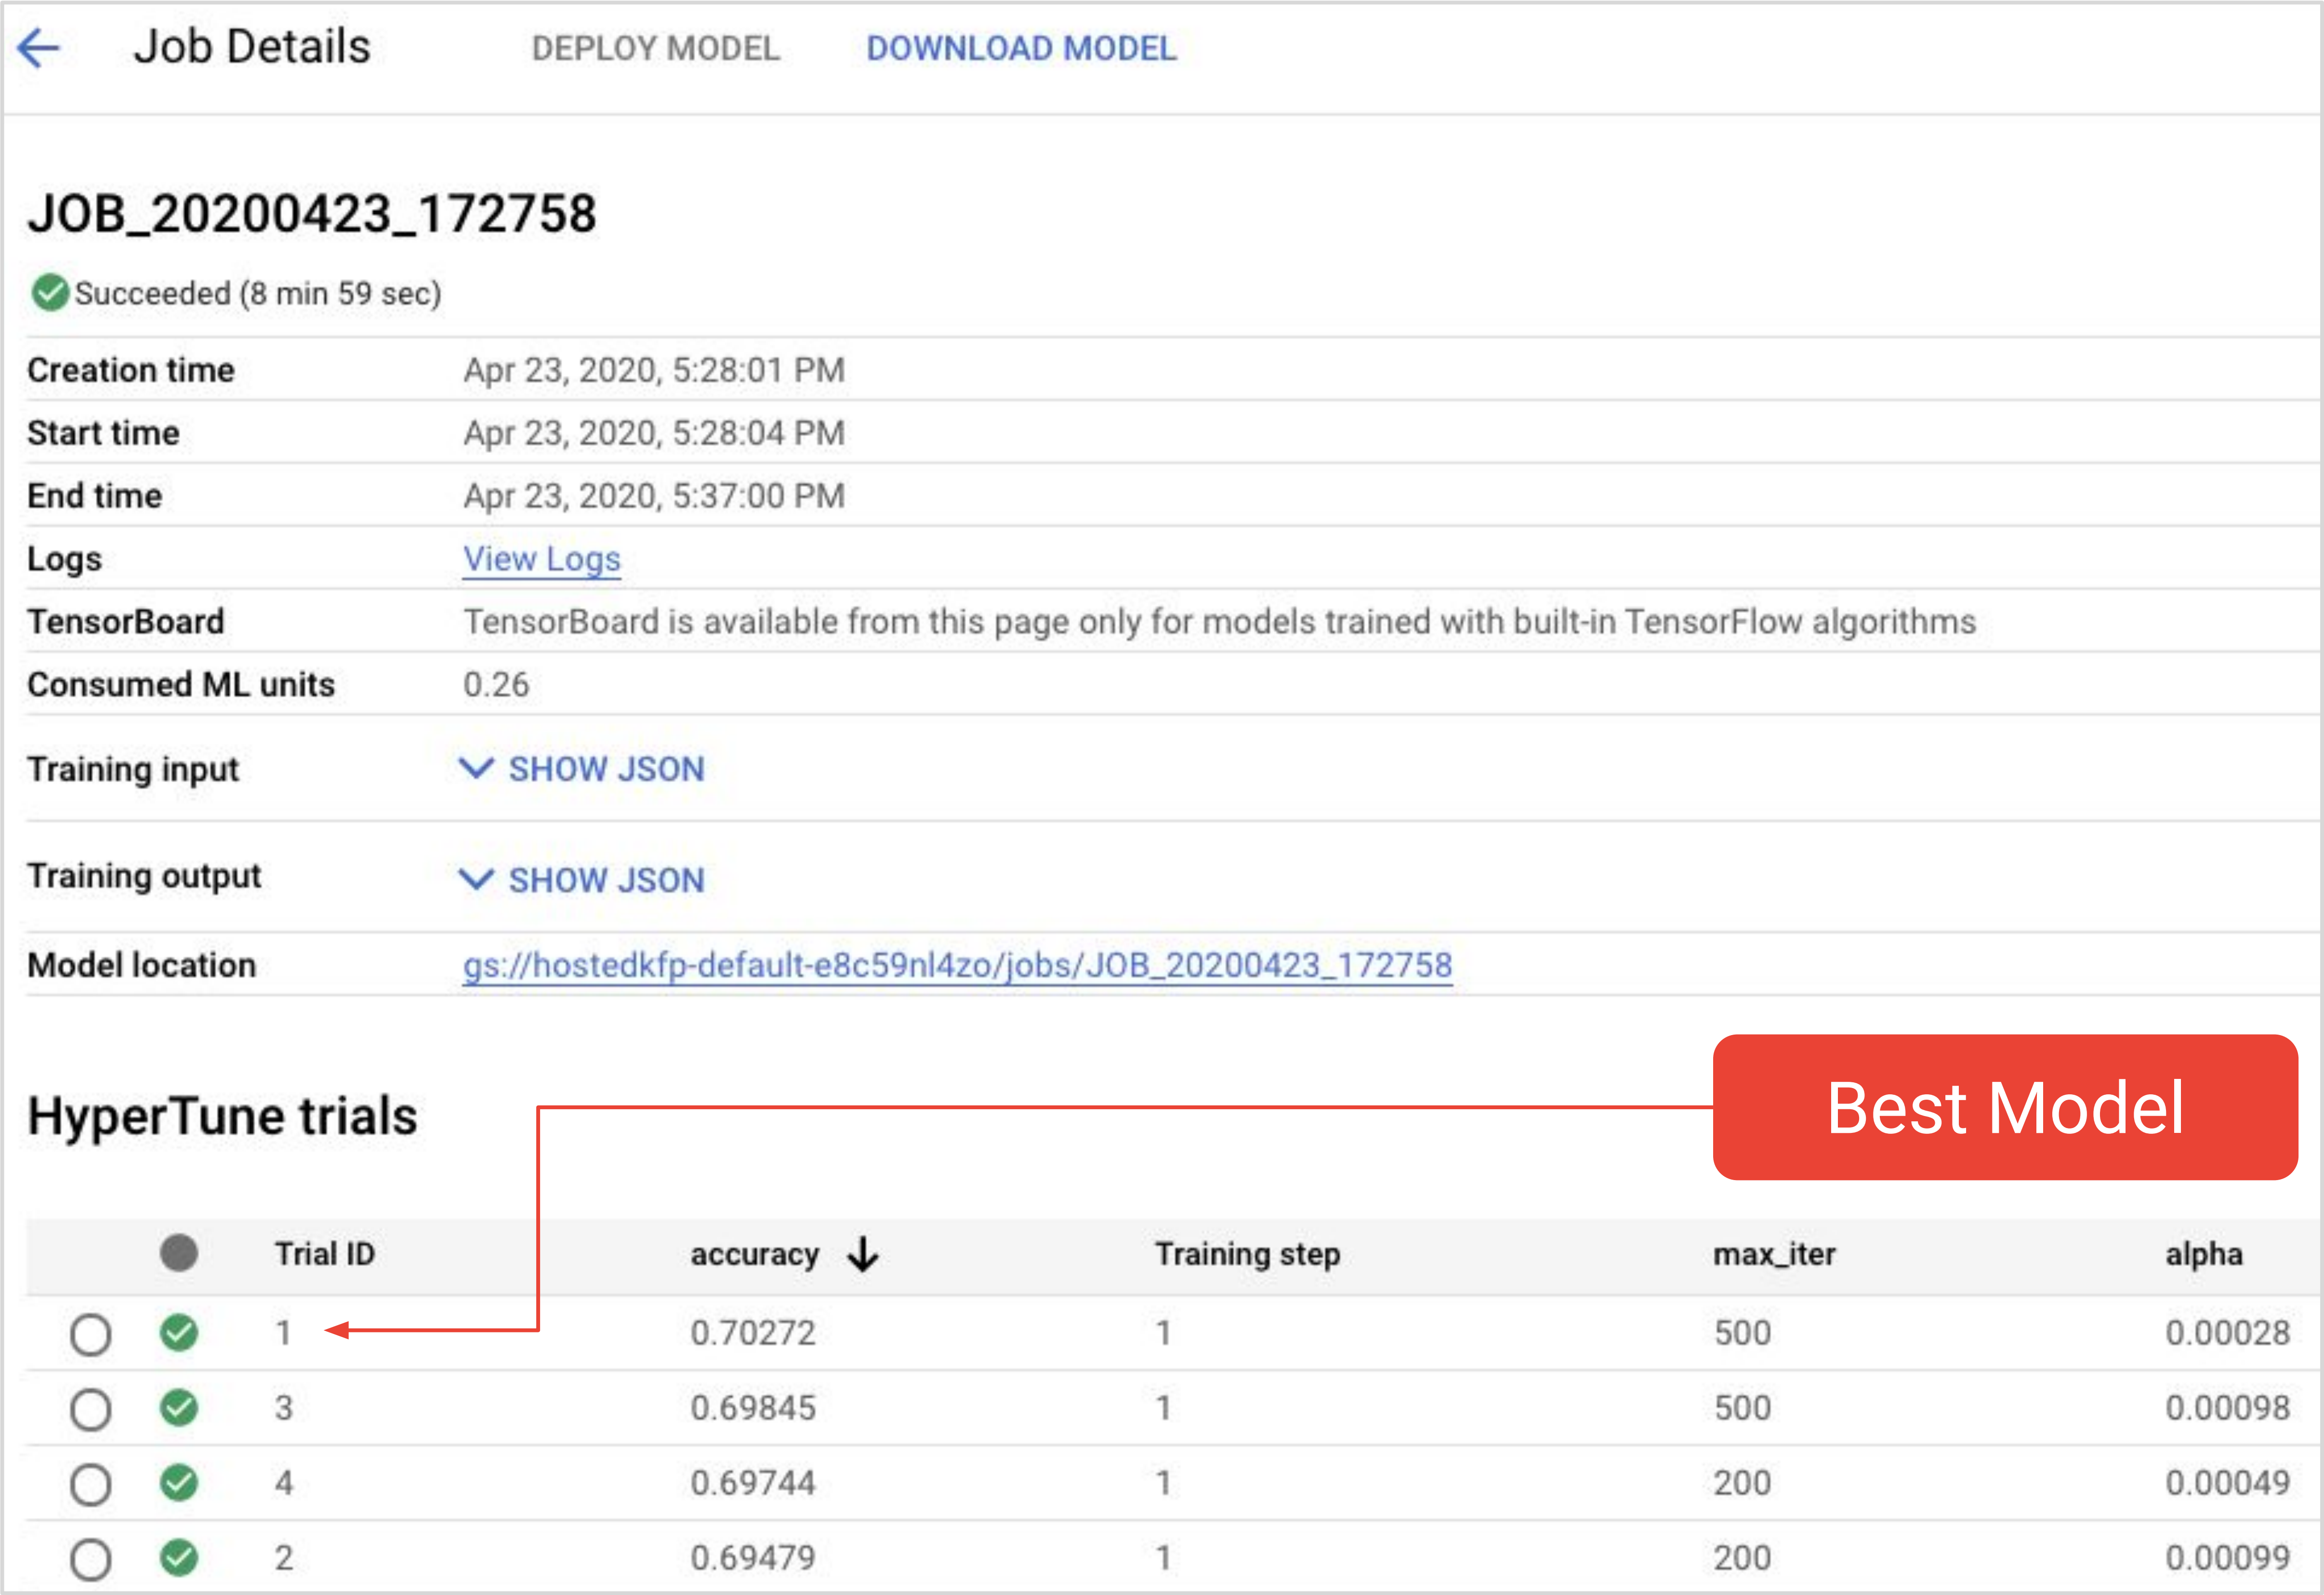

Job Details DEPLOY MODEL DOWNLOAD MODEL

JOB_20200423_172758

Succeeded (8 min 59 sec)

Creation time	Apr 23, 2020, 5:28:01 PM
Start time	Apr 23, 2020, 5:28:04 PM
End time	Apr 23, 2020, 5:37:00 PM
Logs	View Logs
TensorBoard	TensorBoard is available from this page only for models trained with built-in TensorFlow algorithms
Consumed ML units	0.26
Training input	SHOW JSON
Training output	SHOW JSON
Model location	gs://hostedkfp-default-e8c59nl4zo/jobs/JOB_20200423_172758

HyperTune trials

Best Model

Trial ID	accuracy ↓	Training step	max_iter	alpha
1	0.70272	1	500	0.00028
3	0.69845	1	500	0.00098
4	0.69744	1	200	0.00049
2	0.69479	1	200	0.00099

Query AI Platform Training for the best hyperparameters

```
from googleapiclient import discovery

ml = discovery.build('ml', 'v1')

job_id = 'projects/{}/jobs/{}'.format(PROJECT_ID, JOB_NAME)
request = ml.projects().jobs().get(name=job_id)

response = request.execute()

alpha = response['trainingOutput']['trials'][0]['hyperparameters']['alpha']
max_iter = response['trainingOutput']['trials'][0]['hyperparameters']['max_iter']
```

Retrain with the best hyperparameters and export

```
gcloud ai-platform jobs submit training $JOB_NAME \
      --region=$REGION \
      --job-dir=$JOB_DIR \
      --master-image-uri=$IMAGE_URI \
      --scale-tier=$SCALE_TIER \
      -- \
      --training_dataset_path=$TRAINING_FILE_PATH \
      --validation_dataset_path=$VALIDATION_FILE_PATH \
      --alpha=$alpha \
      --max_iter=$max_iter \
      --nohptune
```

Retrain with the best hyperparameters and export

```
gcloud ai-platform jobs submit training $JOB_NAME \
        --region=$REGION \
        --job-dir=$JOB_DIR \
        --master-image-uri=$IMAGE_URI \
        --scale-tier=$SCALE_TIER \
        -- \
        --training_dataset_path=$TRAINING_FILE_PATH \
        --validation_dataset_path=$VALIDATION_FILE_PATH \
        --alpha=$alpha \
        --max_iter=$max_iter \
        --nohptune
```

No more reference to config.yaml

Retrain with the best hyperparameters and export

```
gcloud ai-platform jobs submit training $JOB_NAME \
      --region=$REGION \
      --job-dir=$JOB_DIR \
      --master-image-uri=$IMAGE_URI \
      --scale-tier=$SCALE_TIER \
      -- \
      --training_dataset_path=$TRAINING_FILE_PATH \
      --validation_dataset_path=$VALIDATION_FILE_PATH \
      --alpha=$alpha \
      --max_iter=$max_iter \
      --nohptune
```

Best hyperparameters

Retrain with the best hyperparameters and export

```
gcloud ai-platform jobs submit training $JOB_NAME \
      --region=$REGION \
      --job-dir=$JOB_DIR \
      --master-image-uri=$IMAGE_URI \
      --scale-tier=$SCALE_TIER \
      -- \
      --training_dataset_path=$TRAINING_FILE_PATH \
      --validation_dataset_path=$VALIDATION_FILE_PATH \
      --alpha=$alpha \
      --max_iter=$max_iter \
      --nohptune
```

Training done with `hptune = False`

Model now exported as `model.pkl` on Cloud Storage

```
if not hptune:
    model_filename = 'model.pkl'
    with open(model_filename, 'wb') as model_file:
        pickle.dump(pipeline, model_file)
    gcs_model_path = "{}/{}".format(job_dir, model_filename)
    subprocess.check_call(['gsutil', 'cp', model_filename,
                            gcs_model_path], stderr=sys.stdout)
```

Agenda

System and Concepts Overview

Create a Reproducible Dataset

Implement a Tunable Model

Build and Push a Training Container

Train and Tune a Model

Serve and Query a Model

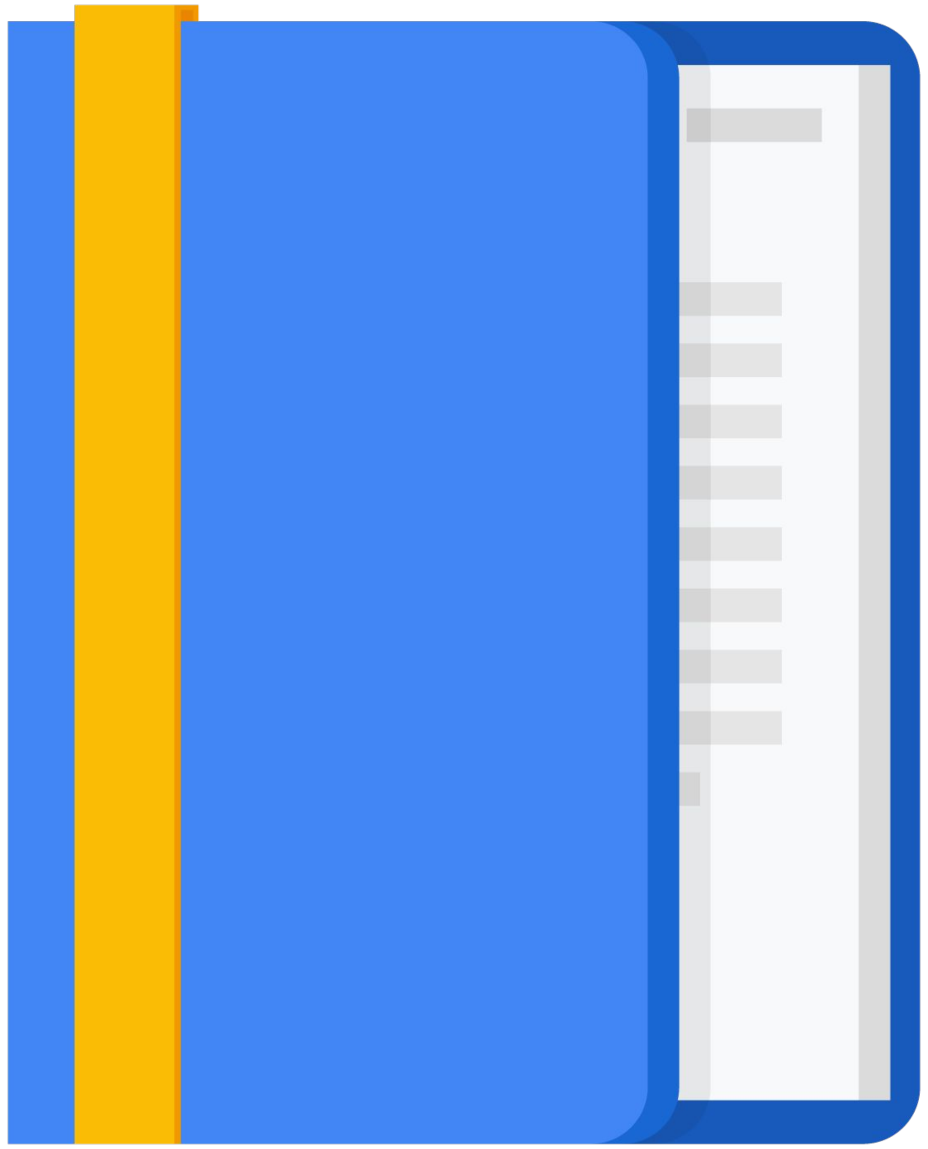

System overview

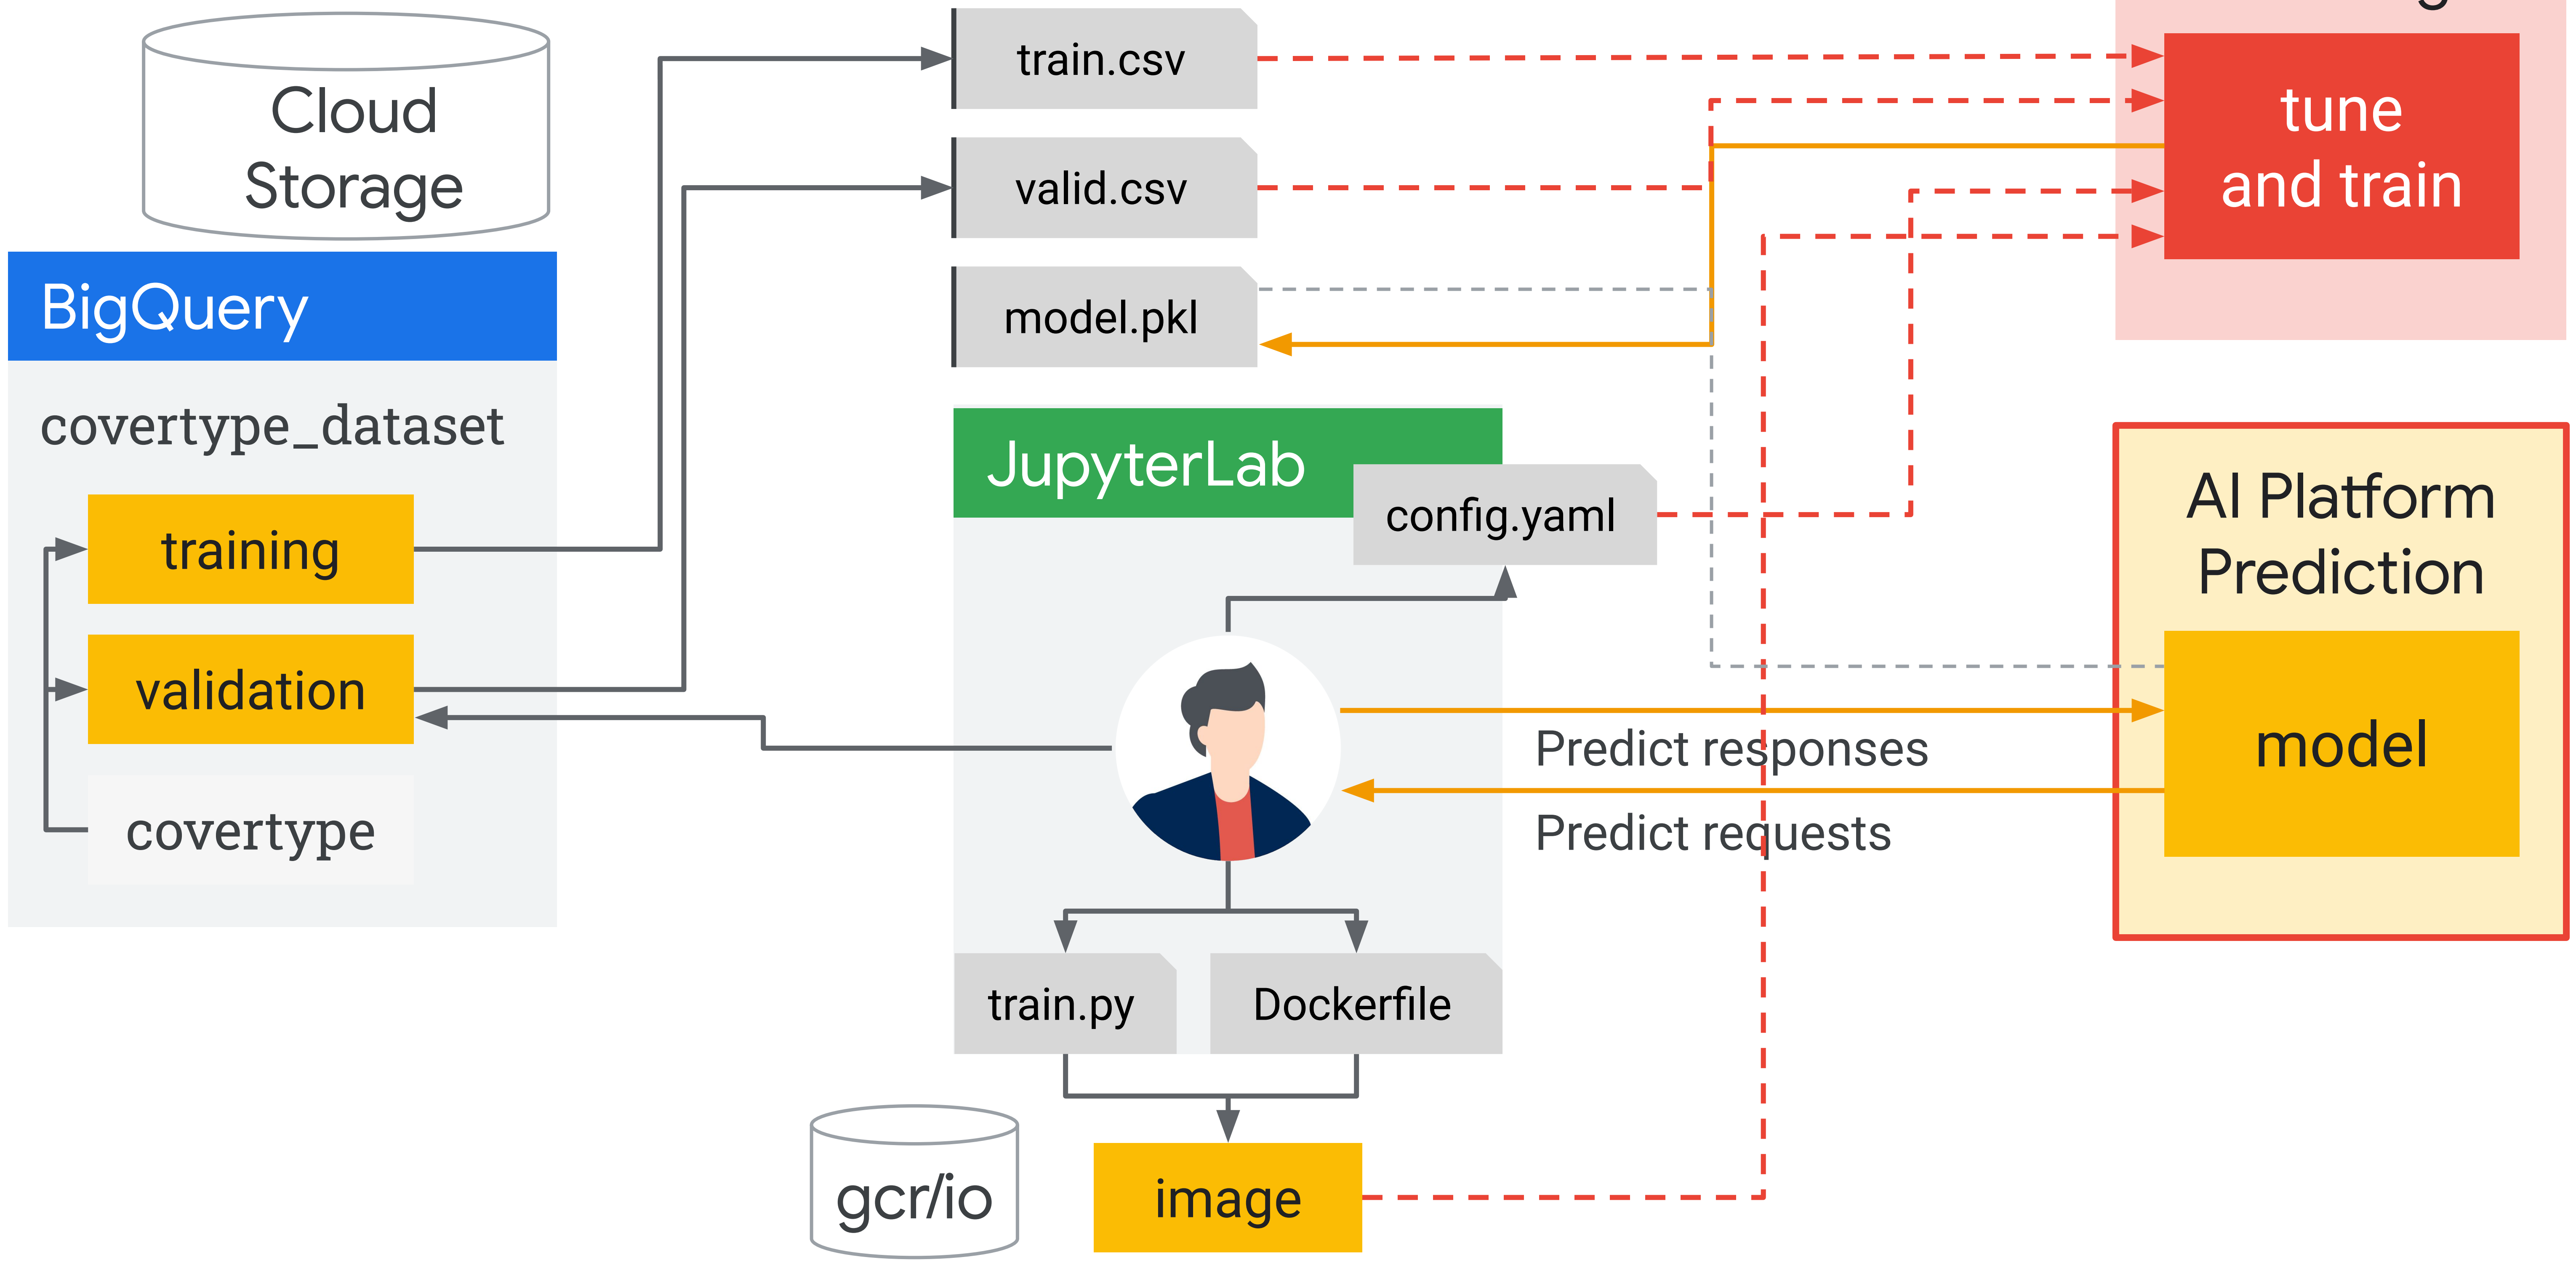

AI Platform Prediction makes deploying models easy

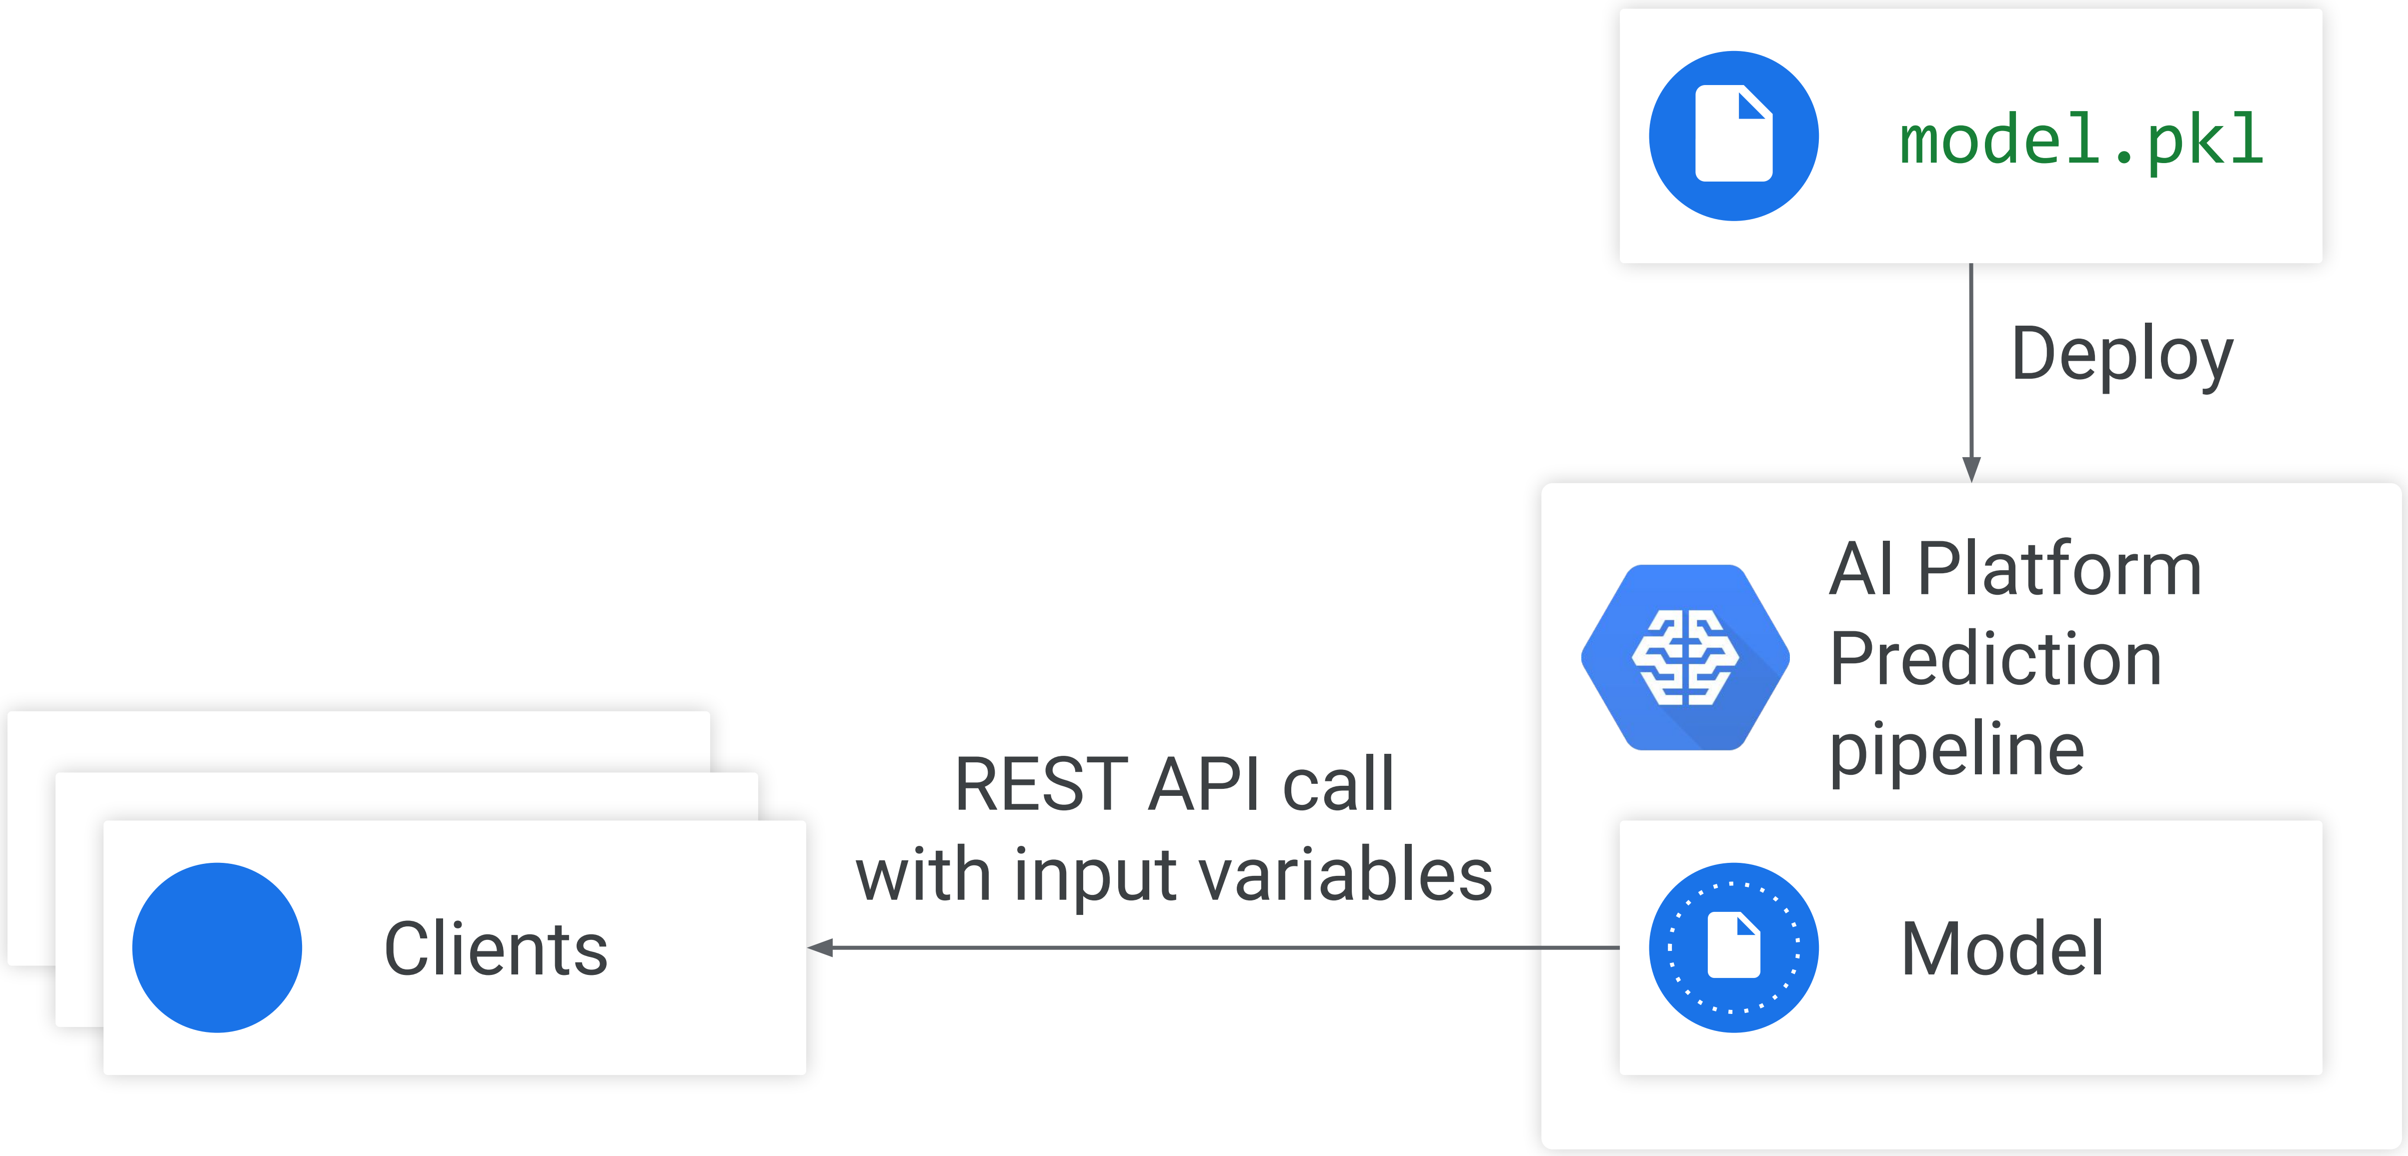

Create a model resource

```
gcloud ai-platform models create $model_name \
  --regions=$REGION \
  --labels=$labels
```

Create a model version

```
gcloud ai-platform versions create {model_version} \
    --model={model_name} \
    --origin=$JOB_DIR \
    --runtime-version=1.15 \
    --framework=scikit-learn \
    --python-version=3.7
```

Create a model version

```
gcloud ai-platform versions create {model_version} \
    --model={model_name} \
    --origin=$JOB_DIR \
    --runtime-version=1.15 \
    --framework=scikit-learn \
    --python-version=3.7
```

where `model.pkl` is exported

Create a model version

```
gcloud ai-platform versions create {model_version} \
    --model={model_name} \
    --origin=$JOB_DIR \
    --runtime-version=1.15 \
    --framework=scikit-learn \
    --python-version=3.7
```

Query the model

```
gcloud ai-platform predict \
  --model $model_name \
  --version $model_version \
  --json-instances $input_file
```

Query the model

```
gcloud ai-platform predict \
  --model $model_name \
  --version $model_version \
  --json-instances $input_file
```

Query the model

```
gcloud ai-platform predict \
  --model $model_name \
  --version $model_version \
  --json-instances $input_file
```

The data we send to the prediction API

Lab

Training, Tuning, and Serving in AI Platform

cloud.google.com