



Kubeflow Pipelines on AI Platform

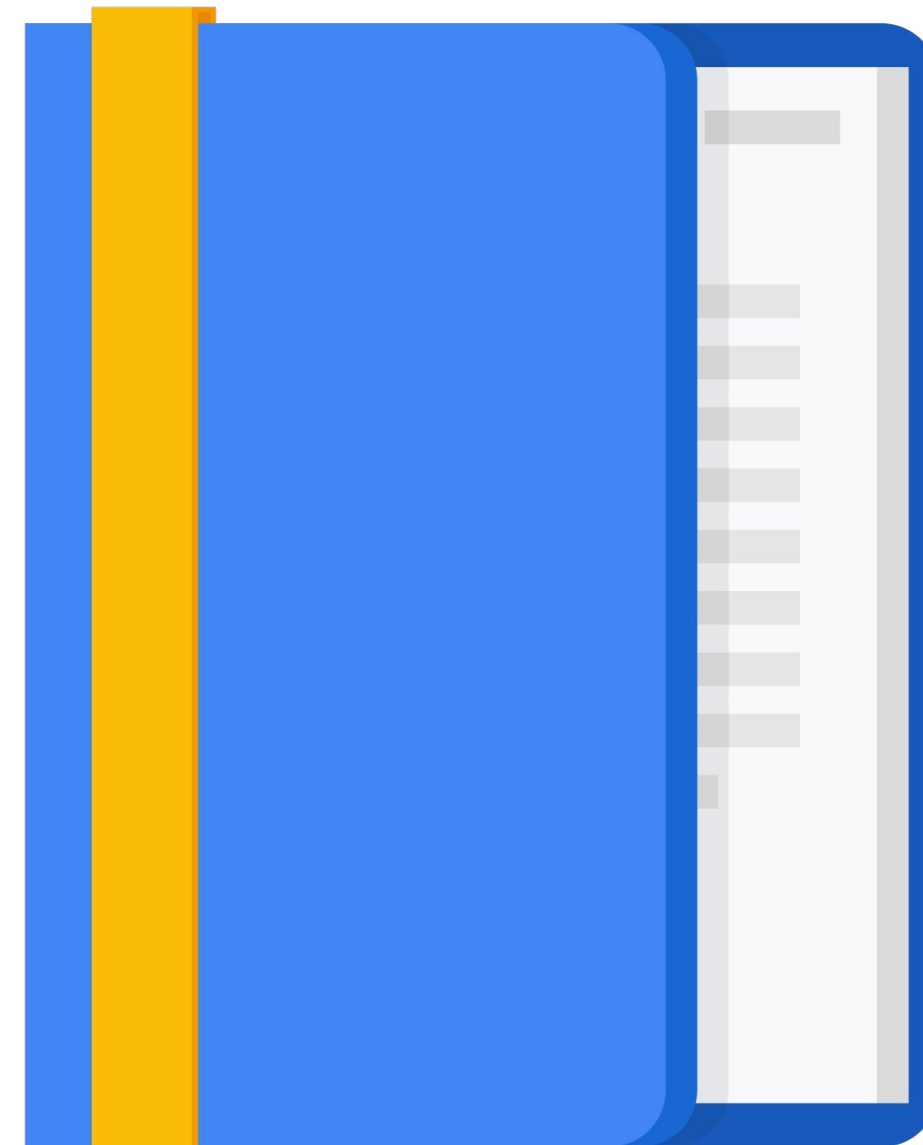
Benoit Dherin

ML Engineer, Google Advanced Solutions Lab

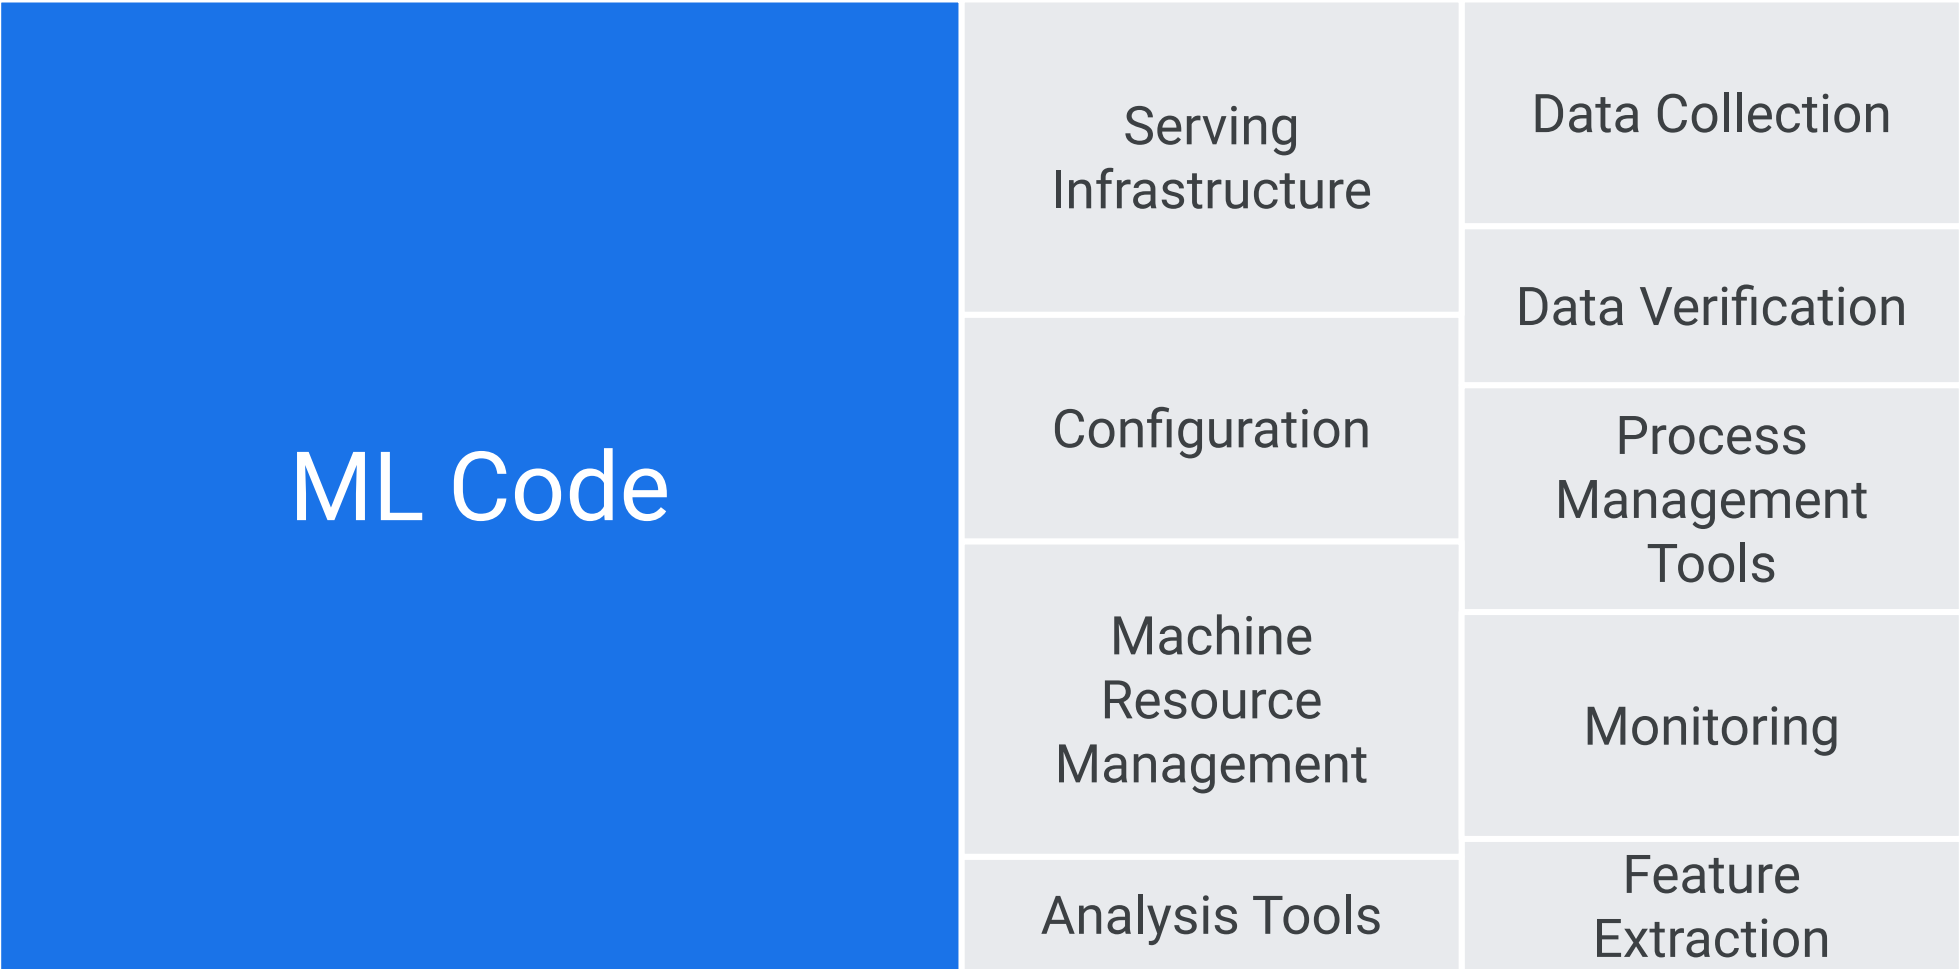


Agenda

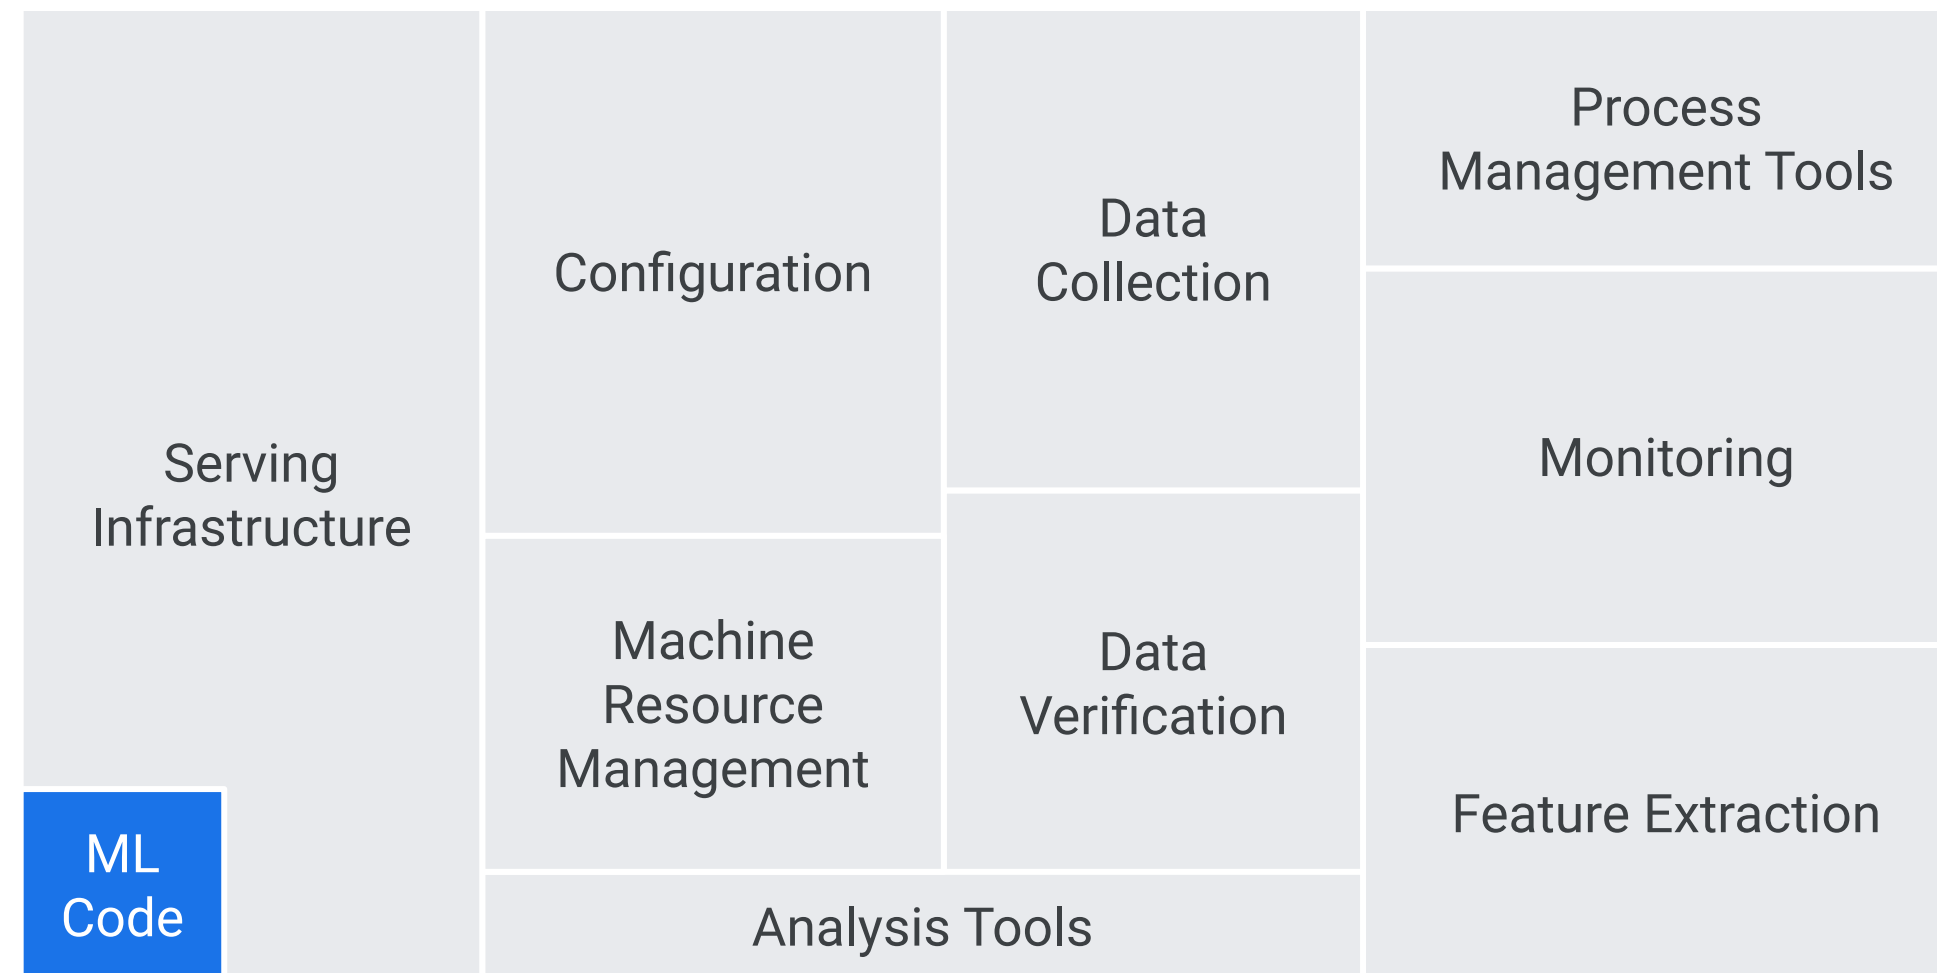
- System and Concept Overview
- Describing a Kubeflow Pipeline with KF DSL
- Pre-built Components
- Lightweight Python Components
- Custom Components
- Compile, Upload, and Run




Perception: ML products are mostly about ML



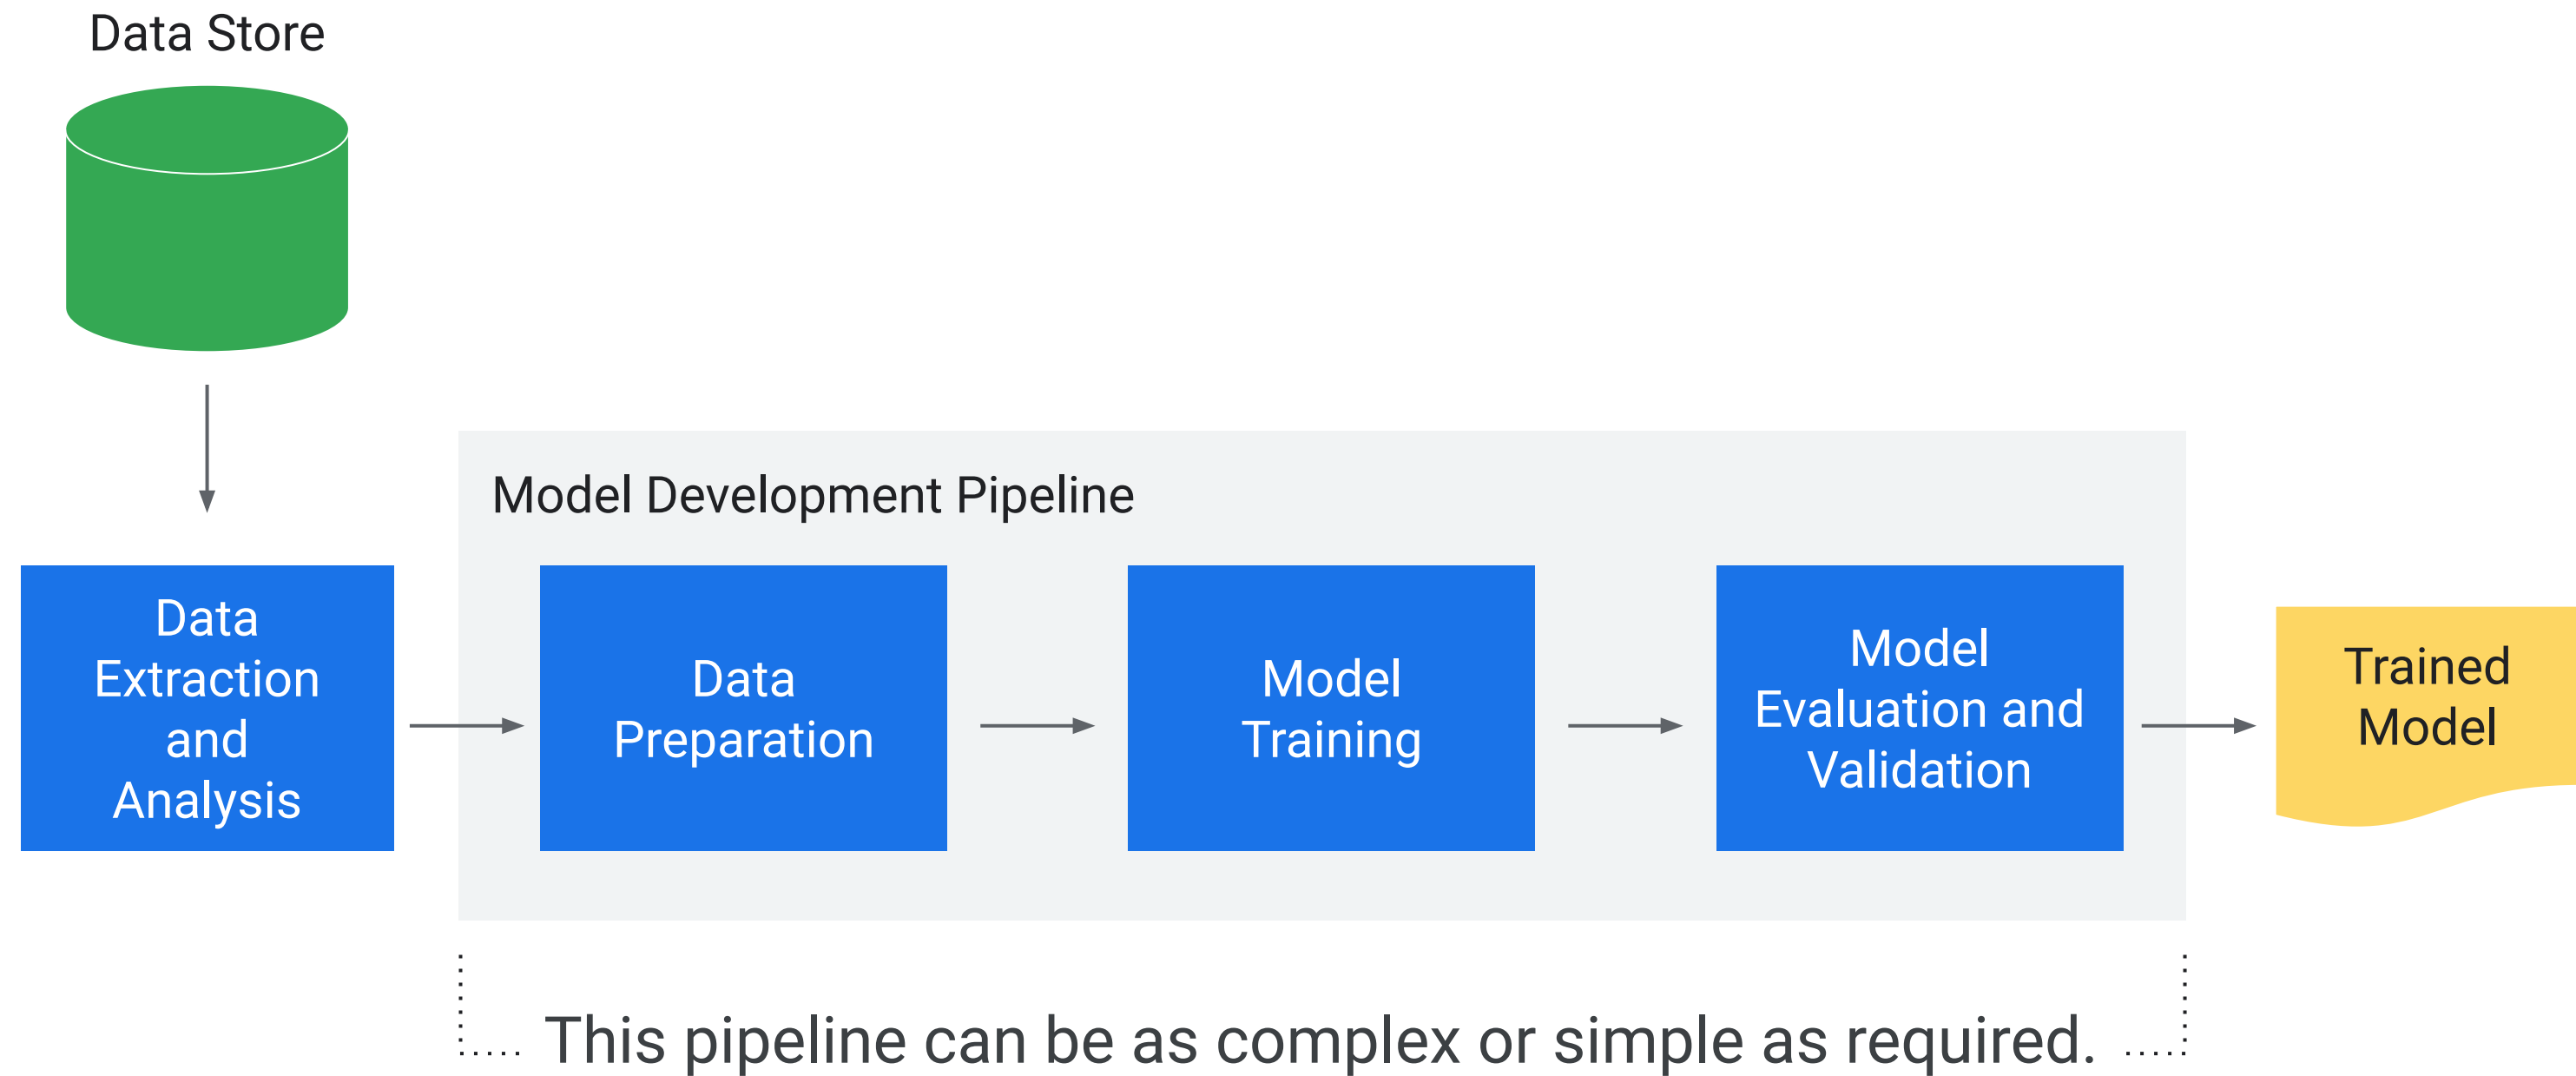
Reality: ML Requires lots of DevOps



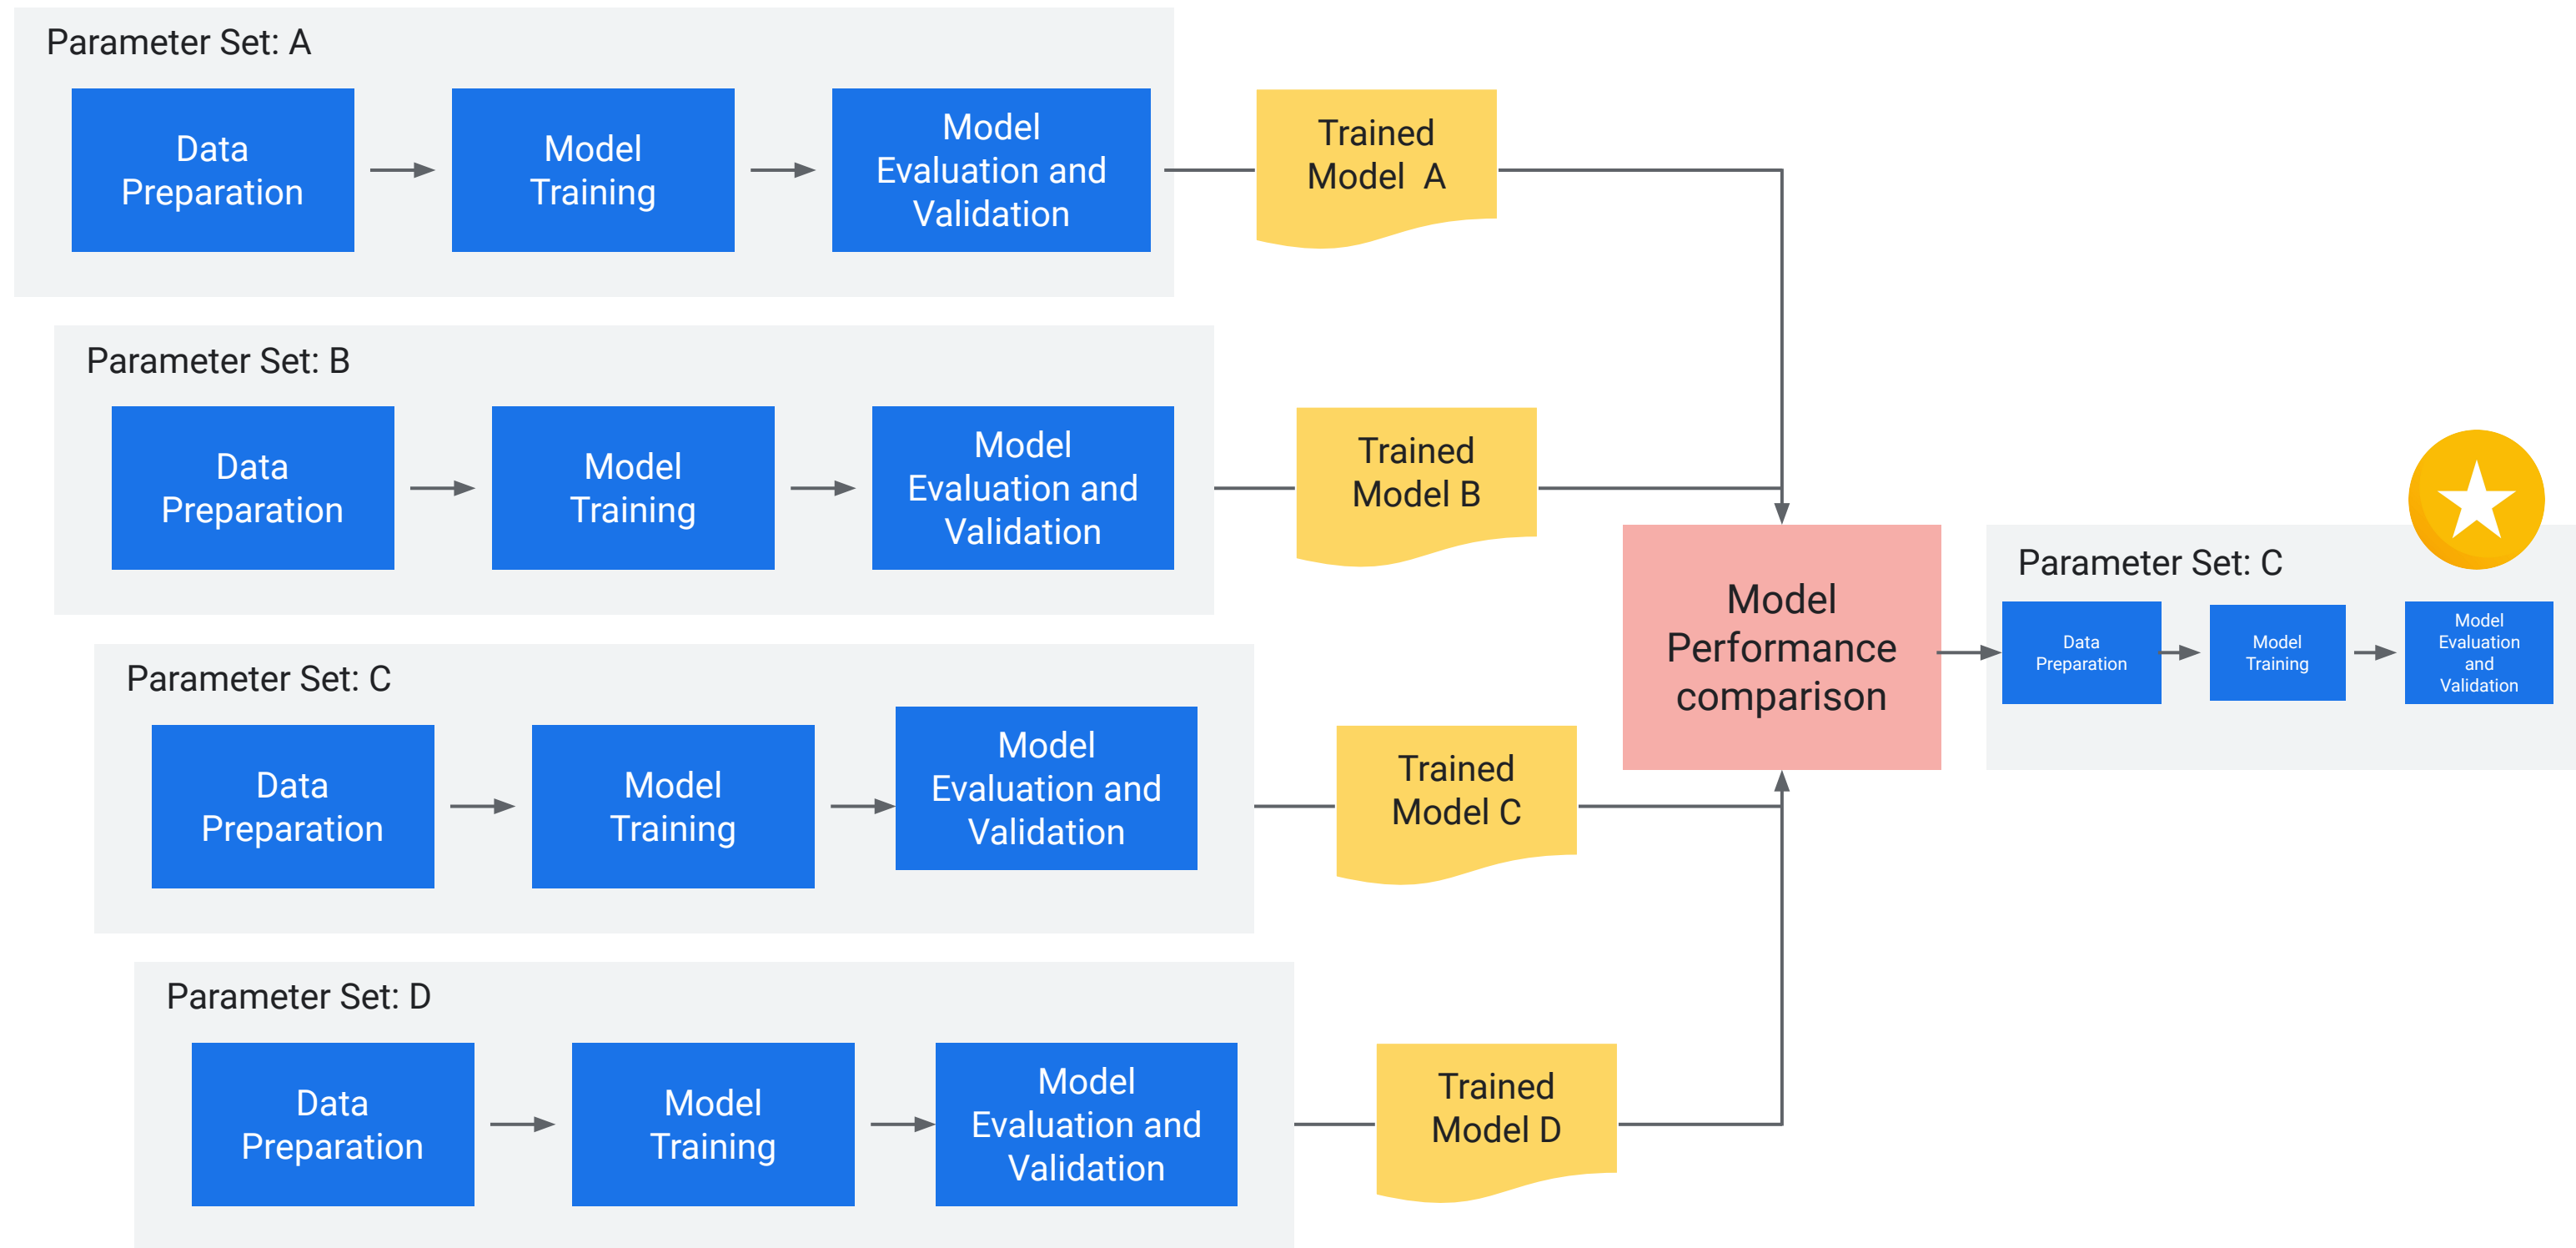
Source: [Sculley et al.: Hidden Technical Debt in Machine Learning Systems](#)



The ML process




Machine learning is all about experimentation!

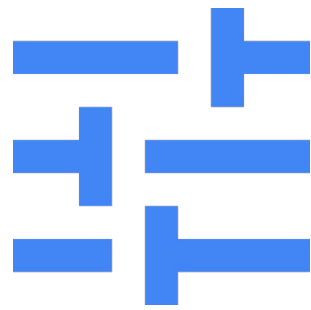


Kubeflow provides a standardized platform for building ML pipelines

- Leverage containers and Kubernetes so that in ML pipelines can be run on a cloud or on-premises with Anthos on GKE.
- Kubeflow is a cloud-native, multi-cloud solution for ML.
- Kubeflow provides a platform for composable, portable, and scalable ML pipelines.
- If you have a Kubernetes-conformant cluster, you can run Kubeflow.



Kubeflow pipelines enable:



ML workflow
orchestration



Share, re-use,
and compose



Rapid, reliable
experimentation

What constitutes a Kubeflow pipeline?

Containerized implementations of ML tasks


- Example of ML tasks: Data import, training, serving, model evaluation
- Containers provide portability, repeatability, and encapsulation.
- A containerized task can invoke other services, such as AI Platform, Dataflow, or Dataproc.

Specification of the sequence of steps

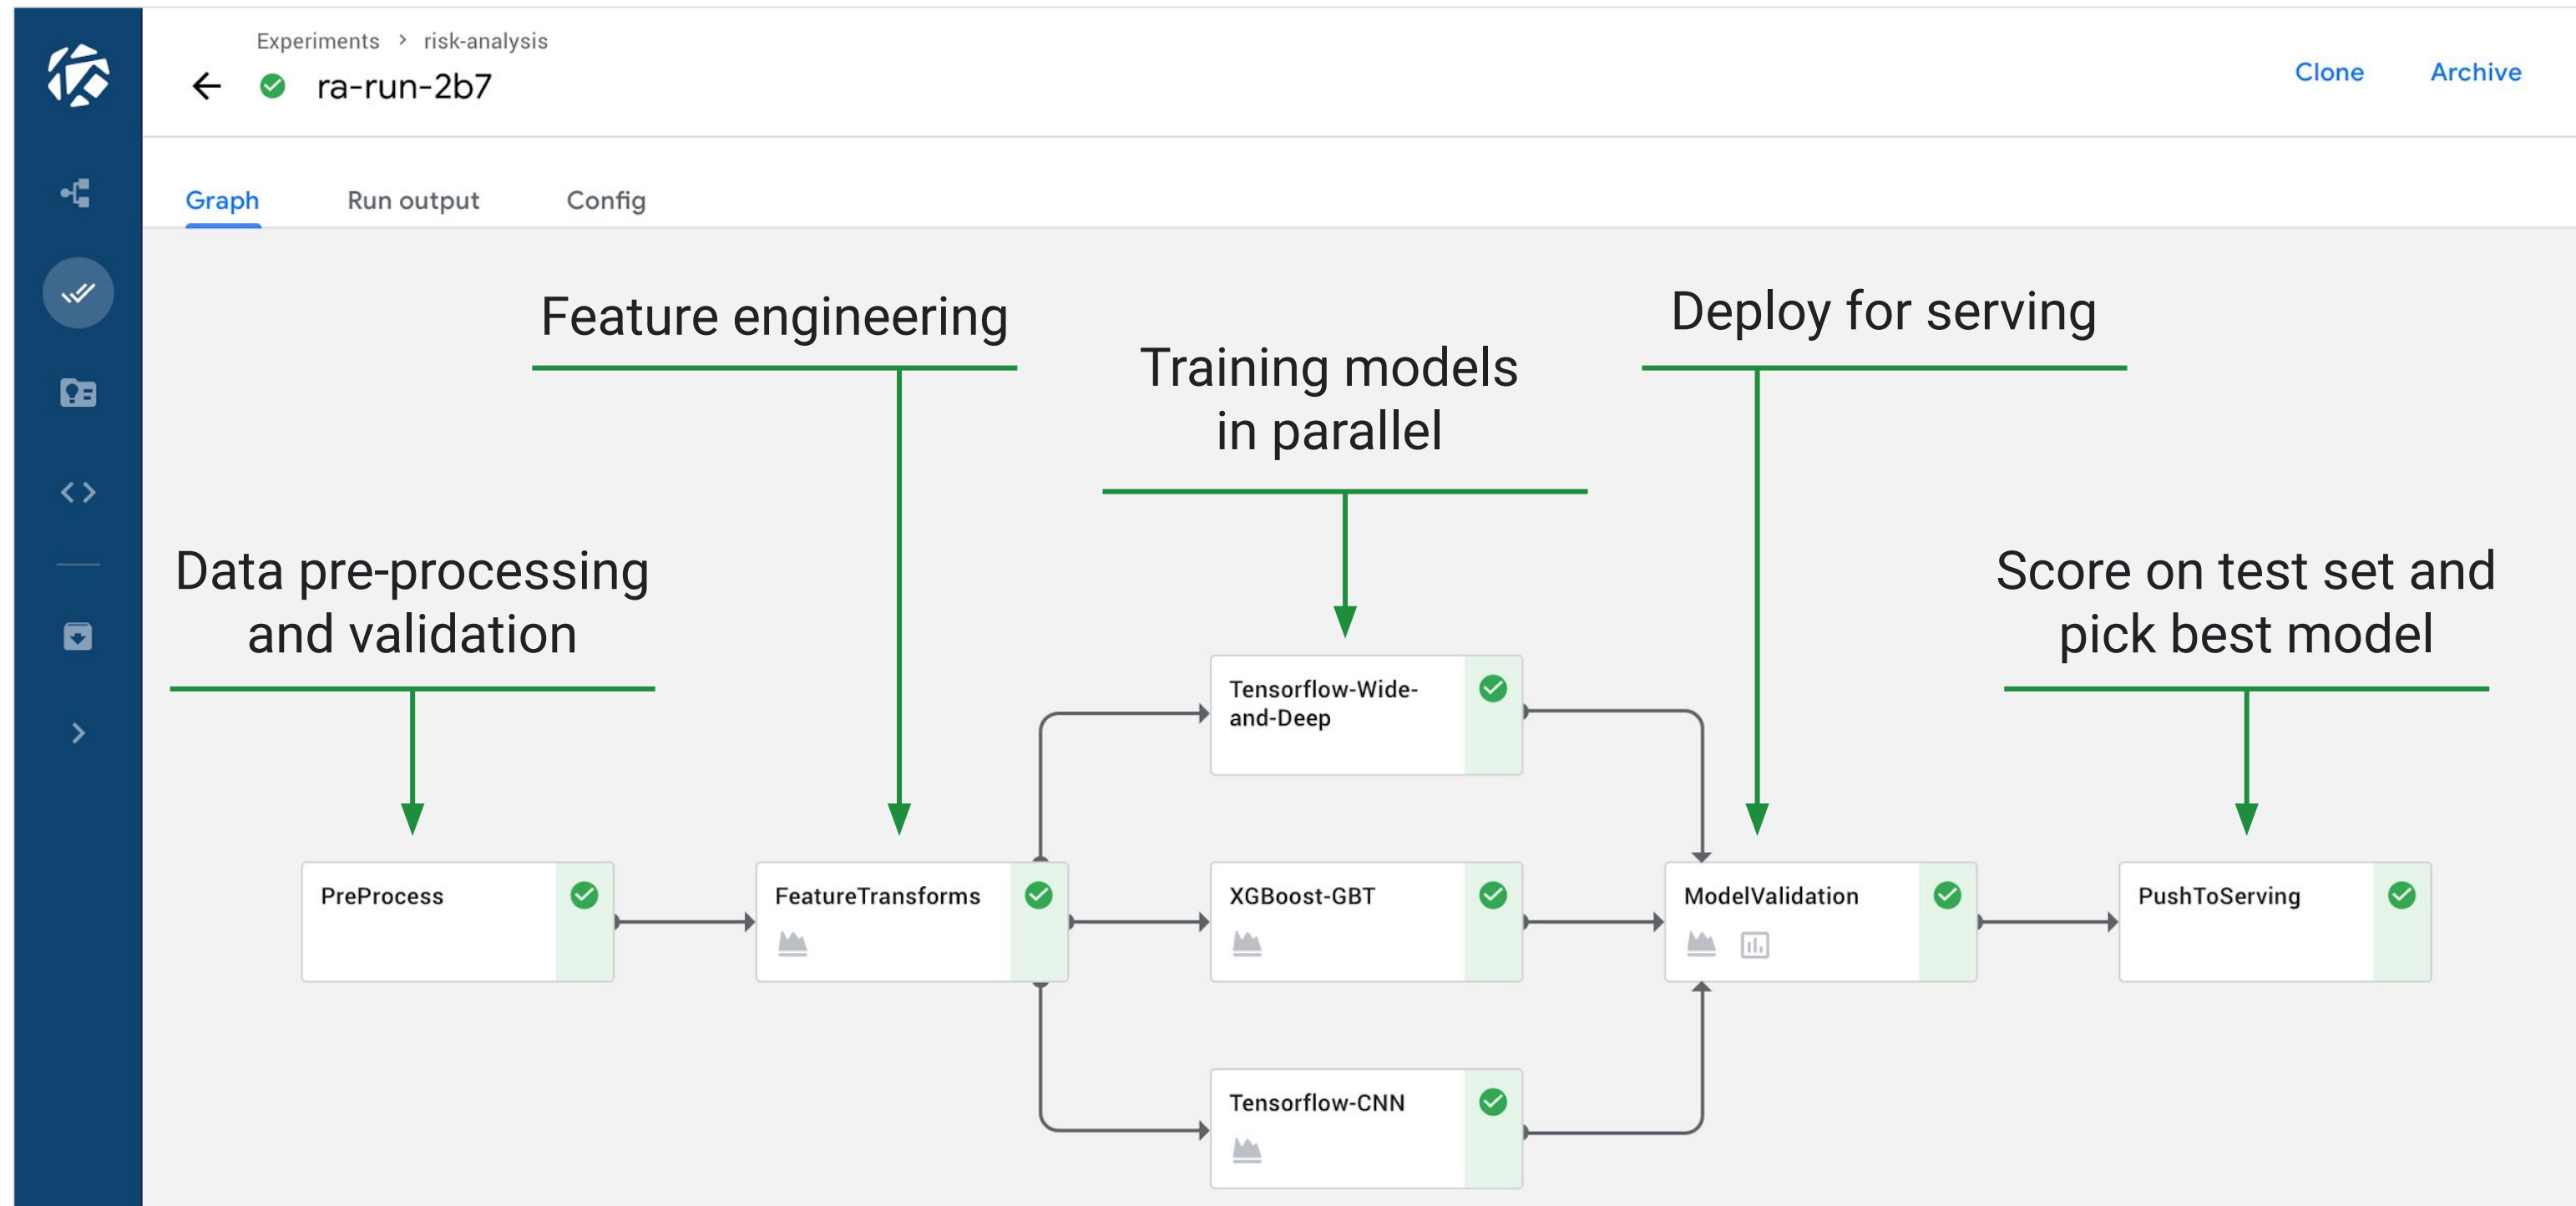
- Specified via Python SDK

Input parameters

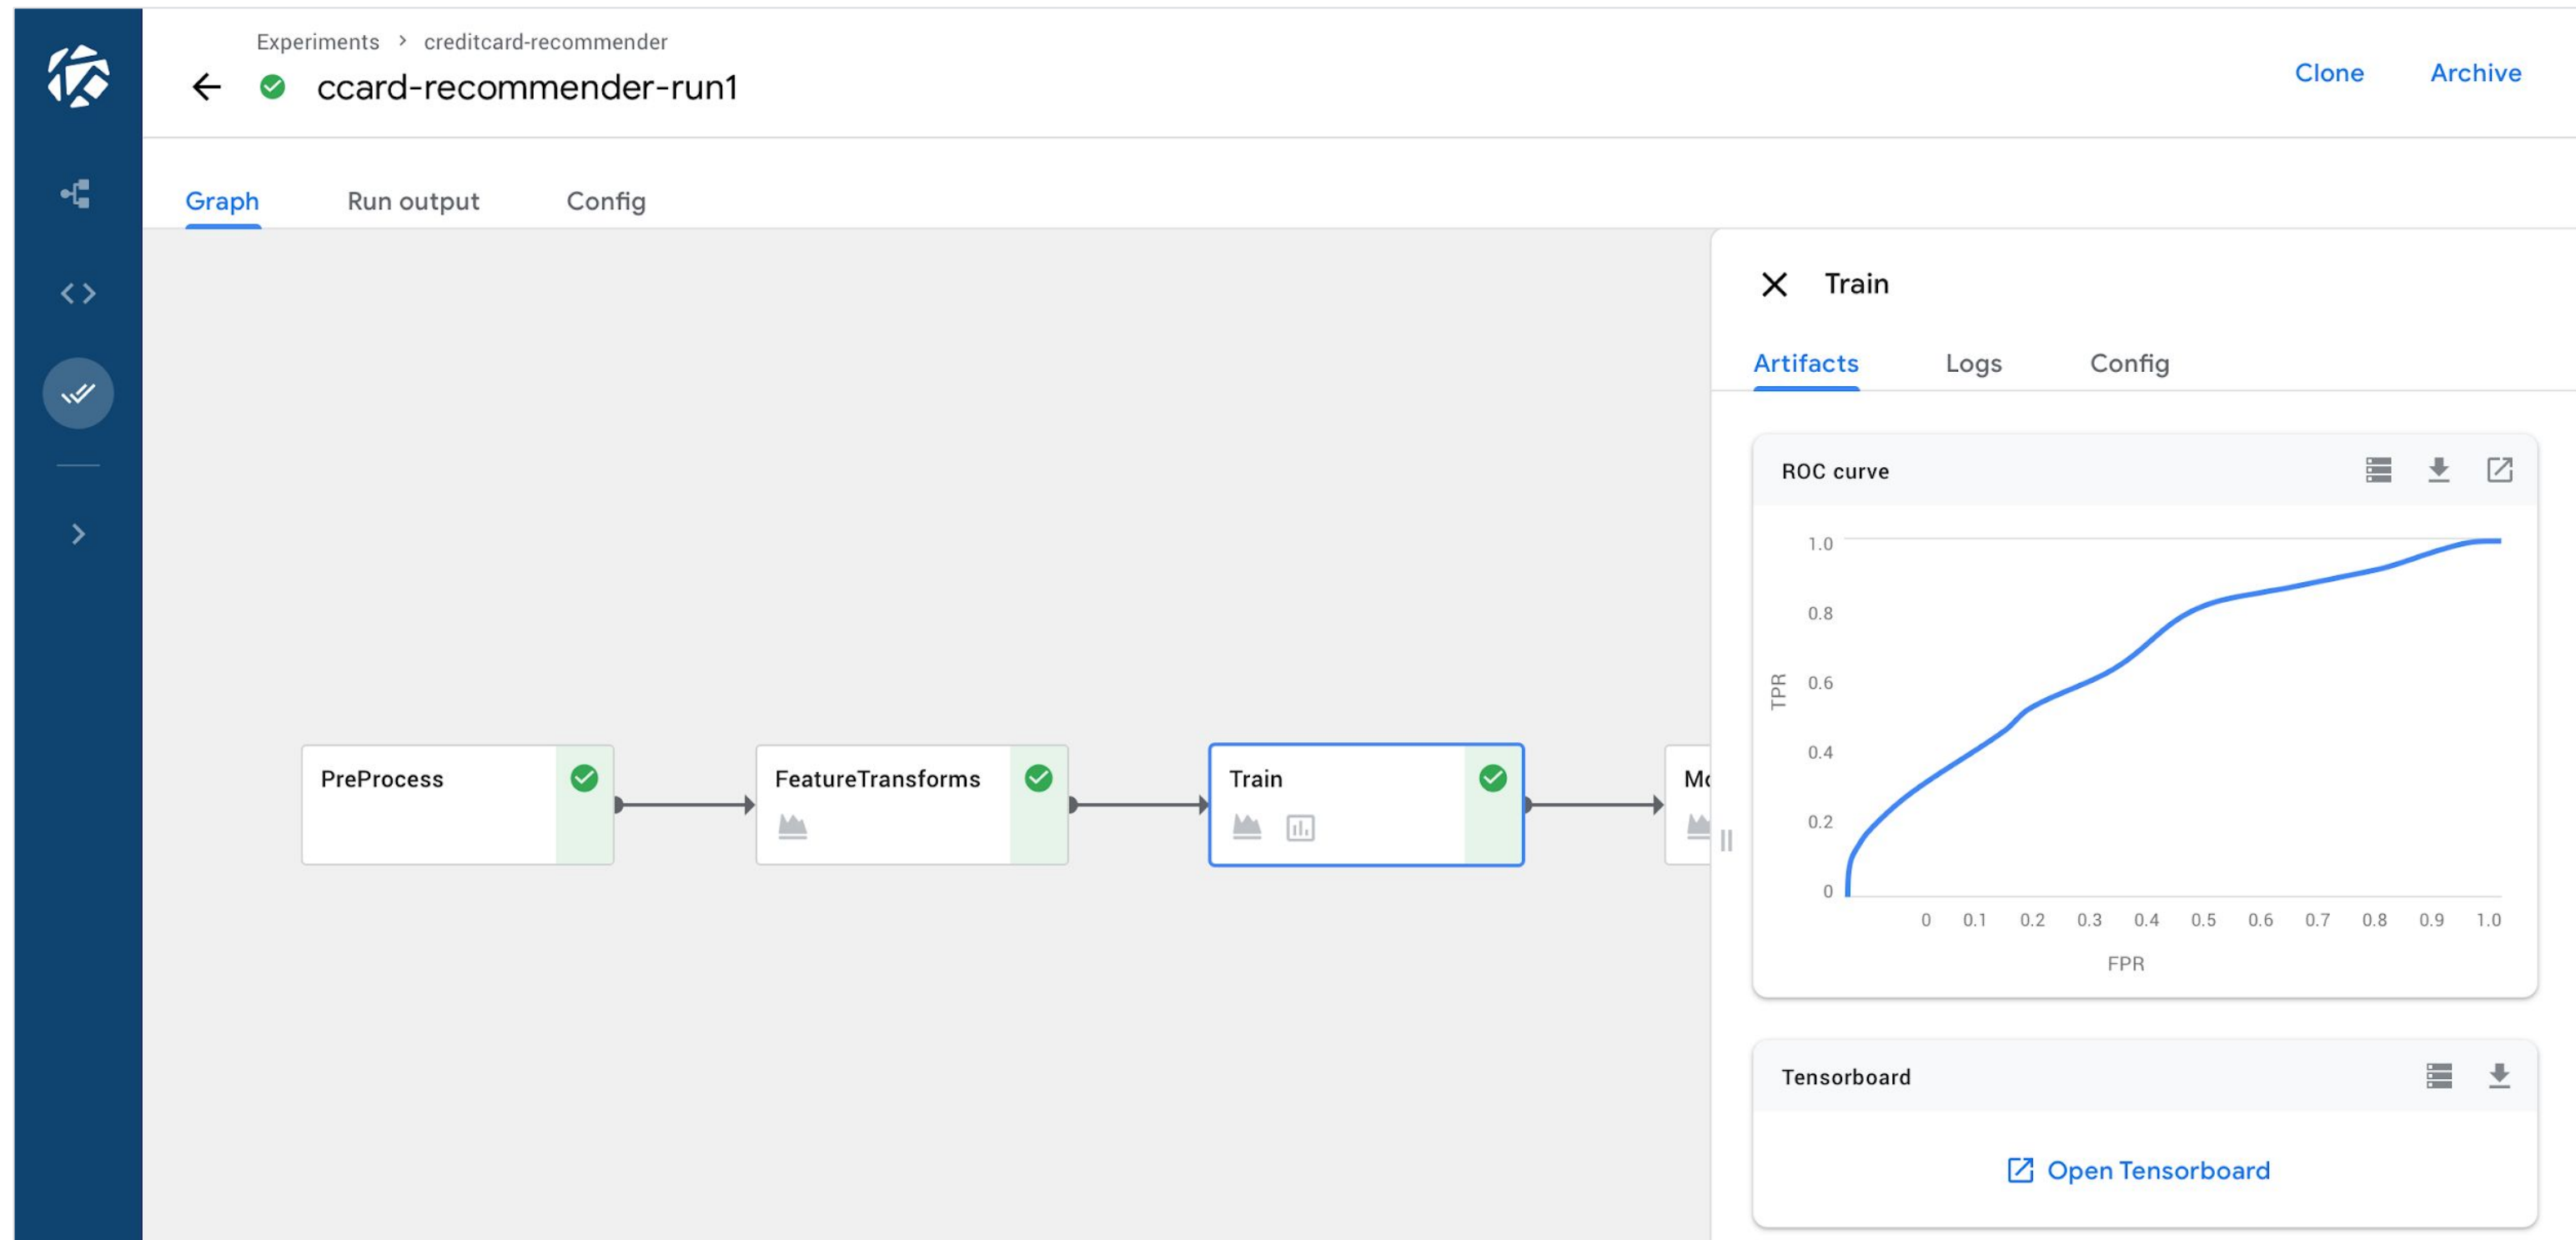
- A “Job” is a pipeline invoked w/specific parameters








Visual depiction of pipeline topology



Rich visualization of metrics



View all configs,
inputs, and outputs



Experiments > Product Image Classification

←

✔ Simple XGBoost Classifier

Graph

Config

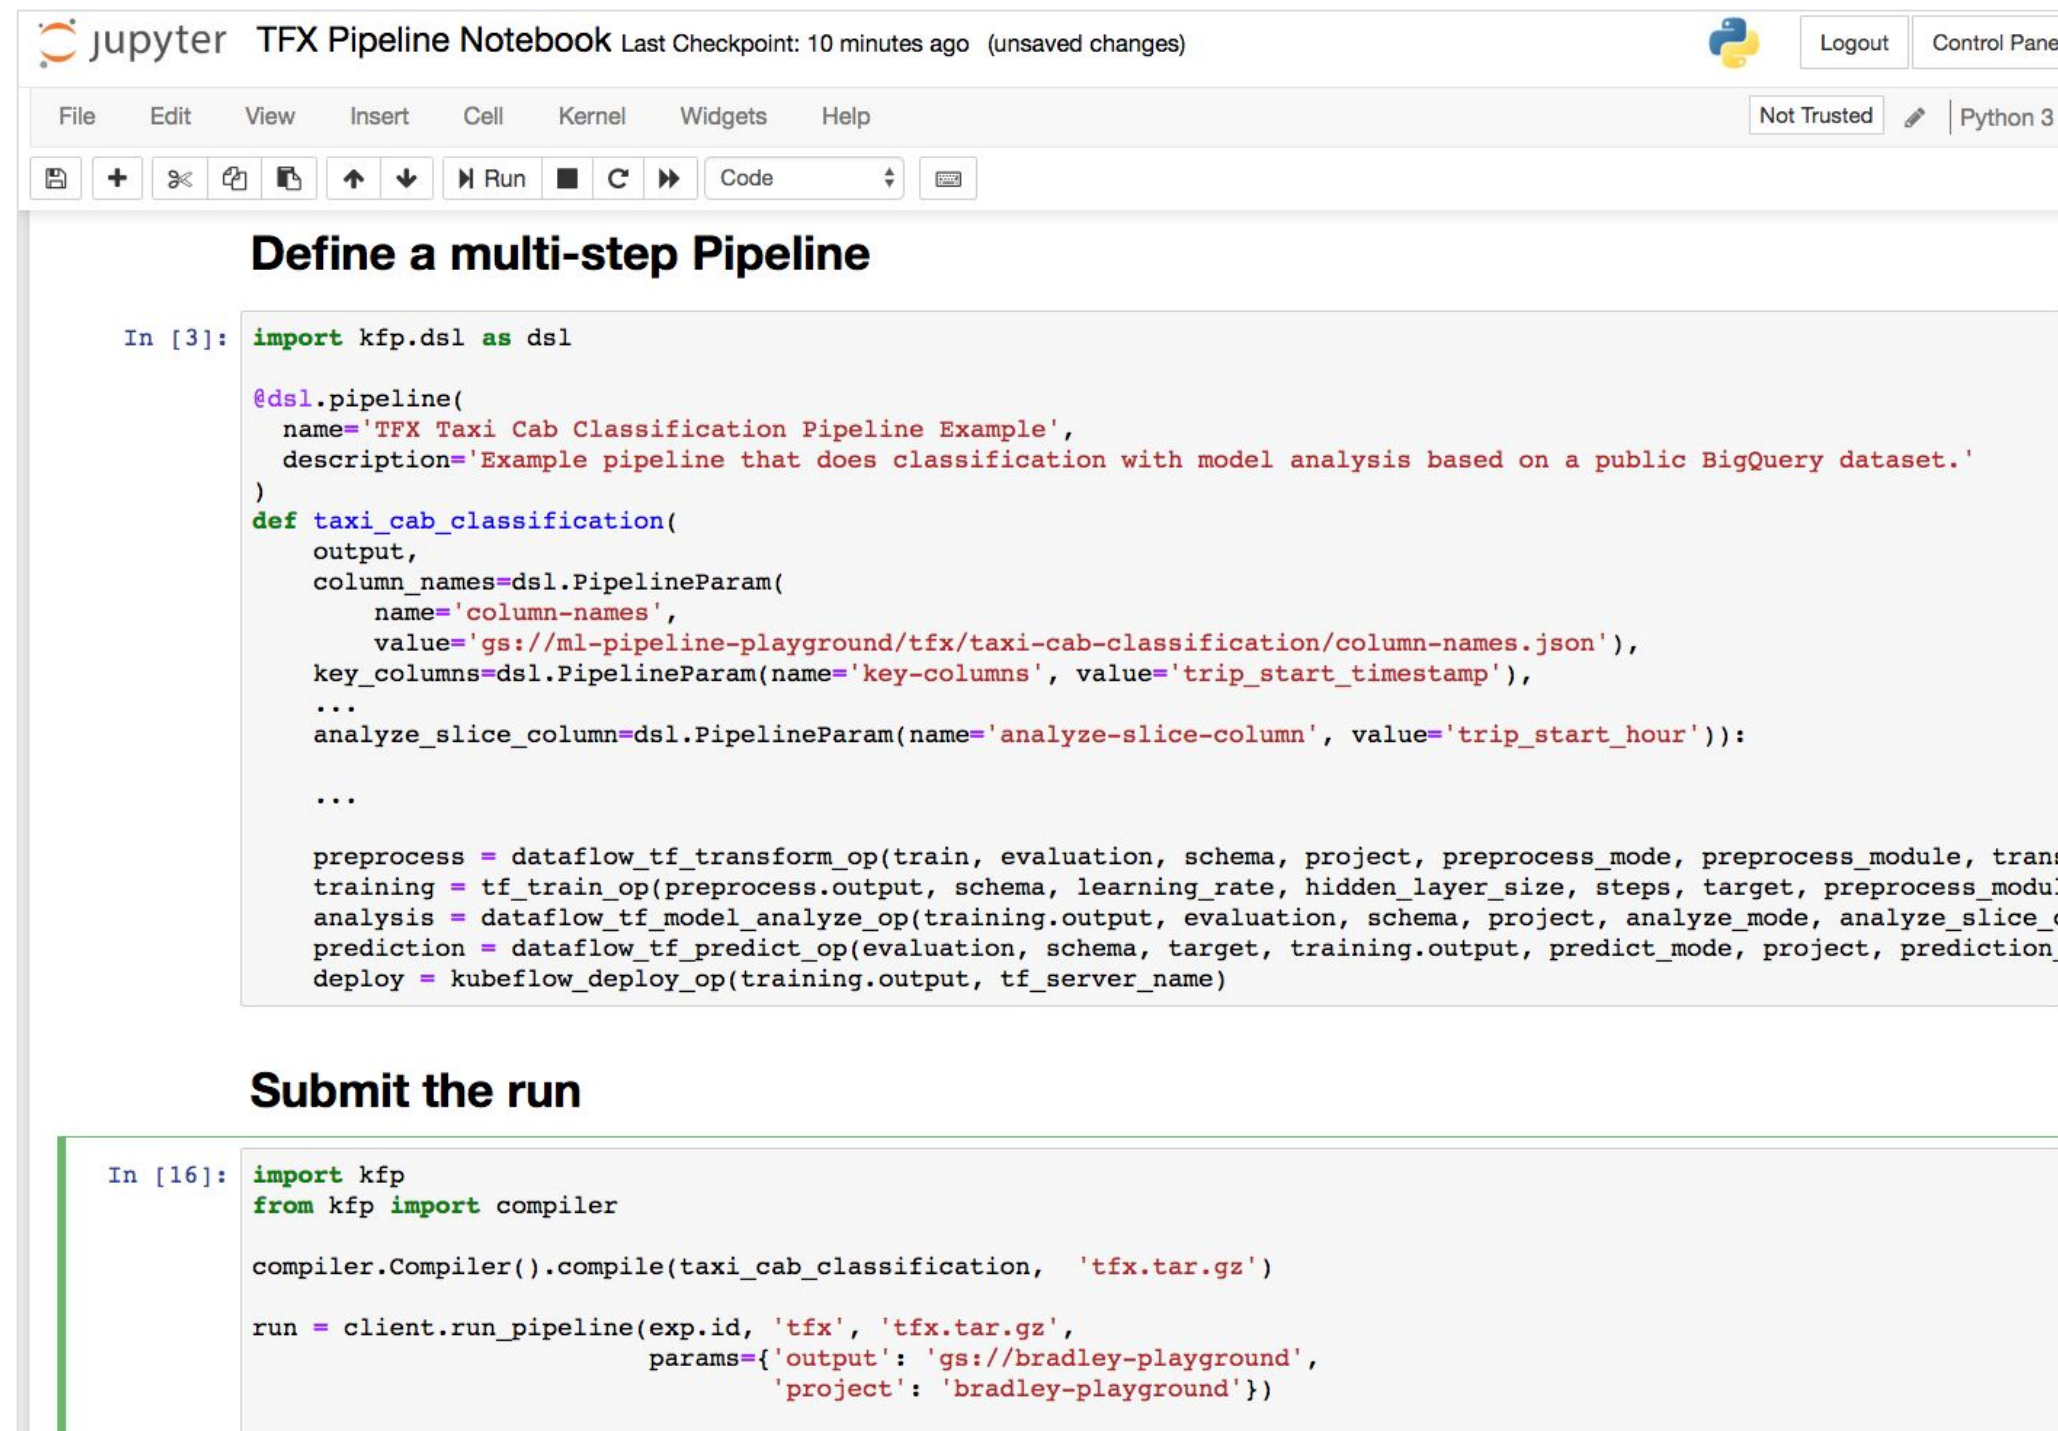
Run details

Status	Succeeded
Description	
Created at	11/25/2018, 12:56:44 PM
Started at	11/25/2018, 12:56:44 PM
Finished	11/25/2018, 12:16:37 PM
Duration	0:19:53

Run parameters

output	gs://mipipelines
project	foo2thebar
region	us-central1
train-data	gs://ml-pipeline-playground/sfpd/train.csv
eval-data	gs://ml-pipeline-playground/sfpd/eval.csv
schema	gs://ml-pipeline-playground/sfpd/schema.json
target	resolution
rounds	200

Author pipelines with an intuitive Python SDK



The screenshot shows a Jupyter Notebook interface with the title 'TFX Pipeline Notebook'. The top bar indicates 'Last Checkpoint: 10 minutes ago' and '(unsaved changes)'. The notebook has a menu bar (File, Edit, View, Insert, Cell, Kernel, Widgets, Help) and a toolbar with icons for saving, adding cells, and running code. The code is written in Python and defines a multi-step pipeline using the KubeFlow Pipeline SDK (kfp.dsl).

Define a multi-step Pipeline

```
In [3]: import kfp.dsl as dsl

@dsl.pipeline(
    name='TFX Taxi Cab Classification Pipeline Example',
    description='Example pipeline that does classification with model analysis based on a public BigQuery dataset.'
)
def taxi_cab_classification(
    output,
    column_names=dsl.PipelineParam(
        name='column-names',
        value='gs://ml-pipeline-playground/tfx/taxi-cab-classification/column-names.json',
    ),
    key_columns=dsl.PipelineParam(name='key-columns', value='trip_start_timestamp'),
    ...
    analyze_slice_column=dsl.PipelineParam(name='analyze-slice-column', value='trip_start_hour')):
    ...


    preprocess = dataflow_tf_transform_op(train, evaluation, schema, project, preprocess_mode, preprocess_module, transform_op_name)
    training = tf_train_op(preprocess.output, schema, learning_rate, hidden_layer_size, steps, target, preprocess_module)
    analysis = dataflow_tf_model_analyze_op(training.output, evaluation, schema, project, analyze_mode, analyze_slice_column)
    prediction = dataflow_tf_predict_op(evaluation, schema, target, training.output, predict_mode, project, prediction_module)
    deploy = kubeflow_deploy_op(training.output, tf_server_name)
```

Submit the run

```
In [16]: import kfp
from kfp import compiler

compiler.Compiler().compile(taxi_cab_classification, 'tfx.tar.gz')

run = client.run_pipeline(exp.id, 'tfx', 'tfx.tar.gz',
    params={'output': 'gs://bradley-playground',
            'project': 'bradley-playground'})
```



Package and share pipelines as zip files

- Upload and execute pipelines via UI (in addition to API/SDK).
- Pipeline steps can be authored as reusable components.

Run details

Pipeline*

xxboost training - confusion matrix

Choose

Run name*

product-recommender-model

Description (optional)

Train XBG model for product recommendation application.

Run parameters

Specify parameters required by the pipeline

output

project

region

us-central1

train-data

gs://ml-pipeline-playground/sfpd/train.csv

eval-data

gs://ml-pipeline-playground/sfpd/eval.csv

schema

gs://ml-pipeline-playground/sfpd/schema.json

target

resolution

rounds

200

workers

2

true-label

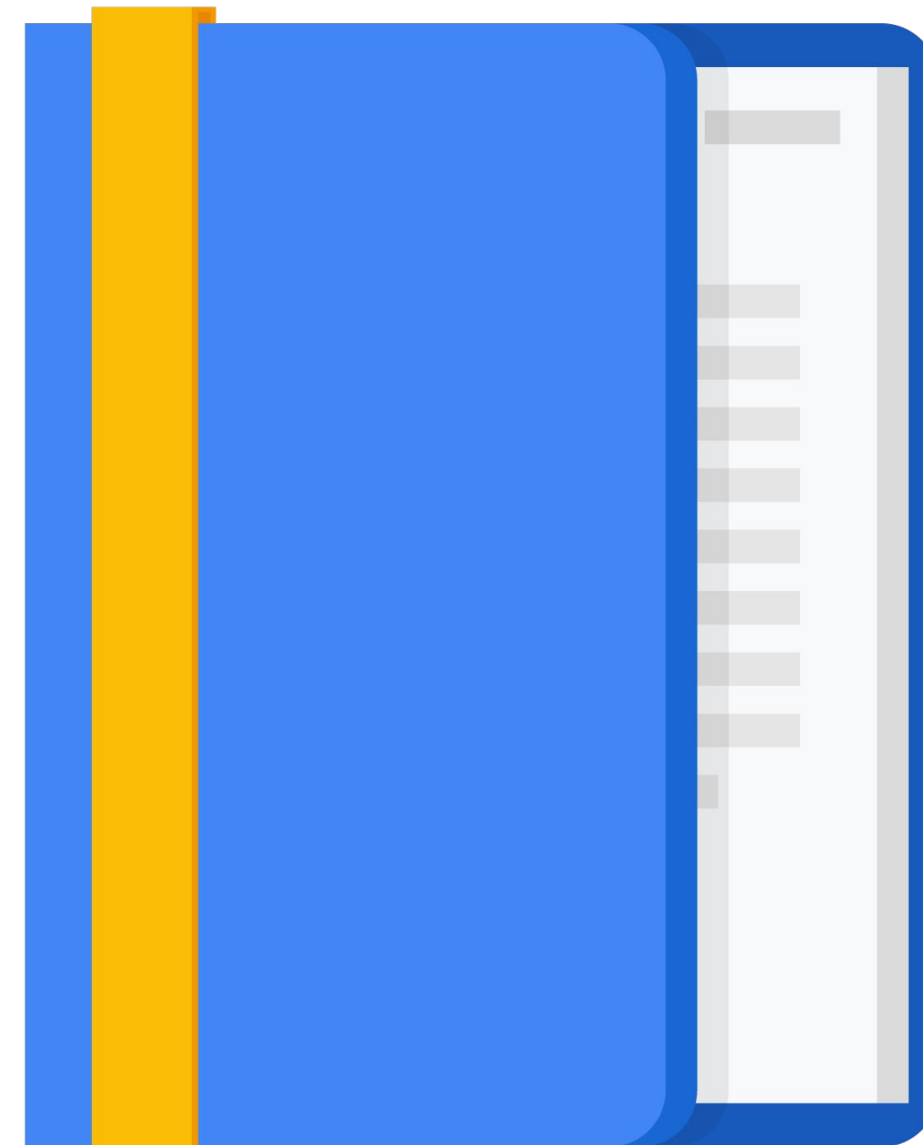
ACTION

Create

Cancel


Agenda

- System and Concept Overview
- Describing a Kubeflow Pipeline with KF DSL
- Pre-built Components
- Lightweight Python Components
- Custom Components
- Compile, Upload, and Run



Kubeflow offers a **Domain Specific Language** (DSL) in Python that allows you to use Python code to describe Kubeflow tasks as they organize themselves in a Directed Acyclic Graph (DAG).

We describe this DSL next...



```
import kfp
```

```
@kfp.dsl.pipeline(  
    name='Covertypes Classifier Training',  
    description='Covertypes training and deployment pipeline',  
)
```

```
def covertypes_train(project_id,  
    region,  
    source_table_name,  
    gcs_root,  
    dataset_id,  
    evaluation_metric_name,  
    evaluation_metric_threshold,  
    model_id,  
    version_id,  
    replace_existing_version,  
    hypertune_settings=HYPERTUNE_SETTINGS,  
    dataset_location='US'):
```

Pipeline
Decorator

Pipeline
Run
Parameters

Run parameters

Specify parameters required by the pipeline

project_id

region

source_table_name

gcs_root

dataset_id

evaluation_metric_name

evaluation_metric_threshold

model_id

version_id

replace_existing_version

hypertune_settings

{ "hyperparameters": { "goal": "MAXIMIZE", "maxTrials": 6, "maxParallelTrials": 3, "hyp

dataset_location

US

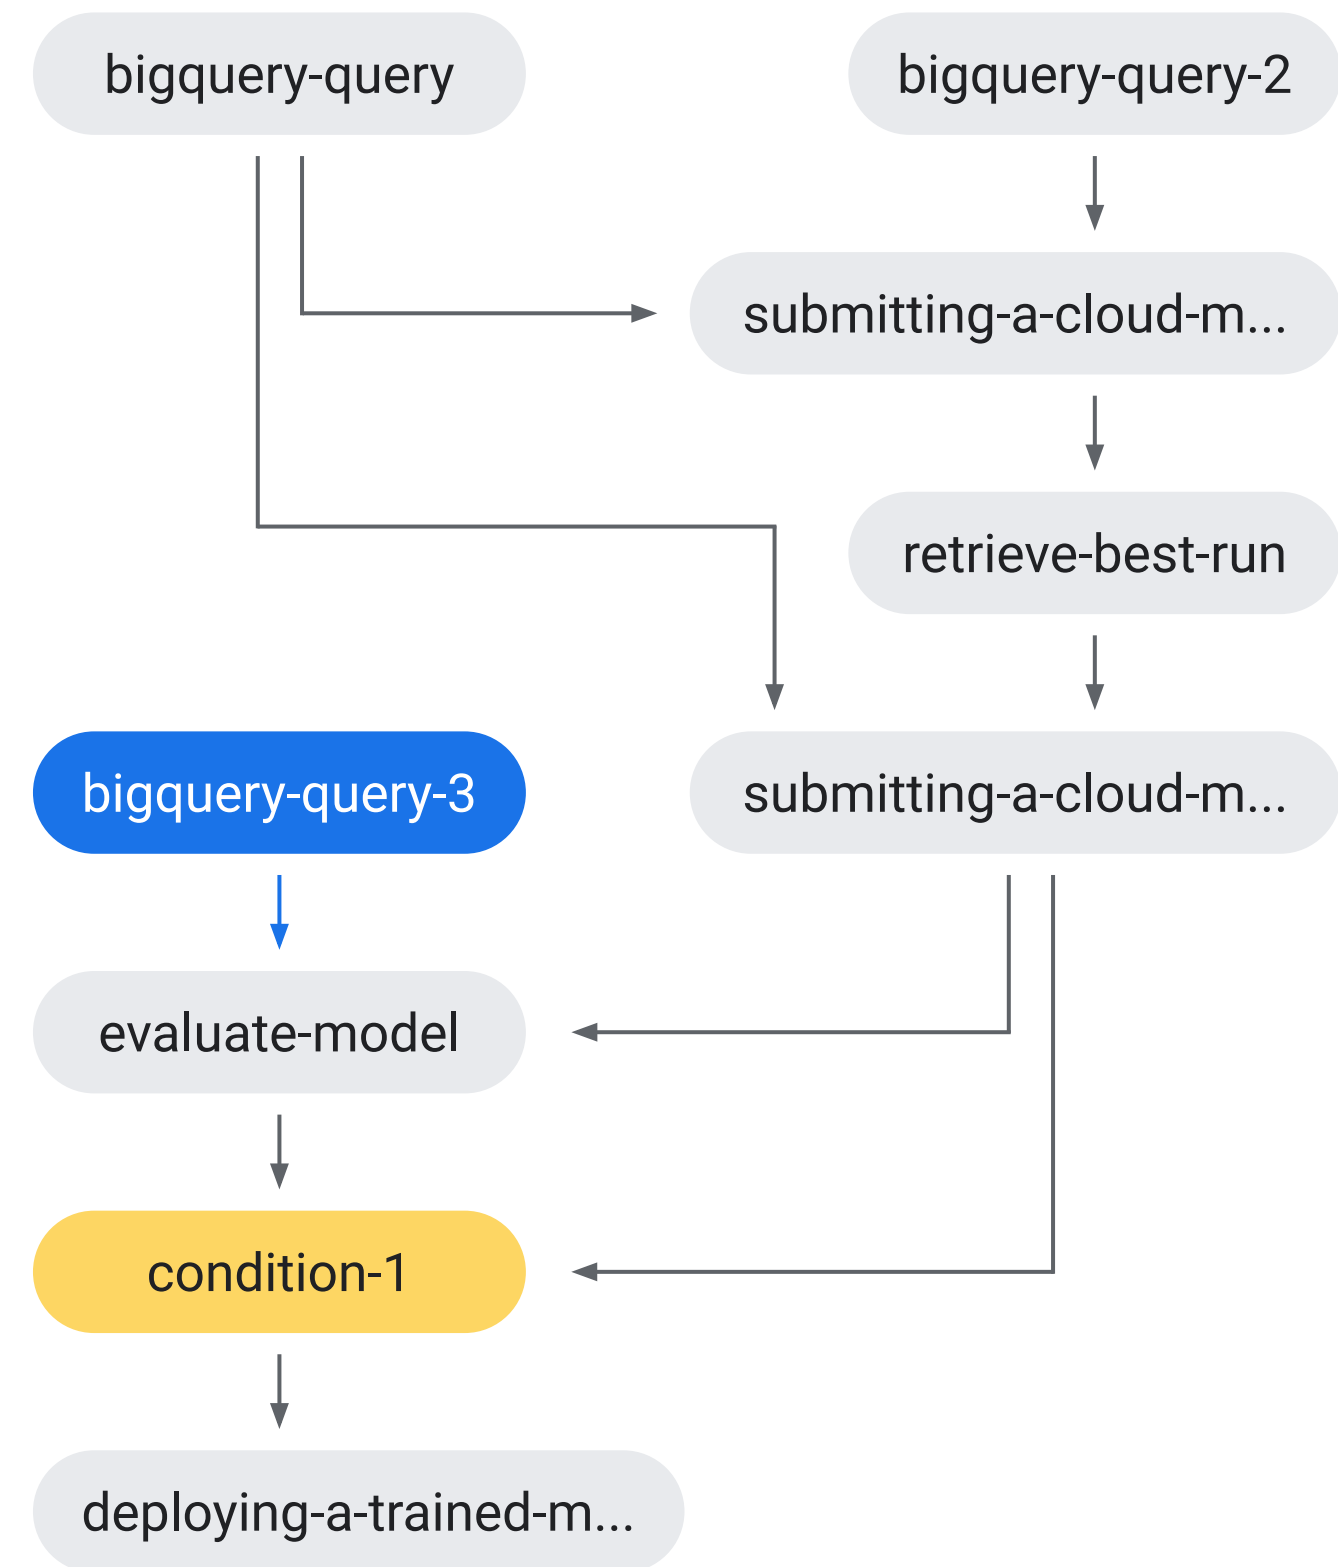
```
def covertime_train(project_id,
                    region,
                    source_table_name,
                    gcs_root,
                    dataset_id,
                    evaluation_metric_name,
                    evaluation_metric_threshold,
                    model_id,
                    version_id,
                    replace_existing_version,
                    hypertune_settings=HYPERTUNE_SETTINGS,
                    dataset_location='US' ):
```

The Run Parameters are supplied at run time.

Define the task DAG within the pipeline function body

```
@kfp.dsl.pipeline(...)
def covertime_train(...):
    # Task DAG defined here
```

1. Create the “ops.”
 2. Compose them into a DAG.
- (OPs = components)




Creation and composition of ops

```
train_model = mlengine_train_op(  
    project_id=project_id,  
    region=region,  
    master_image_uri=TRAINER_IMAGE,  
    job_dir=job_dir,  
    args=train_args)
```

1. Ops creation

2. Ops composition


```
eval_model = evaluate_model_op(  
    dataset_path=str(create_testing_split.outputs['output_gcs_path']),  
    model_path=str(train_model.outputs['job_dir']),  
    metric_name=evaluation_metric_name)
```



Some ops can be triggered **conditionally** to other ops output

```
# Deploy the model if the primary metric is higher than a given threshold
```

```
with kfp.dsl.Condition(eval_model.outputs['metric_value'] >
evaluation_metric_threshold):
    deploy_model = mlengine_deploy_op(
        model_uri=train_model.outputs['job_dir'],
        project_id=project_id,
        model_id=model_id,
        version_id=version_id,
        runtime_version=RUNTIME_VERSION,
        python_version=PYTHON_VERSION,
        replace_existing_version=replace_existing_version)
```



3 main types of Kubeflow components we will look at

01

Pre-built components

- Just load the component from its description and compose.

02


Lightweight Python components

- Implement the component code.

03

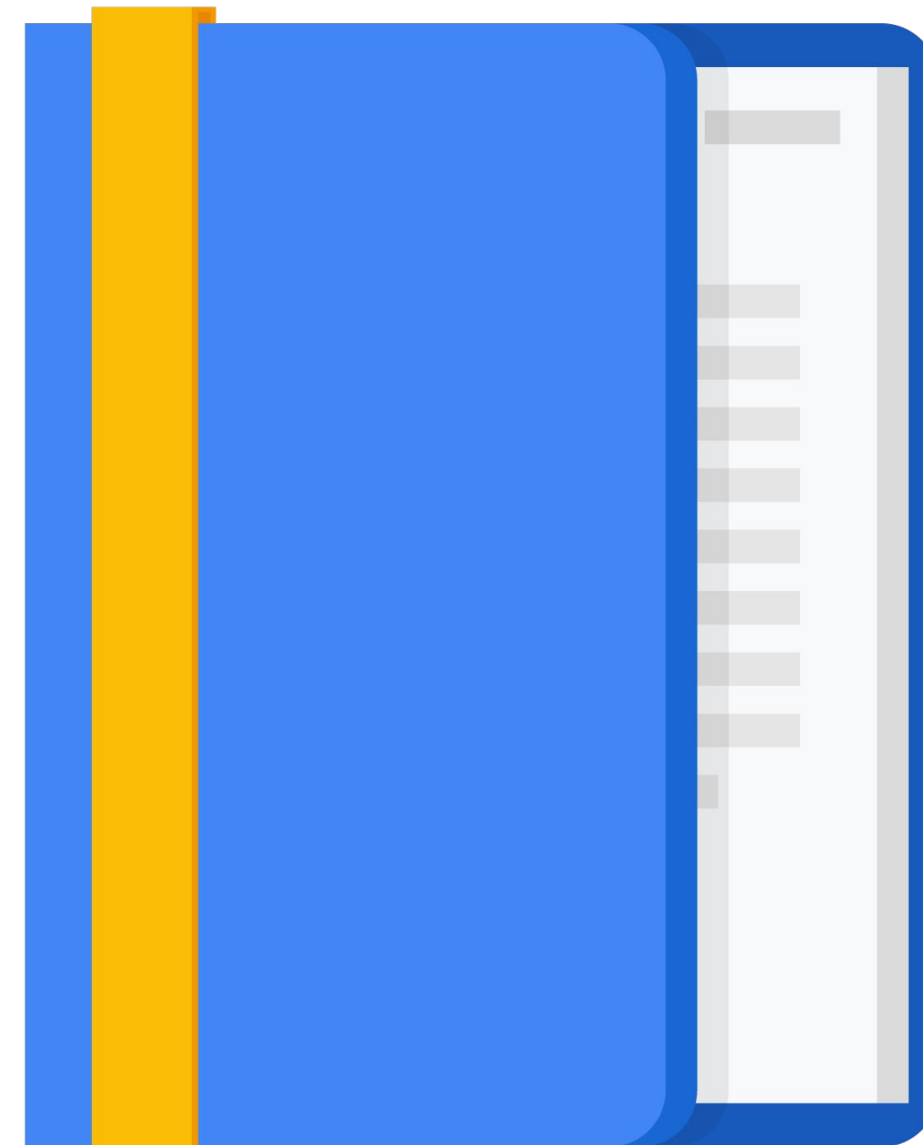
Custom components

- Implement the component code.
- Package it into a Docker container.
- Write the component description.



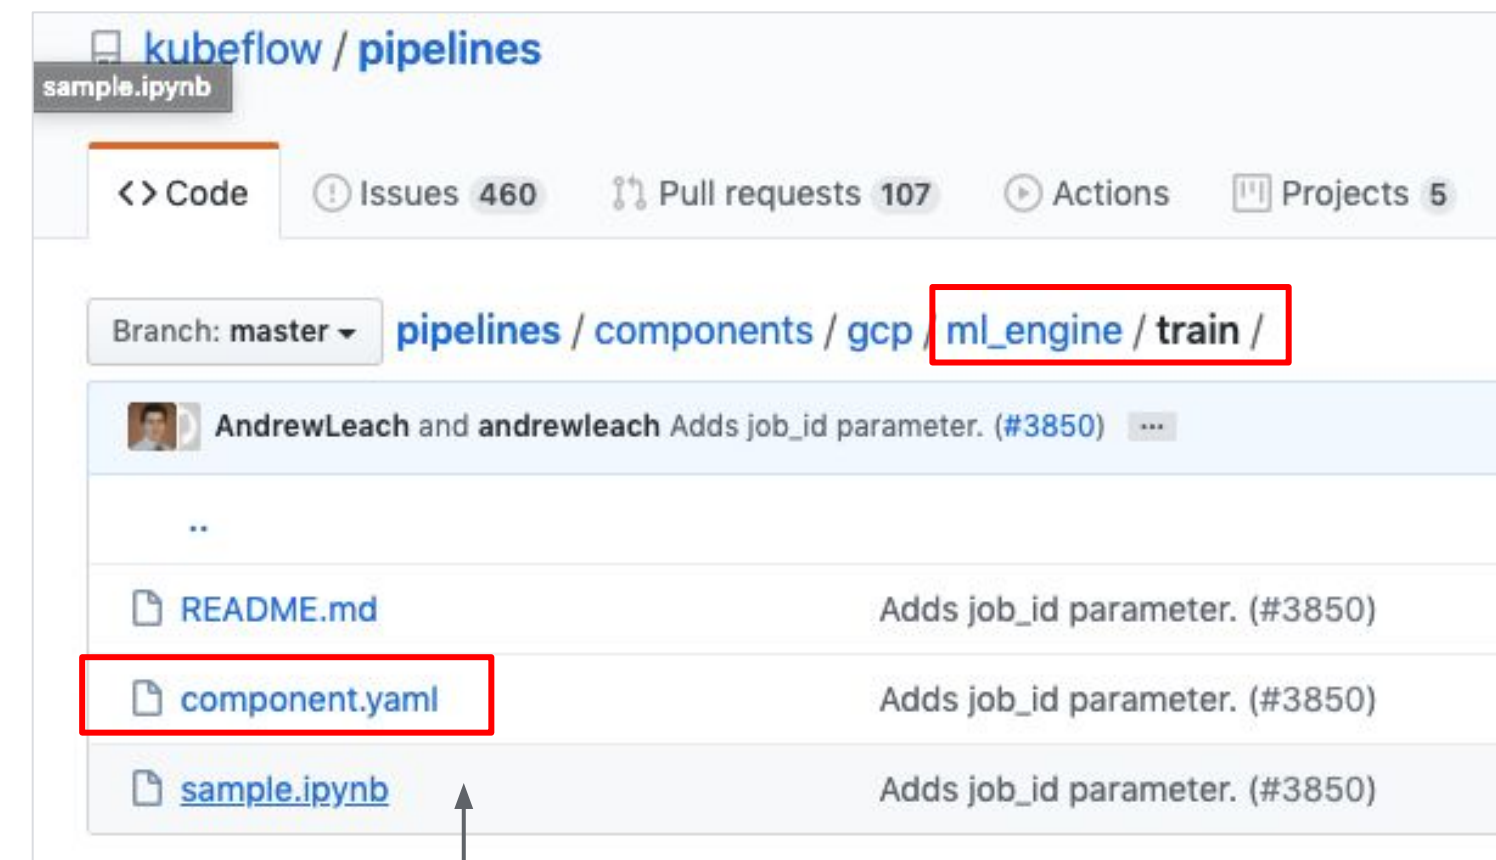
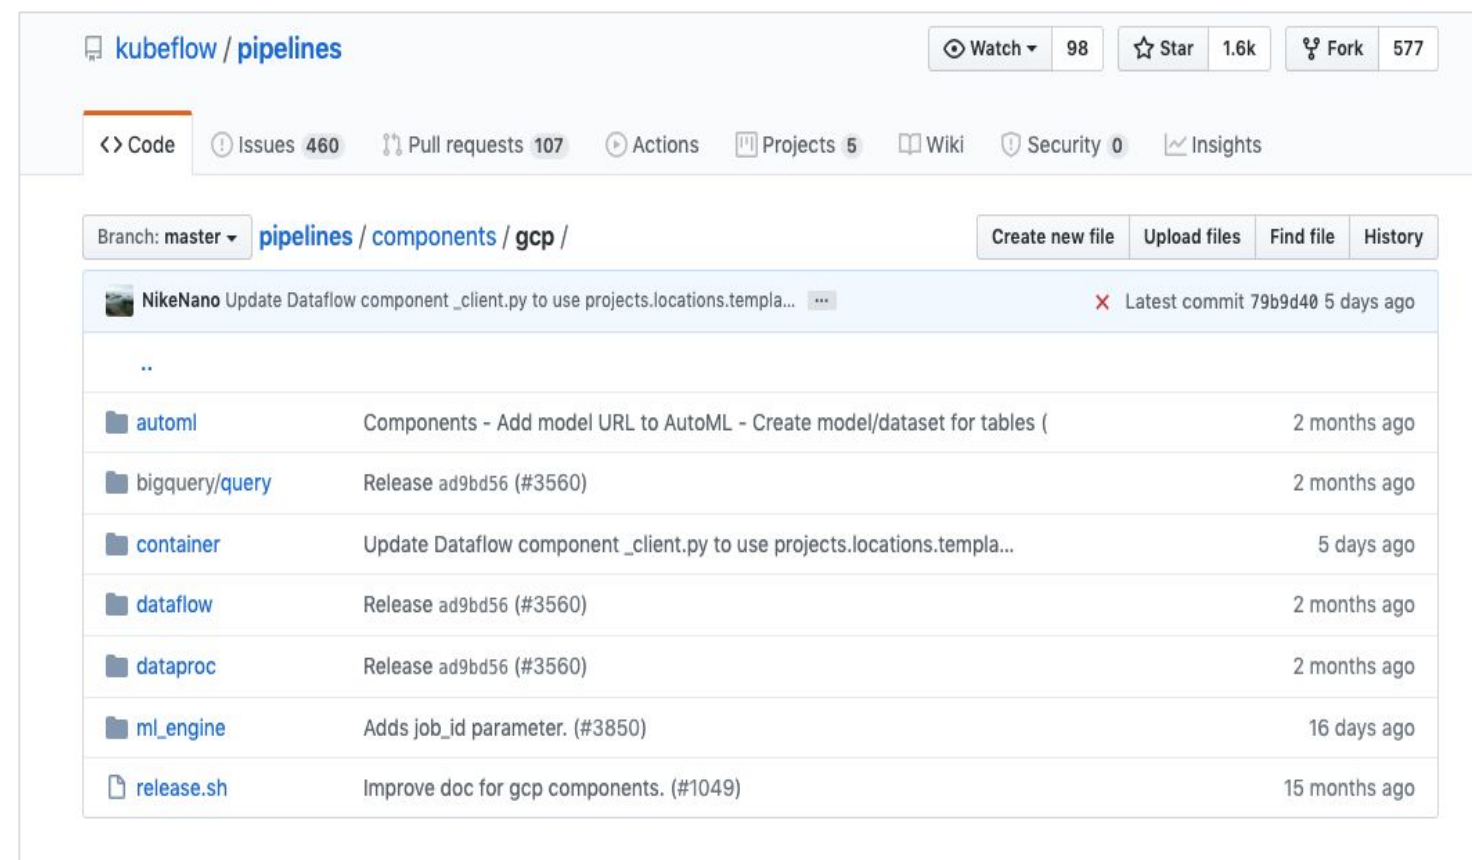
Agenda

- System and Concept Overview
- Describing a Kubeflow Pipeline with KF DSL
- Pre-built Components
- Lightweight Python Components
- Custom Components
- Compile, Upload, and Run



Github repo for pre-built Kubeflow components

<https://github.com/kubeflow/pipelines/blob/master/components>




Component description

component.yaml

```
110 implementation:
111   container:
112     image: gcr.io/ml-pipeline/ml-pipeline-gcp:ad9bd5648dd0453005225779f25d8cebebc7ca00
113     args: [
114       --ui_metadata_path, {outputPath: MLPipeline UI metadata},
115       kfp_component.google.ml_engine, train,
116       --project_id, {inputValue: project_id},
117       --python_module, {inputValue: python_module},
118       --package_uris, {inputValue: package_uris},
119       --region, {inputValue: region},
120       --args, {inputValue: args},
121       --job_dir, {inputValue: job_dir},
122       --python_version, {inputValue: python_version},
123       --runtime_version, {inputValue: runtime_version},
124       --master_image_uri, {inputValue: master_image_uri},
125       --worker_image_uri, {inputValue: worker_image_uri},
126       --training_input, {inputValue: training_input},
127       --job_id_prefix, {inputValue: job_id_prefix},
128       --job_id, {inputValue: job_id},
129       --wait_interval, {inputValue: wait_interval},
130     ]
131   env:
132     KFP_POD_NAME: "{{pod.name}}"
133   fileOutputs:
```

→ Container image URI

→ Run parameters



Loading a pre-built component

```
import kfp
```

```
URI = 'https://raw.githubusercontent.com/kubeflow/pipelines/0.2.5/components/gcp/'
```

```
component_store = kfp.components.ComponentStore(  
    local_search_paths=None, url_search_prefixes=[URI])
```

```
bigquery_query_op = component_store.load_component('bigquery/query')  
mlengine_train_op = component_store.load_component('ml_engine/train')  
mlengine_deploy_op = component_store.load_component('ml_engine/deploy')
```

Using pre-built bigquery/query

```
create_training_split = bigquery_query_op(  
    query=query,  
    project_id=project_id,  
    dataset_id=dataset_id,  
    table_id='',  
    output_gcs_path=training_file_path,  
    dataset_location=dataset_location)
```

Runtime arguments:

Argument	Description	Optional	Data type	Accepted values	Default
query	The query used by BigQuery to fetch the results.	No	String		
project_id	The project ID of the Google Cloud Platform (GCP) project to use to execute the query.	No	GCPProjectID		
dataset_id	The ID of the persistent BigQuery dataset to store the results of the query. If the dataset does not exist, the operation will create a new one.	Yes	String		None

Output:

Name	Description	Type
output_gcs_path	The path to the Cloud Storage bucket containing the query output in CSV format.	GCSPath

<https://github.com/kubeflow/pipelines/tree/master/components/gcp/bigquery/query>

Using pre-built ml_engine/train

```
train_model = mlengine_train_op(  
    project_id=project_id,  
    region=region,  
    master_image_uri=TRAINER_IMAGE,  
    job_dir=job_dir,  
    args=train_args)
```

https://github.com/kubeflow/pipelines/tree/master/components/gcp/ml_engine/train

Runtime arguments:

Argument	Description	Optional	Data type	Accepted values	Default
project_id	The Google Cloud Platform (GCP) project ID of the job.	No	GCPProjectID	-	-
python_module	The name of the Python module to run after installing the training program.	Yes	String	-	None
package_uris	The Cloud Storage location of the packages that contain the training program and any additional dependencies. The maximum number of package URIs is 100.	Yes	List	-	None

Output:

Name	Description	Type
job_id	The ID of the created job.	String
job_dir	The Cloud Storage path that contains the output files with the trained model.	GCSPath

Using pre-built ml_engine/deploy

```
deploy_model = mlengine_deploy_op(  
  
model_uri=train_model.outputs['job_dir'],  
    project_id=project_id,  
    model_id=model_id,  
    version_id=version_id,  
    runtime_version=RUNTIME_VERSION,  
    python_version=PYTHON_VERSION,  
  
replace_existing_version=replace_existing_ve  
rsion)
```

https://github.com/kubeflow/pipelines/tree/master/components/gcp/ml_engine/deploy

Runtime arguments:

Argument	Description	Optional	Data type	Accepted values	Default
model_uri	The URI of a Cloud Storage directory that contains a trained model file. Or An Estimator export base directory that contains a list of subdirectories named by timestamp. The directory with the latest timestamp is used to load the trained model file.	No	GCSPath		
project_id	The ID of the Google Cloud Platform (GCP) project of the serving model.	No	GCPProjectID		
model_id	The name of the trained model.	Yes	String		None
version_id	The name of the version of the model. If it is not provided, the operation uses a random name.	Yes	String		None

Output:

Name	Description	Type
model_uri	The Cloud Storage URI of the trained model.	GCSPath
model_name	The name of the deployed model.	String
version_name	The name of the deployed version.	String


Composing pre-built components: Hyper tuning

```
tune_args = [  
    '--training_dataset_path',  
    create_training_split.outputs['output_gcs_path'],  
    '--validation_dataset_path',  
    create_validation_split.outputs['output_gcs_path'],  
    '--hptune', 'True'  
]
```

```
hypertune = mlengine_train_op(  
    project_id=project_id,  
    region=region,  
    master_image_uri=TRAINER_IMAGE,  
    job_dir=job_dir,  
    args=tune_args,  
    training_input=HYPERTUNE_SETTINGS)
```

Composing pre-built components: Hypertuning

```
HYPERTUNE_SETTINGS = """
{
    "hyperparameters": {
        "goal": "MAXIMIZE",
        "maxTrials": 6,
        "maxParallelTrials": 3,
        "hyperparameterMetricTag": "accuracy",
        "enableTrialEarlyStopping": True,
        "params": [
            {
                "parameterName": "max_iter",
                "type": "DISCRETE",
                "discreteValues": [500, 1000]
            },
            etc.
        ]
    }
}
"""
```



Composing pre-built components: Training best run

```
train_args = [  
    '--training_dataset_path', create_training_split.outputs['output_gcs_path'],  
    '--validation_dataset_path', create_validation_split.outputs['output_gcs_path'],  
    '--alpha', get_best_trial.outputs['alpha'],  
    '--max_iter', get_best_trial.outputs['max_iter'],  
    '--hptune', 'False'  
]
```

```
train_model = mlengine_train_op(  
    project_id=project_id,  
    region=region,  
    master_image_uri=TRAINER_IMAGE,  
    job_dir=job_dir,  
    args=train_args,  
)
```

Components composed through
their input/output



← ✓ Run_001

Graph

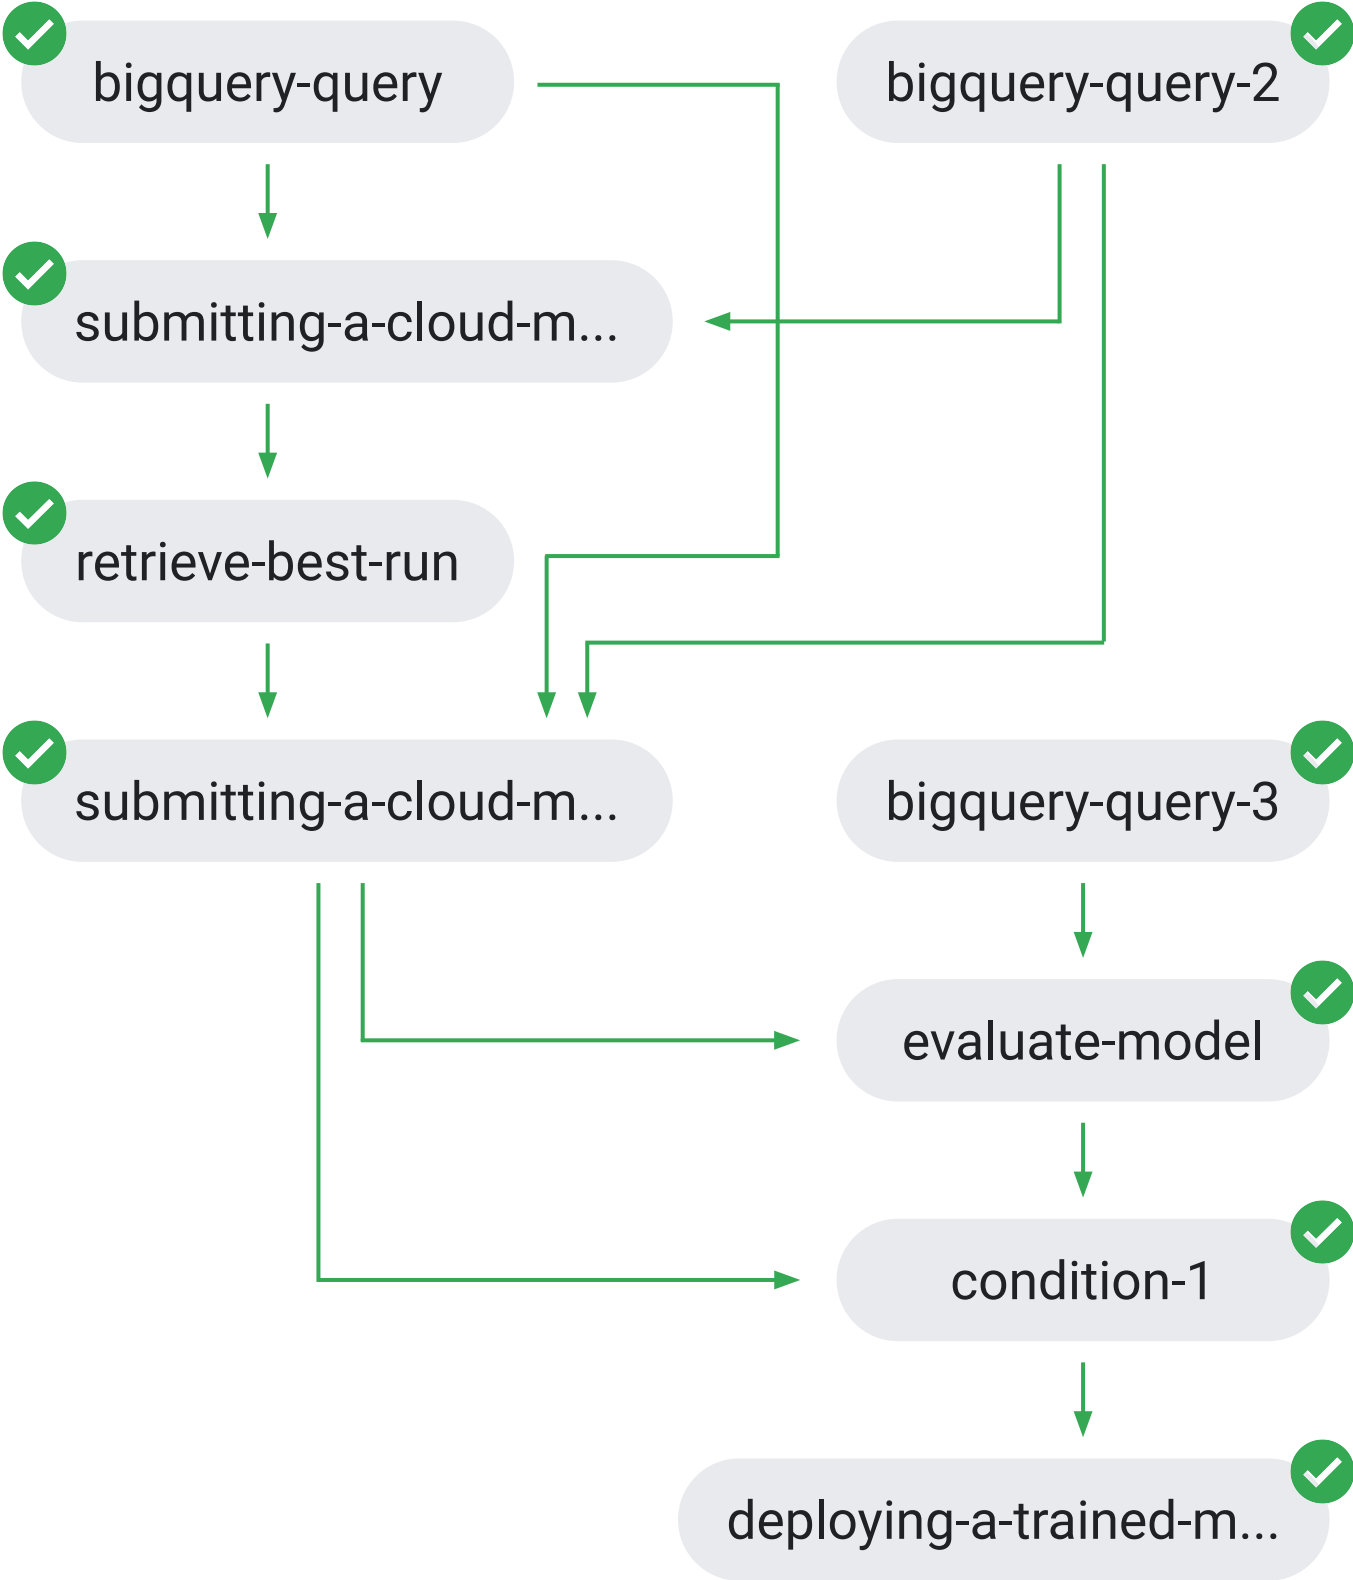
Run output

Config

Metrics

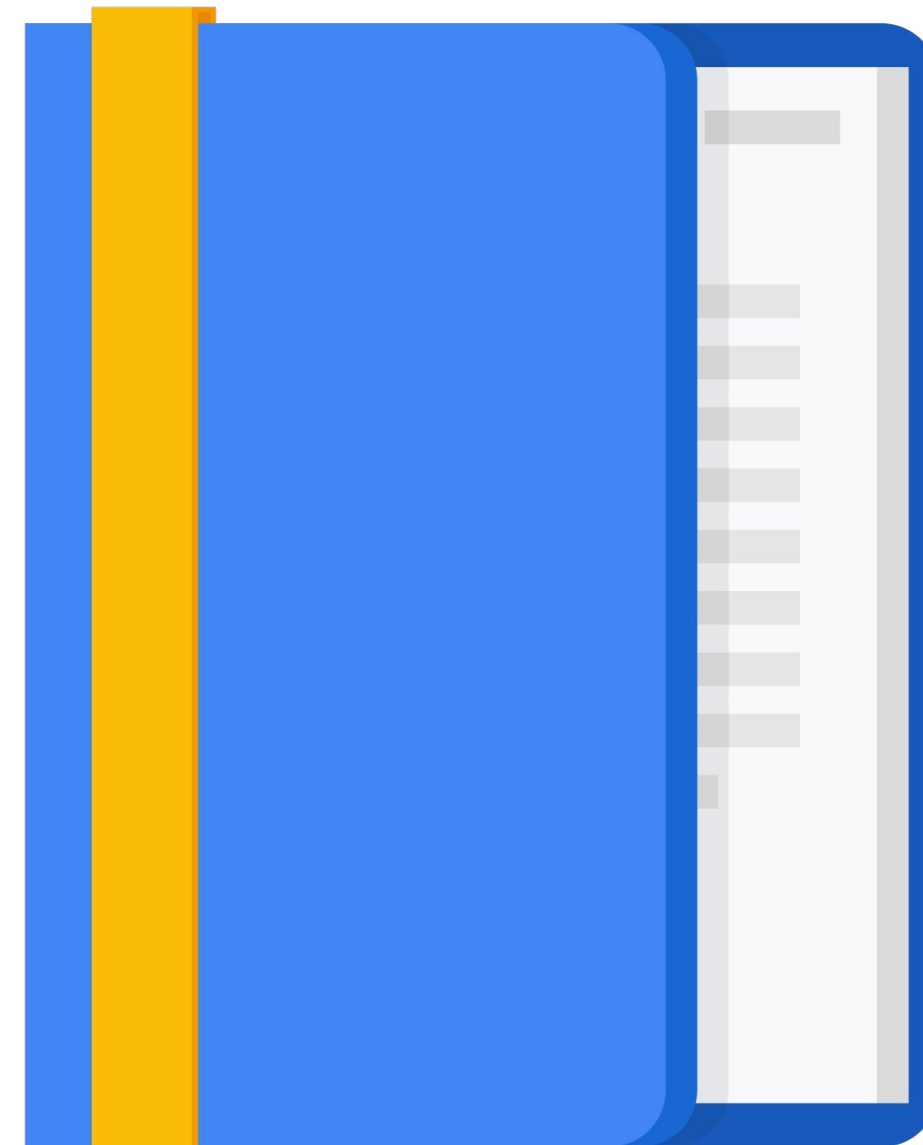
	accuracy
evaluate-model	0.721

No output artifacts found for this run.



Agenda

- System and Concept Overview
- Describing a Kubeflow Pipeline with KF DSL
- Pre-built Components
- **Lightweight Python Components**
- Custom Components
- Compile, Upload, and Run



Wrap Python functions into KF components

helper_components.py

```
def retrieve_best_run(project_id, job_id):  
    """Retrieves the parameters of the best Hypertune run."""  
    # [...]  
    return (metric_value, alpha, max_iter)  
  
def evaluate_model(dataset_path, model_path, metric_name):  
    """Evaluates a trained sklearn model."""  
    # [...]  
    return (metric_name, metric_value, json.dumps(metrics))
```

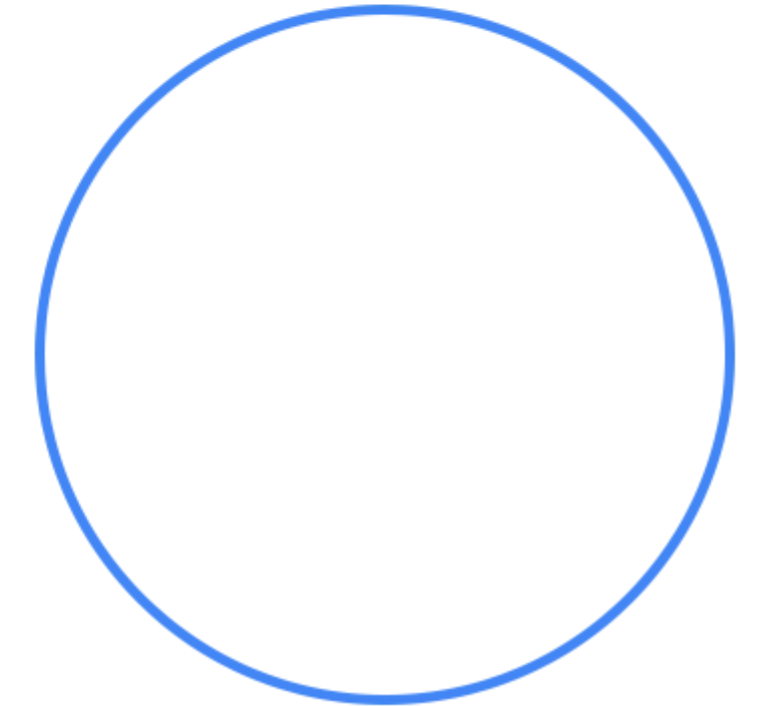
func_to_container_op

```
from helper_components import evaluate_model
from helper_components import retrieve_best_run

from kfp.components import func_to_container_op

retrieve_best_run_op = func_to_container_op(
    retrieve_best_run, base_image=BASE_IMAGE)


evaluate_model_op = func_to_container_op(
    evaluate_model, base_image=BASE_IMAGE)
```



Use and compose the lightweight components as usual


```
get_best_trial = retrieve_best_run_op(  
    project_id, hypertune.outputs['job_id'])
```

```
eval_model = evaluate_model_op(  
    dataset_path=str(create_testing_split.outputs['output_gcs_path']),  
    model_path=str(train_model.outputs['job_dir']),  
    metric_name=evaluation_metric_name)
```



Agenda

- System and Concept Overview
- Describing a Kubeflow Pipeline with KF DSL
- Pre-built Components
- Lightweight Python Components
- Custom Components
- Compile, Upload, and Run



Custom components

Steps

1. Write your own code (in any language).
2. Write a Dockerfile to package it into a Docker container.
3. Build and push your Docker container image to gcr/io.
4. Describe your component in a yaml description file.
5. Use the description file to load the component into the pipeline.

Step 1: Write your custom code (example)

main.py

```
from google.cloud import bigquery
from google.cloud.bigquery.job import ExtractJobConfig

if __name__ == "__main__":

    bq = bigquery.Client()

    dataset_ref = bigquery.Dataset(bq.dataset("taxifare"))
    bq.create_dataset(dataset_ref)

    bq.query(TRAIN_SQL).result()
    bq.query(VALID_SQL).result()

    export_table_to_gcs(dataset_ref, TRAIN_TABLE, train_export_path)
```

Step 2: Package it into a Docker container

Dockerfile

```
FROM google/cloud-sdk:latest

RUN apt-get update && \
    apt-get install --yes python3-pip

COPY . /code
WORKDIR /code

RUN pip3 install google-cloud-bigquery

ENTRYPOINT ["python3", "./main.py"]
```

Step 3: Write the component description

`bq2gcs.yaml`

```
name: bq2gcs
```

```
description: |
```

```
  This component creates the training and  
  validation datasets as BiqQuery tables and exports  
  them into Cloud Storage at gs://<BUCKET>/taxifare/data.
```

```
inputs:
```

```
  - {name: Input Bucket , type: String, description: 'GCS directory path.'}
```

```
implementation:
```

```
  container:
```

```
    image: gcr.io/<YOUR PROJECT>/taxifare-bq2gcs
```

```
    args: ["--bucket", {inputValue: Input Bucket}]
```

Step 4: Load the component into the pipeline

```
@dsl.pipeline(  
    name='Taxifare',  
    description='Train an ML model to predict the taxi fare in NY')  
def pipeline(gcs_bucket_name='<bucket where data and model will be exported>'):  
  
    bq2gcs_op = comp.load_component_from_file(BQ2GCS_YAML)  
    bq2gcs = bq2gcs_op(  
        input_bucket=gcs_bucket_name,  
    )  
  
    [...]
```

Run the pipeline with custom components!

```
import kfp

client = kfp.Client(host=HOST)

client.list_experiments()


exp = client.create_experiment(name='taxifare')

kfp.compiler.Compiler().compile(pipeline, PIPELINE_TAR)

run = client.run_pipeline(
    experiment_id=exp.id,
    job_name='taxifare',
    pipeline_package_path='taxifare.tar.gz',
    params={
        'gcs_bucket_name': BUCKET,
    },
)
```

Agenda

- System and Concept Overview
- Describing a Kubeflow Pipeline with KF DSL
- Pre-built Components
- Lightweight Python Components
- Custom Components
- Compile, Upload, and Run



Step 1: Build and push the trainer container

trainer_image/Dockerfile

Required by ml_engine/train

```
FROM gcr.io/deeplearning-platform-release/base-cpu
RUN pip install -U fire cloudml-hypertune scikit-learn==0.20.4 pandas==0.24.2

WORKDIR /app

COPY train.py .

ENTRYPOINT ["python", "train.py"]
```

```
TRAINER_IMAGE='gcr.io/PROJECT_ID/TRAINER_IMAGE_NAME:TAG'
```

```
gcloud builds submit --timeout 15m --tag $TRAINER_IMAGE trainer_image
```

Step 2: Build and push the base container

base_image/Dockerfile

Required by the Python lightweight ops

```
FROM gcr.io/deeplearning-platform-release/base-cpu
```

```
RUN pip install -U fire scikit-learn==0.20.4 pandas==0.24.2 kfp==0.2.5
```

```
BASE_IMAGE='gcr.io/PROJECT_ID/BASE_IMAGE_NAME:TAG'
```

```
gcloud builds submit --timeout 15m --tag $BASE_IMAGE base_image
```

Step 3: Compile the Kubeflow pipeline

`dsl-compile --py pipeline/covertime_training_pipeline.py --output covertime_training_pipeline.yaml`


```
[16]: !head covertime_training_pipeline.yaml
```

```
"apiVersion": |-
  argoproj.io/v1alpha1
"kind": |-
  Workflow
"metadata":
  "annotations":
    "pipelines.kubeflow.org/pipeline_spec": |-
      {"description": "The pipeline training and deploying the Covertime classifierpipeline_yaml", "inputs": [{"name": "project_id"}, {"name": "region"}, {"name": "source_table_name"}, {"name": "gcs_root"}, {"name": "dataset_id"}, {"name": "evaluation_metric_name"}, {"name": "evaluation_metric_threshold"}, {"name": "model_id"}, {"name": "version_id"}, {"name": "replace_existing_version"}, {"default": "\n{\n  \"hyperparameters\": {\n    \"goal\": \"MAXIMIZE\", \n    \"maxTrials\": 6, \n    \"maxParallelTrials\": 3, \n    \"hyperparameterMetricTag\": \"accuracy\", \n    \"enableTrialEarlyStopping\": True, \n    \"params\": [\n      {\n        \"parameterName\": \"max_iter\", \n        \"type\": \"DISCRETE\", \n        \"discreteValues\": [500, 1000] \n      }, \n      {\n        \"parameterName\": \"alpha\", \n        \"type\": \"DOUBLE\", \n        \"minValue\": 0.0001, \n        \"maxValue\": 0.001, \n        \"scaleType\": \"UNIT_LINEAR_SCALE\" \n      } \n    ] \n  } \n} \n\", \"name\": \"hypertune_settings\", \"optional\": true}, {\"default\": \"US\", \"name\": \"dataset_location\", \"optional\": true}], \"name\": \"Covertime Classifier Training\"}
    "generateName": |-
      covertime-classifier-training-
```

Step 4: Upload the pipeline to the KF cluster

```
kfp --endpoint $ENDPOINT pipeline upload -p $PIPELINE_NAME \
covertime_training_pipeline.yaml
```

```
kfp --endpoint $ENDPOINT pipeline list
```



Step 5: Run the pipeline

```
kfp --endpoint $ENDPOINT run submit \  
-e $EXPERIMENT_NAME \  
-r $RUN_ID \  
-p $PIPELINE_ID \  
project_id=$PROJECT_ID \  
gcs_root=$GCS_STAGING_PATH \  
region=$REGION \  
source_table_name=$SOURCE_TABLE \  
dataset_id=$DATASET_ID \  
evaluation_metric_name=$EVALUATION_METRIC \  
evaluation_metric_threshold=$EVALUATION_METRIC_THRESHOLD \  
model_id=$MODEL_ID \  
version_id=$VERSION_ID \  
replace_existing_version=$REPLACE_EXISTING_VERSION
```

Run
parameters


Lab

Kubeflow Pipelines on CAIP

In this lab, you will learn how to use AI Platform Pipelines to build a Kubeflow pipeline to train, tune, and serve a model automatically.



<https://github.com/GoogleCloudPlatform/mlops-on-gcp/blob/master/workshops/kfp-caip-sklearn/lab-02-kfp-pipeline/exercises/lab-02.ipynb>



cloud.google.com